

Problem statement

predicting the house price in USA.To create a model to help him estimate of what the house would sell for.

```
In [1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [2]: df=pd.read_csv("Sales")
```

To display top 10 rows

```
In [3]: df.head(10)
```

Out[3]:

	MonthYear	Time index	Country	StoreID	City	Dept_ID	Dept. Name	HoursOwn	HoursLease	Sales units	Turnover	Customer	Area (m2)
0	10.2016	1.0	United Kingdom	88253.0	London (I)	1.0	Dry	3184.764	0.0	398560.0	1226244.0	NaN	953.04
1	10.2016	1.0	United Kingdom	88253.0	London (I)	2.0	Frozen	1582.941	0.0	82725.0	387810.0	NaN	720.48
2	10.2016	1.0	United Kingdom	88253.0	London (I)	3.0	other	47.205	0.0	438400.0	654657.0	NaN	966.72
3	10.2016	1.0	United Kingdom	88253.0	London (I)	4.0	Fish	1623.852	0.0	309425.0	499434.0	NaN	1053.36
4	10.2016	1.0	United Kingdom	88253.0	London (I)	5.0	Fruits & Vegetables	1759.173	0.0	165515.0	329397.0	NaN	1053.36
5	10.2016	1.0	United Kingdom	88253.0	London (I)	6.0	Meat	8270.316	0.0	1713310.0	5617137.0	NaN	11735.16
6	10.2016	1.0	United Kingdom	88253.0	London (I)	13.0	Food	16468.251	0.0	3107935.0	8714679.0	NaN	19865.64
7	10.2016	1.0	United Kingdom	88253.0	London (I)	7.0	Clothing	4698.471	0.0	213680.0	1615341.0	NaN	8513.52
8	10.2016	1.0	United Kingdom	88253.0	London (I)	8.0	Household	1183.272	0.0	54915.0	290400.0	NaN	4842.72
9	10.2016	1.0	United Kingdom	88253.0	London (I)	9.0	Hardware	2029.815	0.0	59260.0	450015.0	NaN	5608.8

Data Cleaning And Pre-Processing

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7658 entries, 0 to 7657
Data columns (total 14 columns):
#   Column          Non-Null Count  Dtype
---  -
0   MonthYear       7658 non-null   object
1   Time index      7650 non-null   float64
2   Country         7650 non-null   object
3   StoreID         7650 non-null   float64
4   City            7650 non-null   object
5   Dept_ID         7650 non-null   float64
6   Dept. Name      7650 non-null   object
7   HoursOwn        7650 non-null   object
8   HoursLease      7650 non-null   float64
9   Sales units     7650 non-null   float64
10  Turnover        7650 non-null   float64
11  Customer        0 non-null      float64
12  Area (m2)       7650 non-null   object
13  Opening hours   7650 non-null   object
dtypes: float64(7), object(7)
memory usage: 837.7+ KB
```

```
In [5]: # Display the statistical summary
df.describe()
```

Out[5]:

	Time index	StoreID	Dept_ID	HoursLease	Sales units	Turnover	Customer
count	7650.000000	7650.000000	7650.000000	7650.000000	7.650000e+03	7.650000e+03	0.0
mean	5.000000	61995.220000	9.470588	22.036078	1.076471e+06	3.721393e+06	NaN
std	2.582158	29924.581631	5.337429	133.299513	1.728113e+06	6.003380e+06	NaN
min	1.000000	12227.000000	1.000000	0.000000	0.000000e+00	0.000000e+00	NaN
25%	3.000000	29650.000000	5.000000	0.000000	5.457125e+04	2.726798e+05	NaN
50%	5.000000	75400.500000	9.000000	0.000000	2.932300e+05	9.319575e+05	NaN
75%	7.000000	87703.000000	14.000000	0.000000	9.175075e+05	3.264432e+06	NaN
max	9.000000	98422.000000	18.000000	3984.000000	1.124296e+07	4.271739e+07	NaN

```
In [6]: # To display the col headings  
df.columns
```

```
Out[6]: Index(['MonthYear', 'Time index', 'Country', 'StoreID', 'City', 'Dept_ID',  
              'Dept. Name', 'HoursOwn', 'HoursLease', 'Sales units', 'Turnover',  
              'Customer', 'Area (m2)', 'Opening hours'],  
             dtype='object')
```

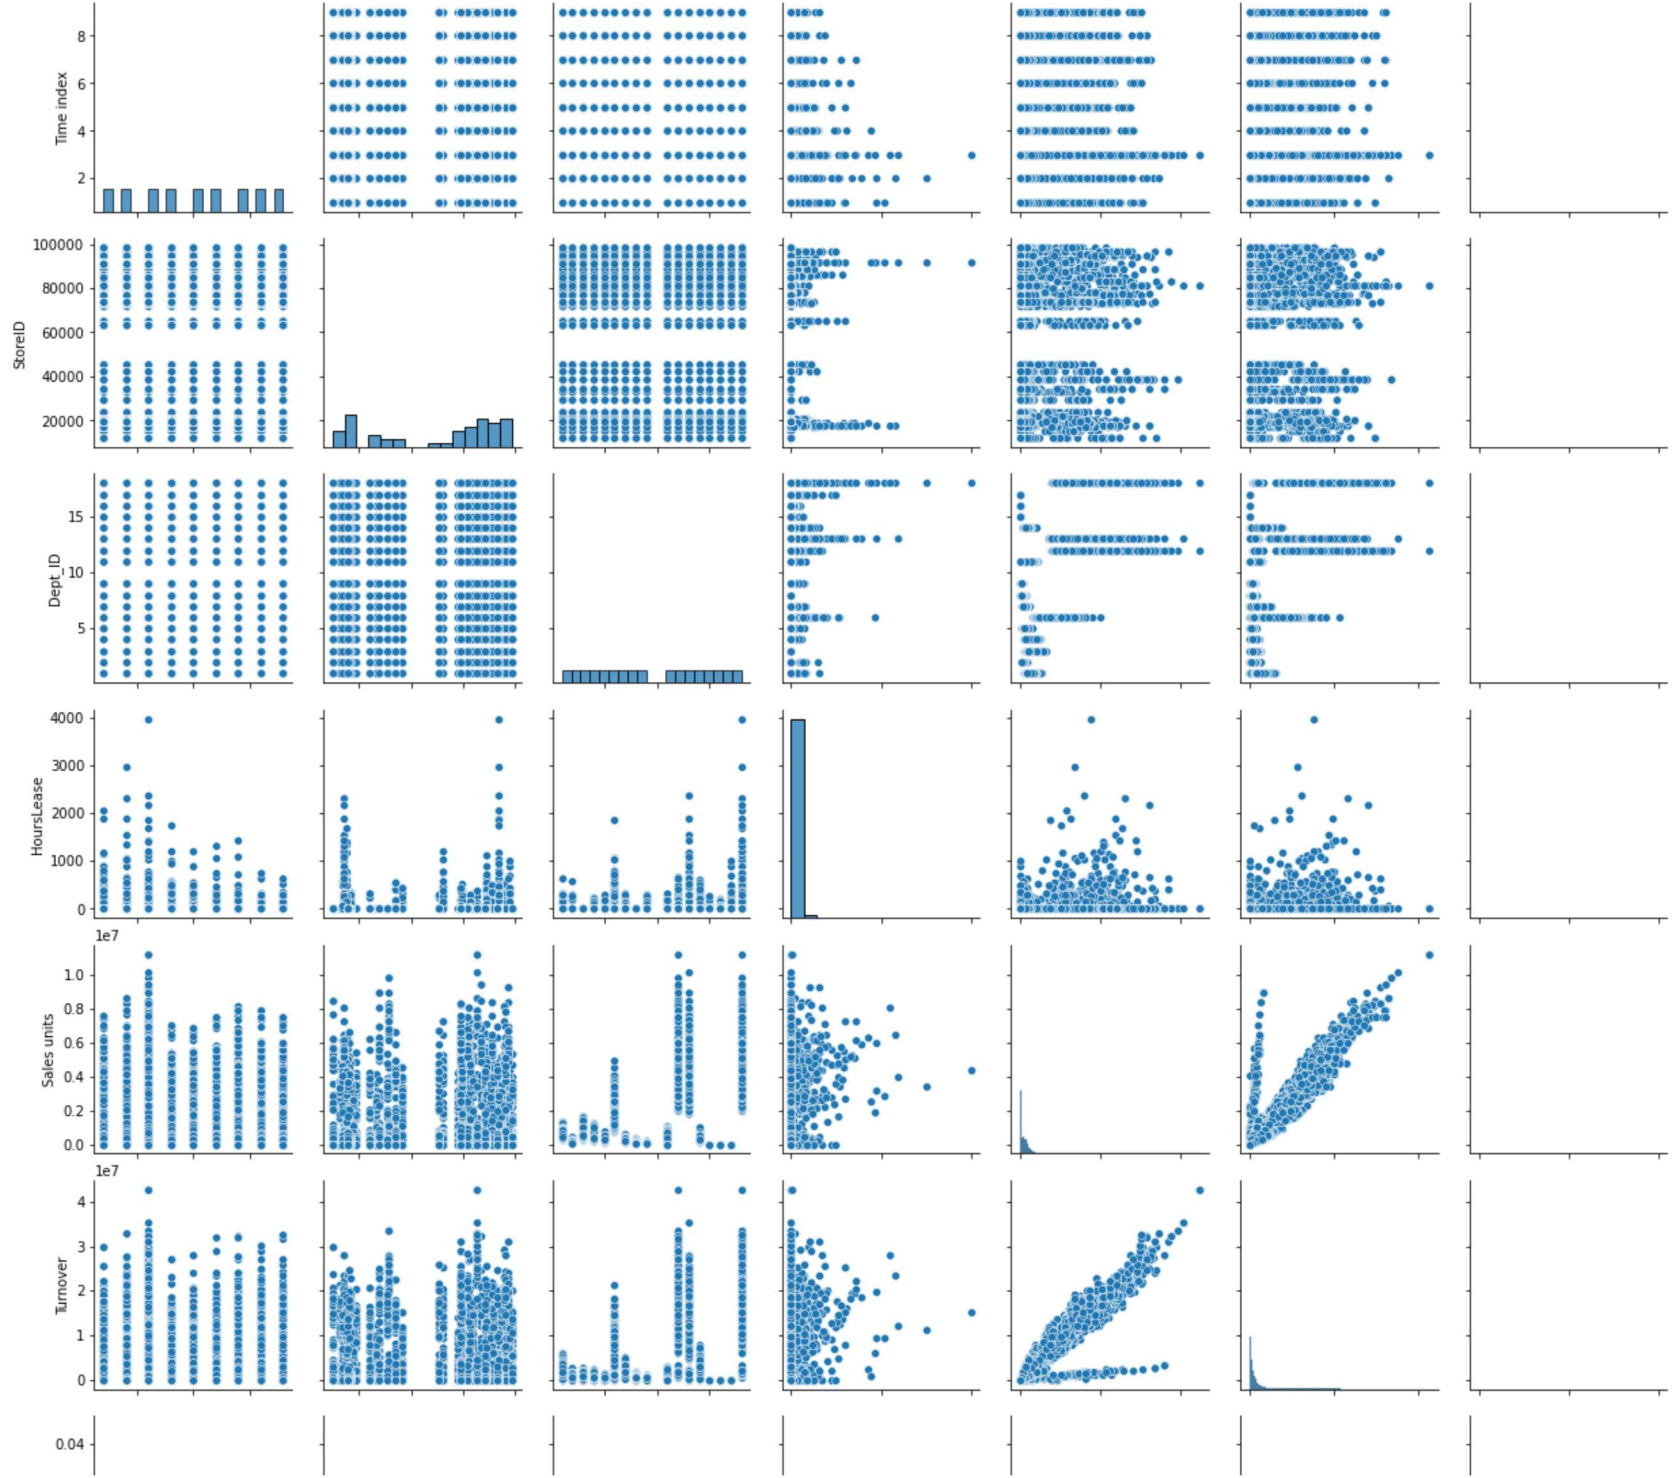
```
In [7]: #cols=df.dropna(axis=1)
```

```
In [8]: #cols.columns
```

EDA and Visualization

```
In [9]: sns.pairplot(df)
```


```
Out[9]: <seaborn.axisgrid.PairGrid at 0x16135d858e0>
```



```
In [10]: sns.displot(df['Sales units'])
```


```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x16139f2dbb0>
```




```
In [11]: # We use displot in older version we get distplot use displot
sns.distplot(df['Sales units'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[11]: <AxesSubplot:xlabel='Sales units', ylabel='Density'>
```



```
In [12]: df1=df[['HoursLease', 'Sales units', 'Turnover',  
               'Customer']]  
df1
```

Out[12]:

	HoursLease	Sales units	Turnover	Customer
0	0.0	398560.0	1226244.0	NaN
1	0.0	82725.0	387810.0	NaN
2	0.0	438400.0	654657.0	NaN
3	0.0	309425.0	499434.0	NaN
4	0.0	165515.0	329397.0	NaN
...
7653	0.0	3886530.0	14538825.0	NaN
7654	0.0	245.0	0.0	NaN
7655	0.0	0.0	0.0	NaN
7656	0.0	245.0	0.0	NaN
7657	0.0	3886530.0	15056214.0	NaN

7658 rows × 4 columns

```
In [13]: sns.heatmap(df1.corr())
```

```
Out[13]: <AxesSubplot:>
```



To train the model - MODEL BUILD

Going to train linear regression model; We split our data into 2 variables x and y where x is independent var(input) and y is dependent on x(output), we could ignore address col as it is not required for our model

```
In [14]: x=df1[['HoursLease', 'Sales units', 'Turnover',  
              'Customer']]  
y=df1[['Sales units']]
```

To split the dataset into test data

```
In [15]: # importing lib for splitting test data  
from sklearn.model_selection import train_test_split
```

```
In [16]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [17]: from sklearn.linear_model import LinearRegression
```

```
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

[illegible]

```
65          # extra_args > 0
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in check_array(array, accept_sparse,
accept_large_sparse, dtype, order, copy, force_all_finite, ensure_2d, allow_nd, ensure_min_samples, ensure_m
in_features, estimator)
```

```
661
662         if force_all_finite:
--> 663             _assert_all_finite(array,
664                               allow_nan=force_all_finite == 'allow-nan')
665
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py in _assert_all_finite(X, allow_nan, m
sg_dtype)
```

```
101         not allow_nan and not np.isfinite(X).all()):
102         type_err = 'infinity' if allow_nan else 'NaN, infinity'
--> 103         raise ValueError(
104             msg_err.format
105             (type_err,
```

ValueError: Input contains NaN, infinity or a value too large for dtype('float64').

```
In [ ]: print(lr.intercept_)
```

```
In [ ]: print(lr.score(x_test,y_test))
```

```
In [ ]: coeff=pd.DataFrame(lr.coef_)
coeff
```

```
In [ ]: pred = lr.predict(x_test)
plt.scatter(y_test,pred)
```

```
In [ ]: cols=df.dropna()
cols
```

```
In [ ]:
```

