

Problem statement

predicting the house price in USA.To create a model to help him estimate of what the house would sell for.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 df=pd.read_csv("placement")
```

To display top 10 rows

```
In [3]: 1 df.head(10)
```

Out[3]:

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0
5	7.30	23.0	1
6	6.69	11.0	0
7	7.12	39.0	1
8	6.45	38.0	0
9	7.75	94.0	1

Data Cleaning And Pre-Processing

In [4]:

```
1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   cgpa                  1000 non-null   float64
1   placement_exam_marks 1000 non-null   float64
2   placed                1000 non-null   int64   
dtypes: float64(2), int64(1)
memory usage: 23.6 KB
```

In [5]:

```
1 # Display the statistical summary
2 df.describe()
```

Out[5]:

	cgpa	placement_exam_marks	placed
count	1000.000000	1000.000000	1000.000000
mean	6.961240	32.225000	0.489000
std	0.615898	19.130822	0.500129
min	4.890000	0.000000	0.000000
25%	6.550000	17.000000	0.000000
50%	6.960000	28.000000	0.000000
75%	7.370000	44.000000	1.000000
max	9.120000	100.000000	1.000000

In [6]:

```
1 # To display the col headings
2 df.columns
```

Out[6]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')

```
In [7]: 1 cols=df.dropna(axis=1)
```

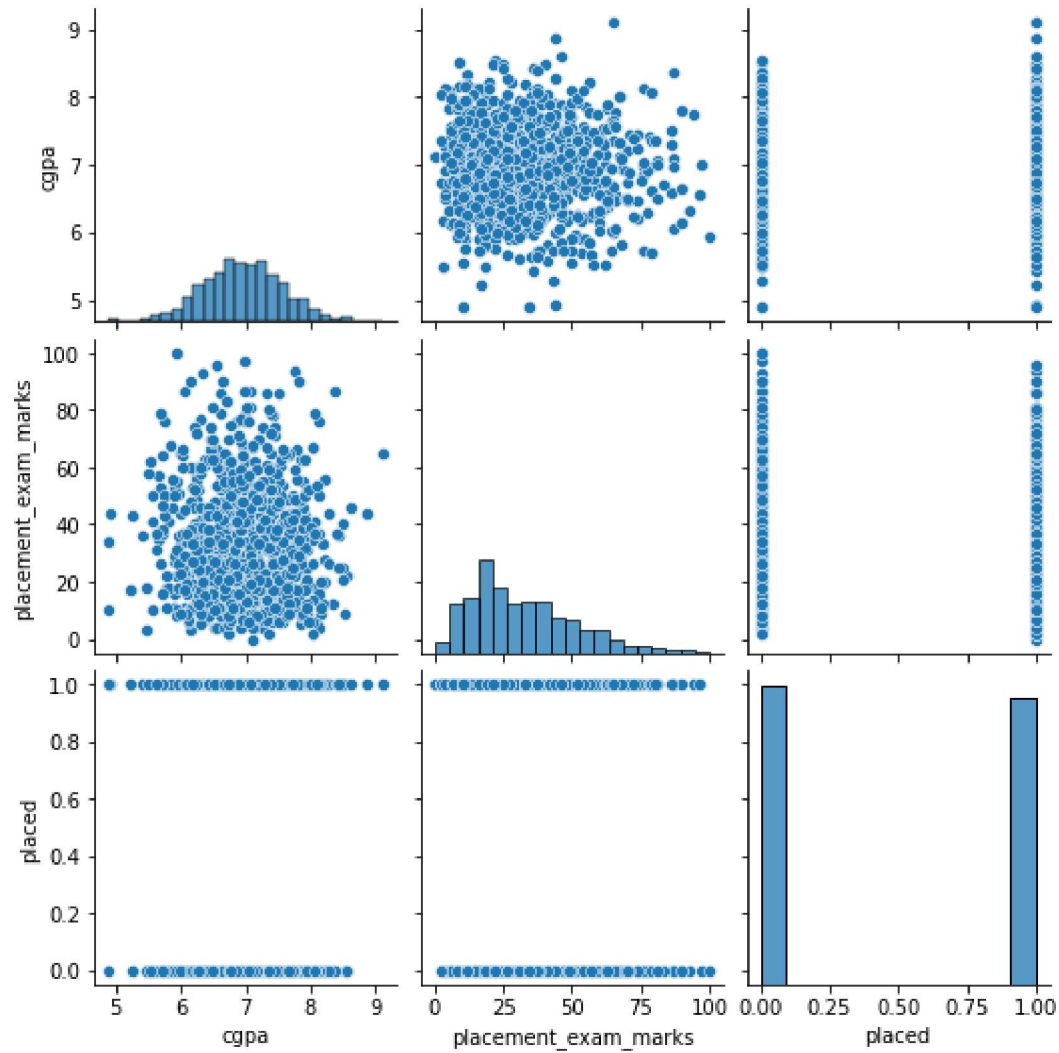
```
In [8]: 1 cols.columns
```

```
Out[8]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')
```

EDA and Visualization


```
In [9]: 1 sns.pairplot(cols)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x239544a6880>
```



```
In [10]: 1 sns.displot(df['placement_exam_marks'])
```


```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x23956525700>
```



```
In [11]: 1 # We use displot in older version we get distplot use displot
        2 sns.distplot(df['placement_exam_marks'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[11]: <AxesSubplot:xlabel='placement_exam_marks', ylabel='Density'>
```



```
In [12]: 1 df1=cols[['cgpa', 'placement_exam_marks', 'placed']]
          2 df1
```


Out[12]:

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0
...
995	8.87	44.0	1
996	9.12	65.0	1
997	4.89	34.0	0
998	8.62	46.0	1
999	4.90	10.0	1

1000 rows × 3 columns

```
In [13]: 1 sns.heatmap(df1.corr())
```

```
Out[13]: <AxesSubplot:>
```



To train the model - MODEL BUILD

Going to train linear regression model; We split our data into 2 variables x and y where x is independent var(input) and y is dependent on x(output), we could ignore address col as it is not required for our model

```
In [14]: 1 x=df1[['cgpa', 'placement_exam_marks', 'placed']]  
2 y=df1[['placed']]
```

To split the dataset into test data

```
In [15]: 1 # importing lib for splitting test data  
2 from sklearn.model_selection import train_test_split
```

```
In [16]: 1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```



```
In [17]: 1 from sklearn.linear_model import LinearRegression
        2
        3 lr=LinearRegression()
        4 lr.fit(x_train,y_train)
```

Out[17]: LinearRegression()

```
In [18]: 1 print(lr.intercept_)

[-1.11022302e-16]
```

```
In [19]: 1 print(lr.score(x_test,y_test))

1.0
```


```
In [20]: 1 coeff=pd.DataFrame(lr.coef_)
        2 coeff
```

Out[20]:

	0	1	2
0	1.508114e-18	-6.017229e-19	1.0

```
In [21]: 1 pred = lr.predict(x_test)
        2 plt.scatter(y_test,pred)
```

Out[21]: <matplotlib.collections.PathCollection at 0x23956dfeee0>



```
In [22]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [23]: 1 rr=Ridge(alpha=20)  
2 rr.fit(x_train,y_train)
```

Out[23]: Ridge(alpha=20)

```
In [24]: 1 rr.score(x_test,y_test)
```

Out[24]: 0.9893931081876479

```
In [25]: 1 la=Lasso(alpha=20)  
2 la.fit(x_train,y_train)
```

Out[25]: Lasso(alpha=20)

```
In [26]: 1 la.score(x_test,y_test)
```

Out[26]: -0.0036014405762305746

```
In [ ]: 1
```