

Problem statement

predicting the house price in USA.To create a model to help him estimate of what the house would sell for.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 df=pd.read_csv("drug")
```

To display top 10 rows

```
In [3]: 1 df.head(10)
```

Out[3]:

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY
5	22	F	NORMAL	HIGH	8.607	drugX
6	49	F	NORMAL	HIGH	16.275	drugY
7	41	M	LOW	HIGH	11.037	drugC
8	60	M	NORMAL	HIGH	15.171	drugY
9	43	M	LOW	NORMAL	19.368	drugY

Data Cleaning And Pre-Processing

In [4]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 200 entries, 0 to 199  
Data columns (total 6 columns):  
#   Column          Non-Null Count  Dtype    
---  ---            -  
0   Age             200 non-null   int64    
1   Sex             200 non-null   object   
2   BP              200 non-null   object   
3   Cholesterol     200 non-null   object   
4   Na_to_K         200 non-null   float64  
5   Drug            200 non-null   object   
dtypes: float64(1), int64(1), object(4)  
memory usage: 9.5+ KB
```

In [5]:

```
1 # Display the statistical summary  
2 df.describe()
```

Out[5]:

	Age	Na_to_K
count	200.000000	200.000000
mean	44.315000	16.084485
std	16.544315	7.223956
min	15.000000	6.269000
25%	31.000000	10.445500
50%	45.000000	13.936500
75%	58.000000	19.380000
max	74.000000	38.247000

```
In [6]: 1 # To display the col headings
        2 df.columns
```

```
Out[6]: Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

```
In [7]: 1 cols=df.dropna(axis=1)
```


```
In [8]: 1 cols.columns
```

```
Out[8]: Index(['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K', 'Drug'], dtype='object')
```

EDA and Visualization


```
In [9]: 1 sns.pairplot(cols)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x1f150e726d0>
```



```
In [10]: 1 sns.displot(df['Na_to_K'])
```


```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x1f151768310>
```



```
In [11]: 1 # We use displot in older version we get distplot use displot
        2 sns.distplot(df['Na_to_K'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[11]: <AxesSubplot:xlabel='Na_to_K', ylabel='Density'>
```



```
In [12]: 1 df1=cols[['Age', 'Na_to_K']]
          2 df1
```


Out[12]:

	Age	Na_to_K
0	23	25.355
1	47	13.093
2	47	10.114
3	28	7.798
4	61	18.043
...
195	56	11.567
196	16	12.006
197	52	9.894
198	23	14.020
199	40	11.349

200 rows × 2 columns

```
In [13]: 1 sns.heatmap(df1.corr())
```

```
Out[13]: <AxesSubplot:>
```



To train the model - MODEL BUILD

Going to train linear regression model; We split our data into 2 variables x and y where x is independent var(input) and y is dependent on x(output), we could ignore address col as it is not required for our model

```
In [14]: 1 x=df1[['Age', 'Na_to_K']]  
2 y=df1[['Na_to_K']]
```

To split the dataset into test data

```
In [15]: 1 # importing lib for splitting test data  
2 from sklearn.model_selection import train_test_split
```

```
In [16]: 1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [17]: 1 from sklearn.linear_model import LinearRegression
2
3 lr=LinearRegression()
4 lr.fit(x_train,y_train)
```

Out[17]: LinearRegression()

```
In [18]: 1 print(lr.intercept_)
```

[7.10542736e-15]

```
In [19]: 1 print(lr.score(x_test,y_test))
```

1.0


```
In [20]: 1 coeff=pd.DataFrame(lr.coef_)
2 coeff
```

Out[20]:

	0	1
0	-4.687824e-18	1.0


```
In [21]: 1 pred = lr.predict(x_test)
         2 plt.scatter(y_test,pred)
```

Out[21]: <matplotlib.collections.PathCollection at 0x1f1531c2760>



```
In [22]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [23]: 1 rr=Ridge(alpha=10)
         2 rr.fit(x_train,y_train)
```

Out[23]: Ridge(alpha=10)

```
In [24]: 1 rr.score(x_test,y_test)
```

Out[24]: 0.9999983447574403

```
In [25]: 1 la=Lasso(alpha=10)
         2 la.fit(x_train,y_train)
```

Out[25]: Lasso(alpha=10)

```
In [26]: 1 la.score(x_test,y_test)
```

Out[26]: 0.9673495603384751

ELASTICNET

```
In [27]: 1 from sklearn.linear_model import ElasticNet
        2 en=ElasticNet()
        3 en.fit(x_train,y_train)
```

Out[27]: ElasticNet()

```
In [28]: 1 print(en.coef_)

        [-0.          0.98209656]
```

```
In [29]: 1 print(en.intercept_)

        [0.288742]
```

```
In [30]: 1 prediction=en.predict(x_test)
        2 prediction
```

Out[30]: array([11.10064299, 10.93466867, 18.25619851, 26.86820322, 13.92417059,
35.28968119, 9.40358014, 18.69224938, 9.59410687, 9.68445976,
18.30824963, 14.05773573, 31.4074535 , 8.74164706, 9.56464398,
14.00568461, 29.62887664, 16.34307441, 18.6568939 , 10.54477634,
12.951895 , 17.20535519, 19.73032544, 10.99163027, 14.01059509,
12.82618664, 9.56267978, 20.85580809, 12.36361917, 20.61126605,
22.34270228, 22.78366364, 11.0692159 , 10.17550804, 18.93973771,
10.05569226, 8.88208687, 7.99328949, 10.22166657, 18.95545126,
16.71332482, 28.40813062, 15.79604663, 13.97425752, 18.76099614,
24.13011802, 15.17339741, 10.9268119 , 13.64329098, 9.21010712,
10.29532382, 11.34911342, 19.16561992, 11.28625924, 19.94344039,
22.57938755, 9.77972312, 32.62132485, 12.37540432, 15.50141766])

```
In [31]: 1 print(en.score(x_test,y_test))

        0.9996793149923281
```

EVALUATION METRICS

```
In [32]: 1 from sklearn import metrics
```

```
In [33]: print("Mean Absolute Error:", metrics.mean_absolute_error(y_test, prediction))
```

Mean Absolute Error: 0.09616337750064127

```
In [34]: print("Mean Squared Error:", metrics.mean_squared_error(y_test, prediction))
```

Mean Squared Error: 0.014079799306715247

```
In [35]: 1 print("Root Mean Squared Error:", np.sqrt(metrics.mean_squared_error(y_test, prediction)))
```

Root Mean Squared Error: 0.11865833011936097