

## Problem statement

predicting the house price in USA.To create a model to help him estimate of what the house would sell for.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 df=pd.read_csv("Horse Racing Results")
```

## To display top 10 rows

In [3]: 1 df.head(10)

Out[3]:

	Dato	Track	Race Number	Distance	Surface	Prize money	Starting position	Jockey	Jockey weight	Country	...	TrainerName	Race time	Path	Final place	f
0	03.09.2017	Sha Tin	10	1400	Gress	1310000	6	K C Leung	52	Sverige	...	CH Yip	83,38	2	9	
1	16.09.2017	Sha Tin	10	1400	Gress	1310000	14	C Y Ho	52	Sverige	...	CH Yip	81,56	3	4	
2	14.10.2017	Sha Tin	10	1400	Gress	1310000	8	C Y Ho	52	Sverige	...	CH Yip	82,36	1	6	
3	11.11.2017	Sha Tin	9	1600	Gress	1310000	13	Brett Prebble	54	Sverige	...	CH Yip	96,53	0	8	
4	26.11.2017	Sha Tin	9	1600	Gress	1310000	9	C Y Ho	52	Sverige	...	CH Yip	94,17	0	3	
5	10.12.2017	Sha Tin	1	1800	Gress	1310000	4	C Y Ho	52	Sverige	...	CH Yip	107,92	2	3	
6	01.01.2018	Sha Tin	9	1800	Gress	1310000	9	C Schofield	54	Sverige	...	CH Yip	108,63	5	1	
7	04.02.2018	Sha Tin	5	1800	Gress	1310000	6	Joao Moreira	57	Sverige	...	CH Yip	108,19	0	4	
8	03.03.2018	Sha Tin	8	1800	Gress	1310000	3	C Y Ho	56	Sverige	...	CH Yip	107,65	1	3	
9	11.03.2018	Sha Tin	10	1600	Gress	1310000	8	C Y Ho	57	Sverige	...	CH Yip	94,46	1	6	

10 rows × 21 columns



## Data Cleaning And Pre-Processing

In [4]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27008 entries, 0 to 27007
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Dato                   27008 non-null  object
1   Track                  27008 non-null  object
2   Race Number            27008 non-null  int64
3   Distance                27008 non-null  int64
4   Surface                 27008 non-null  object
5   Prize money             27008 non-null  int64
6   Starting position       27008 non-null  int64
7   Jockey                  27008 non-null  object
8   Jockey weight           27008 non-null  int64
9   Country                 27008 non-null  object
10  Horse age               27008 non-null  int64
11  TrainerName             27008 non-null  object
12  Race time                27008 non-null  object
13  Path                    27008 non-null  int64
14  Final place              27008 non-null  int64
15  FGrating                27008 non-null  int64
16  Odds                    27008 non-null  object
17  RaceType                27008 non-null  object
18  HorseId                 27008 non-null  int64
19  JockeyId                 27008 non-null  int64
20  TrainerID               27008 non-null  int64
dtypes: int64(12), object(9)
memory usage: 4.3+ MB
```

```
In [5]: 1 # Display the statistical summary
        2 df.describe()
```

Out[5]:

	Race Number	Distance	Prize money	Starting position	Jockey weight	Horse age	Path	Final place	FGrating
count	27008.000000	27008.000000	2.700800e+04	27008.000000	27008.000000	27008.000000	27008.000000	27008.000000	27008.000000
mean	5.268624	1401.666173	1.479445e+06	6.741447	55.867373	5.246408	1.678021	6.685834	113.428318
std	2.780088	276.065045	2.162109e+06	3.691071	2.737006	1.519880	1.631784	3.664551	13.314949
min	1.000000	1000.000000	6.600000e+05	1.000000	47.000000	2.000000	0.000000	1.000000	-5.000000
25%	3.000000	1200.000000	9.200000e+05	4.000000	54.000000	4.000000	0.000000	4.000000	106.000000
50%	5.000000	1400.000000	9.670000e+05	7.000000	56.000000	5.000000	1.000000	7.000000	114.000000
75%	8.000000	1650.000000	1.450000e+06	10.000000	58.000000	6.000000	3.000000	10.000000	122.000000
max	11.000000	2400.000000	2.800000e+07	14.000000	63.000000	12.000000	11.000000	14.000000	157.000000

```
In [6]: 1 # To display the col headings
        2 df.columns
```

Out[6]: Index(['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money', 'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age', 'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odds', 'RaceType', 'HorseId', 'JockeyId', 'TrainerID'], dtype='object')

```
In [7]: 1 cols=df.dropna(axis=1)
```

```
In [8]: 1 cols.columns
```

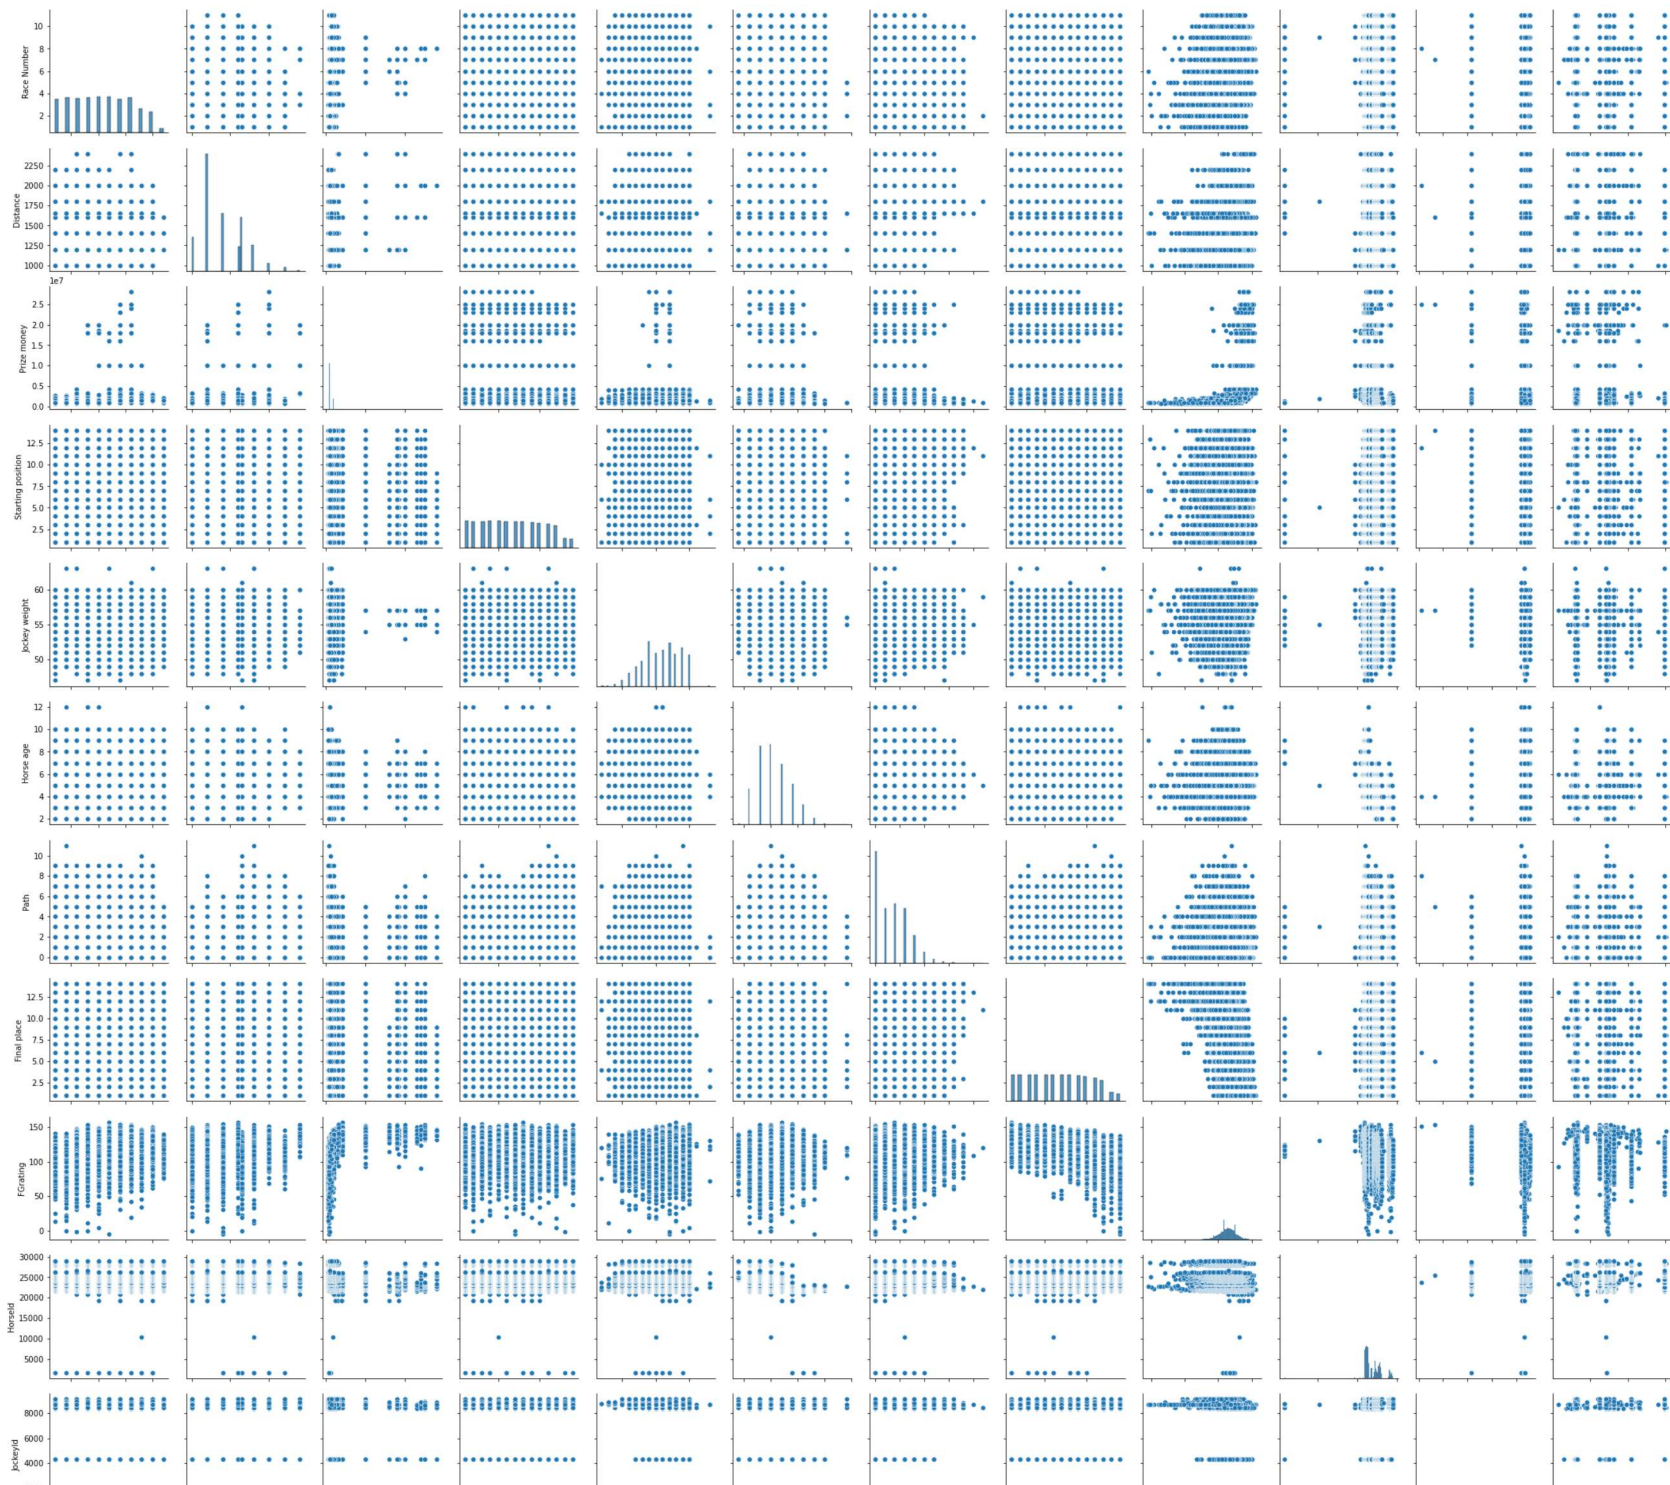
Out[8]: Index(['Dato', 'Track', 'Race Number', 'Distance', 'Surface', 'Prize money', 'Starting position', 'Jockey', 'Jockey weight', 'Country', 'Horse age', 'TrainerName', 'Race time', 'Path', 'Final place', 'FGrating', 'Odds', 'RaceType', 'HorseId', 'JockeyId', 'TrainerID'], dtype='object')

# EDA and Visualization


In [9]: 1 sns.pairplot(cols)

Out[9]: <seaborn.axisgrid.PairGrid at 0x1b6182f23d0>










```
In [10]: 1 sns.displot(df['Jockey weight'])
```

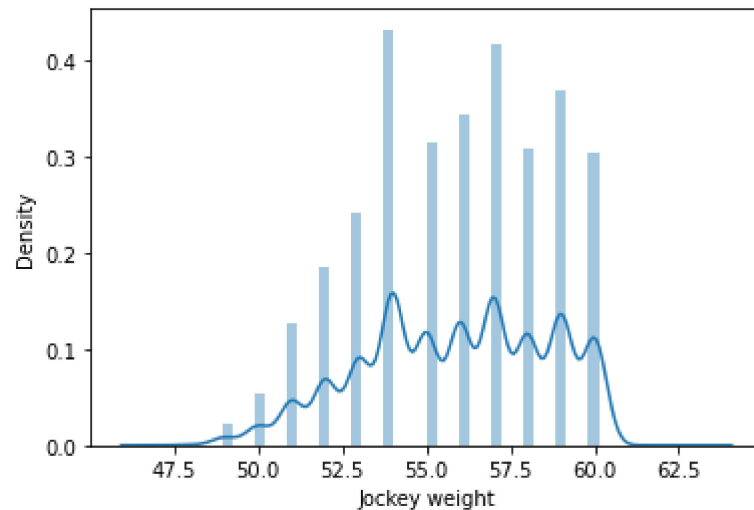
```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x1b62e14eee0>
```



```
In [11]: 1 # We use displot in older version we get distplot use displot
        2 sns.distplot(df['Jockey weight'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

```
Out[11]: <AxesSubplot:xlabel='Jockey weight', ylabel='Density'>
```



In [12]:

```
1 df1=cols[['Race Number', 'Distance',  
2          'Jockey weight', 'Horse age']]  
3 df1
```


Out[12]:

	Race Number	Distance	Jockey weight	Horse age
0	10	1400	52	7
1	10	1400	52	7
2	10	1400	52	7
3	9	1600	54	7
4	9	1600	52	7
...	...	...	...	...
27003	11	1200	59	3
27004	2	1200	57	3
27005	4	1200	57	3
27006	5	1200	57	3
27007	11	1200	55	4

27008 rows × 4 columns

```
In [13]: 1 sns.heatmap(df1.corr())
```

```
Out[13]: <AxesSubplot:>
```



## To train the model - MODEL BUILD

Going to train linear regression model; We split our data into 2 variables x and y where x is independent var(input) and y is dependent on x(output), we could ignore address col as it is not required for our model

```
In [14]: 1 x=df1[['Race Number', 'Distance',  
2         'Jockey weight', 'Horse age']]  
3 y=df1[['Distance']]
```

## To split the dataset into test data

```
In [15]: 1 # importing lib for splitting test data  
2 from sklearn.model_selection import train_test_split
```

```
In [16]: 1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [17]: 1 from sklearn.linear_model import LinearRegression  
2  
3 lr=LinearRegression()  
4 lr.fit(x_train,y_train)
```

Out[17]: LinearRegression()

```
In [18]: 1 print(lr.intercept_)
```

[-6.82121026e-13]

```
In [19]: 1 print(lr.score(x_test,y_test))
```

1.0


```
In [20]: 1 coeff=pd.DataFrame(lr.coef_)  
2 coeff
```

Out[20]:

	0	1	2	3
0	-1.293701e-15	1.0	8.004529e-17	7.520177e-16

```
In [21]: 1 pred = lr.predict(x_test)
         2 plt.scatter(y_test,pred)
```

Out[21]: <matplotlib.collections.PathCollection at 0x1b62ffecd90>



```
In [22]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [23]: 1 rr=Ridge(alpha=10)
         2 rr.fit(x_train,y_train)
```

Out[23]: Ridge(alpha=10)

```
In [24]: 1 rr.score(x_test,y_test)
```

Out[24]: 1.0

```
In [25]: 1 la=Lasso(alpha=10)
         2 la.fit(x_train,y_train)
```

Out[25]: Lasso(alpha=10)

```
In [26]: 1 la.score(x_test,y_test)
```

Out[26]: 0.9999999831678777

# Elastic net

```
In [27]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

Out[27]: ElasticNet()

```
In [28]: 1 print(en.coef_)
```

[0. 0.99998703 0. 0. ]

```
In [29]: 1 print(en.intercept_)
```

[0.01818225]

```
In [30]: 1 prediction=en.predict(x_test)
          2 prediction
```

Out[30]: array([1200.00261376, 1649.99677557, 1200.00261376, ..., 1649.99677557,  
 1400.00001901, 1200.00261376])

```
In [31]: 1 print(en.score(x_test,y_test))
```

0.999999999831681

## EVALUATION METRICS

```
In [32]: 1 from sklearn import metrics
```

```
In [33]: 1 print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

Mean Absolute Error: 0.002922189737464215

In [34]: 1 `print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))`

Mean Squared Error: 1.2486062600118047e-05

In [35]: 1 `print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))`

Root Mean Squared Error: 0.0035335623102073703