

# Problem statement

predicting the house price in USA.To create a model to help him estimate of what the house would sell for.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 df=pd.read_csv("placement")
```

## To display top 10 rows

```
In [3]: 1 df.head(10)
```

Out[3]:

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0
5	7.30	23.0	1
6	6.69	11.0	0
7	7.12	39.0	1
8	6.45	38.0	0
9	7.75	94.0	1

# Data Cleaning And Pre-Processing

In [4]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   cgpa                  1000 non-null  float64
1   placement_exam_marks 1000 non-null  float64
2   placed                1000 non-null  int64  
dtypes: float64(2), int64(1)
memory usage: 23.6 KB
```

In [5]:

```
1 # Display the statistical summary
2 df.describe()
```

Out[5]:

	cgpa	placement_exam_marks	placed
count	1000.000000	1000.000000	1000.000000
mean	6.961240	32.225000	0.489000
std	0.615898	19.130822	0.500129
min	4.890000	0.000000	0.000000
25%	6.550000	17.000000	0.000000
50%	6.960000	28.000000	0.000000
75%	7.370000	44.000000	1.000000
max	9.120000	100.000000	1.000000

In [6]:

```
1 # To display the col headings
2 df.columns
```

Out[6]: Index(['cgpa', 'placement\_exam\_marks', 'placed'], dtype='object')

```
In [7]: 1 cols=df.dropna(axis=1)
```

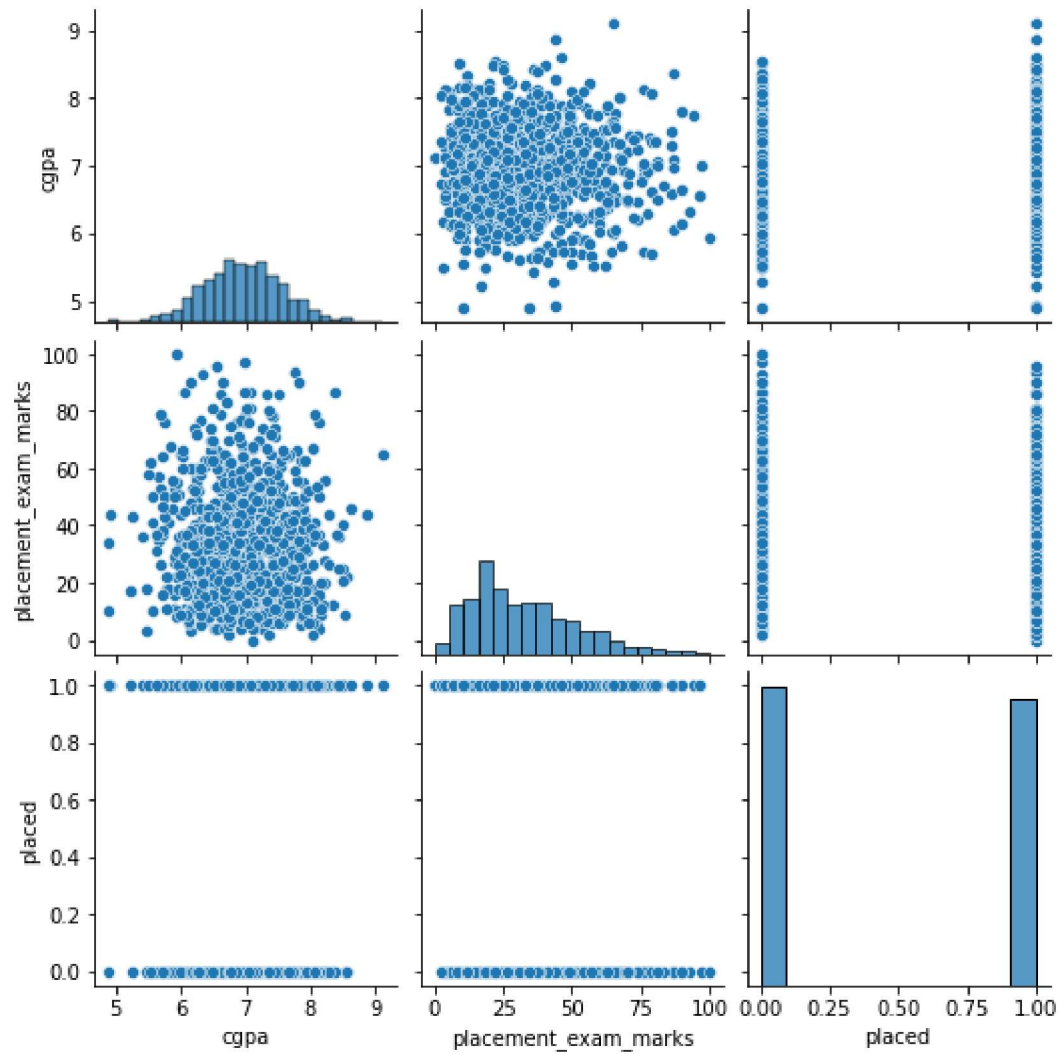
```
In [8]: 1 cols.columns
```

```
Out[8]: Index(['cgpa', 'placement_exam_marks', 'placed'], dtype='object')
```

## EDA and Visualization


```
In [9]: 1 sns.pairplot(cols)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x2bf6f49e280>
```



```
In [10]: 1 sns.displot(df['placement_exam_marks'])
```

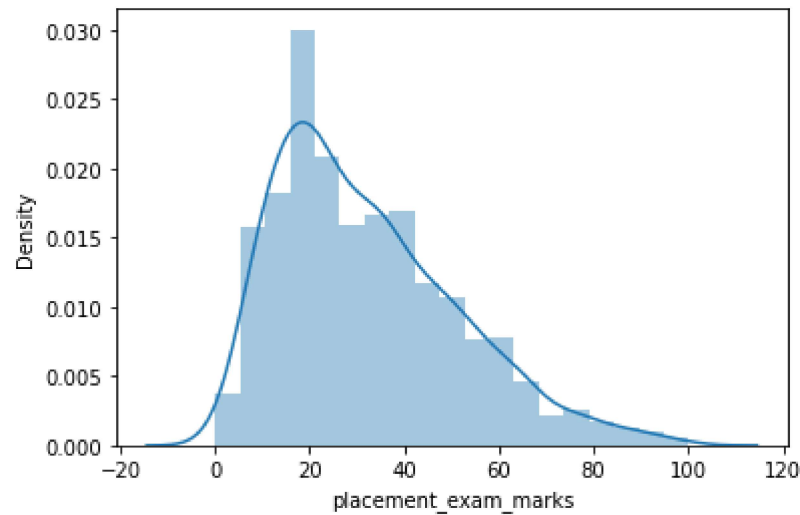
```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x2bf711198b0>
```



```
In [11]: 1 # We use displot in older version we get distplot use displot
        2 sns.distplot(df['placement_exam_marks'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

Out[11]: <AxesSubplot:xlabel='placement\_exam\_marks', ylabel='Density'>



```
In [12]: 1 df1=cols[['cgpa', 'placement_exam_marks', 'placed']]
          2 df1
```


Out[12]:

	cgpa	placement_exam_marks	placed
0	7.19	26.0	1
1	7.46	38.0	1
2	7.54	40.0	1
3	6.42	8.0	1
4	7.23	17.0	0
...	...	...	...
995	8.87	44.0	1
996	9.12	65.0	1
997	4.89	34.0	0
998	8.62	46.0	1
999	4.90	10.0	1

1000 rows × 3 columns

```
In [13]: 1 sns.heatmap(df1.corr())
```

```
Out[13]: <AxesSubplot:>
```



## To train the model - MODEL BUILD

Going to train linear regression model; We split our data into 2 variables x and y where x is independent var(input) and y is dependent on x(output), we could ignore address col as it is not required for our model

```
In [14]: 1 x=df1[['cgpa', 'placement_exam_marks', 'placed']]  
2 y=df1[['placed']]
```

## To split the dataset into test data

```
In [15]: 1 # importing lib for splitting test data  
2 from sklearn.model_selection import train_test_split
```



```
In [16]: 1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [17]: 1 from sklearn.linear_model import LinearRegression
2
3 lr=LinearRegression()
4 lr.fit(x_train,y_train)
```

Out[17]: LinearRegression()

```
In [18]: 1 print(lr.intercept_)
```

[-3.33066907e-16]

```
In [19]: 1 print(lr.score(x_test,y_test))
```

1.0


```
In [20]: 1 coeff=pd.DataFrame(lr.coef_)
2 coeff
```

Out[20]:

	0	1	2
0	1.196068e-16	-6.880975e-18	1.0

```
In [21]: 1 pred = lr.predict(x_test)
         2 plt.scatter(y_test,pred)
```

Out[21]: <matplotlib.collections.PathCollection at 0x2bf71ce5f70>



```
In [22]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [23]: 1 rr=Ridge(alpha=20)
         2 rr.fit(x_train,y_train)
```

Out[23]: Ridge(alpha=20)

```
In [24]: 1 rr.score(x_test,y_test)
```

Out[24]: 0.9888145970795182

```
In [25]: 1 la=Lasso(alpha=20)
         2 la.fit(x_train,y_train)
```

Out[25]: Lasso(alpha=20)

```
In [26]: 1 la.score(x_test,y_test)
```

Out[26]: -0.019886363636363535

# ELASTIC NET

```
In [30]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

Out[30]: ElasticNet()

```
In [31]: 1 print(en.coef_)
          [ 0.          -0.00021605  0.          ]
```

```
In [32]: 1 print(en.intercept_)
          [0.51692275]
```

```
In [33]: 1 prediction=en.predict(x_test)
          2 prediction
```

```
Out[33]: array([0.51497833, 0.51303391, 0.51022531, 0.513466 , 0.51497833,
0.51022531, 0.51130554, 0.51130554, 0.51216972, 0.51195368,
0.51324996, 0.50698461, 0.50568833, 0.51368205, 0.50892903,
0.513466 , 0.50871298, 0.51368205, 0.50914507, 0.51216972,
0.51238577, 0.50439205, 0.5074167 , 0.50093531, 0.51000926,
0.50568833, 0.49899089, 0.513466 , 0.51130554, 0.51022531,
0.51454624, 0.51303391, 0.51433019, 0.51411414, 0.50439205,
0.51195368, 0.50676856, 0.51541042, 0.51130554, 0.51216972,
0.51238577, 0.51303391, 0.50590438, 0.50892903, 0.50698461,
0.50806484, 0.51216972, 0.50892903, 0.51108949, 0.51303391,
0.50395996, 0.50482414, 0.51130554, 0.51195368, 0.50633647,
0.51087345, 0.50331182, 0.51108949, 0.50828089, 0.51411414,
0.51411414, 0.50201554, 0.50720065, 0.51216972, 0.50331182,
0.51519437, 0.50936112, 0.51562647, 0.51476228, 0.51303391,
0.51260182, 0.50287973, 0.51324996, 0.50892903, 0.50849693,
0.51605856, 0.5106574 , 0.51519437, 0.51173763, 0.50698461,
0.50849693, 0.50849693, 0.50525624, 0.51519437, 0.51368205,
0.51368205, 0.51108949, 0.51303391, 0.50266368, 0.51152158,
0.4981267 , 0.50504019, 0.51303391, 0.50331182, 0.51411414,
0.513466 , 0.51324996, 0.51238577, 0.50957717, 0.51324996,
0.51130554, 0.51238577, 0.51476228, 0.5106574 , 0.51000926,
0.50763275, 0.51108949, 0.51152158, 0.5106574 , 0.50914507,
0.51368205, 0.51000926, 0.51368205, 0.5074167 , 0.51476228,
0.51324996, 0.51562647, 0.51195368, 0.50504019, 0.51173763,
0.51541042, 0.50871298, 0.50936112, 0.50763275, 0.51022531,
0.49963903, 0.51497833, 0.513466 , 0.51087345, 0.51044135,
0.51130554, 0.50504019, 0.50676856, 0.50914507, 0.51216972,
0.50892903, 0.51454624, 0.50892903, 0.50806484, 0.51152158,
0.5074167 , 0.50504019, 0.50914507, 0.51368205, 0.50698461,
0.51022531, 0.50784879, 0.51324996, 0.504176 , 0.51130554,
0.50007112, 0.50936112, 0.50979321, 0.50244763, 0.50676856,
0.51022531, 0.50871298, 0.51173763, 0.50504019, 0.49618228,
0.50914507, 0.51519437, 0.50849693, 0.51195368, 0.50007112,
0.51605856, 0.5138981 , 0.51130554, 0.51454624, 0.50828089,
0.50698461, 0.5046081 , 0.51087345, 0.50871298, 0.51044135,
0.51541042, 0.51152158, 0.51281786, 0.50914507, 0.51497833,
0.50957717, 0.51238577, 0.51368205, 0.51497833, 0.51497833,
0.51497833, 0.50568833, 0.50439205, 0.51411414, 0.51152158,
0.5013674 , 0.51216972, 0.50892903, 0.51476228, 0.51519437,
0.51216972, 0.50806484, 0.50784879, 0.51173763, 0.51000926,
0.51454624, 0.51087345, 0.51281786, 0.50568833, 0.5074167 ,
0.50374391, 0.50568833, 0.50957717, 0.50849693, 0.51087345,
0.50784879, 0.51324996, 0.50892903, 0.50936112, 0.51454624,
```

```
0.50936112, 0.51260182, 0.51108949, 0.50763275, 0.50957717,  
0.51108949, 0.50179949, 0.51411414, 0.51173763, 0.51087345,  
0.50784879, 0.50287973, 0.50957717, 0.50892903, 0.51087345,  
0.50504019, 0.50979321, 0.51324996, 0.50028717, 0.51000926,  
0.51368205, 0.50957717, 0.50547228, 0.51411414, 0.50612042,  
0.50914507, 0.51000926, 0.51044135, 0.50849693, 0.51216972,  
0.51303391, 0.50892903, 0.51281786, 0.51022531, 0.51000926,  
0.504176 , 0.50979321, 0.51368205, 0.51433019, 0.51303391,  
0.51324996, 0.51195368, 0.5013674 , 0.50655252, 0.51497833,  
0.51281786, 0.51260182, 0.51130554, 0.51044135, 0.4981267 ,  
0.50806484, 0.50655252, 0.50395996, 0.50590438, 0.51022531,  
0.50806484, 0.49834275, 0.50957717, 0.51044135, 0.51303391,  
0.51000926, 0.50504019, 0.51433019, 0.51476228, 0.50093531,  
0.50676856, 0.50763275, 0.51260182, 0.49985507, 0.50720065,  
0.51303391, 0.51216972, 0.51195368, 0.50568833, 0.51216972,  
0.51281786, 0.51497833, 0.51368205, 0.50568833, 0.50504019,  
0.51368205, 0.50979321, 0.51087345, 0.50612042, 0.50806484])
```

```
In [34]: 1 print(en.score(x_test,y_test))
```

```
-0.02088205440129043
```

## EVALUATION METRICS

```
In [35]: 1 from sklearn import metrics
```

```
In [36]: 1 print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
```

```
Mean Absolute Error: 0.501432280822711
```

```
In [37]: 1 print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

```
Mean Squared Error: 0.251545338204478
```

In [38]: 1 `print("Root Mean Squared Error:", np.sqrt(metrics.mean_squared_error(y_test, prediction)))`

Root Mean Squared Error: 0.5015429574866723