

## Problem statement

predicting the house price in USA.To create a model to help him estimate of what the house would sell for.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 df=pd.read_csv("Sleep_health")
```

## To display top 10 rows

In [3]: 1 df.head(10)

Out[3]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate	Daily Steps	Sleep Disorder
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	77	4200	None
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75	10000	None
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75	10000	None
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	3000	Sleep Apnea
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	3000	Sleep Apnea
5	6	Male	28	Software Engineer	5.9	4	30	8	Obese	140/90	85	3000	Insomnia
6	7	Male	29	Teacher	6.3	6	40	7	Obese	140/90	82	3500	Insomnia
7	8	Male	29	Doctor	7.8	7	75	6	Normal	120/80	70	8000	None
8	9	Male	29	Doctor	7.8	7	75	6	Normal	120/80	70	8000	None
9	10	Male	29	Doctor	7.8	7	75	6	Normal	120/80	70	8000	None

## Data Cleaning And Pre-Processing

In [4]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374 entries, 0 to 373
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype  
---  -
0   Person ID                            374 non-null    int64  
1   Gender                                374 non-null    object  
2   Age                                    374 non-null    int64  
3   Occupation                            374 non-null    object  
4   Sleep Duration                        374 non-null    float64 
5   Quality of Sleep                      374 non-null    int64  
6   Physical Activity Level               374 non-null    int64  
7   Stress Level                          374 non-null    int64  
8   BMI Category                          374 non-null    object  
9   Blood Pressure                        374 non-null    object  
10  Heart Rate                            374 non-null    int64  
11  Daily Steps                           374 non-null    int64  
12  Sleep Disorder                        374 non-null    object  
dtypes: float64(1), int64(7), object(5)
memory usage: 38.1+ KB
```

In [5]:

```
1 # Display the statistical summary
2 df.describe()
```

Out[5]:

	Person ID	Age	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	Heart Rate	Daily Steps
<b>count</b>	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
<b>mean</b>	187.500000	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775	6816.844920
<b>std</b>	108.108742	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676	1617.915679
<b>min</b>	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000	3000.000000
<b>25%</b>	94.250000	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000	5600.000000
<b>50%</b>	187.500000	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000	7000.000000
<b>75%</b>	280.750000	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000	8000.000000
<b>max</b>	374.000000	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000	10000.000000

```
In [6]: 1 # To display the col headings
        2 df.columns
```

```
Out[6]: Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration',
              'Quality of Sleep', 'Physical Activity Level', 'Stress Level',
              'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps',
              'Sleep Disorder'],
              dtype='object')
```

```
In [7]: 1 cols=df.dropna(axis=1)
```

```
In [8]: 1 cols.columns
```


```
Out[8]: Index(['Person ID', 'Gender', 'Age', 'Occupation', 'Sleep Duration',
              'Quality of Sleep', 'Physical Activity Level', 'Stress Level',
              'BMI Category', 'Blood Pressure', 'Heart Rate', 'Daily Steps',
              'Sleep Disorder'],
              dtype='object')
```


## EDA and Visualization

```
In [9]: 1 sns.pairplot(cols)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x20092c3d610>
```








```
In [10]: 1 sns.displot(df['Quality of Sleep'])
```

```
Out[10]: <seaborn.axisgrid.FacetGrid at 0x20095ead3a0>
```

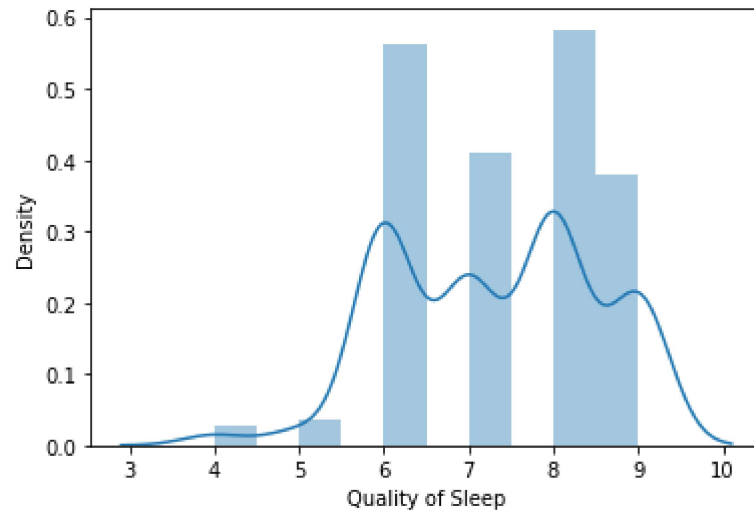




```
In [11]: 1 # We use displot in older version we get distplot use displot
        2 sns.distplot(df['Quality of Sleep'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

Out[11]: <AxesSubplot:xlabel='Quality of Sleep', ylabel='Density'>



```
In [12]: 1 df1=cols[['Quality of Sleep', 'Physical Activity Level', 'Stress Level']]
          2 df1
```


Out[12]:

	Quality of Sleep	Physical Activity Level	Stress Level
0	6	42	6
1	6	60	8
2	6	60	8
3	4	30	8
4	4	30	8
...	...	...	...
369	9	75	3
370	9	75	3
371	9	75	3
372	9	75	3
373	9	75	3

374 rows × 3 columns

```
In [13]: 1 sns.heatmap(df1.corr())
```

```
Out[13]: <AxesSubplot:>
```



## To train the model - MODEL BUILD

Going to train linear regression model; We split our data into 2 variables x and y where x is independent var(input) and y is dependent on x(output), we could ignore address col as it is not required for our model

```
In [14]: 1 x=df1[['Quality of Sleep', 'Physical Activity Level', 'Stress Level']]  
2 y=df1[['Quality of Sleep']]
```

## To split the dataset into test data

```
In [15]: 1 # importing lib for splitting test data
         2 from sklearn.model_selection import train_test_split
```

```
In [16]: 1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [17]: 1 from sklearn.linear_model import LinearRegression
         2
         3 lr=LinearRegression()
         4 lr.fit(x_train,y_train)
```

Out[17]: LinearRegression()

```
In [18]: 1 print(lr.intercept_)
```

[-1.77635684e-15]

```
In [19]: 1 print(lr.score(x_test,y_test))
```

1.0


```
In [20]: 1 coeff=pd.DataFrame(lr.coef_)
         2 coeff
```

Out[20]:

	0	1	2
0	1.0	8.111872e-17	-1.707167e-16

```
In [21]: 1 pred = lr.predict(x_test)
          2 plt.scatter(y_test,pred)
```

Out[21]: <matplotlib.collections.PathCollection at 0x2009828f6d0>



```
In [22]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [23]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

Out[23]: Ridge(alpha=10)

```
In [24]: 1 rr.score(x_test,y_test)
```

Out[24]: 0.9969900659175175

```
In [25]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

Out[25]: Lasso(alpha=10)

```
In [26]: 1 la.score(x_test,y_test)
```

Out[26]: -0.0032932238920933

# ELASTIC NET

```
In [27]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

Out[27]: ElasticNet()

```
In [28]: 1 print(en.coef_)
          [ 0.17198479  0.00581034 -0.28979061]
```

```
In [29]: 1 print(en.intercept_)
          [7.25607742]
```

```
In [30]: 1 prediction=en.predict(x_test)
         2 prediction
```

```
Out[30]: array([7.531623 , 8.37034408, 7.04079592, 7.70593316, 6.14397145,
                7.82141361, 8.37034408, 8.37034408, 6.31988066, 8.37034408,
                5.81162255, 7.531623 , 7.56227374, 8.37034408, 7.82141361,
                7.82141361, 7.15700269, 6.49259177, 7.18372901, 6.14397145,
                6.52091714, 6.31988066, 8.10887884, 8.37034408, 6.14397145,
                8.37034408, 7.82141361, 6.14397145, 7.531623 , 8.37034408,
                7.15700269, 7.82141361, 7.70593316, 7.15700269, 8.45749916,
                6.52091714, 8.42844747, 7.531623 , 8.10887884, 7.82141361,
                7.56227374, 7.70593316, 7.70593316, 8.10887884, 6.49259177,
                7.531623 , 8.37034408, 8.37034408, 6.49259177, 6.69290192,
                7.15700269, 8.02404914, 7.04079592, 6.29082896, 6.14397145,
                7.82141361, 8.37034408, 6.14397145, 6.14397145, 8.37034408,
                7.56227374, 7.15700269, 6.14397145, 7.56227374, 7.15700269,
                8.37034408, 7.82141361, 6.14397145, 8.10887884, 8.10887884,
                8.10887884, 7.15700269, 7.531623 , 7.82141361, 6.14397145,
                7.531623 , 7.56227374, 7.2127807 , 7.531623 , 7.70593316,
                6.29082896, 6.31828161, 6.49259177, 7.15700269, 7.70593316,
                6.52091714, 6.52091714, 7.70593316, 8.37034408, 6.14397145,
                7.35963822, 5.80000188, 6.14397145, 7.531623 , 7.15700269,
                6.69290192, 7.70593316, 7.70593316, 7.15700269, 7.15700269,
                7.38868991, 6.49259177, 7.33058652, 6.31828161, 8.37034408,
                8.10887884, 7.531623 , 6.49259177, 6.89553745, 6.14397145,
                7.15700269, 6.49259177, 8.39939578])
```

```
In [31]: 1 print(en.score(x_test,y_test))
```

```
0.7950486692003498
```

## EVALUATION METRICS

```
In [32]: 1 from sklearn import metrics
```

In [33]: 1 `print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))`

Mean Absolute Error: 0.45495704210174887

In [34]: 1 `print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))`

Mean Squared Error: 0.3194409417788893

In [35]: 1 `print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))`

Root Mean Squared Error: 0.5651910666127777

In [ ]:

1