

Problem statement

predicting the house price in USA.To create a model to help him estimate of what the house would sell for.

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```

```
In [2]: 1 df=pd.read_csv("states")
```

To display top 10 rows

```
In [3]: 1 df.head(10)
```

Out[3]:

| | id | name | country_id | country_code | country_name | state_code | type | latitude | longitude |
|---|------|------------|------------|--------------|--------------|------------|------|-----------|-----------|
| 0 | 3901 | Badakhshan | 1 | AF | Afghanistan | BDS | NaN | 36.734772 | 70.811995 |
| 1 | 3871 | Badghis | 1 | AF | Afghanistan | BDG | NaN | 35.167134 | 63.769538 |
| 2 | 3875 | Baghlan | 1 | AF | Afghanistan | BGL | NaN | 36.178903 | 68.745306 |
| 3 | 3884 | Balkh | 1 | AF | Afghanistan | BAL | NaN | 36.755060 | 66.897537 |
| 4 | 3872 | Bamyan | 1 | AF | Afghanistan | BAM | NaN | 34.810007 | 67.821210 |
| 5 | 3892 | Daykundi | 1 | AF | Afghanistan | DAY | NaN | 33.669495 | 66.046353 |
| 6 | 3899 | Farah | 1 | AF | Afghanistan | FRA | NaN | 32.495328 | 62.262663 |
| 7 | 3889 | Faryab | 1 | AF | Afghanistan | FYB | NaN | 36.079561 | 64.905955 |
| 8 | 3870 | Ghazni | 1 | AF | Afghanistan | GHA | NaN | 33.545059 | 68.417397 |
| 9 | 3888 | Ghōr | 1 | AF | Afghanistan | GHO | NaN | 34.099578 | 64.905955 |

Data Cleaning And Pre-Processing

In [4]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5077 entries, 0 to 5076
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   id               5077 non-null   int64
1   name             5077 non-null   object
2   country_id       5077 non-null   int64
3   country_code     5063 non-null   object
4   country_name     5077 non-null   object
5   state_code       5072 non-null   object
6   type             1597 non-null   object
7   latitude         5008 non-null   float64
8   longitude        5008 non-null   float64
dtypes: float64(2), int64(2), object(5)
memory usage: 357.1+ KB
```

In [5]:

```
1 # Display the statistical summary
2 df.describe()
```

Out[5]:

| | id | country_id | latitude | longitude |
|-------|-------------|-------------|-------------|-------------|
| count | 5077.000000 | 5077.000000 | 5008.000000 | 5008.000000 |
| mean | 2609.765413 | 133.467599 | 27.576415 | 17.178713 |
| std | 1503.376799 | 72.341160 | 22.208161 | 61.269334 |
| min | 1.000000 | 1.000000 | -54.805400 | -178.116500 |
| 25% | 1324.000000 | 74.000000 | 11.399747 | -3.943859 |
| 50% | 2617.000000 | 132.000000 | 34.226432 | 17.501792 |
| 75% | 3905.000000 | 201.000000 | 45.802822 | 41.919647 |
| max | 5220.000000 | 248.000000 | 77.874972 | 179.852222 |

```
In [6]: 1 # To display the col headings
        2 df.columns
```

```
Out[6]: Index(['id', 'name', 'country_id', 'country_code', 'country_name',
              'state_code', 'type', 'latitude', 'longitude'],
              dtype='object')
```

```
In [7]: 1 cols=df.dropna(axis=1)
```


```
In [8]: 1 cols.columns
```

```
Out[8]: Index(['id', 'name', 'country_id', 'country_name'], dtype='object')
```

EDA and Visualization


```
In [9]: 1 sns.pairplot(cols)
```

```
Out[9]: <seaborn.axisgrid.PairGrid at 0x2053e1c3730>
```



```
In [11]: 1 sns.displot(df['country_id'])
```


```
Out[11]: <seaborn.axisgrid.FacetGrid at 0x2053942e760>
```



```
In [12]: 1 # We use distplot in older version we get distplot use displot
        2 sns.distplot(df['country_id'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[12]: <AxesSubplot:xlabel='country_id', ylabel='Density'>
```



```
In [15]: 1 df1=cols[['id', 'country_id']]
          2 df1
```


Out[15]:

| | id | country_id |
|------|------|------------|
| 0 | 3901 | 1 |
| 1 | 3871 | 1 |
| 2 | 3875 | 1 |
| 3 | 3884 | 1 |
| 4 | 3872 | 1 |
| ... | ... | ... |
| 5072 | 1953 | 247 |
| 5073 | 1960 | 247 |
| 5074 | 1954 | 247 |
| 5075 | 1952 | 247 |
| 5076 | 1957 | 247 |

5077 rows × 2 columns

```
In [16]: 1 sns.heatmap(df1.corr())
```

```
Out[16]: <AxesSubplot:>
```



To train the model - MODEL BUILD

Going to train linear regression model; We split our data into 2 variables x and y where x is independent var(input) and y is dependent on x(output), we could ignore address col as it is not required for our model

```
In [18]: 1 x=df1[['id', 'country_id']]  
2 y=df1[['id']]
```

To split the dataset into test data

```
In [19]: 1 # importing lib for splitting test data  
2 from sklearn.model_selection import train_test_split
```

```
In [20]: 1 x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```



```
In [21]: 1 from sklearn.linear_model import LinearRegression
         2
         3 lr=LinearRegression()
         4 lr.fit(x_train,y_train)
```

Out[21]: LinearRegression()

```
In [22]: 1 print(lr.intercept_)

         [0.]
```

```
In [23]: 1 print(lr.score(x_test,y_test))

         1.0
```


```
In [24]: 1 coeff=pd.DataFrame(lr.coef_)
         2 coeff
```

Out[24]:

| | 0 | 1 |
|---|-----|--------------|
| 0 | 1.0 | 1.327549e-17 |

```
In [25]: 1 pred = lr.predict(x_test)
         2 plt.scatter(y_test,pred)
```

Out[25]: <matplotlib.collections.PathCollection at 0x2053f6c78b0>



```
In [26]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [27]: 1 rr=Ridge(alpha=10)
         2 rr.fit(x_train,y_train)
```

Out[27]: Ridge(alpha=10)

```
In [28]: 1 rr.score(x_test,y_test)
```

Out[28]: 1.0

```
In [29]: 1 la=Lasso(alpha=10)
         2 la.fit(x_train,y_train)
```

Out[29]: Lasso(alpha=10)

```
In [30]: 1 la.score(x_test,y_test)
```

Out[30]: 0.9999999999806893

ELASTIC NET

```
In [31]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

Out[31]: ElasticNet()

```
In [32]: 1 print(en.coef_)
          [0.99999956 0.          ]
```

```
In [33]: 1 print(en.intercept_)
          [0.00114742]
```

```
In [34]: 1 prediction=en.predict(x_test)
          2 prediction
```

Out[34]: array([3316.99968981, 3096.99978649, 2767.99993106, ..., 4380.99922225,
 66.00111842, 2944.99985328])

```
In [35]: 1 print(en.score(x_test,y_test))
          0.9999999999998069
```

EVALUATION METRICS

```
In [36]: 1 from sklearn import metrics
```

```
In [37]: 1 print("Mean Absolute Error:",metrics.mean_absolute_error(y_test,prediction))
          Mean Absolute Error: 0.000568406353368604
```

```
In [38]: 1 print("Mean Squared Error:",metrics.mean_squared_error(y_test,prediction))
```

Mean Squared Error: 4.291842482179795e-07

```
In [39]: 1 print("Root Mean Squared Error:",np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

Root Mean Squared Error: 0.0006551215522465885

MODEL SAVING

```
In [40]: 1 import pickle
```

```
In [43]: 1 filename='prediction2'  
2 pickle.dump(lr,open(filename,'wb'))
```

```
In [ ]: 1
```