

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sn
```

```
In [2]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [3]: 1 df=pd.read_csv(r"C5_health care diabetes")
```

```
In [4]: 1 df
```

Out[4]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

In [5]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Pregnancies            768 non-null   int64
1   Glucose                768 non-null   int64
2   BloodPressure          768 non-null   int64
3   SkinThickness          768 non-null   int64
4   Insulin                768 non-null   int64
5   BMI                   768 non-null   float64
6   DiabetesPedigreeFunction 768 non-null   float64
7   Age                   768 non-null   int64
8   Outcome                768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

In [6]: 1 df.dropna(inplace=True)

In [7]: 1 df.columns

Out[7]: Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
 'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
 dtype='object')

In [8]:

```
1 df
```

Out[8]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

In [9]:

```
1 from sklearn.linear_model import LogisticRegression
```

In [10]:

```
1 logr =LogisticRegression()
```

In [11]:

```
1 feature_matrix=df[['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',  
2                     'BMI', 'DiabetesPedigreeFunction', 'Age']]  
3 target_vector=df['Outcome']
```

In [12]:

```
1 feature_matrix.shape
```

Out[12]: (768, 8)

In [13]:

```
1 target_vector.shape
```

Out[13]: (768,)

```
In [14]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [15]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [16]: 1 logr=LogisticRegression()  
2 logr.fit(fs,target_vector)
```

```
Out[16]: LogisticRegression()
```

```
In [19]: 1 observation=[[1,2,3,4,5,6,7,8]]
```

```
In [20]: 1 prediction = logr.predict(observation)  
2 print(prediction)
```

```
[1]
```

```
In [21]: 1 logr.classes_
```

```
Out[21]: array([0, 1], dtype=int64)
```

```
In [22]: 1 logr.predict_proba(observation)[0][1]
```

```
Out[22]: 0.9997076305131244
```

```
In [23]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[23]: 0.00029236948687560993
```


Linear regression 2

```
In [24]: 1 import re
2 from sklearn.datasets import load_digits
3 import numpy as np
4 import pandas as pd
5 import matplotlib.pyplot as plt
6 import seaborn as sns
7 from sklearn.linear_model import LogisticRegression
8 from sklearn.model_selection import train_test_split
```

```
In [25]: 1 digits =load_digits()
2 digits
```

```
Out[25]: {'data': array([[ 0.,  0.,  5., ...,  0.,  0.,  0.],
        [ 0.,  0.,  0., ..., 10.,  0.,  0.],
        [ 0.,  0.,  0., ..., 16.,  9.,  0.],
        ...,
        [ 0.,  0.,  1., ...,  6.,  0.,  0.],
        [ 0.,  0.,  2., ..., 12.,  0.,  0.],
        [ 0.,  0., 10., ..., 12.,  1.,  0.])),
 'target': array([0, 1, 2, ..., 8, 9, 8]),
 'frame': None,
 'feature_names': ['pixel_0_0',
 'pixel_0_1',
 'pixel_0_2',
 'pixel_0_3',
 'pixel_0_4',
 'pixel_0_5',
 'pixel_0_6',
 'pixel_0_7',
 'pixel_1_0',
 'pixel_1_1',
 'pixel_1_2',
 'pixel_1_3',
 'pixel_1_4',
 'pixel_1_5',
 'pixel_1_6',
 'pixel_1_7',
 'pixel_2_0',
 'pixel_2_1',
 'pixel_2_2',
 'pixel_2_3',
 'pixel_2_4',
 'pixel_2_5',
 'pixel_2_6',
 'pixel_2_7',
 'pixel_3_0',
 'pixel_3_1',
 'pixel_3_2',
 'pixel_3_3',
 'pixel_3_4',
 'pixel_3_5',
 'pixel_3_6',
 'pixel_3_7',
 'pixel_4_0',
 'pixel_4_1',
 'pixel_4_2',
 'pixel_4_3',
 'pixel_4_4',
 'pixel_4_5',
 'pixel_4_6',
 'pixel_4_7',
 'pixel_5_0',
 'pixel_5_1',
 'pixel_5_2',
 'pixel_5_3',
 'pixel_5_4',
 'pixel_5_5',
 'pixel_5_6',
 'pixel_5_7',
 'pixel_6_0',
 'pixel_6_1',
 'pixel_6_2',
 'pixel_6_3',
 'pixel_6_4',
 'pixel_6_5',
 'pixel_6_6',
 'pixel_6_7',
 'pixel_7_0',
 'pixel_7_1',
 'pixel_7_2',
 'pixel_7_3',
 'pixel_7_4',
 'pixel_7_5',
 'pixel_7_6',
 'pixel_7_7',
 'pixel_8_0',
 'pixel_8_1',
 'pixel_8_2',
 'pixel_8_3',
 'pixel_8_4',
 'pixel_8_5',
 'pixel_8_6',
 'pixel_8_7',
 'pixel_9_0',
 'pixel_9_1',
 'pixel_9_2',
 'pixel_9_3',
 'pixel_9_4',
 'pixel_9_5',
 'pixel_9_6',
 'pixel_9_7']},
 'data': array([[ 0.,  0.,  5., ...,  0.,  0.,  0.],
        [ 0.,  0.,  0., ..., 10.,  0.,  0.],
        [ 0.,  0.,  0., ..., 16.,  9.,  0.],
        ...,
        [ 0.,  0.,  1., ...,  6.,  0.,  0.],
        [ 0.,  0.,  2., ..., 12.,  0.,  0.],
        [ 0.,  0., 10., ..., 12.,  1.,  0.]])
```

```
In [26]: 1 plt.figure(figsize=(20,4))
2 for index,(image,label) in enumerate(zip(digits.data[0:5],digits.target[0:5])):
3     plt.subplot(1,5,index+1)
4     plt.imshow(np.reshape(image,(8,8)),cmap=plt.cm.gray)
5     plt.title("Number:%i\n"%label,fontsize=15)
```



```
In [27]: 1 x_train,x_test,y_train,y_test=train_test_split(digits.data,digits.target,test_size=0.30)
```

```
In [28]: 1 print(x_train.shape)
2 print(x_test.shape)
3 print(y_train.shape)
4 print(y_test.shape)
```

```
(1257, 64)
(540, 64)
(1257,)
(540,)
```

```
In [29]: 1 logre=LogisticRegression(max_iter=10000)
2 logre.fit(x_train,y_train)
```

```
Out[29]: LogisticRegression(max_iter=10000)
```

In [30]:

```
1 print(logre.predict(x_test))
```

```
[9 8 2 9 3 0 0 1 7 5 1 2 2 0 0 1 0 3 0 6 1 8 9 6 1 5 0 3 7 2 9 2 2 0 4 5 9
 3 5 1 1 3 6 4 6 3 8 4 5 8 4 2 9 4 0 6 2 2 6 7 9 3 0 9 2 5 7 4 0 3 1 7 2 1
 9 3 5 5 4 8 7 1 2 8 7 2 6 6 6 5 2 8 6 8 9 4 1 0 0 8 6 5 1 5 6 9 3 6 3 5 5
 7 4 4 2 2 5 2 8 9 1 4 3 3 0 6 5 3 6 5 8 2 6 5 6 6 6 0 7 8 1 2 4 9 7 0 1 7
 8 8 5 5 6 0 5 4 9 4 8 2 8 9 7 3 9 7 8 7 0 6 8 8 8 5 0 5 4 5 7 9 5 3 1 3 1
 1 2 8 9 8 3 0 7 7 5 7 3 7 4 9 4 7 0 9 0 2 8 0 1 3 3 5 4 8 8 2 6 5 9 6 0 9
 7 5 9 0 2 2 5 5 3 6 8 5 5 3 1 4 9 4 3 2 1 0 1 3 5 8 3 4 0 7 2 8 8 8 2 7 4
 9 2 5 4 6 3 2 1 3 9 8 6 8 2 0 3 1 6 8 3 6 2 6 6 2 3 7 7 9 9 4 1 7 9 6 2 2
 0 4 9 7 1 1 4 8 3 1 8 0 7 1 1 6 2 4 5 5 7 7 9 4 1 4 6 8 0 8 1 4 4 1 3 7 6
 2 5 4 4 1 3 2 1 1 1 6 8 7 1 6 5 4 4 3 3 6 2 5 4 1 3 7 2 7 8 4 2 5 0 3 5 0
 1 8 0 6 5 4 0 4 5 1 5 5 4 2 1 4 2 0 6 5 4 2 3 6 4 2 7 4 0 4 4 7 5 3 3 1 0
 9 1 4 1 9 4 0 9 8 9 2 6 4 3 9 3 3 2 6 2 8 8 6 1 9 4 5 2 8 0 8 0 4 7 3 4 2
 5 3 1 7 3 3 4 0 2 4 8 1 2 6 4 0 4 7 0 6 3 3 5 4 2 4 1 3 7 3 6 7 2 8 3 9 3
 7 7 4 2 6 9 2 6 2 6 0 3 4 6 5 6 4 3 3 5 3 6 6 9 3 7 6 9 9 2 1 0 7 2 7 3 3
 1 2 0 9 2 6 7 9 9 3 6 0 3 6 2 5 1 1 9 7 2 4]
```

In [31]:

```
1 print(logre.score(x_test,y_test))
```

0.95