

Importing Libraries

```
In [1]: import numpy as np  
import pandas as pd  
import seaborn as sns  
import matplotlib.pyplot as plt
```

Importing Datasets

```
In [2]: df=pd.read_csv("innovation_and_development_database.csv")  
df
```

Out[2]:

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO
0	Aruba	ABW	1960	0	0	1	0	0	0	0.0	...	NaN	NaN
1	Aruba	ABW	1961	0	0	1	0	0	0	0.0	...	NaN	NaN
2	Aruba	ABW	1962	0	0	1	0	0	0	0.0	...	NaN	NaN
3	Aruba	ABW	1963	0	0	1	0	0	0	0.0	...	NaN	NaN
4	Aruba	ABW	1964	0	0	1	0	0	0	0.0	...	NaN	NaN
...
8290	Zimbabwe	ZWE	1998	0	0	0	0	1	0	0.0	...	8.290000e+09	8.0
8291	Zimbabwe	ZWE	1999	0	0	0	0	1	0	0.0	...	8.230000e+09	8.0
8292	Zimbabwe	ZWE	2000	0	0	0	0	1	0	0.0	...	7.830000e+09	8.0
8293	Zimbabwe	ZWE	2001	0	0	0	0	1	0	0.0	...	NaN	8.0
8294	Zimbabwe	ZWE	2002	0	0	0	0	1	0	0.0	...	NaN	NaN

8295 rows × 33 columns



Data Cleaning and Data Preprocessing

In [3]:

```
df=df.fillna(1)
df
```

Out[3]:

	country	code	year	eap	eca	lac	mena	sha	sa	hi	...	y	stockpatEPO
0	Aruba	ABW	1960	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0
1	Aruba	ABW	1961	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0
2	Aruba	ABW	1962	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0
3	Aruba	ABW	1963	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0
4	Aruba	ABW	1964	0	0	1	0	0	0	0.0	...	1.000000e+00	1.0
...
8290	Zimbabwe	ZWE	1998	0	0	0	0	1	0	0.0	...	8.290000e+09	8.0
8291	Zimbabwe	ZWE	1999	0	0	0	0	1	0	0.0	...	8.230000e+09	8.0
8292	Zimbabwe	ZWE	2000	0	0	0	0	1	0	0.0	...	7.830000e+09	8.0
8293	Zimbabwe	ZWE	2001	0	0	0	0	1	0	0.0	...	1.000000e+00	8.0
8294	Zimbabwe	ZWE	2002	0	0	0	0	1	0	0.0	...	1.000000e+00	1.0

8295 rows × 33 columns



In [4]:

```
df.columns
```

Out[4]:

```
Index(['country', 'code', 'year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa',
       'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdfinabro',
       'rdfinprod', 'rdperfprod', 'rdperfhe', 'rdperfpub', 'lowrdexp',
       'lowrdfinprod', 'lowrdperfprod', 'y', 'stockpatEPO', 'poptotal',
       'labor', 'rdexpgdp', 'patgrantedstock', 'plantpatstock',
       'designpatstock', 'plantpat', 'designpat'],
      dtype='object')
```

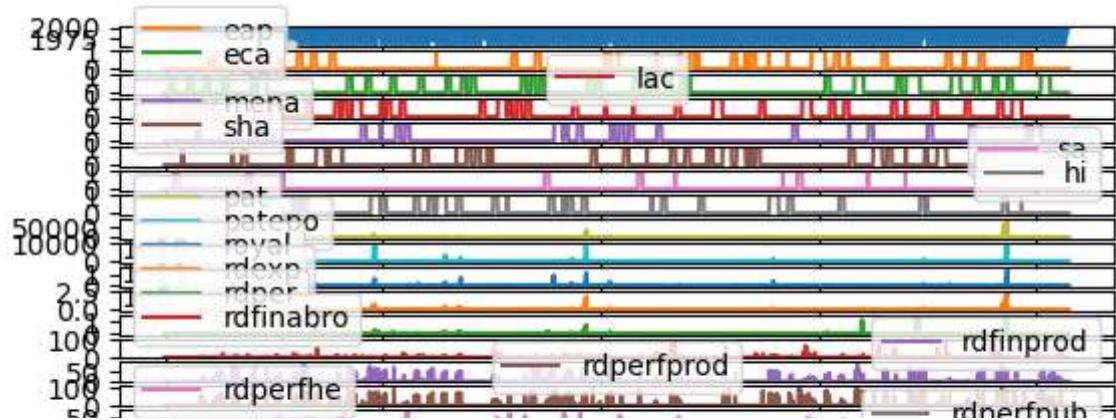
In [5]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8295 entries, 0 to 8294
Data columns (total 33 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   country          8295 non-null    object  
 1   code              8295 non-null    object  
 2   year              8295 non-null    int64  
 3   eap               8295 non-null    int64  
 4   eca               8295 non-null    int64  
 5   lac               8295 non-null    int64  
 6   mena              8295 non-null    int64  
 7   sha               8295 non-null    int64  
 8   sa                8295 non-null    int64  
 9   hi                8295 non-null    float64 
 10  pat               8295 non-null    float64 
 11  patepo            8295 non-null    float64 
 12  royal              8295 non-null    float64 
 13  rdexp              8295 non-null    float64 
 14  rdper              8295 non-null    float64 
 15  rdfinabro          8295 non-null    float64 
 16  rdfinprod           8295 non-null    float64 
 17  rdperfprod          8295 non-null    float64 
 18  rdperfhe            8295 non-null    float64 
 19  rdperfpub           8295 non-null    float64 
 20  lowrdexp             8295 non-null    float64 
 21  lowrdfinprod          8295 non-null    float64 
 22  lowrdperfprod         8295 non-null    float64 
 23  y                  8295 non-null    float64 
 24  stockpatEPO           8295 non-null    float64 
 25  poptotal            8295 non-null    float64 
 26  labor              8295 non-null    float64 
 27  rdexpgdp             8295 non-null    float64 
 28  patgrantedstock        8295 non-null    float64 
 29  plantpatstock          8295 non-null    float64 
 30  designpatstock          8295 non-null    float64 
 31  plantpat              8295 non-null    float64 
 32  designpat              8295 non-null    float64 
dtypes: float64(24), int64(7), object(2)
memory usage: 2.1+ MB
```

Line chart

```
In [6]: df.plot.line(subplots=True)
```

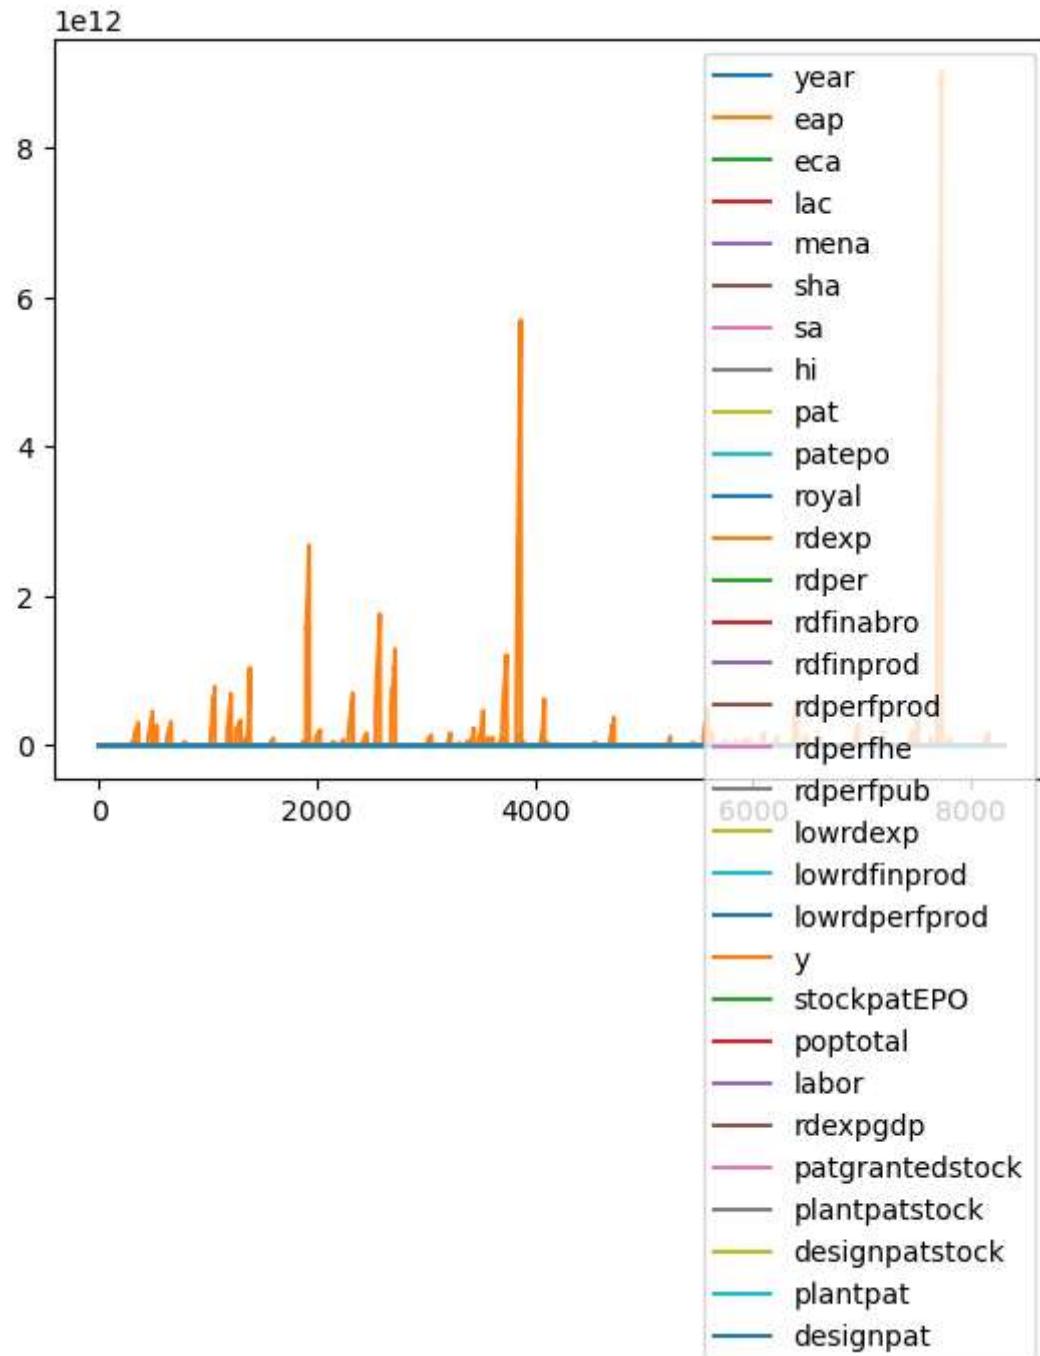
```
Out[6]: array([<Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >,
   <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >,
   <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >,
   <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >, <Axes: >,
   <Axes: >], dtype=object)
```



Line chart

In [7]: `df.plot.line()`

Out[7]: <Axes: >

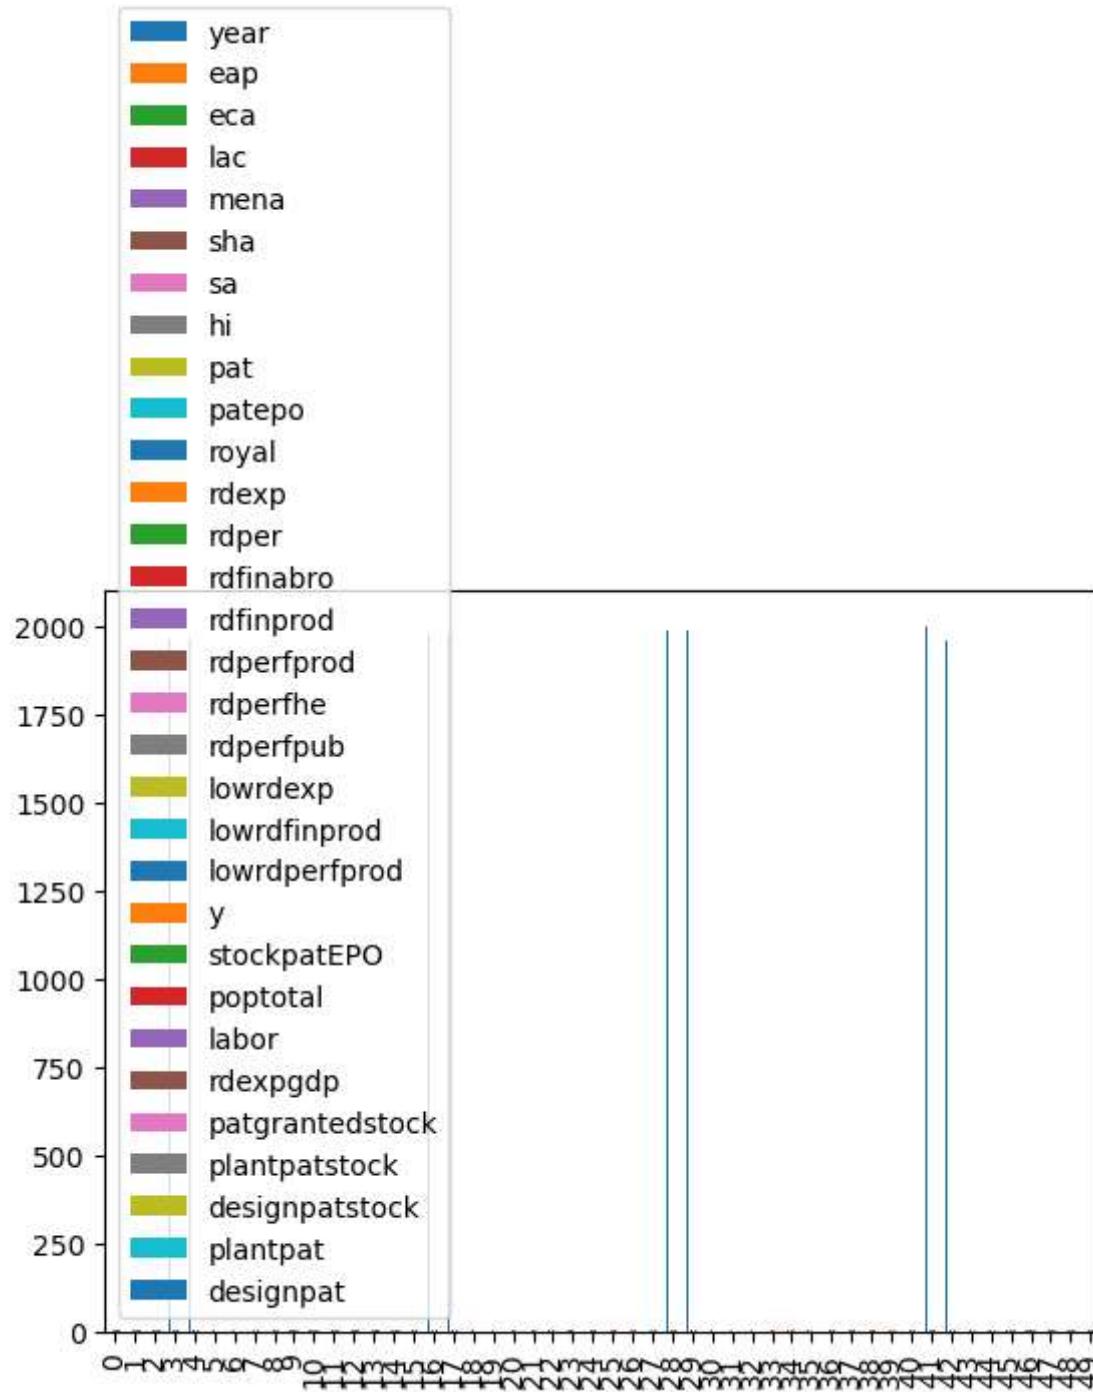


Bar chart

In [8]: `b=df[0:50]`

In [9]: `b.plot.bar()`

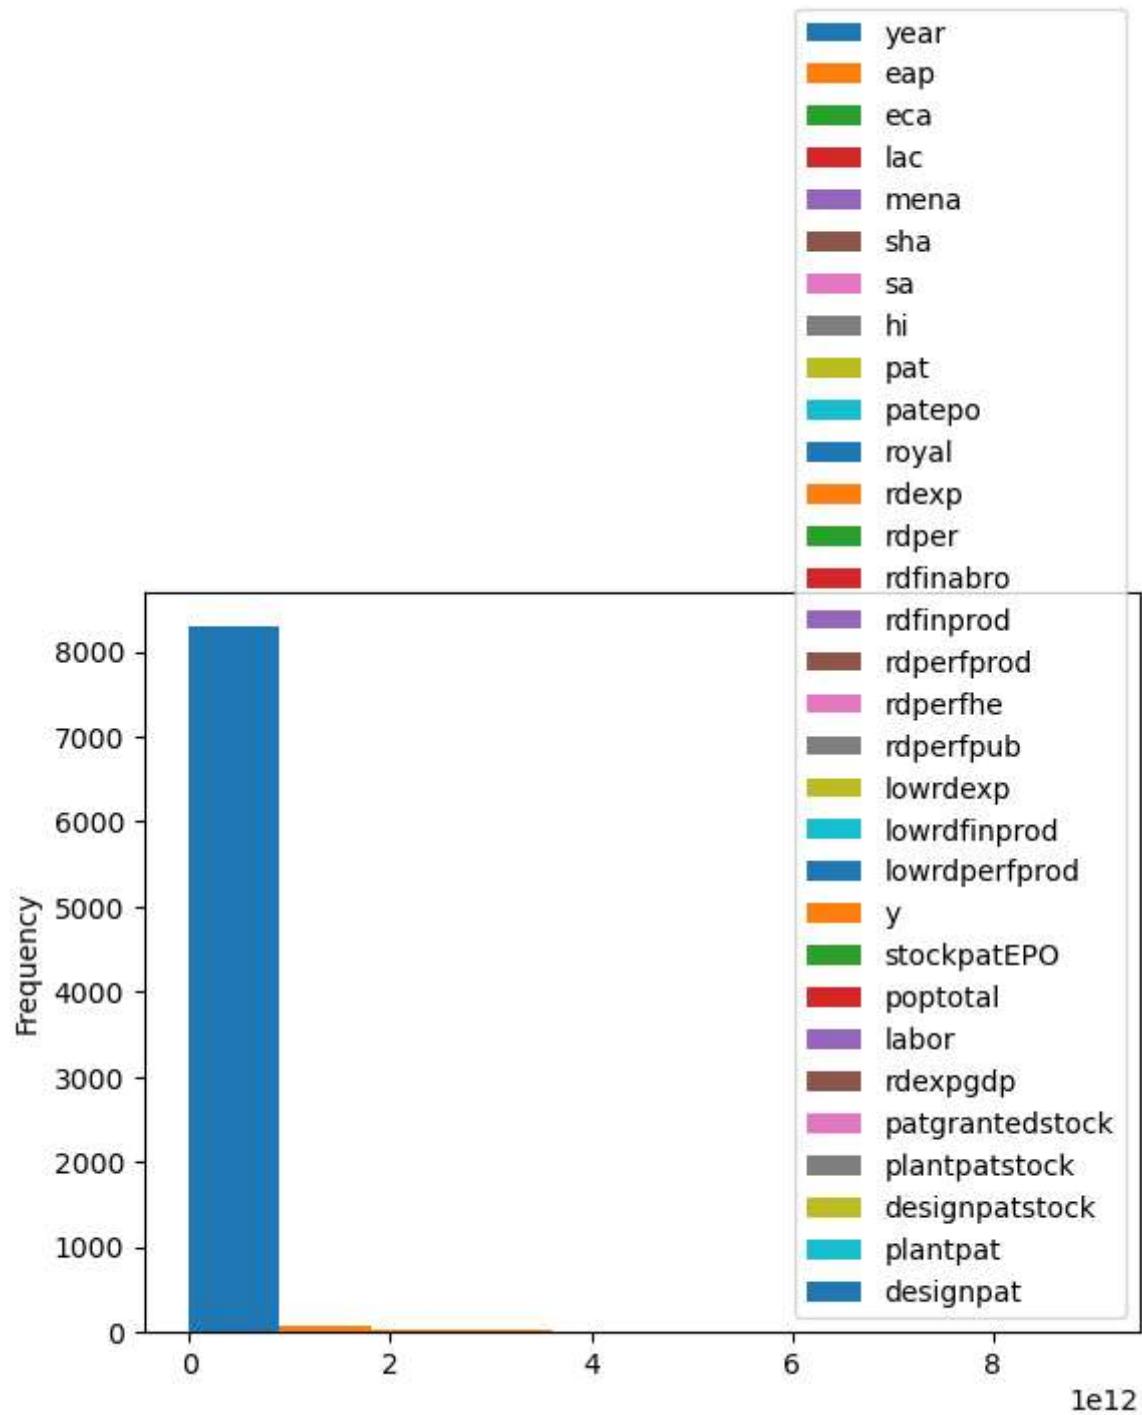
Out[9]: <Axes: >



Histogram

```
In [10]: df.plot.hist()
```

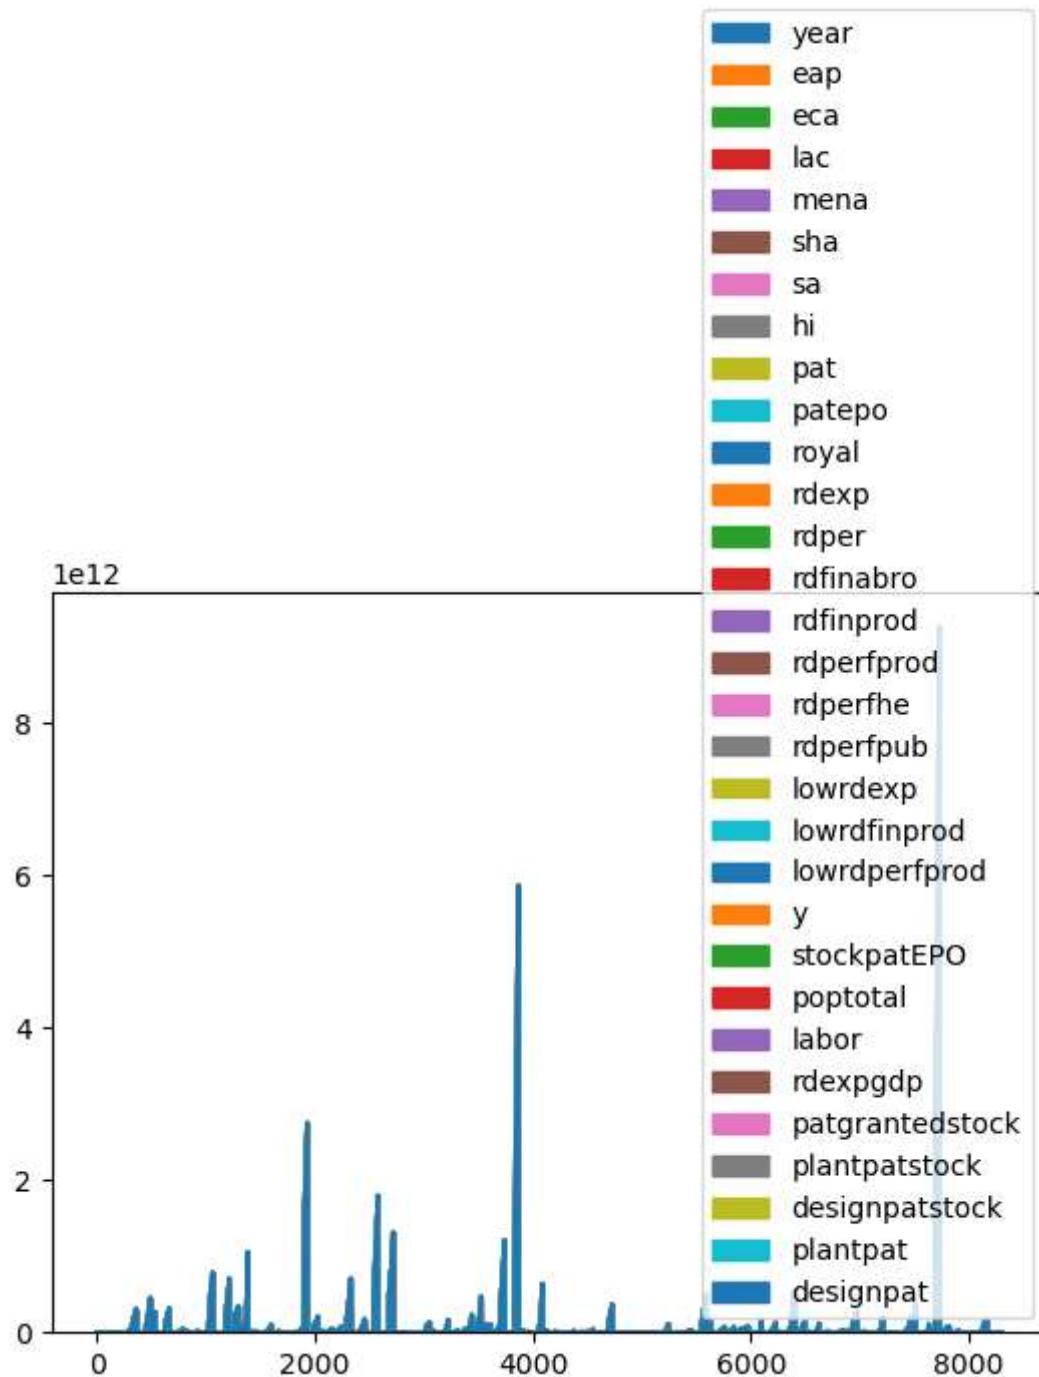
```
Out[10]: <Axes: ylabel='Frequency'>
```



Area chart

```
In [11]: df.plot.area()
```

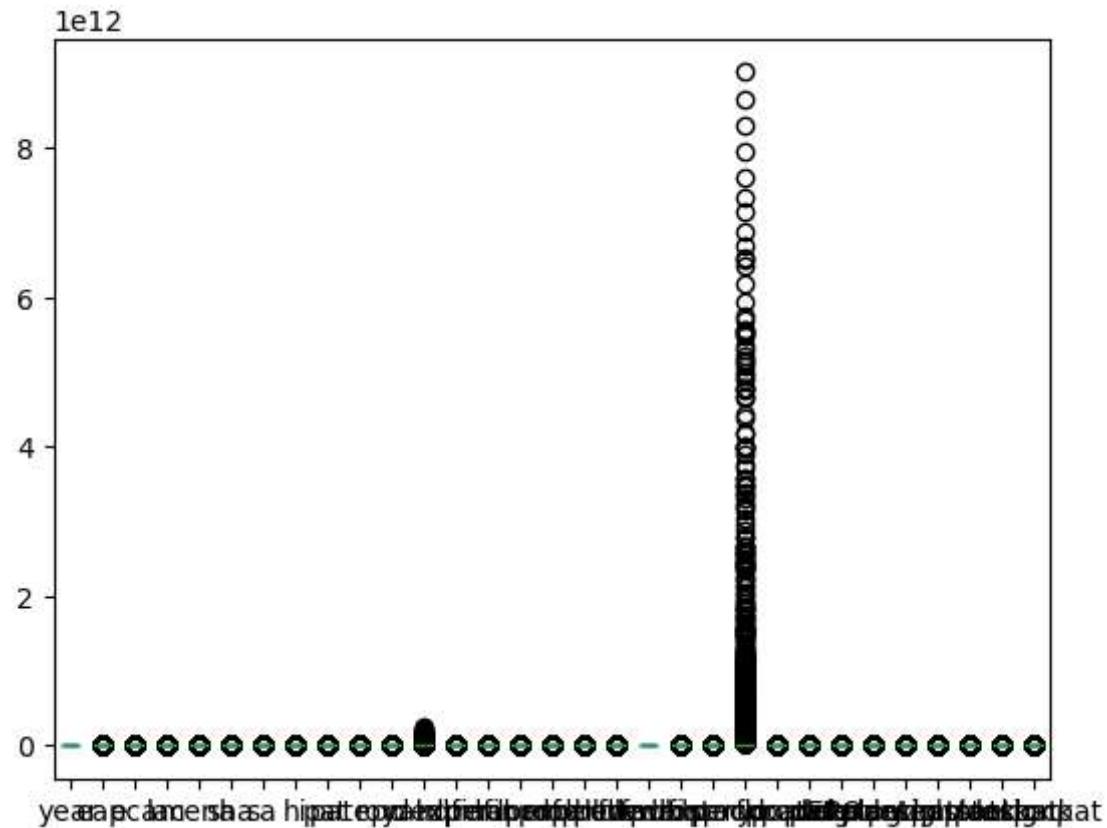
```
Out[11]: <Axes: >
```



Box chart

In [12]: df.plot.box()

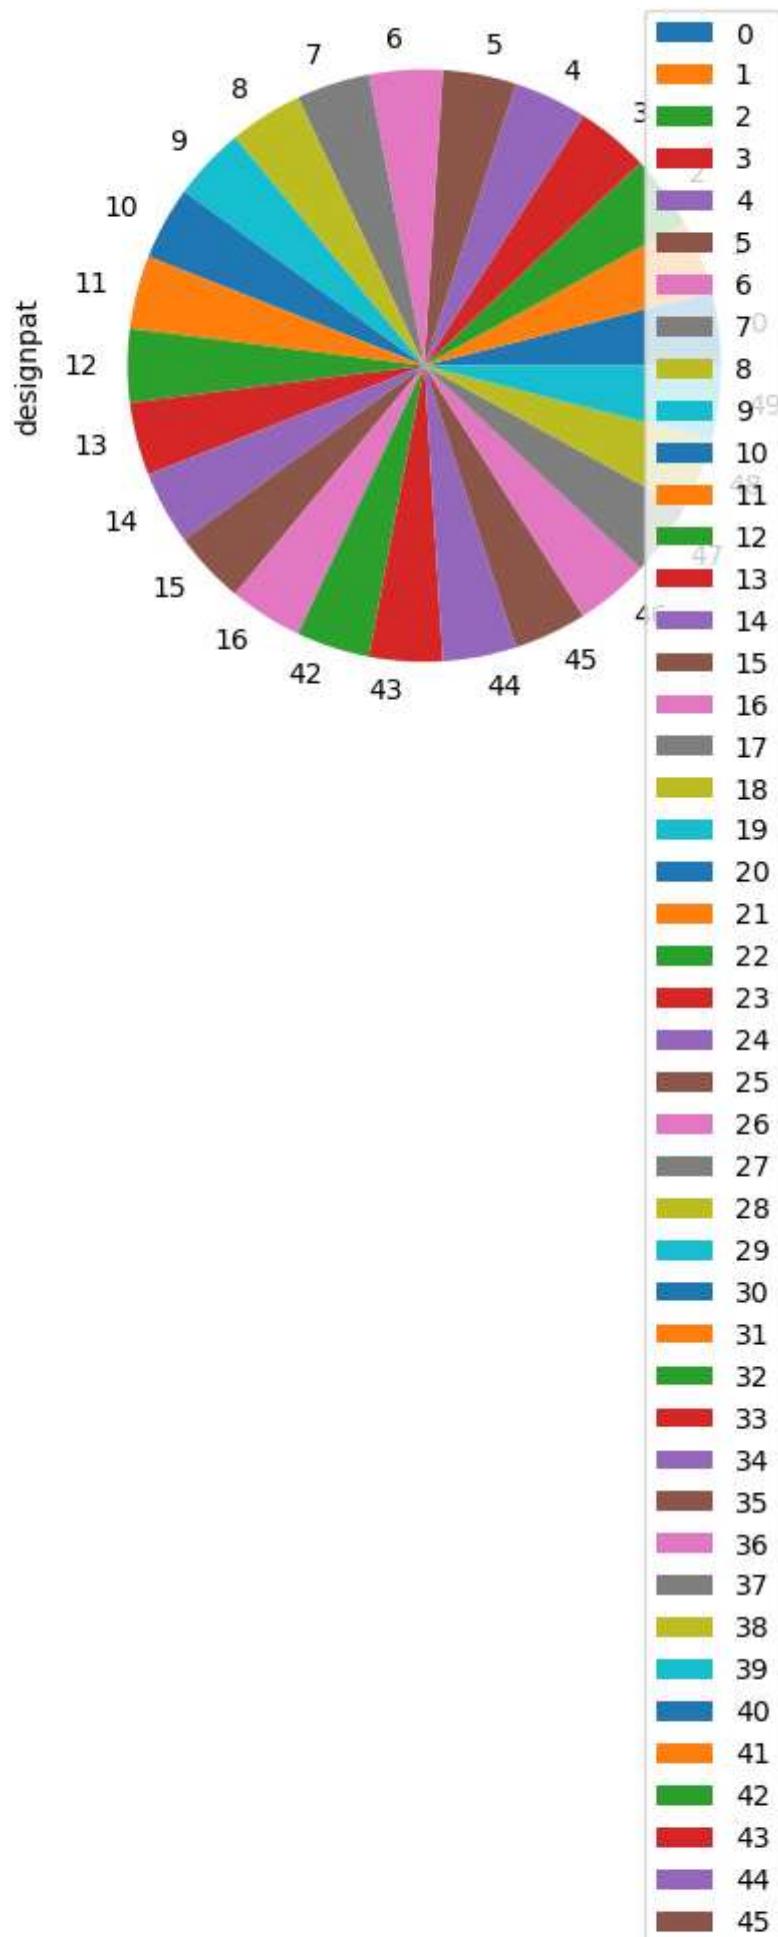
Out[12]: <Axes: >



Pie chart

```
In [13]: b.plot.pie(y='designpat' )
```

```
Out[13]: <Axes: ylabel='designpat'>
```

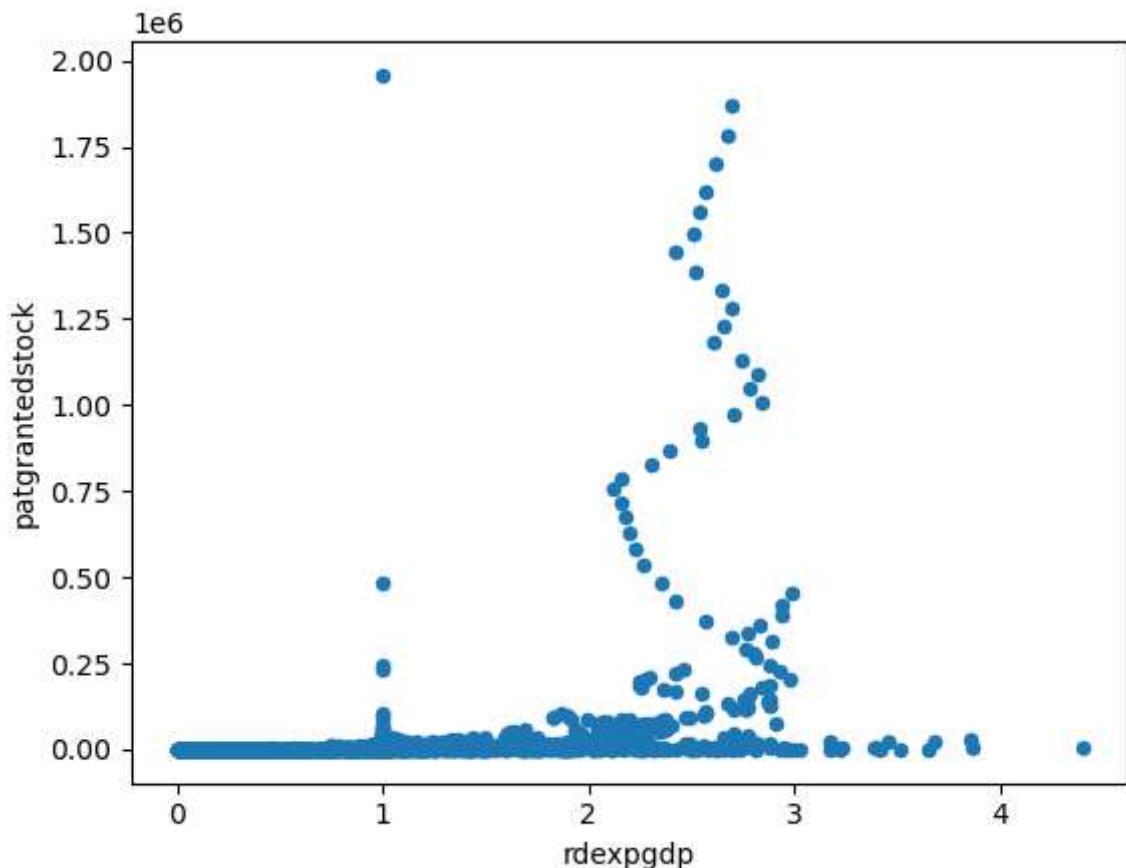





Scatter chart

```
In [14]: df.plot.scatter( x='rdexpgdp',y= 'patgrantedstock')
```

```
Out[14]: <Axes: xlabel='rdexpgdp', ylabel='patgrantedstock'>
```



In [15]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8295 entries, 0 to 8294
Data columns (total 33 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   country          8295 non-null    object  
 1   code              8295 non-null    object  
 2   year              8295 non-null    int64  
 3   eap               8295 non-null    int64  
 4   eca               8295 non-null    int64  
 5   lac               8295 non-null    int64  
 6   mena              8295 non-null    int64  
 7   sha               8295 non-null    int64  
 8   sa                8295 non-null    int64  
 9   hi                8295 non-null    float64 
 10  pat               8295 non-null    float64 
 11  patepo            8295 non-null    float64 
 12  royal              8295 non-null    float64 
 13  rdexp              8295 non-null    float64 
 14  rdper              8295 non-null    float64 
 15  rdfinabro          8295 non-null    float64 
 16  rdfinprod          8295 non-null    float64 
 17  rdperfprod          8295 non-null    float64 
 18  rdperfhe            8295 non-null    float64 
 19  rdperfpub            8295 non-null    float64 
 20  lowrdexp             8295 non-null    float64 
 21  lowrdfinprod          8295 non-null    float64 
 22  lowrdperfprod          8295 non-null    float64 
 23  y                  8295 non-null    float64 
 24  stockpatEPO           8295 non-null    float64 
 25  poptotal            8295 non-null    float64 
 26  labor              8295 non-null    float64 
 27  rdexpgdp             8295 non-null    float64 
 28  patgrantedstock          8295 non-null    float64 
 29  plantpatstock           8295 non-null    float64 
 30  designpatstock           8295 non-null    float64 
 31  plantpat              8295 non-null    float64
```

In [16]: df.describe()

Out[16]:

	year	eap	eca	lac	mena	sha	hi
count	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000	8295.000000
mean	1981.203014	0.094515	0.195901	0.176974	0.098614	0.150090	0.026524
std	12.421590	0.292561	0.396917	0.381670	0.298161	0.357182	0.160691
min	1960.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1970.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	1981.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	1992.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	2002.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000

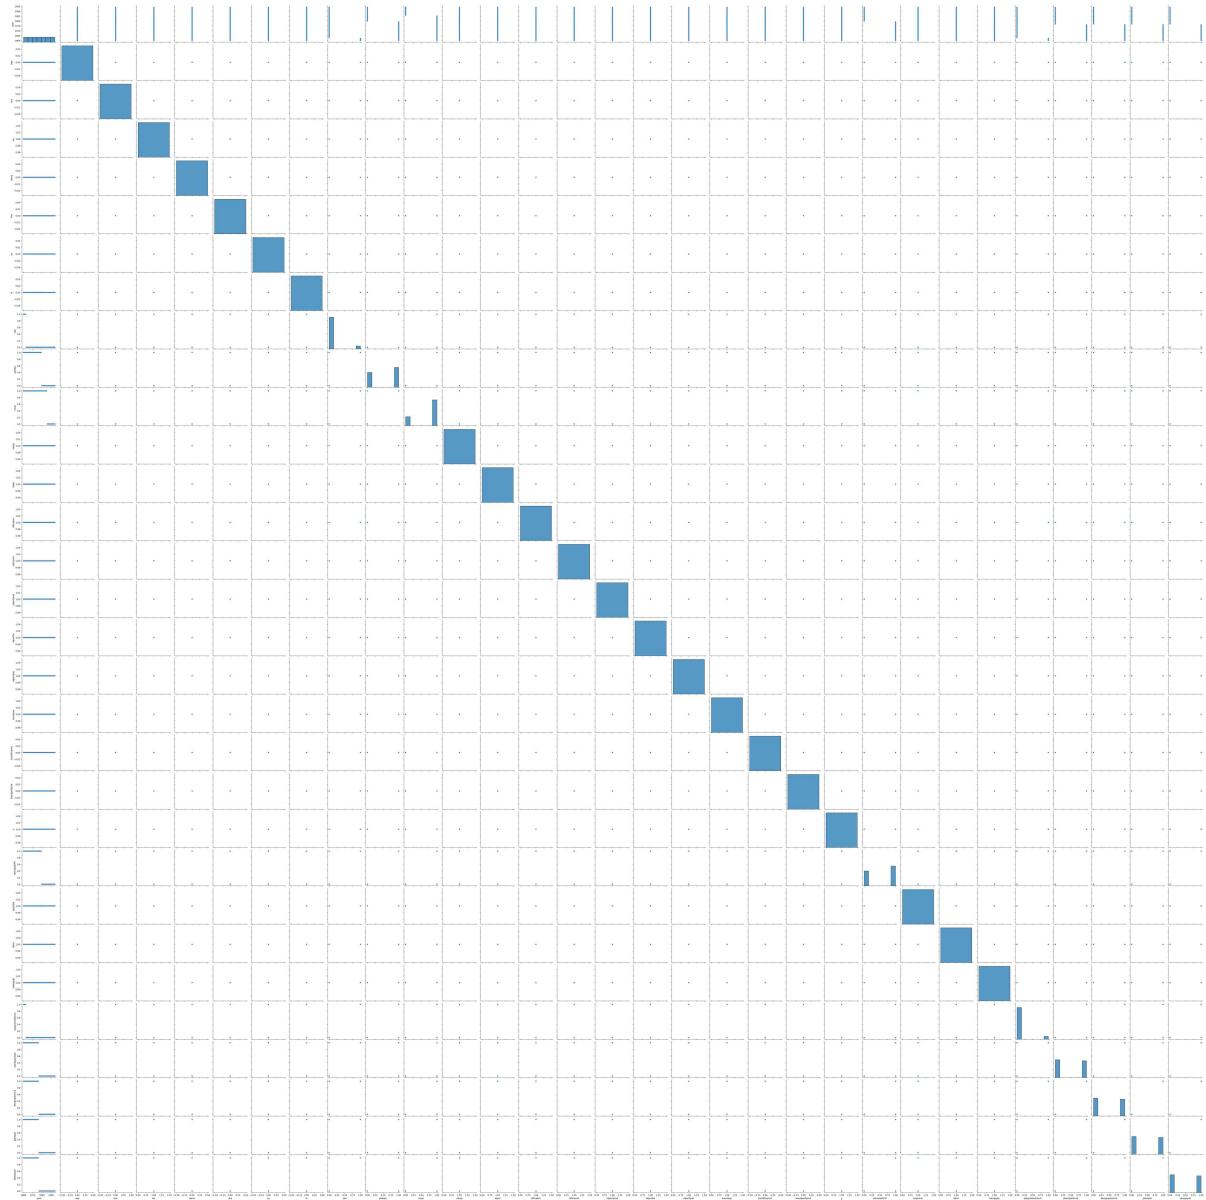
8 rows × 31 columns

In [17]: df1=df[['year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa', 'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdfinabro', 'rdfinprod', 'rdperfprod', 'rdperfhe', 'rdperfpub', 'lowrdexp', 'lowrdfinprod', 'lowrdperfprod', 'y', 'stockpatEPO', 'poptotal', 'labor', 'rdexpgdp', 'patgrantedstock', 'plantpatstock', 'designpatstock', 'plantpat', 'designpat']]

EDA AND VISUALIZATION

In [18]: `sns.pairplot(df1[0:35])`

Out[18]: <seaborn.axisgrid.PairGrid at 0x238624632e0>



```
In [20]: sns.distplot(df1['designpat'])
```

C:\Users\USER\AppData\Local\Temp\ipykernel_13864\3705146837.py:1: UserWarning:

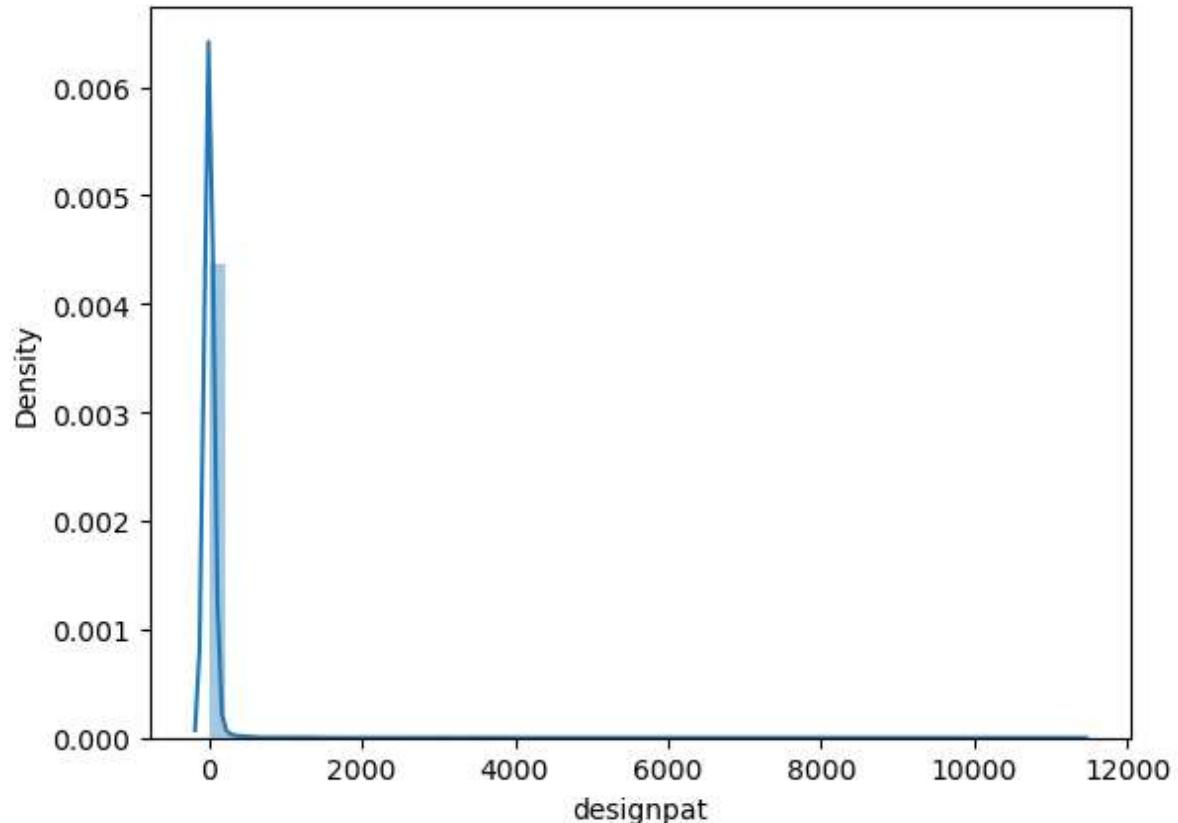
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751> (<https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>)

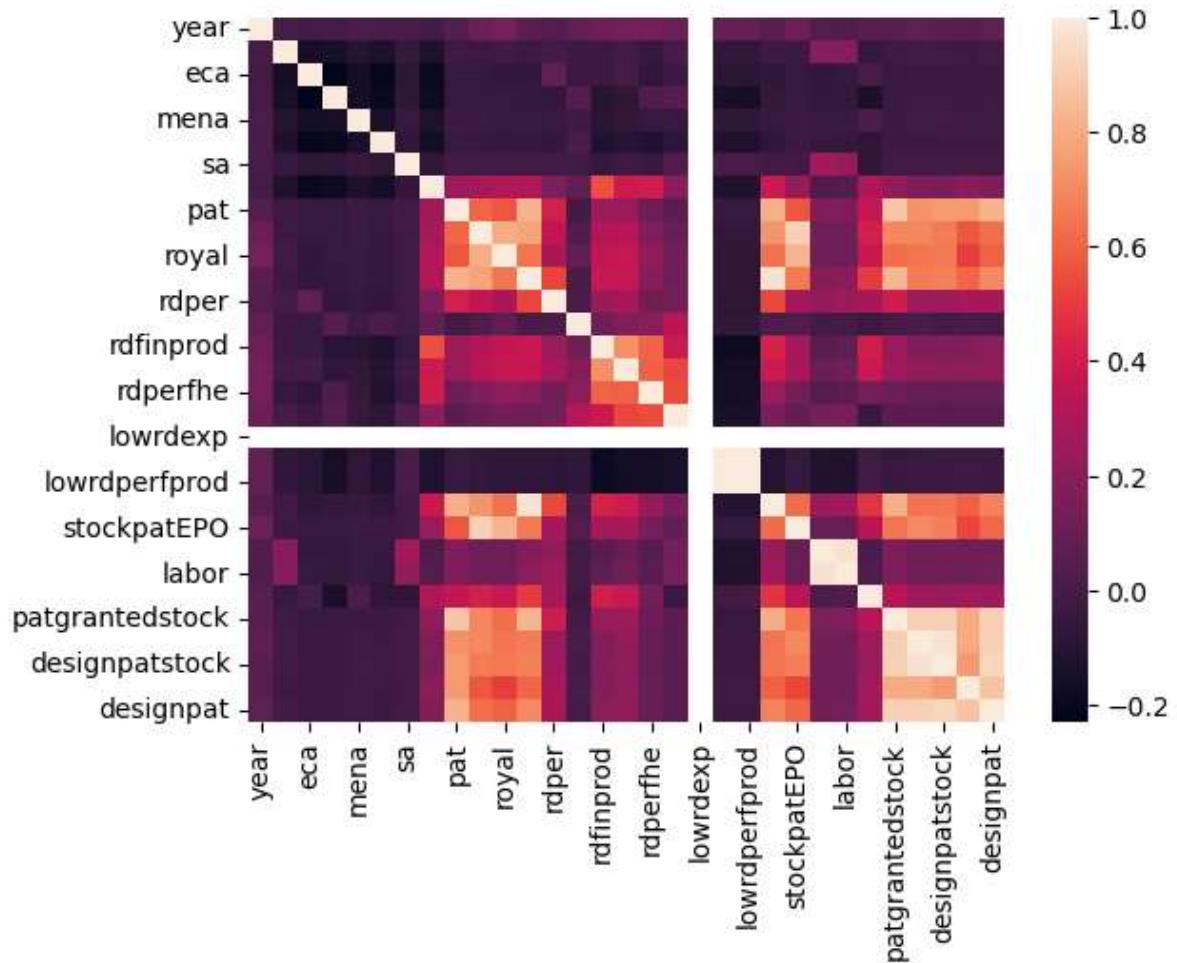
```
sns.distplot(df1['designpat'])
```

Out[20]: <Axes: xlabel='designpat', ylabel='Density'>



In [21]: `sns.heatmap(df1.corr())`

Out[21]: <Axes: >



TO TRAIN THE MODEL AND MODEL BUILDING

In [22]: `x=df[['year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa',
 'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdfinabro',
 'rdfinprod', 'rdperfprod', 'rdperfhe', 'rdperfpub', 'lowrdexp',
 'lowrdfinprod', 'lowrdperfprod', 'y', 'stockpatEPO', 'poptotal',
 'labor', 'rdexpgdp', 'patgrantedstock', 'plantpatstock',
 'designpatstock', 'plantpat']]
 y=df['designpat']`

In [23]: `from sklearn.model_selection import train_test_split
 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)`

Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]:
```

```
  ▾ LinearRegression  
  LinearRegression()
```

```
In [25]: lr.intercept_
```

```
Out[25]: -785.1242838967256
```

In [26]: `coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])
coeff`

Out[26]:

	Co-efficient
year	3.957525e-01
eap	7.187019e+00
eca	-2.280144e+00
lac	6.026328e-01
mena	7.132861e-01
sha	1.122245e+00
sa	1.387282e+00
hi	-1.348329e+01
pat	1.222491e-02
patepo	1.058544e-02
royal	-2.256661e-08
rdexp	-8.631242e-10
rdper	-1.941006e-05
rdfinabro	-5.749730e-02
rdfinprod	4.299086e-01
rdperfprod	8.627065e-02
rdperfhe	1.013243e-02
rdperfpub	-8.539015e-02
lowrdexp	-1.594280e-13
lowrdfinprod	1.421951e+00
lowrdperfprod	4.295657e-01
y	1.183526e-11
stockpatEPO	-1.202728e-03
poptotal	-1.664159e-08
labor	1.937966e-08
rdexpgdp	-5.340751e+00
patgrantedstock	-7.851185e-04
plantpatstock	-9.294575e-01
designpatstock	1.018973e-01
plantpat	1.275950e+01

```
In [27]: prediction = lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x2382d5e5f00>
```



ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.9132963219687931
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.966292141886354
```

Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
C:\ProgramData\anaconda3\lib\site-packages\sklearn\linear_model\_ridge.py:21  
6: LinAlgWarning: Ill-conditioned matrix (rcond=6.65421e-27): result may not  
be accurate.  
    return linalg.solve(A, Xy, assume_a="pos", overwrite_a=True).T
```

```
Out[31]:
```

```
▼      Ridge  
Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.9132838297272209
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.966292058632804
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

```
Out[34]: 
▼ Lasso
Lasso(alpha=10)
```

```
In [35]: la.score(x_train,y_train)
```

```
Out[35]: 0.9660737773836561
```

Accuracy(Lasso)

```
In [36]: la.score(x_test,y_test)
```

```
Out[36]: 0.9130973620296752
```

ElasticNet

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

```
Out[37]: 
▼ ElasticNet
ElasticNet()
```

```
In [38]: en.coef_
```

```
Out[38]: array([ 4.20549621e-01,  3.76106260e-01, -0.00000000e+00,  0.00000000e+00,
   0.00000000e+00,  0.00000000e+00,  0.00000000e+00, -8.52201281e-01,
   1.24638348e-02,  1.12625937e-02, -2.36533482e-08, -6.89583803e-10,
  -1.78042545e-05, -2.64369890e-02,  2.79827944e-01,  6.54884840e-02,
  -2.58576248e-02, -1.76266534e-02,  0.00000000e+00,  0.00000000e+00,
   0.00000000e+00,  2.35796724e-12, -1.22671515e-03, -9.61306119e-09,
   2.41098171e-08, -0.00000000e+00, -7.83809412e-04, -9.24271102e-01,
   1.02248713e-01,  1.25877033e+01])
```

```
In [39]: en.intercept_
```

```
Out[39]: -838.9298157129731
```

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

```
Out[41]: 0.9128762036800263
```

Evaluation Metrics

```
In [42]: from sklearn import metrics  
print(metrics.mean_absolute_error(y_test,prediction))  
print(metrics.mean_squared_error(y_test,prediction))  
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
14.263756514423774
```

```
13251.195703099474
```

```
115.1138380174142
```

Logistic Regression

```
In [43]: from sklearn.linear_model import LogisticRegression
```

```
In [44]: feature_matrix=df[['year', 'eap', 'eca', 'lac', 'mena', 'sha', 'sa',  
'hi', 'pat', 'patepo', 'royal', 'rdexp', 'rdper', 'rdfinabro',  
'rdfinprod', 'rdperfprod', 'rdperfhe', 'rdperfpub', 'lowrdexp',  
'lowrdfinprod', 'lowrdperfprod', 'y', 'stockpatEPO', 'poptotal',  
'labor', 'rdexpgdp', 'patgrantedstock', 'plantpatstock',  
'designpatstock', 'plantpat']]  
target_vector=df['designpat']
```

```
In [45]: feature_matrix.shape
```

```
Out[45]: (8295, 30)
```

```
In [46]: target_vector.shape
```

```
Out[46]: (8295,)
```

```
In [47]: from sklearn.preprocessing import StandardScaler
```

```
In [48]: fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression
          LogisticRegression(max_iter=10000)
```

```
In [50]: observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14,15,16,17,18,19,20,21,22,23,24,2
```

```
In [51]: prediction=logr.predict(observation)
          print(prediction)
```

```
[212.]
```

In [52]: `logr.classes_`

Out[52]: `array([0.0000e+00, 1.0000e+00, 2.0000e+00, 3.0000e+00, 4.0000e+00, 5.0000e+00, 6.0000e+00, 7.0000e+00, 8.0000e+00, 9.0000e+00, 1.0000e+01, 1.1000e+01, 1.2000e+01, 1.3000e+01, 1.4000e+01, 1.5000e+01, 1.6000e+01, 1.7000e+01, 1.8000e+01, 1.9000e+01, 2.0000e+01, 2.1000e+01, 2.2000e+01, 2.3000e+01, 2.4000e+01, 2.5000e+01, 2.6000e+01, 2.7000e+01, 2.8000e+01, 2.9000e+01, 3.0000e+01, 3.2000e+01, 3.3000e+01, 3.4000e+01, 3.7000e+01, 3.8000e+01, 3.9000e+01, 4.0000e+01, 4.1000e+01, 4.2000e+01, 4.3000e+01, 4.4000e+01, 4.5000e+01, 4.6000e+01, 4.7000e+01, 4.8000e+01, 4.9000e+01, 5.0000e+01, 5.4000e+01, 5.5000e+01, 5.6000e+01, 5.7000e+01, 5.8000e+01, 5.9000e+01, 6.0000e+01, 6.1000e+01, 6.2000e+01, 6.3000e+01, 6.4000e+01, 6.5000e+01, 6.6000e+01, 6.7000e+01, 6.9000e+01, 7.0000e+01, 7.1000e+01, 7.2000e+01, 7.3000e+01, 7.4000e+01, 7.6000e+01, 7.7000e+01, 7.8000e+01, 8.0000e+01, 8.1000e+01, 8.2000e+01, 8.3000e+01, 8.4000e+01, 8.5000e+01, 8.6000e+01, 8.7000e+01, 8.8000e+01, 8.9000e+01, 9.0000e+01, 9.1000e+01, 9.3000e+01, 9.4000e+01, 9.5000e+01, 9.6000e+01, 9.7000e+01, 9.8000e+01, 9.9000e+01, 1.0000e+02, 1.0100e+02, 1.0200e+02, 1.0300e+02, 1.0500e+02, 1.0600e+02, 1.0700e+02, 1.0900e+02, 1.1000e+02, 1.1200e+02, 1.1300e+02, 1.1400e+02, 1.1500e+02, 1.1700e+02, 1.1800e+02, 1.1900e+02, 1.2000e+02, 1.2100e+02, 1.2300e+02, 1.2400e+02, 1.2500e+02, 1.2600e+02, 1.2700e+02, 1.2900e+02, 1.3200e+02, 1.3300e+02, 1.3400e+02, 1.3600e+02, 1.3800e+02, 1.3900e+02, 1.4200e+02, 1.4300e+02, 1.4400e+02, 1.4700e+02, 1.5000e+02, 1.5200e+02, 1.5300e+02, 1.5600e+02, 1.5900e+02, 1.6000e+02, 1.6200e+02, 1.6300e+02, 1.6500e+02, 1.6900e+02, 1.7200e+02, 1.7300e+02, 1.7600e+02, 1.7900e+02, 1.8000e+02, 1.8100e+02, 1.8200e+02, 1.8400e+02, 1.8500e+02, 1.8600e+02, 1.9000e+02, 1.9300e+02, 1.9500e+02, 1.9600e+02, 1.9700e+02, 2.0100e+02, 2.0200e+02, 2.0500e+02, 2.0800e+02, 2.1100e+02, 2.1200e+02, 2.1300e+02, 2.1500e+02, 2.1600e+02, 2.2100e+02, 2.2200e+02, 2.2700e+02, 2.2800e+02, 2.3000e+02, 2.3100e+02, 2.3400e+02, 2.3700e+02, 2.3900e+02, 2.4000e+02, 2.4300e+02, 2.4700e+02, 2.5000e+02, 2.5300e+02, 2.5400e+02, 2.5700e+02, 2.5800e+02, 2.6000e+02, 2.6500e+02, 2.7500e+02, 2.8200e+02, 3.0000e+02, 3.0600e+02, 3.2000e+02, 3.3000e+02, 3.3800e+02, 3.4100e+02, 3.5000e+02, 3.5600e+02, 3.6000e+02, 3.6800e+02, 3.7000e+02, 3.7200e+02, 3.8200e+02, 3.9000e+02, 3.9600e+02, 4.0100e+02, 4.1000e+02, 4.1800e+02, 4.3800e+02, 4.3900e+02, 4.6600e+02, 4.8200e+02, 4.8400e+02, 4.8500e+02, 5.0300e+02, 5.0500e+02, 5.0900e+02, 5.2200e+02, 5.3900e+02, 5.4700e+02, 5.7600e+02, 5.8800e+02, 6.2000e+02, 6.9800e+02, 7.0500e+02, 7.9500e+02, 8.3300e+02, 8.6100e+02, 9.3600e+02, 9.4800e+02, 1.0260e+03, 1.0370e+03, 1.0560e+03, 1.1350e+03, 1.1370e+03, 1.1490e+03, 1.1680e+03, 1.2070e+03, 1.2940e+03, 1.3070e+03, 1.3100e+03, 1.3640e+03, 1.4970e+03, 1.5460e+03, 2.4690e+03, 3.0520e+03, 3.0550e+03, 3.0650e+03, 3.2780e+03, 3.4280e+03, 3.4460e+03, 3.4750e+03, 3.5460e+03, 3.5700e+03, 3.6450e+03, 3.8830e+03, 3.9020e+03, 5.0690e+03, 6.0130e+03, 6.0750e+03, 7.4160e+03, 7.6970e+03, 7.7470e+03, 7.8630e+03, 8.2510e+03, 9.3250e+03, 9.6540e+03, 9.9130e+03, 1.0346e+04, 1.1285e+04])`

```
In [53]: logr.score(fs,target_vector)
```

```
Out[53]: 0.8763110307414105
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 0.0
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([[0.00000000e+00, 7.37298243e-85, 3.91060558e-86, 4.57591854e-80,
 8.17391709e-83, 2.12264191e-56, 3.81522172e-68, 2.46900847e-74,
3.76048442e-68, 7.25475358e-60, 4.73052191e-85, 5.13680225e-55,
9.53516840e-74, 1.06274517e-59, 1.03519482e-50, 8.02008700e-68,
1.76246137e-50, 2.80704013e-61, 4.29949277e-50, 6.82350359e-49,
3.28463252e-52, 7.46248036e-50, 3.69048735e-44, 1.01089687e-70,
1.95377914e-32, 1.37129799e-51, 2.97734078e-25, 1.08619424e-42,
3.64918135e-36, 3.61000410e-58, 2.36698757e-38, 4.09631532e-33,
4.21068349e-49, 1.95886535e-43, 1.45589951e-54, 1.26567419e-50,
1.56091794e-22, 2.09150015e-29, 1.19950895e-34, 8.70042058e-30,
2.08847713e-44, 3.92847541e-15, 5.11190552e-36, 3.79636173e-12,
1.63336791e-26, 5.79073514e-31, 1.85014185e-18, 1.58628939e-35,
4.00701759e-25, 6.94094671e-09, 1.77509224e-36, 7.10277816e-19,
9.14293882e-31, 3.19498459e-26, 4.23231442e-32, 3.36602209e-26,
3.11543195e-38, 2.03118856e-46, 1.57576773e-24, 5.50419983e-29,
4.35270183e-14, 4.99698872e-35, 5.40475297e-13, 3.87058945e-43,
1.82955140e-38, 5.64465160e-41, 1.71473076e-22, 9.97950063e-29,
4.17987001e-32, 2.21514878e-21, 1.32505012e-13, 2.87354482e-17,
1.90987336e-27, 4.38501647e-37, 5.79577486e-11, 2.43559393e-04,
9.68305316e-34, 8.89984336e-25, 1.55495412e-19, 1.38727283e-26,
3.41249500e-25, 3.64647789e-25, 5.26127881e-21, 1.06382046e-32,
4.09959239e-30, 2.69988301e-14, 1.77895816e-16, 2.29319128e-31,
4.69209523e-20, 1.08043494e-31, 6.66081511e-18, 7.97780936e-15,
9.06834607e-20, 4.89729025e-12, 1.40296280e-26, 2.56888330e-28,
3.50040540e-23, 1.40505898e-20, 1.44768405e-35, 2.25809627e-19,
2.78514612e-28, 3.05925885e-25, 2.01444825e-25, 2.77986579e-47,
3.90179831e-38, 1.95789134e-26, 5.13561823e-29, 2.90465495e-43,
2.46749241e-24, 1.59644437e-22, 1.42446443e-44, 4.50938335e-24,
1.13018689e-27, 1.27201921e-12, 4.04835136e-18, 1.45199763e-36,
6.06801672e-25, 9.84779095e-17, 2.39086869e-31, 3.03801479e-28,
2.71070452e-19, 1.02077347e-20, 8.25400919e-15, 1.02235612e-23,
9.82253276e-19, 3.52672575e-23, 5.80639277e-11, 1.43391307e-29,
1.98005831e-26, 3.00611013e-35, 4.11542311e-41, 9.90008623e-43,
2.10702094e-18, 4.66506556e-11, 2.26197054e-22, 1.52022135e-19,
1.06402978e-23, 1.22412990e-17, 5.58098881e-41, 1.07447174e-12,
2.68853014e-12, 4.93151010e-06, 1.29962630e-37, 5.01745991e-07,
3.47147071e-26, 3.26754889e-11, 6.33815450e-18, 3.97823858e-11,
6.23061618e-12, 4.88047433e-10, 4.60046666e-16, 2.06050476e-29,
2.67597016e-22, 1.49443276e-23, 9.99242277e-01, 1.79190853e-40,
1.43453397e-08, 3.87186752e-33, 1.28855360e-18, 6.11715517e-17,
1.55178169e-21, 2.30500300e-13, 3.87067774e-09, 3.34342145e-20,
1.03550559e-13, 2.50075847e-16, 2.66556954e-19, 5.15625777e-10,
3.95312702e-18, 2.16203465e-23, 8.26129216e-15, 1.83303171e-10,
1.98082923e-34, 9.93381574e-11, 2.34895819e-34, 6.62295795e-16,
8.44466966e-24, 1.38346301e-10, 9.11918021e-05, 1.32167615e-13,
6.06196749e-16, 2.33862024e-26, 4.71430452e-15, 1.40132304e-13,
2.95020036e-18, 4.31926632e-17, 2.35229003e-26, 3.20306264e-11,
1.15341742e-31, 2.33867238e-24, 3.19561667e-07, 1.34557006e-40,
1.42392518e-16, 1.61781243e-25, 5.13663060e-18, 3.09949948e-04,
1.26366802e-10, 3.52922290e-12, 2.06905300e-21, 5.20685146e-23,
2.65047516e-17, 4.33950755e-10, 1.64277819e-07, 4.27560643e-20,
1.05451971e-04, 1.53962720e-19, 3.28532575e-21, 1.91426528e-20,
7.03698449e-12, 2.06743782e-22, 5.20658652e-22, 9.44102390e-20,
4.43029008e-19, 2.39421742e-18, 5.94203832e-17, 1.27354304e-15,
5.02692361e-17, 8.68296241e-19, 2.16594941e-17, 8.41980378e-15,
6.20409912e-18, 8.48302492e-21, 1.96932159e-13, 9.92060527e-15,
5.67647798e-15, 7.41013080e-26, 1.54541595e-11, 1.57470260e-24,
```

```
2.28307783e-18, 2.96395503e-12, 5.19565122e-26, 1.79030842e-10,  
6.35598379e-20, 1.15201253e-18, 6.86160819e-21, 1.76777588e-18,  
8.34429323e-13, 1.25286156e-08, 2.09997113e-13, 2.78330570e-08,  
1.49224723e-11, 1.02061972e-11, 2.52582135e-13, 5.76603530e-07,  
3.02732364e-09, 6.01603889e-11, 3.66949947e-13, 2.68530288e-09,  
5.50477419e-07, 9.89413877e-09, 2.03030087e-08, 4.16722719e-07,  
2.75858662e-11, 2.55808065e-11, 2.22765734e-21, 1.30234666e-13,  
3.36151853e-09, 5.95469588e-12, 1.20970424e-09]])
```

Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier  
| RandomForestClassifier()
```

```
In [58]: parameters={  
    'max_depth':[1,2,3,4,5],  
    'min_samples_leaf':[5,10,15,20,25],  
    'n_estimators':[10,20,30,40,50]  
}
```

```
In [59]: from sklearn.model_selection import GridSearchCV  
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="acc")  
grid_search.fit(x_train,y_train)
```

```
C:\ProgramData\anaconda3\lib\site-packages\sklearn\model_selection\_split.py:  
700: UserWarning: The least populated class in y has only 1 members, which is  
less than n_splits=2.  
    warnings.warn(
```

```
Out[59]: GridSearchCV  
| estimator: RandomForestClassifier  
| | RandomForestClassifier()
```

```
In [60]: grid_search.best_score_
```

```
Out[60]: 0.8797795384085428
```

```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree

plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5], feature_names=x.columns, class_names=['variable1', 'variable11', 'variable12', 'variable13', 'variable14', 'variable15', 'variable2', 'variable21', 'variable22', 'variable23', 'variable24', 'variable25', 'variable3', 'variable31', 'variable32', 'variable33', 'variable34', 'variable35', 'variable4', 'variable41', 'variable42', 'variable43', 'variable44', 'variable45', 'variable5', 'variable51', 'variable52', 'variable53', 'variable54', 'variable55', 'variable6', 'variable61', 'variable62', 'variable63', 'variable64', 'variable65', 'variable7', 'variable71', 'variable72', 'variable73', 'variable74', 'variable75', 'variable8', 'variable81', 'variable82', 'variable83', 'variable84', 'variable85', 'variable9', 'variable91', 'variable92', 'variable93', 'variable94', 'variable95', 'variable10'])
```

Conclusion

Accuracy

```
In [63]: print("Linear Regression:",lr.score(x_test,y_test))
print("Ridge Regression:",rr.score(x_test,y_test))
print("Lasso Regression",la.score(x_test,y_test))
print("ElasticNet Regression:",en.score(x_test,y_test))
print("Logistic Regression:",logr.score(fs,target_vector))
print("Random Forest:",grid_search.best_score_)
```

```
Linear Regression: 0.9132963219687931
Ridge Regression: 0.9132838297272209
Lasso Regression: 0.9130973620296752
ElasticNet Regression: 0.9128762036800263
Logistic Regression: 0.8763110307414105
Random Forest: 0.8797795384085428
```

Linear Regression is suitable for this dataset

