

Importing Libraries

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```

1 df=pd.read_csv("madrid_2001")
2 df

```

Out[2]:

| | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PXY | SO_2 | TCH | TOL | ... |
|--------|---------------------|-------|------|------|-------|------|-----------|------------|------|-----------|------------|------|-----------|------|-------|-----|
| 0 | 2001-08-01 01:00:00 | NaN | 0.37 | NaN | NaN | NaN | 58.400002 | 87.150002 | NaN | 34.529999 | 105.000000 | NaN | 6.340000 | NaN | NaN | 280 |
| 1 | 2001-08-01 01:00:00 | 1.50 | 0.34 | 1.49 | 4.10 | 0.07 | 56.250000 | 75.169998 | 2.11 | 42.160000 | 100.599998 | 1.73 | 8.110000 | 1.24 | 10.82 | 280 |
| 2 | 2001-08-01 01:00:00 | NaN | 0.28 | NaN | NaN | NaN | 50.660000 | 61.380001 | NaN | 46.310001 | 100.099998 | NaN | 7.850000 | NaN | NaN | 280 |
| 3 | 2001-08-01 01:00:00 | NaN | 0.47 | NaN | NaN | NaN | 69.790001 | 73.449997 | NaN | 40.650002 | 69.779999 | NaN | 6.460000 | NaN | NaN | 280 |
| 4 | 2001-08-01 01:00:00 | NaN | 0.39 | NaN | NaN | NaN | 22.830000 | 24.799999 | NaN | 66.309998 | 75.180000 | NaN | 8.800000 | NaN | NaN | 280 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 217867 | 2001-04-01 00:00:00 | 10.45 | 1.81 | NaN | NaN | NaN | 73.000000 | 264.399994 | NaN | 5.200000 | 47.880001 | NaN | 39.910000 | NaN | 28.35 | 280 |
| 217868 | 2001-04-01 00:00:00 | 5.20 | 0.69 | 4.56 | NaN | 0.13 | 71.080002 | 129.300003 | NaN | 13.460000 | 26.809999 | NaN | 13.450000 | 1.32 | 16.08 | 280 |
| 217869 | 2001-04-01 00:00:00 | 0.49 | 1.09 | NaN | 1.00 | 0.19 | 76.279999 | 128.399994 | 0.35 | 5.020000 | 40.770000 | 0.61 | 14.700000 | 1.40 | 1.55 | 280 |
| 217870 | 2001-04-01 00:00:00 | 5.62 | 1.01 | 5.04 | 11.38 | NaN | 80.019997 | 197.000000 | 2.58 | 5.840000 | 37.889999 | 4.31 | 39.919998 | NaN | 20.75 | 280 |
| 217871 | 2001-04-01 00:00:00 | 8.09 | 1.62 | 6.66 | 13.04 | 0.18 | 76.809998 | 206.300003 | 5.20 | 8.340000 | 35.369999 | 4.95 | 27.340000 | 1.41 | 22.27 | 280 |

217872 rows × 16 columns



Data Cleaning and Data Preprocessing

```
In [3]: 1 df=df.dropna()
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
               dtype='object')
```

```
In [5]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 29669 entries, 1 to 217871
Data columns (total 16 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      29669 non-null   object 
 1   BEN       29669 non-null   float64
 2   CO        29669 non-null   float64
 3   EBE       29669 non-null   float64
 4   MXY       29669 non-null   float64
 5   NMHC      29669 non-null   float64
 6   NO_2      29669 non-null   float64
 7   NOx       29669 non-null   float64
 8   OXY       29669 non-null   float64
 9   O_3        29669 non-null   float64
 10  PM10      29669 non-null   float64
 11  PXY       29669 non-null   float64
 12  SO_2      29669 non-null   float64
 13  TCH       29669 non-null   float64
 14  TOL       29669 non-null   float64
 15  station    29669 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 3.8+ MB
```

```
In [6]: 1 data=df[['CO' , 'station']]  
2 data
```

Out[6]:

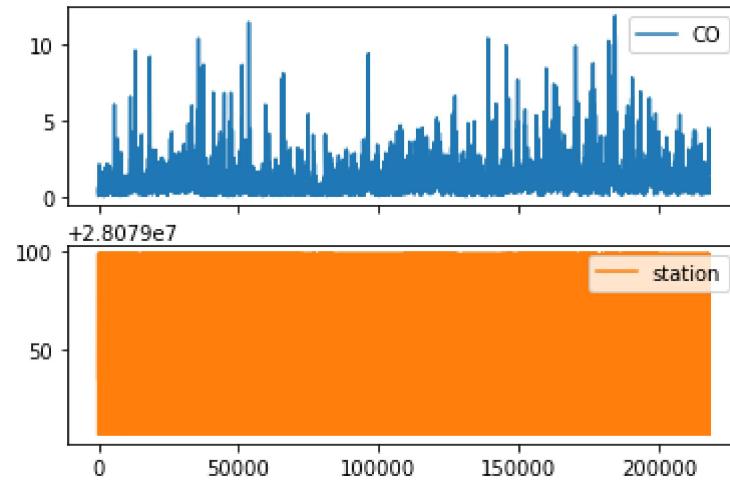
| | CO | station |
|---------------|------|----------|
| 1 | 0.34 | 28079035 |
| 5 | 0.63 | 28079006 |
| 21 | 0.43 | 28079024 |
| 23 | 0.34 | 28079099 |
| 25 | 0.06 | 28079035 |
| ... | ... | ... |
| 217829 | 4.48 | 28079006 |
| 217847 | 2.65 | 28079099 |
| 217849 | 1.22 | 28079035 |
| 217853 | 1.83 | 28079006 |
| 217871 | 1.62 | 28079099 |

29669 rows × 2 columns

Line chart

```
In [7]: 1 data.plot.line(subplots=True)
```

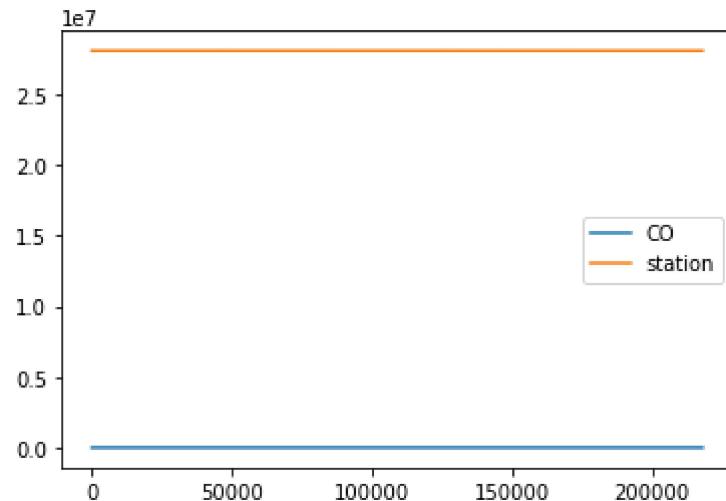
```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



Line chart

```
In [8]: 1 data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

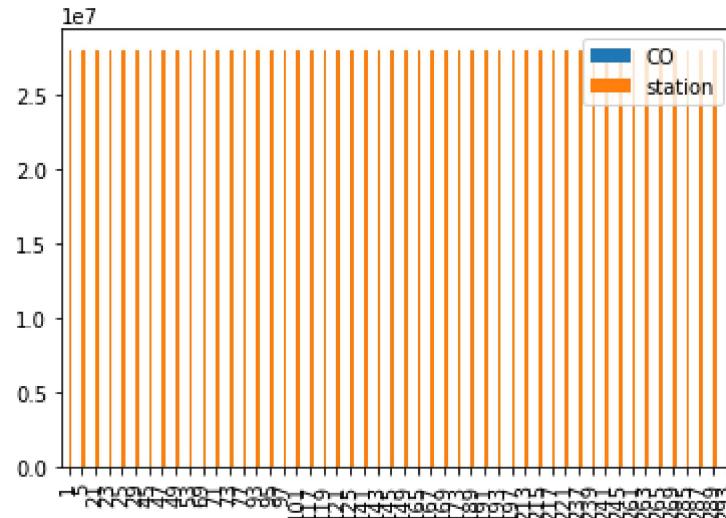


Bar chart

```
In [9]: 1 b=data[0:50]
```

```
In [10]: 1 b.plot.bar()
```

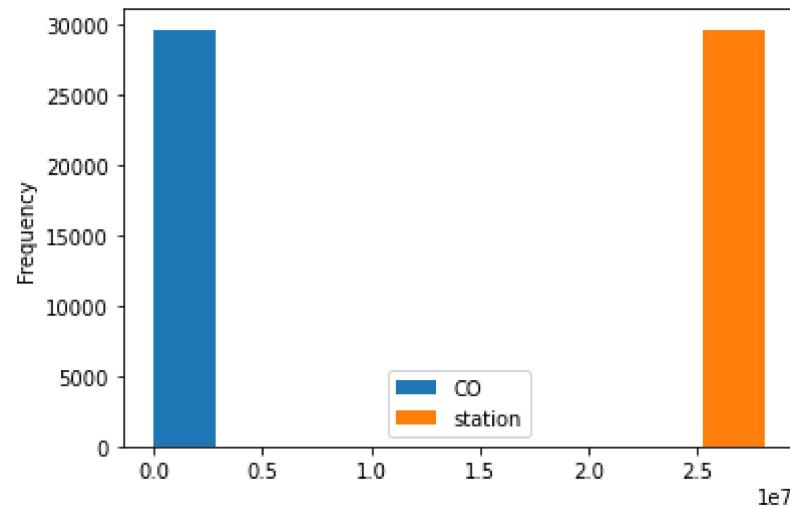
```
Out[10]: <AxesSubplot:>
```



Histogram

```
In [11]: 1 data.plot.hist()
```

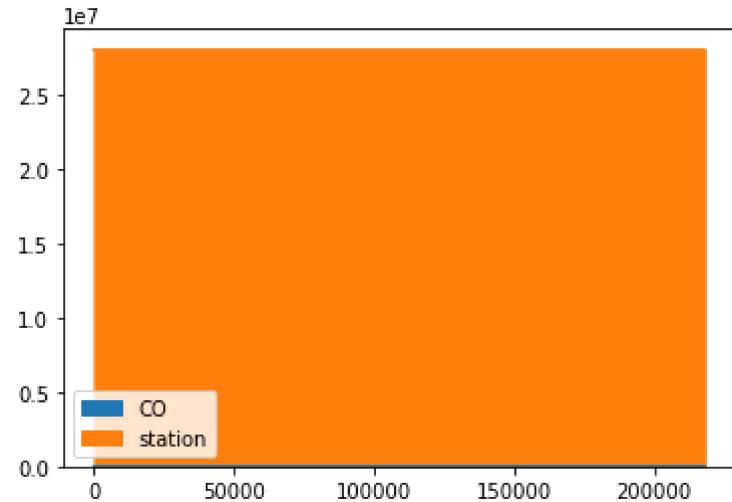
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

In [12]: 1 data.plot.area()

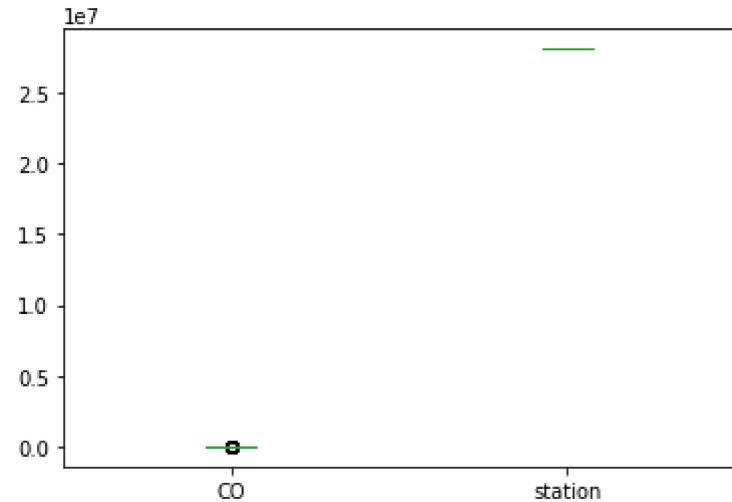
Out[12]: <AxesSubplot:>



Box chart

In [13]: 1 data.plot.box()

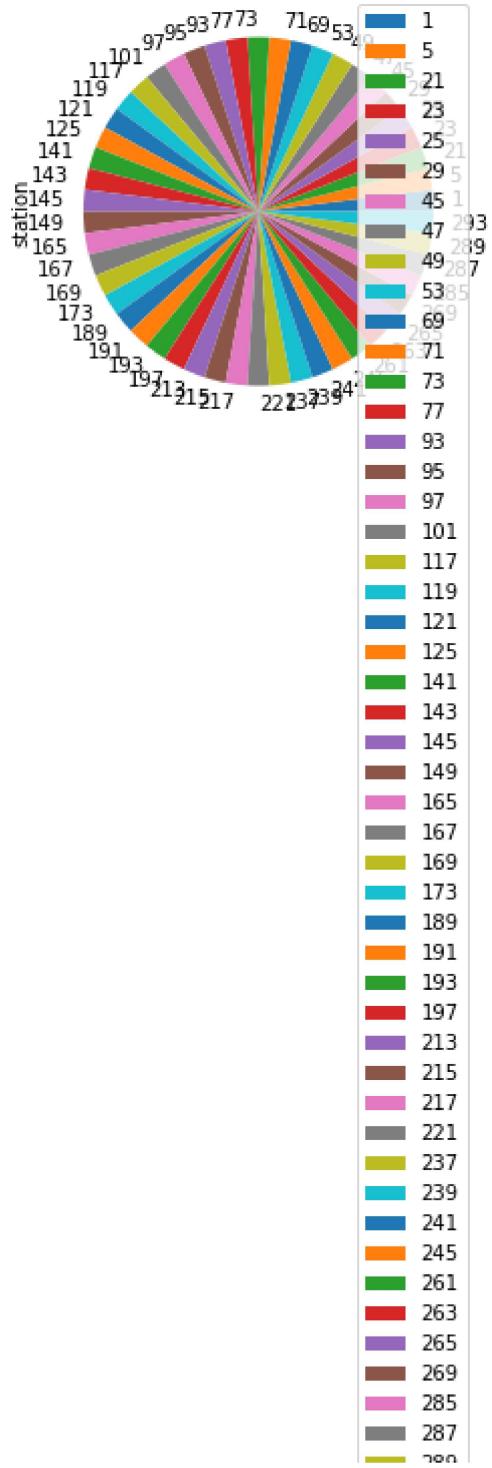
Out[13]: <AxesSubplot:>



Pie chart

```
In [14]: 1 b.plot.pie(y='station' )
```

```
Out[14]: <AxesSubplot:ylabel='station'>
```

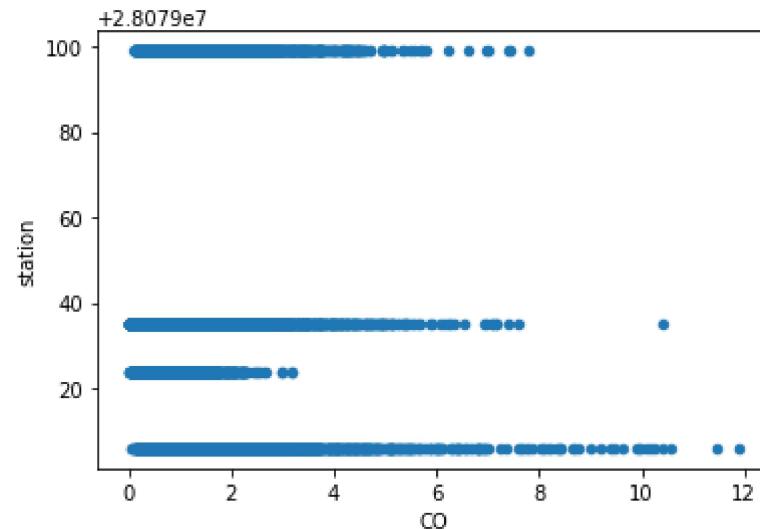





Scatter chart

In [15]: 1 data.plot.scatter(x='CO' ,y='station')

Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>



```
In [16]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 29669 entries, 1 to 217871
Data columns (total 16 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   date      29669 non-null    object  
 1   BEN        29669 non-null    float64 
 2   CO         29669 non-null    float64 
 3   EBE        29669 non-null    float64 
 4   MXY        29669 non-null    float64 
 5   NMHC       29669 non-null    float64 
 6   NO_2       29669 non-null    float64 
 7   NOx        29669 non-null    float64 
 8   OXY        29669 non-null    float64 
 9   O_3         29669 non-null    float64 
 10  PM10       29669 non-null    float64 
 11  PXY        29669 non-null    float64 
 12  SO_2       29669 non-null    float64 
 13  TCH        29669 non-null    float64 
 14  TOL        29669 non-null    float64
```

```
In [17]: 1 df.describe()
```

Out[17]:

| | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 29669.000000 | 29669.000000 | 29669.000000 | 29669.000000 | 29669.000000 | 29669.000000 | 29669.000000 | 29669.000000 | 29669.000000 |
| mean | 3.361895 | 1.005413 | 3.580229 | 8.113086 | 0.195222 | 67.652292 | 163.004858 | 3.736694 | 29.801946 |
| std | 3.176669 | 0.863135 | 3.744496 | 7.909701 | 0.192585 | 34.003120 | 147.491777 | 3.702559 | 24.099665 |
| min | 0.100000 | 0.000000 | 0.140000 | 0.210000 | 0.000000 | 1.180000 | 1.280000 | 0.190000 | 0.290000 |
| 25% | 1.280000 | 0.470000 | 1.390000 | 3.040000 | 0.080000 | 44.299999 | 68.089996 | 1.480000 | 9.310000 |
| 50% | 2.510000 | 0.760000 | 2.600000 | 5.830000 | 0.140000 | 64.449997 | 123.699997 | 2.730000 | 23.830000 |
| 75% | 4.420000 | 1.270000 | 4.580000 | 10.640000 | 0.250000 | 86.540001 | 213.199997 | 4.830000 | 43.980000 |
| max | 54.560001 | 11.890000 | 77.260002 | 150.600006 | 2.880000 | 292.700012 | 1940.000000 | 89.510002 | 173.600006 |

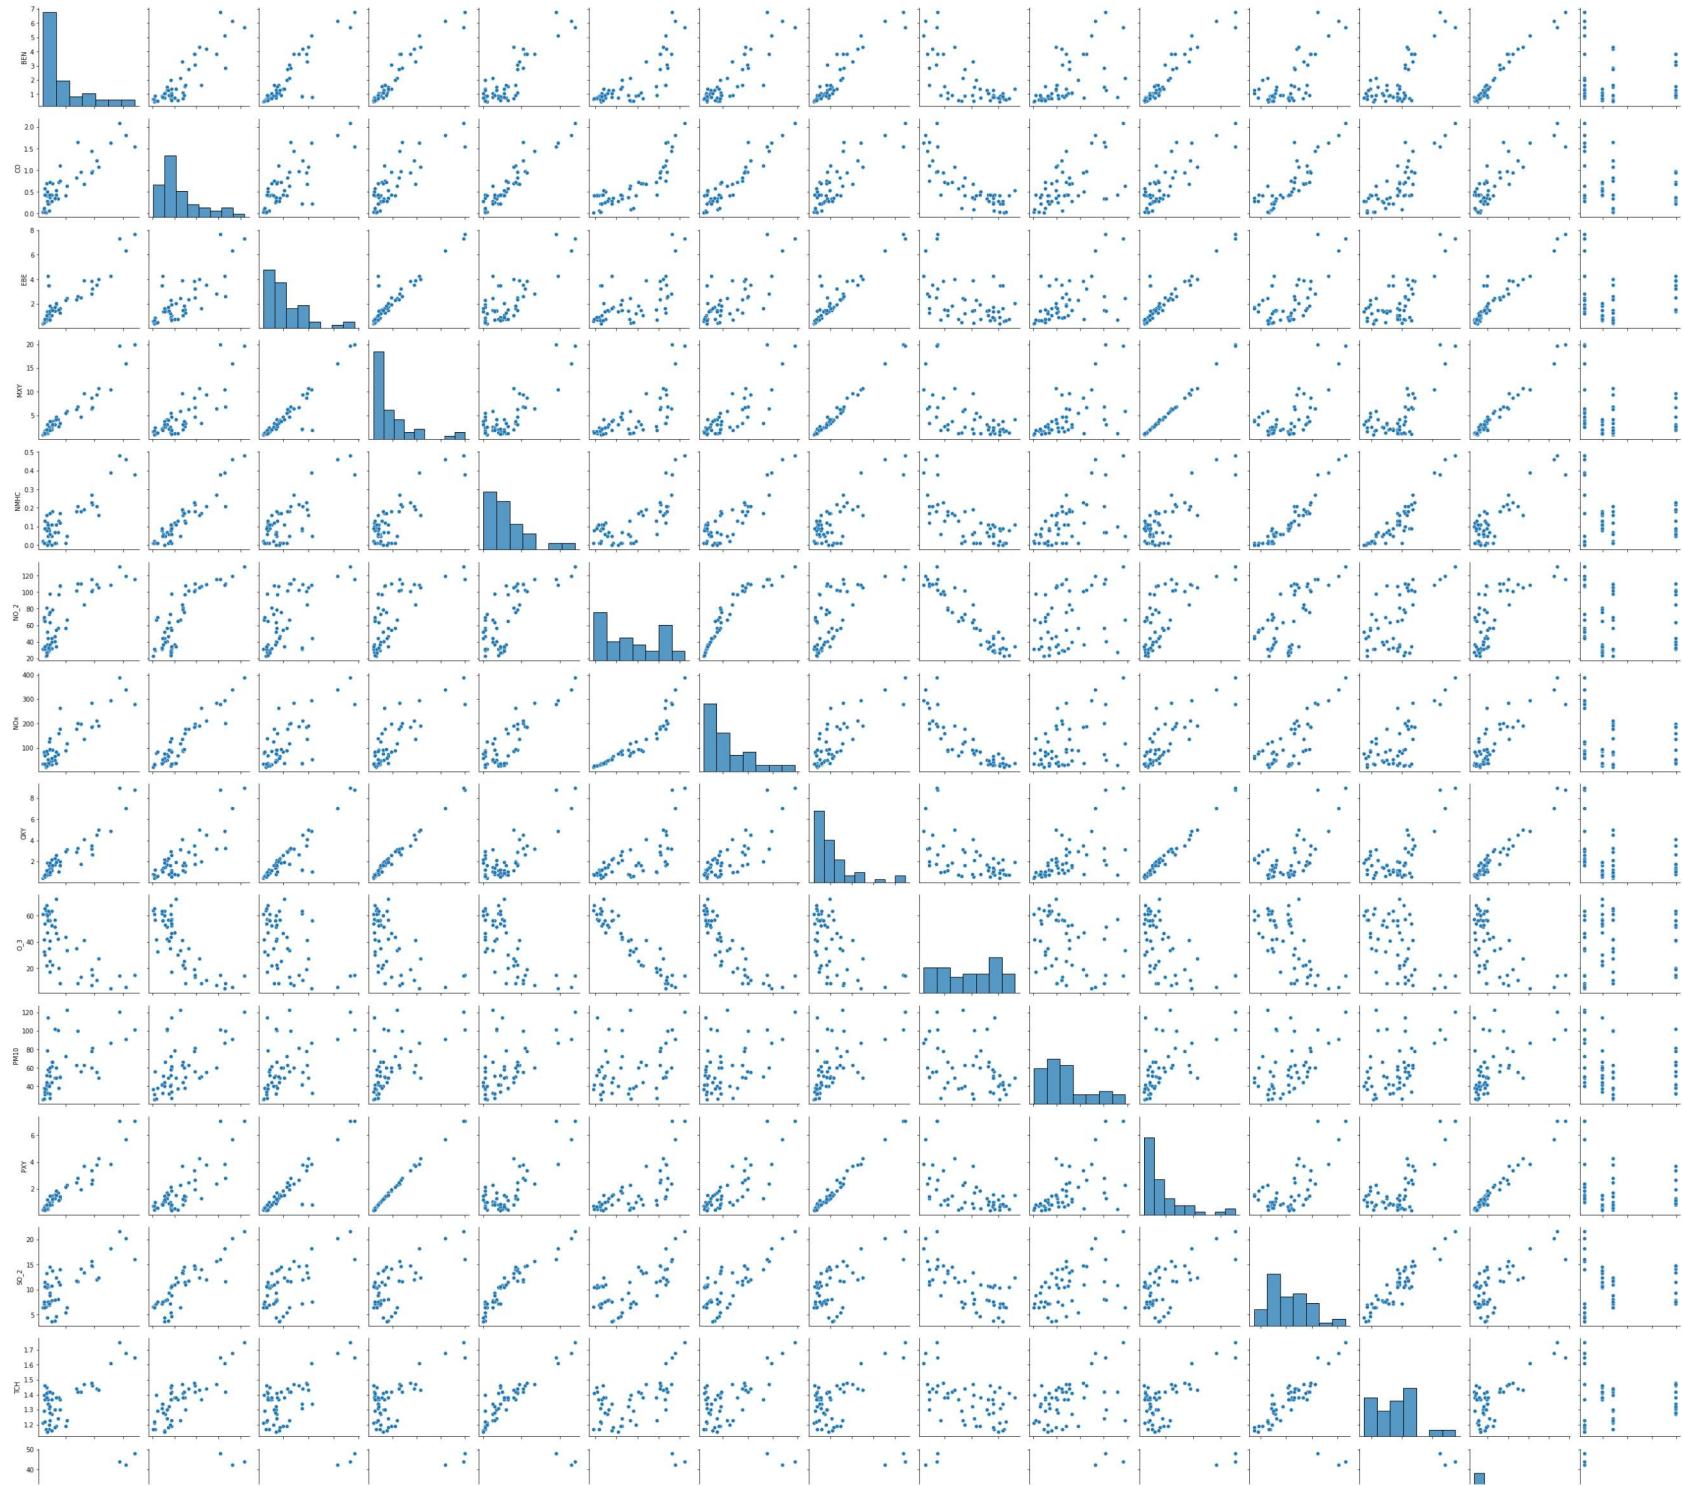
In [18]:

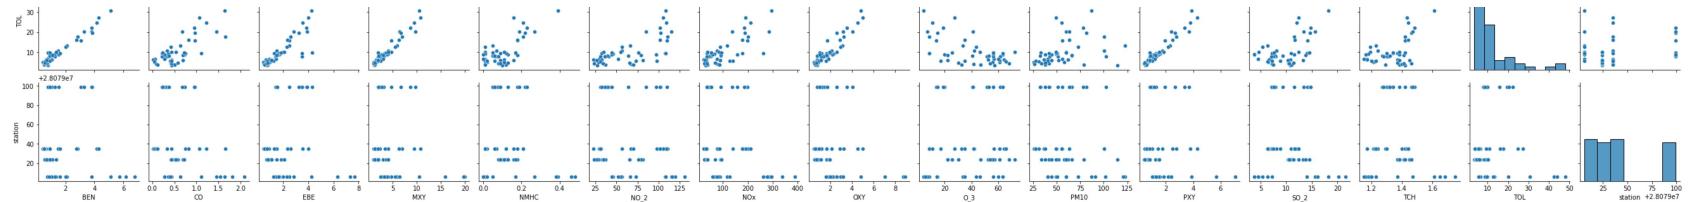
```
1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2         'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

```
In [19]: 1 sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x1c7cf7600a0>
```

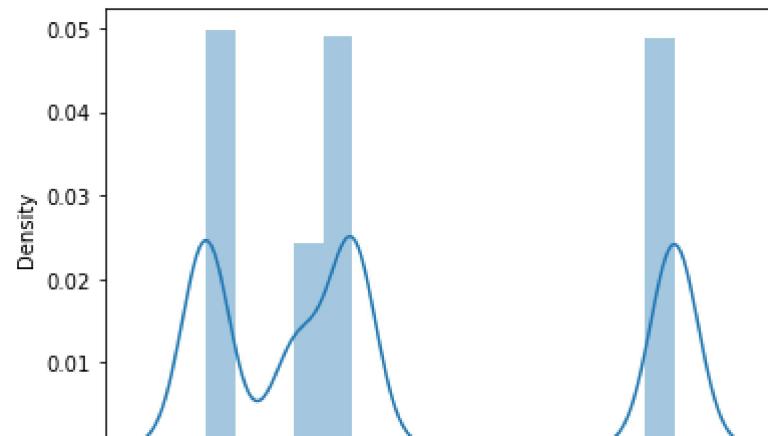





```
In [20]: 1 sns.distplot(df1['station'])
```

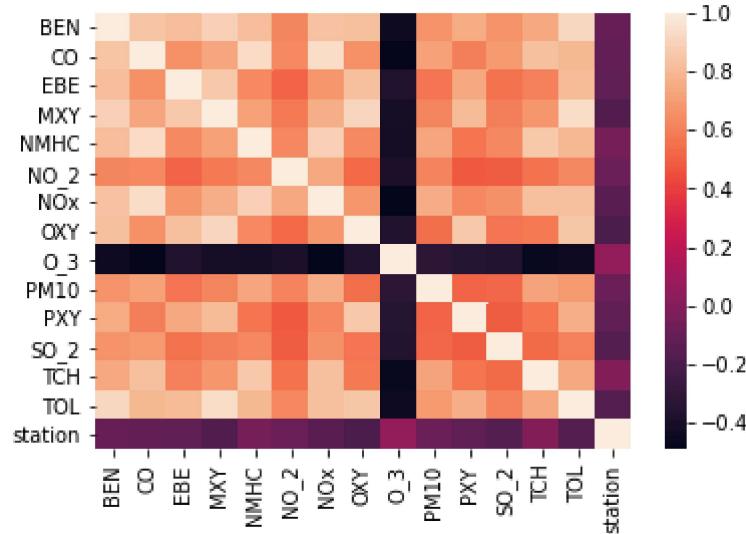
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: 1 sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
```

```
In [23]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: 1 from sklearn.linear_model import LinearRegression  
2 lr=LinearRegression()  
3 lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: 1 lr.intercept_
```

```
Out[25]: 28079006.920903727
```

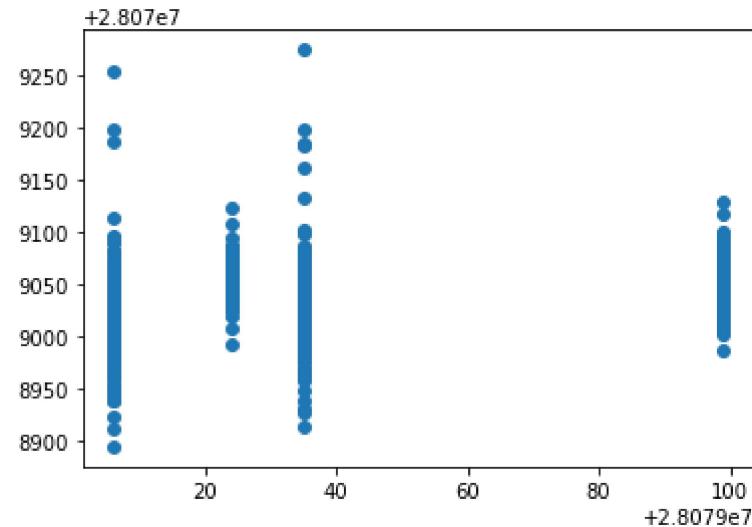
```
In [26]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
2 coeff
```

```
Out[26]:
```

| | Co-efficient |
|------|--------------|
| BEN | 7.167385 |
| CO | -17.219286 |
| EBE | 0.695649 |
| MXY | -0.312471 |
| NMHC | 84.856456 |
| NO_2 | 0.121145 |
| NOx | -0.081024 |
| OXY | -3.089740 |
| O_3 | -0.029332 |
| PM10 | -0.053788 |
| PXY | 1.724651 |
| SO_2 | -0.280828 |
| TCH | 37.211048 |
| TOL | -1.210974 |

```
In [27]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x1c7deef35b0>
```



ACCURACY

```
In [28]: 1 lr.score(x_test,y_test)
```

```
Out[28]: 0.1625629724433405
```

```
In [29]: 1 lr.score(x_train,y_train)
```

```
Out[29]: 0.1653930183128306
```

Ridge and Lasso

```
In [30]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: 1 rr.score(x_test,y_test)
```

```
Out[32]: 0.16291711658435537
```

```
In [33]: 1 rr.score(x_train,y_train)
```

```
Out[33]: 0.16511999698249946
```

```
In [34]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

```
In [35]: 1 la.score(x_train,y_train)
```

```
Out[35]: 0.03882024770011494
```

Accuracy(Lasso)

```
In [36]: 1 la.score(x_test,y_test)
```

```
Out[36]: 0.0405390347370842
```

```
In [37]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: 1 en.coef_
```

```
Out[38]: array([ 4.86193102,  0.          ,  0.68085738, -0.45047435,  0.07055131,
 0.07159909, -0.03470946, -2.43261621, -0.03199313,  0.07877159,
 1.06512407, -0.31461248,  1.20230068, -0.64051274])
```

```
In [39]: 1 en.intercept_
```

```
Out[39]: 28079048.57322989
```

```
In [40]: 1 prediction=en.predict(x_test)
```

```
In [41]: 1 en.score(x_test,y_test)
```

```
Out[41]: 0.1044284862453626
```

Evaluation Metrics

```
In [42]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
30.382443453690524
```

```
1213.0170622202952
```

```
34.828394482380254
```

Logistic Regression

```
In [43]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [44]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df['station']
```

```
In [45]: 1 feature_matrix.shape
```

```
Out[45]: (29669, 14)
```

```
In [46]: 1 target_vector.shape
```

```
Out[46]: (29669,)
```

```
In [47]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [48]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: 1 logr=LogisticRegression(max_iter=10000)  
2 logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: 1 prediction=logr.predict(observation)  
2 print(prediction)
```

```
[28079035]
```

```
In [52]: 1 logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079035, 28079099], dtype=int64)
```

```
In [53]: 1 logr.score(fs,target_vector)
```

```
Out[53]: 0.8087229094340894
```

```
In [54]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[54]: 1.724527777144498e-43
```

```
In [55]: 1 logr.predict_proba(observation)
```

```
Out[55]: array([[1.72452778e-43, 2.43756289e-56, 9.99998565e-01, 1.43537418e-06]])
```

Random Forest

```
In [56]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: 1 rfc=RandomForestClassifier()
2 rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: 1 parameters={'max_depth':[1,2,3,4,5],
2                 'min_samples_leaf':[5,10,15,20,25],
3                 'n_estimators':[10,20,30,40,50]
4 }
```

```
In [59]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')
```

```
In [60]: 1 grid_search.best_score_
```

```
Out[60]: 0.7245762711864407
```

```
In [61]: 1 rfc_best=grid_search.best_estimator_
```

In [62]:

```
1 from sklearn.tree import plot_tree
2
3 plt.figure(figsize=(80,40))
4 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

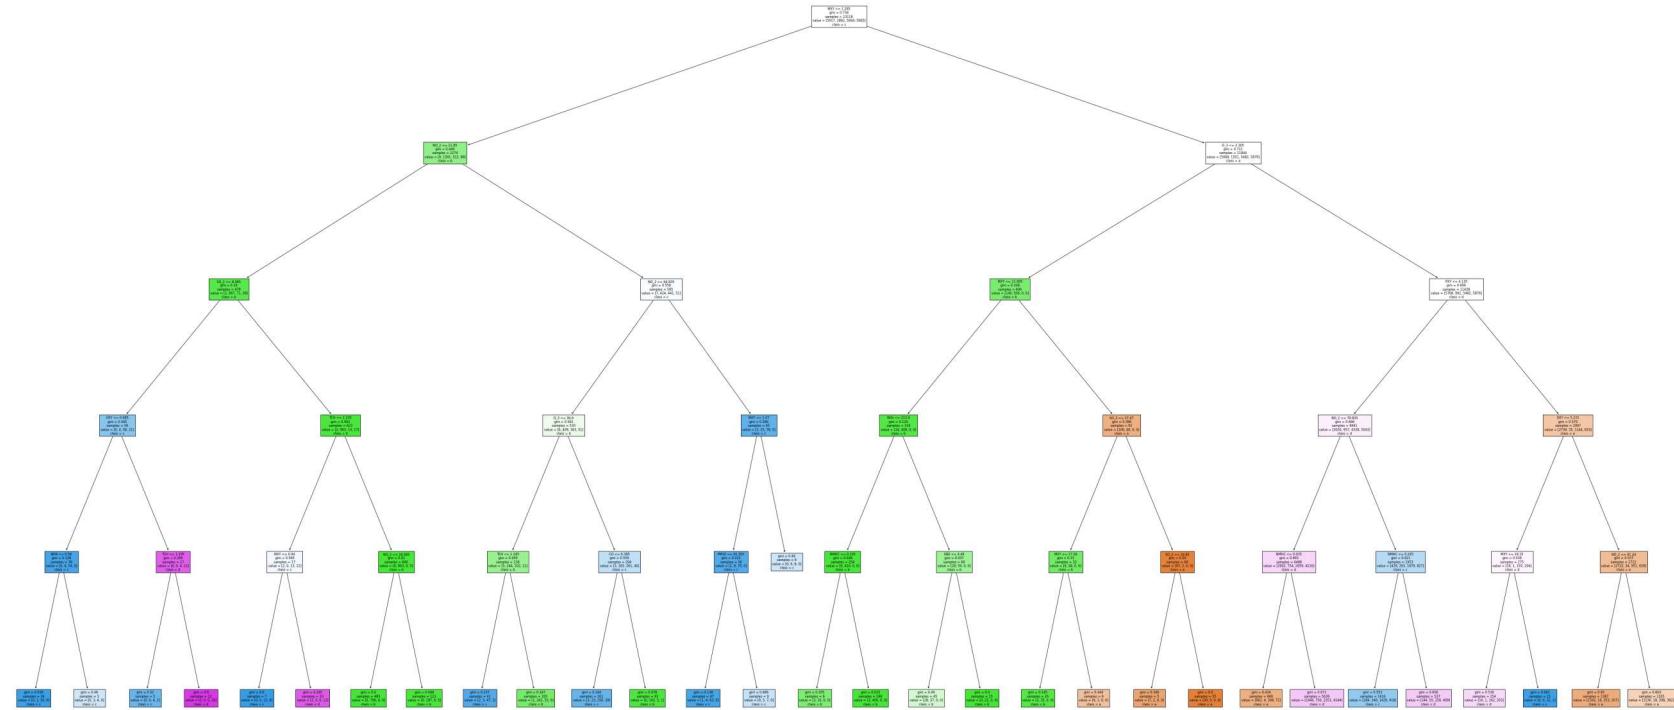
```
Out[62]: [Text(2222.700000000003, 1993.2, 'MXY <= 1.265\ngini = 0.734\nsamples = 13118\nvalue = [5917, 2892, 5994, 5965]\nklass = c'),  
Text(1171.800000000002, 1630.800000000002, 'NO_2 <= 21.85\ngini = 0.449\nsamples = 1274\nvalue = [9, 139, 1, 512, 89]\nklass = b'),  
Text(595.2, 1268.4, 'SO_2 <= 8.095\ngini = 0.19\nsamples = 679\nvalue = [2, 967, 71, 38]\nklass = b'),  
Text(297.6, 906.0, 'OXY <= 0.685\ngini = 0.445\nsamples = 56\nvalue = [0, 4, 58, 21]\nklass = c'),  
Text(148.8, 543.599999999999, 'BEN <= 0.58\ngini = 0.128\nsamples = 39\nvalue = [0, 4, 54, 0]\nklass = c'),  
Text(74.4, 181.1999999999982, 'gini = 0.038\nsamples = 34\nvalue = [0, 1, 50, 0]\nklass = c'),  
Text(223.200000000002, 181.1999999999982, 'gini = 0.49\nsamples = 5\nvalue = [0, 3, 4, 0]\nklass = c'),  
Text(446.400000000003, 543.599999999999, 'TCH <= 1.195\ngini = 0.269\nsamples = 17\nvalue = [0, 0, 4, 21]\nklass = d'),  
Text(372.0, 181.1999999999982, 'gini = 0.32\nsamples = 5\nvalue = [0, 0, 4, 1]\nklass = c'),  
Text(520.800000000001, 181.1999999999982, 'gini = 0.0\nsamples = 12\nvalue = [0, 0, 0, 20]\nklass = d'),  
Text(892.800000000001, 906.0, 'TCH <= 1.235\ngini = 0.063\nsamples = 623\nvalue = [2, 963, 13, 17]\nklass = b'),  
Text(744.0, 543.599999999999, 'MXY <= 0.94\ngini = 0.565\nsamples = 17\nvalue = [2, 0, 13, 12]\nklass = c'),  
Text(669.6, 181.1999999999982, 'gini = 0.0\nsamples = 7\nvalue = [0, 0, 13, 0]\nklass = c'),  
Text(818.400000000001, 181.1999999999982, 'gini = 0.245\nsamples = 10\nvalue = [2, 0, 0, 12]\nklass = d'),  
Text(1041.600000000001, 543.599999999999, 'NO_2 <= 16.565\ngini = 0.01\nsamples = 606\nvalue = [0, 963, 0, 5]\nklass = b'),  
Text(967.2, 181.1999999999982, 'gini = 0.0\nsamples = 483\nvalue = [0, 766, 0, 0]\nklass = b'),  
Text(1116.0, 181.1999999999982, 'gini = 0.048\nsamples = 123\nvalue = [0, 197, 0, 5]\nklass = b'),  
Text(1748.4, 1268.4, 'NO_2 <= 64.835\ngini = 0.558\nsamples = 595\nvalue = [7, 424, 441, 51]\nklass = c'),  
Text(1488.0, 906.0, 'O_3 <= 36.6\ngini = 0.561\nsamples = 530\nvalue = [6, 409, 363, 51]\nklass = b'),  
Text(1339.2, 543.599999999999, 'TCH <= 1.245\ngini = 0.459\nsamples = 236\nvalue = [3, 244, 102, 11]\nklass = b'),  
Text(1264.800000000002, 181.1999999999982, 'gini = 0.237\nsamples = 41\nvalue = [2, 3, 47, 2]\nklass = c'),  
Text(1413.600000000001, 181.1999999999982, 'gini = 0.347\nsamples = 195\nvalue = [1, 241, 55, 9]\nklass = b'),  
Text(1636.800000000002, 543.599999999999, 'CO <= 0.385\ngini = 0.559\nsamples = 294\nvalue = [3, 165, 261, 40]\nklass = c'),  
Text(1562.4, 181.1999999999982, 'gini = 0.344\nsamples = 201\nvalue = [3, 23, 256, 39]\nklass = c'),  
Text(1711.2, 181.1999999999982, 'gini = 0.078\nsamples = 93\nvalue = [0, 142, 5, 1]\nklass = b'),  
Text(2008.800000000002, 906.0, 'MXY <= 1.07\ngini = 0.286\nsamples = 65\nvalue = [1, 15, 78, 0]\nklass = c'),  
Text(1934.4, 543.599999999999, 'PM10 <= 61.205\ngini = 0.222\nsamples = 56\nvalue = [1, 9, 70, 0]\nklass = c'),  
Text(1860.000000000002, 181.1999999999982, 'gini = 0.138\nsamples = 47\nvalue = [1, 4, 63, 0]\nklass = c'),
```

Text(2008.800000000002, 181.1999999999982, 'gini = 0.486\nsamples = 9\nvalue = [0, 5, 7, 0]\nclass = c'),
Text(2083.200000000003, 543.5999999999999, 'gini = 0.49\nsamples = 9\nvalue = [0, 6, 8, 0]\nclass = c'),
Text(3273.600000000004, 1630.800000000002, 'O_3 <= 3.305\ngini = 0.711\nsamples = 11844\nvalue = [5908, 1
501, 5482, 5876]\nclass = a'),
Text(2678.4, 1268.4, 'MXY <= 11.995\ngini = 0.338\nsamples = 406\nvalue = [140, 509, 0, 0]\nclass = b'),
Text(2380.8, 906.0, 'NOx <= 213.6\ngini = 0.126\nsamples = 314\nvalue = [34, 469, 0, 0]\nclass = b'),
Text(2232.0, 543.5999999999999, 'NMHC <= 0.105\ngini = 0.028\nsamples = 254\nvalue = [6, 410, 0, 0]\nclass
= b'),
Text(2157.600000000004, 181.1999999999982, 'gini = 0.355\nsamples = 6\nvalue = [3, 10, 0, 0]\nclass =
b'),
Text(2306.4, 181.1999999999982, 'gini = 0.015\nsamples = 248\nvalue = [3, 400, 0, 0]\nclass = b'),
Text(2529.600000000004, 543.5999999999999, 'EBE <= 4.49\ngini = 0.437\nsamples = 60\nvalue = [28, 59, 0,
0]\nclass = b'),
Text(2455.200000000003, 181.1999999999982, 'gini = 0.49\nsamples = 45\nvalue = [28, 37, 0, 0]\nclass =
b'),
Text(2604.0, 181.1999999999982, 'gini = 0.0\nsamples = 15\nvalue = [0, 22, 0, 0]\nclass = b'),
Text(2976.0, 906.0, 'SO_2 <= 17.47\ngini = 0.398\nsamples = 92\nvalue = [106, 40, 0, 0]\nclass = a'),
Text(2827.200000000003, 543.5999999999999, 'MXY <= 17.26\ngini = 0.31\nsamples = 32\nvalue = [9, 38, 0, 0]
\nclass = b'),
Text(2752.8, 181.1999999999982, 'gini = 0.145\nsamples = 26\nvalue = [3, 35, 0, 0]\nclass = b'),
Text(2901.600000000004, 181.1999999999982, 'gini = 0.444\nsamples = 6\nvalue = [6, 3, 0, 0]\nclass = a'),
Text(3124.8, 543.5999999999999, 'SO_2 <= 19.85\ngini = 0.04\nsamples = 60\nvalue = [97, 2, 0, 0]\nclass =
a'),
Text(3050.4, 181.1999999999982, 'gini = 0.346\nsamples = 5\nvalue = [7, 2, 0, 0]\nclass = a'),
Text(3199.200000000003, 181.1999999999982, 'gini = 0.0\nsamples = 55\nvalue = [90, 0, 0, 0]\nclass = a'),
Text(3868.8, 1268.4, 'PXY <= 4.135\ngini = 0.699\nsamples = 11438\nvalue = [5768, 992, 5482, 5876]\nclass =
d'),
Text(3571.200000000003, 906.0, 'NO_2 <= 79.835\ngini = 0.696\nsamples = 8441\nvalue = [3030, 957, 4338, 50
43]\nclass = d'),
Text(3422.4, 543.5999999999999, 'NMHC <= 0.035\ngini = 0.693\nsamples = 6488\nvalue = [2602, 754, 2659, 421
6]\nclass = d'),
Text(3348.000000000005, 181.1999999999982, 'gini = 0.434\nsamples = 849\nvalue = [962, 4, 308, 72]\nclass
= a'),
Text(3496.8, 181.1999999999982, 'gini = 0.671\nsamples = 5639\nvalue = [1640, 750, 2351, 4144]\nclass =
d'),
Text(3720.000000000005, 543.5999999999999, 'NMHC <= 0.245\ngini = 0.621\nsamples = 1953\nvalue = [428, 20
3, 1679, 827]\nclass = c'),
Text(3645.600000000004, 181.1999999999982, 'gini = 0.553\nsamples = 1416\nvalue = [284, 148, 1429, 418]\n
class = c'),
Text(3794.4, 181.1999999999982, 'gini = 0.656\nsamples = 537\nvalue = [144, 55, 250, 409]\nclass = d'),
Text(4166.400000000001, 906.0, 'OXY <= 5.215\ngini = 0.579\nsamples = 2997\nvalue = [2738, 35, 1144, 833]\n
class = a'),
Text(4017.600000000004, 543.5999999999999, 'MXY <= 18.15\ngini = 0.538\nsamples = 275\nvalue = [16, 1, 19

```

3, 204]\nclass = d'),
Text(3943.200000000003, 181.19999999999982, 'gini = 0.536\nsamples = 254\nvalue = [16, 1, 162, 203]\nclass
= d'),
Text(4092.000000000005, 181.19999999999982, 'gini = 0.061\nsamples = 21\nvalue = [0, 0, 31, 1]\nclass =
c'),
Text(4315.200000000001, 543.5999999999999, 'NO_2 <= 91.24\ngini = 0.537\nsamples = 2722\nvalue = [2722, 34,
951, 629]\nclass = a'),
Text(4240.8, 181.19999999999982, 'gini = 0.45\nsamples = 1387\nvalue = [1592, 16, 353, 267]\nclass = a'),
Text(4389.6, 181.19999999999982, 'gini = 0.603\nsamples = 1335\nvalue = [1130, 18, 598, 362]\nclass = a'])

```



Conclusion

Accuracy

Linear Regression: 0.1653930183128306

Ridge Regression: 0.16511999698249946

Lasso Regression: 0.03882024770011494

ElasticNet Regression: 0.1044284862453626

Logistic Regression: 0.8087229094340894

Random Forest: 0.7245762711864407

Logistic Regression is suitable for this dataset