

## Importing Libraries

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

## Importing Datasets

In [2]:

```

1 df=pd.read_csv("madrid_2002")
2 df

```

Out[2]:

		date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PXY	SO_2	TCH	TOL	s
0		2002-04-01 01:00:00	NaN	1.39	NaN	NaN	NaN	145.100006	352.100006	NaN	6.54	41.990002	NaN	21.320000	NaN	NaN	280
1		2002-04-01 01:00:00	1.93	0.71	2.33	6.20	0.15	98.150002	153.399994	2.67	6.85	20.980000	2.53	11.660000	1.82	10.980000	280
2		2002-04-01 01:00:00	NaN	0.80	NaN	NaN	NaN	103.699997	134.000000	NaN	13.01	28.440001	NaN	13.670000	NaN	NaN	280
3		2002-04-01 01:00:00	NaN	1.61	NaN	NaN	NaN	97.599998	268.000000	NaN	5.12	42.180000	NaN	16.990000	NaN	NaN	280
4		2002-04-01 01:00:00	NaN	1.90	NaN	NaN	NaN	92.089996	237.199997	NaN	7.28	76.330002	NaN	15.260000	NaN	NaN	280
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
217291		2002-11-01 00:00:00	4.16	1.14	NaN	NaN	NaN	81.080002	265.700012	NaN	7.21	36.750000	NaN	13.210000	NaN	NaN	280
217292		2002-11-01 00:00:00	3.67	1.73	2.89	NaN	0.38	113.900002	373.100006	NaN	5.66	63.389999	NaN	15.640000	1.78	15.690000	280
217293		2002-11-01 00:00:00	1.37	0.58	1.17	2.37	0.15	65.389999	107.699997	1.30	9.11	9.640000	0.94	5.620000	1.43	4.330000	280
217294		2002-11-01 00:00:00	4.51	0.91	4.83	10.99	NaN	149.800003	202.199997	1.00	5.75	NaN	5.52	24.219999	NaN	22.129999	280
217295		2002-11-01 00:00:00	3.11	1.17	3.00	7.77	0.26	80.110001	180.300003	2.25	7.38	29.240000	3.35	12.910000	1.54	15.510000	280

217296 rows × 16 columns



# Data Cleaning and Data Preprocessing

```
In [3]: 1 df=df.dropna()
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
               dtype='object')
```

```
In [5]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 32381 entries, 1 to 217295
Data columns (total 16 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      32381 non-null   object 
 1   BEN       32381 non-null   float64
 2   CO        32381 non-null   float64
 3   EBE       32381 non-null   float64
 4   MXY       32381 non-null   float64
 5   NMHC      32381 non-null   float64
 6   NO_2      32381 non-null   float64
 7   NOx       32381 non-null   float64
 8   OXY       32381 non-null   float64
 9   O_3        32381 non-null   float64
 10  PM10      32381 non-null   float64
 11  PXY       32381 non-null   float64
 12  SO_2      32381 non-null   float64
 13  TCH       32381 non-null   float64
 14  TOL       32381 non-null   float64
 15  station    32381 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.2+ MB
```

```
In [6]: 1 data=df[['CO' , 'station']]  
2 data
```

Out[6]:

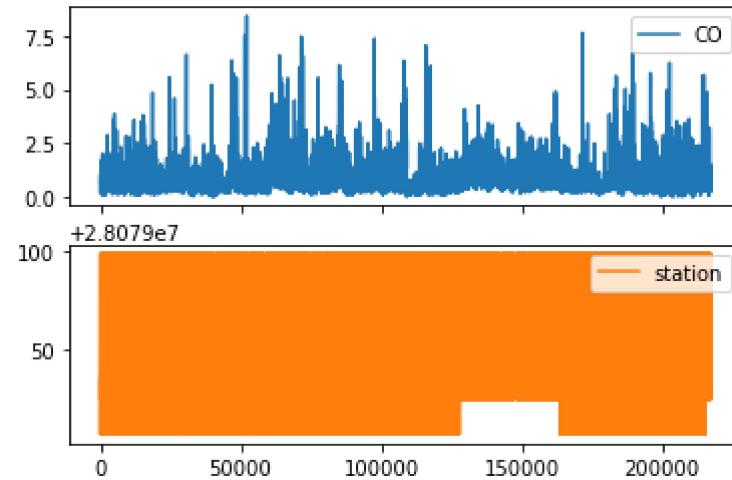
	CO	station
1	0.71	28079035
5	0.72	28079006
22	0.80	28079024
24	1.04	28079099
26	0.53	28079035
...	...	...
217269	0.28	28079024
217271	1.30	28079099
217273	0.97	28079035
217293	0.58	28079024
217295	1.17	28079099

32381 rows × 2 columns

## Line chart

```
In [7]: 1 data.plot.line(subplots=True)
```

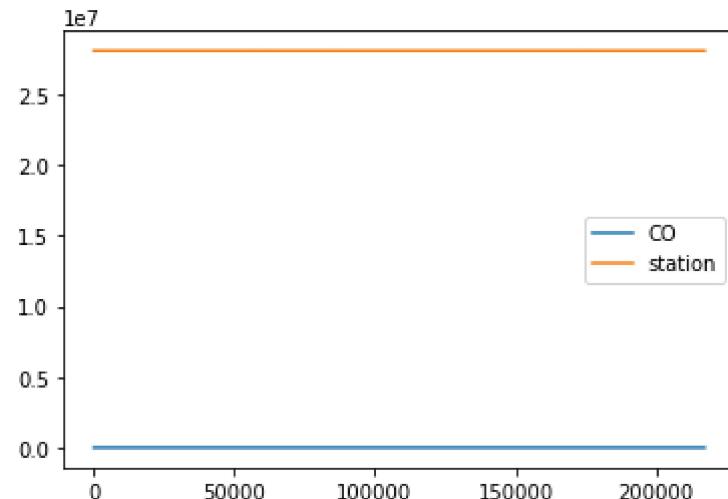
```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



## Line chart

```
In [8]: 1 data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

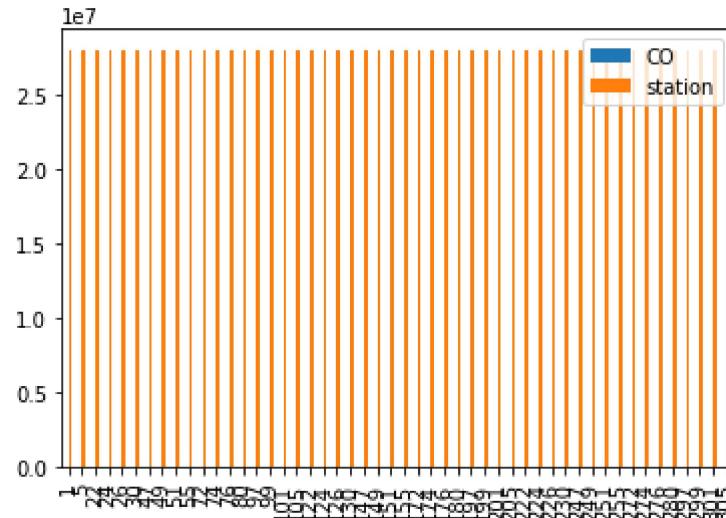


## Bar chart

```
In [9]: 1 b=data[0:50]
```

```
In [10]: 1 b.plot.bar()
```

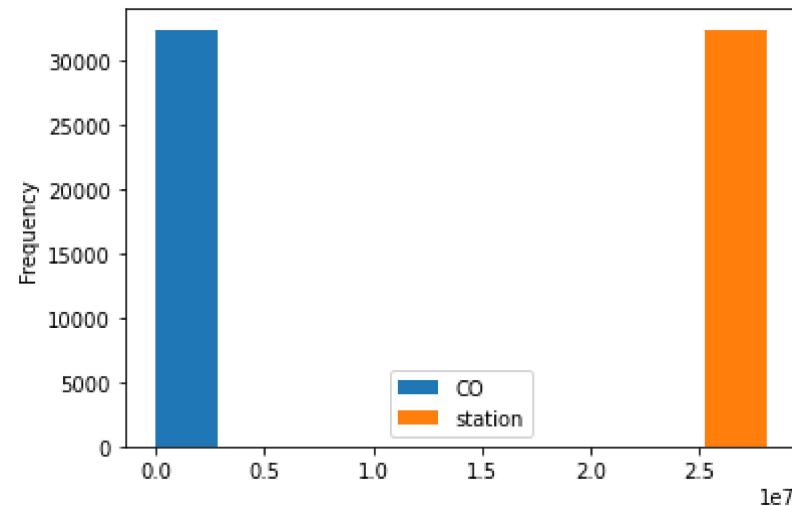
```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: 1 data.plot.hist()
```

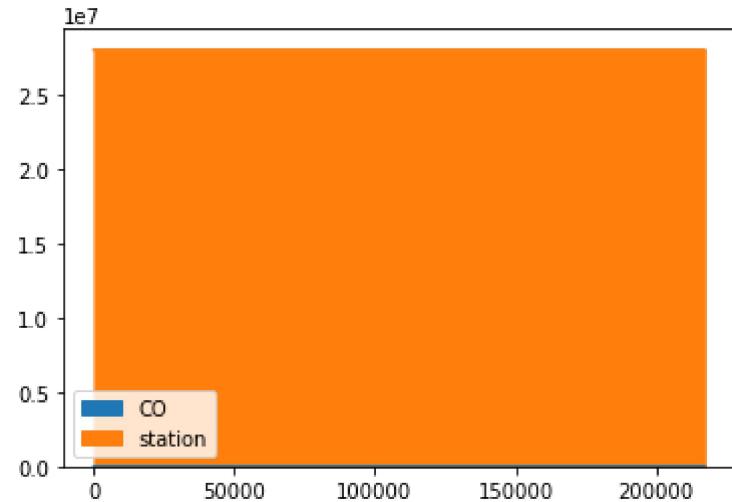
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

In [12]: 1 data.plot.area()

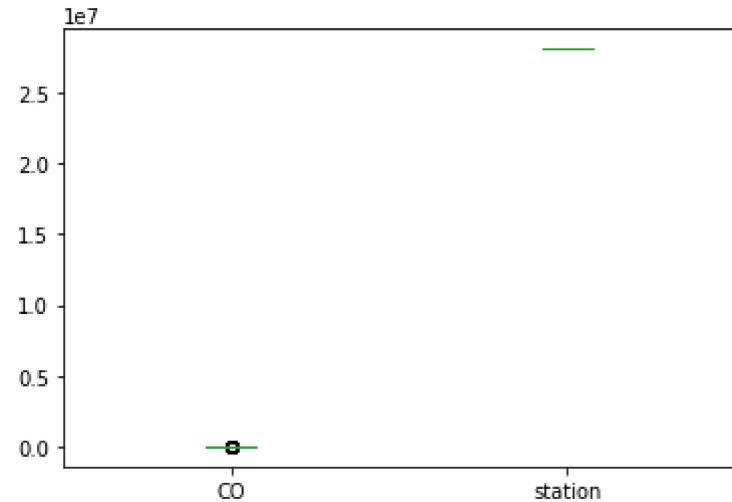
Out[12]: <AxesSubplot:>



## Box chart

In [13]: 1 data.plot.box()

Out[13]: <AxesSubplot:>

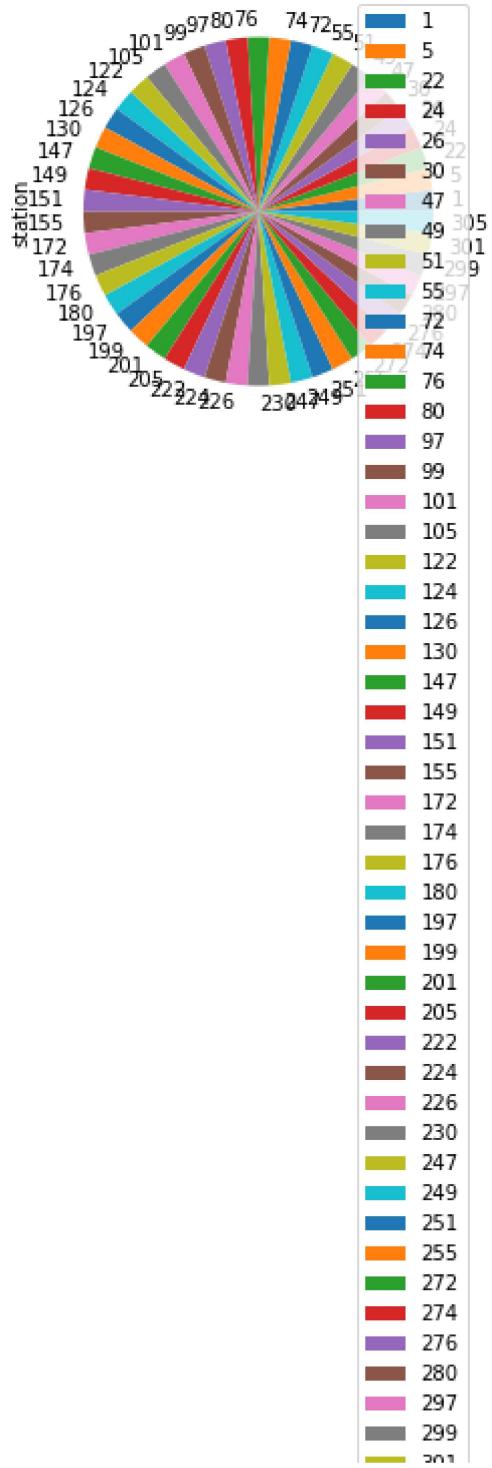


## Pie chart

```
In [14]: 1 b.plot.pie(y='station' )
```

```
Out[14]: <AxesSubplot:ylabel='station'>
```



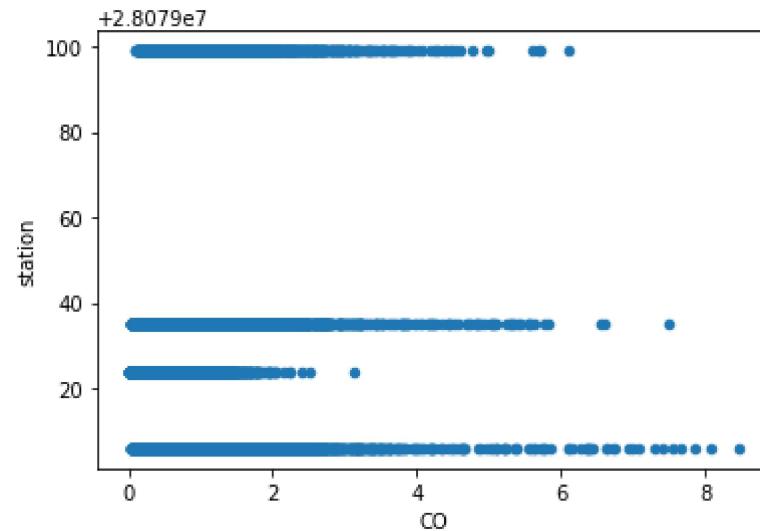




## Scatter chart

```
In [15]: 1 data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 32381 entries, 1 to 217295
Data columns (total 16 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   date      32381 non-null    object  
 1   BEN        32381 non-null    float64 
 2   CO         32381 non-null    float64 
 3   EBE        32381 non-null    float64 
 4   MXY        32381 non-null    float64 
 5   NMHC       32381 non-null    float64 
 6   NO_2       32381 non-null    float64 
 7   NOx        32381 non-null    float64 
 8   OXY        32381 non-null    float64 
 9   O_3         32381 non-null    float64 
 10  PM10       32381 non-null    float64 
 11  PXY        32381 non-null    float64 
 12  SO_2       32381 non-null    float64 
 13  TCH        32381 non-null    float64 
 14  TOL        32381 non-null    float64
```

```
In [17]: 1 df.describe()
```

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
count	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000	32381.000000
mean	2.479155	0.787323	2.914004	7.013636	0.155827	58.936796	126.009340	3.169093	33.128518
std	2.280959	0.610810	2.667881	6.774365	0.135731	31.472733	114.035078	2.950771	25.679650
min	0.180000	0.000000	0.180000	0.190000	0.000000	0.890000	1.710000	0.180000	1.880000
25%	0.970000	0.420000	1.140000	2.420000	0.080000	35.660000	48.720001	1.190000	10.290000
50%	1.840000	0.620000	2.130000	5.140000	0.130000	57.160000	96.830002	2.340000	27.290001
75%	3.250000	0.980000	3.830000	9.420000	0.200000	78.769997	167.500000	4.130000	50.160000
max	32.660000	8.460000	41.740002	99.879997	2.700000	263.600006	1336.000000	76.339996	151.399994

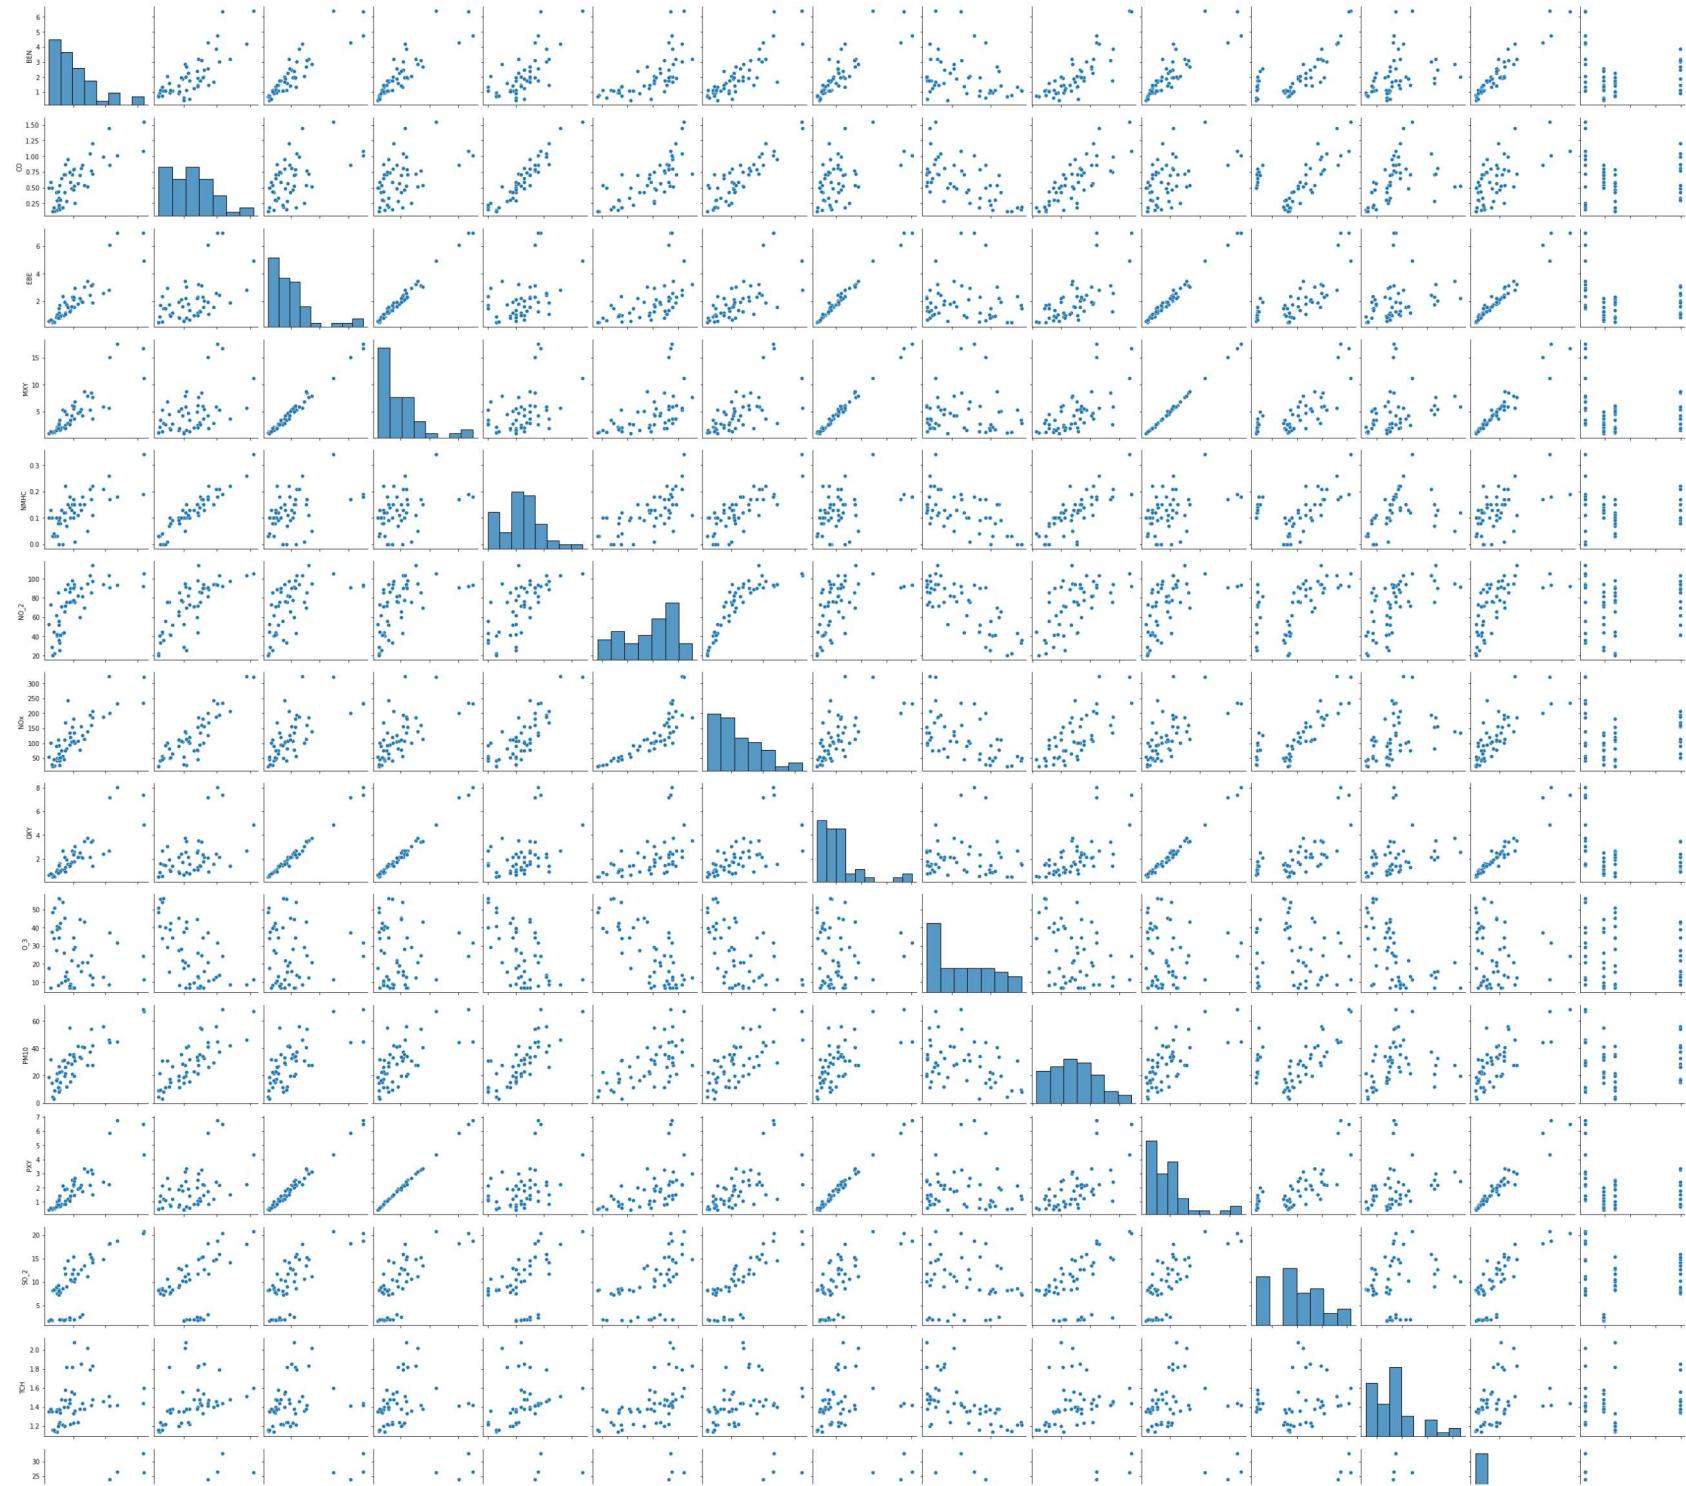
```
In [18]: 1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

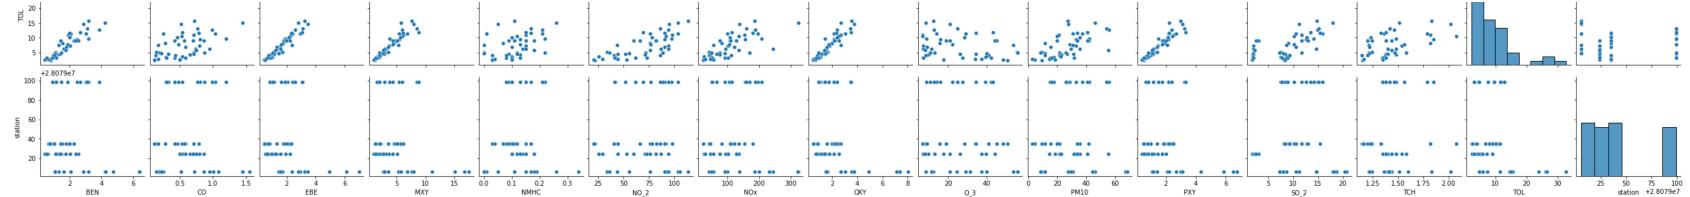
## EDA AND VISUALIZATION

```
In [19]: 1 sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x195dde002e0>
```



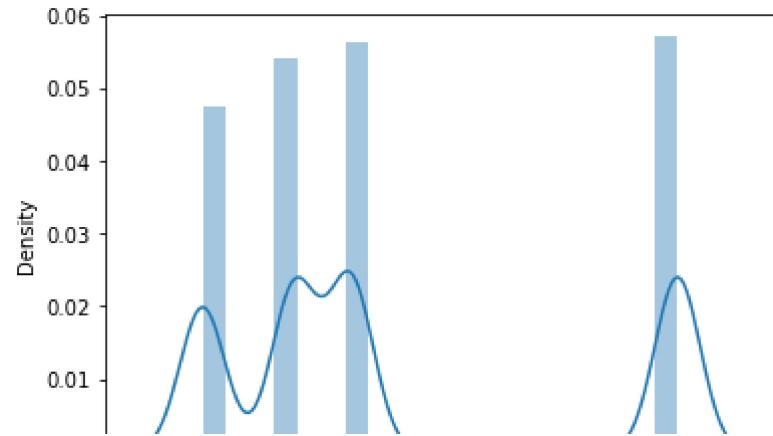




```
In [20]: 1 sns.distplot(df1['station'])
```

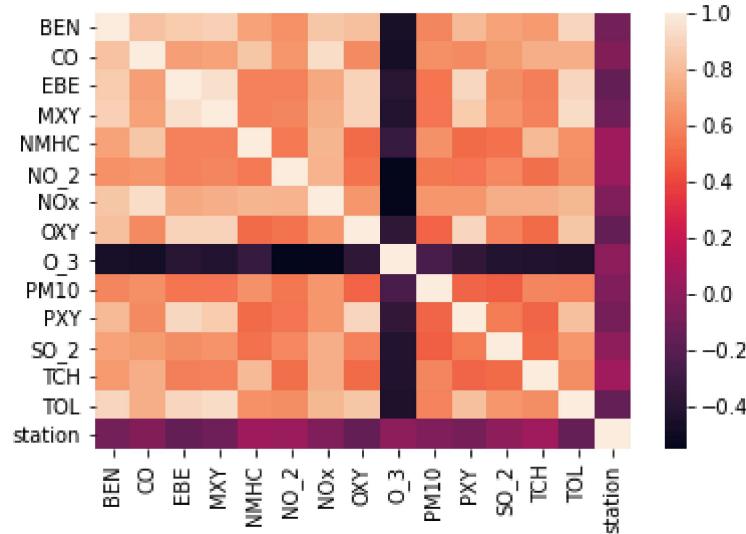
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: 1 sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
```

```
In [23]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

```
In [24]: 1 from sklearn.linear_model import LinearRegression  
2 lr=LinearRegression()  
3 lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: 1 lr.intercept_
```

```
Out[25]: 28078988.04789076
```

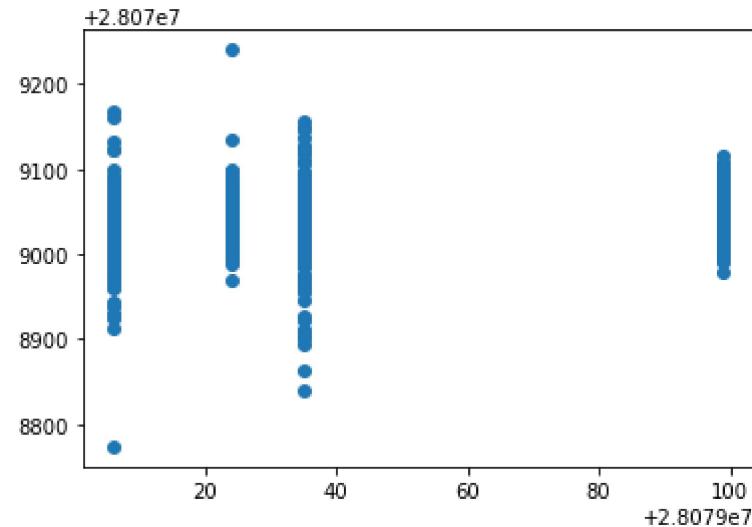
```
In [26]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
2 coeff
```

```
Out[26]:
```

	Co-efficient
BEN	2.127880
CO	-13.377859
EBE	-11.898731
MXY	4.213690
NMHC	85.350840
NO_2	0.249613
NOx	-0.102582
OXY	-5.063797
O_3	-0.023072
PM10	-0.126225
PXY	7.819506
SO_2	0.631406
TCH	43.085702
TOL	-1.502361

```
In [27]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x195ed610580>
```



## ACCURACY

```
In [28]: 1 lr.score(x_test,y_test)
```

```
Out[28]: 0.19084160353876733
```

```
In [29]: 1 lr.score(x_train,y_train)
```

```
Out[29]: 0.20226473515319232
```

## Ridge and Lasso

```
In [30]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: 1 rr.score(x_test,y_test)
```

```
Out[32]: 0.19070124926960652
```

```
In [33]: 1 rr.score(x_train,y_train)
```

```
Out[33]: 0.20204408096897053
```

```
In [34]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

```
In [35]: 1 la.score(x_train,y_train)
```

```
Out[35]: 0.05783026530167035
```

## Accuracy(Lasso)

```
In [36]: 1 la.score(x_test,y_test)
```

```
Out[36]: 0.05896723554942673
```

```
In [37]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: 1 en.coef_
```

```
Out[38]: array([ 1.05363840e+00,  0.00000000e+00, -2.94374581e+00,  1.54056083e+00,
   2.08377419e-01,  2.26935966e-01, -2.78379496e-02, -2.34722815e+00,
  -1.84086853e-02,  5.92739824e-04,  2.21824675e+00,  4.20780309e-01,
  1.12599389e+00, -1.17689144e+00])
```

```
In [39]: 1 en.intercept_
```

```
Out[39]: 28079038.032892514
```

```
In [40]: 1 prediction=en.predict(x_test)
```

```
In [41]: 1 en.score(x_test,y_test)
```

```
Out[41]: 0.09906557806958294
```

## Evaluation Metrics

```
In [42]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
28.665579995440105
```

```
1128.1578624152505
```

```
33.58806130778093
```

## Logistic Regression

```
In [43]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [44]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df[ 'station']
```

```
In [45]: 1 feature_matrix.shape
```

```
Out[45]: (32381, 14)
```

```
In [46]: 1 target_vector.shape
```

```
Out[46]: (32381,)
```

```
In [47]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [48]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: 1 logr=LogisticRegression(max_iter=10000)
2 logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: 1 prediction=logr.predict(observation)
2 print(prediction)
```

```
[28079035]
```

```
In [52]: 1 logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079035, 28079099], dtype=int64)
```

```
In [53]: 1 logr.score(fs,target_vector)
```

```
Out[53]: 0.8480899292795158
```

```
In [54]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[54]: 2.5638972732451705e-10
```

```
In [55]: 1 logr.predict_proba(observation)
```

```
Out[55]: array([[2.56389727e-10, 3.44199742e-71, 1.00000000e+00, 1.43898646e-13]])
```

## Random Forest

```
In [56]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: 1 rfc=RandomForestClassifier()
2 rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: 1 parameters={'max_depth':[1,2,3,4,5],
2                 'min_samples_leaf':[5,10,15,20,25],
3                 'n_estimators':[10,20,30,40,50]
4 }
```

```
In [59]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')
```

```
In [60]: 1 grid_search.best_score_
```

```
Out[60]: 0.7758757610517957
```

```
In [61]: 1 rfc_best=grid_search.best_estimator_
```

In [62]:

```
1 from sklearn.tree import plot_tree
2
3 plt.figure(figsize=(80,40))
4 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

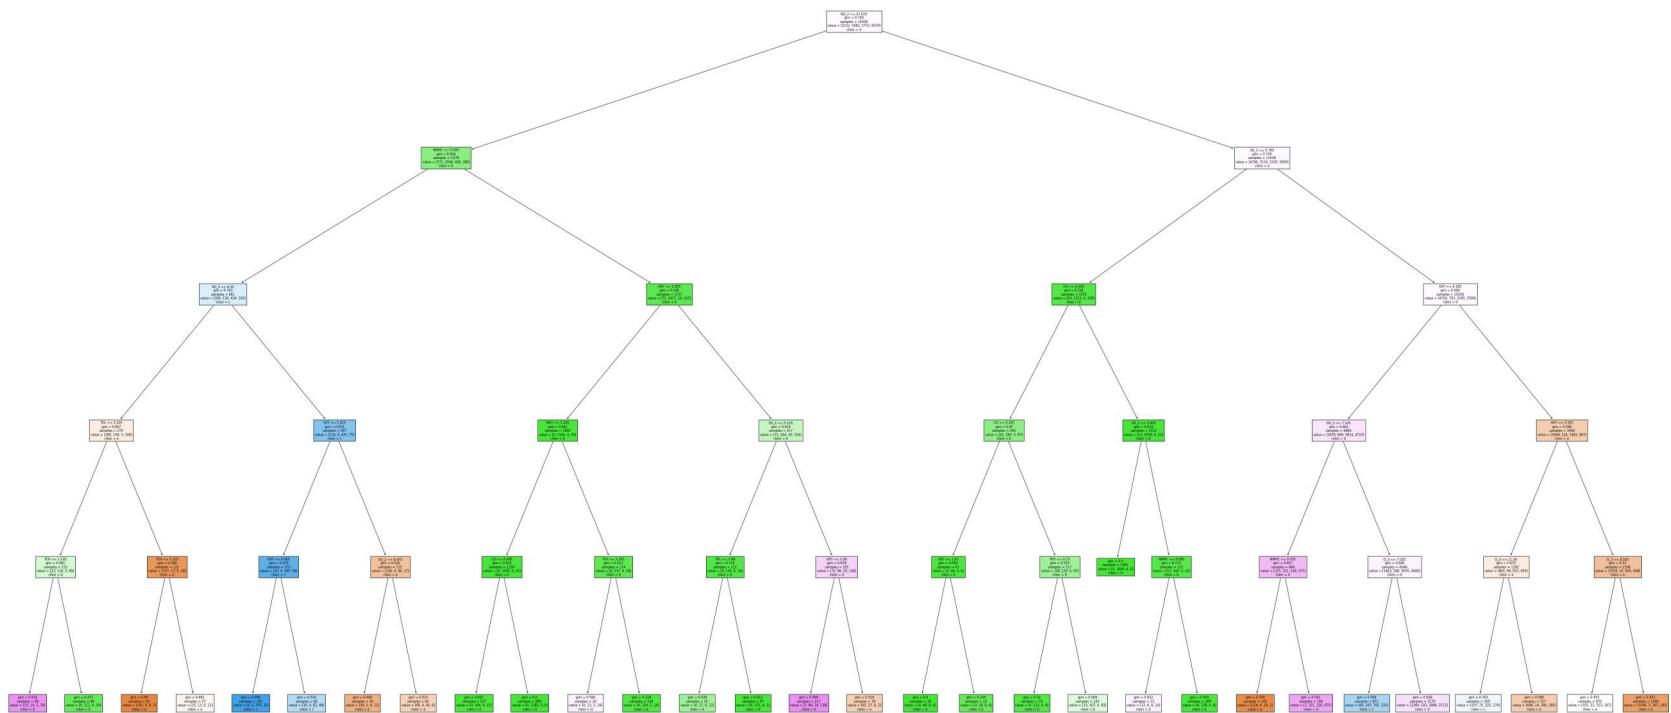
```
Out[62]: [Text(2278.5, 1993.2, 'NO_2 <= 27.015\ngini = 0.749\nsamples = 14306\nvalue = [5152, 5682, 5753, 6079]\nclas  
s = d'),  
Text(1190.4, 1630.800000000002, 'NMHC <= 0.055\ngini = 0.504\nsamples = 2378\nvalue = [372, 2566, 458, 38  
0]\nklass = b'),  
Text(595.2, 1268.4, 'SO_2 <= 6.33\ngini = 0.703\nsamples = 661\nvalue = [300, 139, 434, 183]\nklass = c'),  
Text(297.6, 906.0, 'TOL <= 3.105\ngini = 0.657\nsamples = 274\nvalue = [180, 139, 3, 108]\nklass = a'),  
Text(148.8, 543.599999999999, 'TCH <= 1.235\ngini = 0.581\nsamples = 152\nvalue = [23, 126, 3, 90]\nklass  
= b'),  
Text(74.4, 181.1999999999982, 'gini = 0.534\nsamples = 68\nvalue = [23, 14, 3, 70]\nklass = d'),  
Text(223.200000000002, 181.1999999999982, 'gini = 0.257\nsamples = 84\nvalue = [0, 112, 0, 20]\nklass =  
b'),  
Text(446.400000000003, 543.599999999999, 'TCH <= 1.225\ngini = 0.289\nsamples = 122\nvalue = [157, 13,  
0, 18]\nklass = a'),  
Text(372.0, 181.1999999999982, 'gini = 0.09\nsamples = 95\nvalue = [142, 0, 0, 7]\nklass = a'),  
Text(520.800000000001, 181.1999999999982, 'gini = 0.661\nsamples = 27\nvalue = [15, 13, 0, 11]\nklass =  
a'),  
Text(892.800000000001, 906.0, 'OXY <= 1.415\ngini = 0.475\nsamples = 387\nvalue = [120, 0, 431, 75]\nklass  
= c'),  
Text(744.0, 543.599999999999, 'OXY <= 0.905\ngini = 0.279\nsamples = 272\nvalue = [16, 0, 385, 58]\nklass  
= c'),  
Text(669.6, 181.1999999999982, 'gini = 0.096\nsamples = 187\nvalue = [6, 0, 303, 10]\nklass = c'),  
Text(818.400000000001, 181.1999999999982, 'gini = 0.534\nsamples = 85\nvalue = [10, 0, 82, 48]\nklass =  
c'),  
Text(1041.600000000001, 543.599999999999, 'SO_2 <= 6.875\ngini = 0.526\nsamples = 115\nvalue = [104, 0, 4  
6, 17]\nklass = a'),  
Text(967.2, 181.1999999999982, 'gini = 0.449\nsamples = 30\nvalue = [36, 0, 4, 11]\nklass = a'),  
Text(1116.0, 181.1999999999982, 'gini = 0.523\nsamples = 85\nvalue = [68, 0, 42, 6]\nklass = a'),  
Text(1785.600000000001, 1268.4, 'OXY <= 1.005\ngini = 0.198\nsamples = 1717\nvalue = [72, 2427, 24, 197]\nklass = b'),  
Text(1488.0, 906.0, 'MXY <= 1.245\ngini = 0.042\nsamples = 1400\nvalue = [0, 2163, 4, 43]\nklass = b'),  
Text(1339.2, 543.599999999999, 'CO <= 0.365\ngini = 0.015\nsamples = 1226\nvalue = [0, 1926, 0, 15]\nklass  
= b'),  
Text(1264.800000000002, 181.1999999999982, 'gini = 0.052\nsamples = 337\nvalue = [0, 544, 0, 15]\nklass =  
b'),  
Text(1413.600000000001, 181.1999999999982, 'gini = 0.0\nsamples = 889\nvalue = [0, 1382, 0, 0]\nklass =  
b'),  
Text(1636.800000000002, 543.599999999999, 'TCH <= 1.245\ngini = 0.213\nsamples = 174\nvalue = [0, 237, 4,  
28]\nklass = b'),  
Text(1562.4, 181.1999999999982, 'gini = 0.584\nsamples = 20\nvalue = [0, 13, 3, 14]\nklass = d'),  
Text(1711.2, 181.1999999999982, 'gini = 0.118\nsamples = 154\nvalue = [0, 224, 1, 14]\nklass = b'),  
Text(2083.200000000003, 906.0, 'SO_2 <= 5.125\ngini = 0.619\nsamples = 317\nvalue = [72, 264, 20, 154]\nklass = b'),  
Text(1934.4, 543.599999999999, 'TOL <= 2.98\ngini = 0.135\nsamples = 122\nvalue = [0, 178, 0, 14]\nklass =
```

```
b'),
Text(1860.000000000002, 181.19999999999982, 'gini = 0.439\nsamples = 25\nvalue = [0, 27, 0, 13]\nclass =
b'),
Text(2008.800000000002, 181.19999999999982, 'gini = 0.013\nsamples = 97\nvalue = [0, 151, 0, 1]\nclass =
b'),
Text(2232.0, 543.599999999999, 'OXY <= 2.06\ngini = 0.678\nsamples = 195\nvalue = [72, 86, 20, 140]\nclass
= d'),
Text(2157.600000000004, 181.19999999999982, 'gini = 0.499\nsamples = 117\nvalue = [7, 49, 14, 138]\nclass
= d'),
Text(2306.4, 181.19999999999982, 'gini = 0.534\nsamples = 78\nvalue = [65, 37, 6, 2]\nclass = a'),
Text(3366.600000000004, 1630.800000000002, 'SO_2 <= 5.795\ngini = 0.739\nsamples = 11928\nvalue = [4780,
3116, 5295, 5699]\nclass = d'),
Text(2864.4, 1268.4, 'CO <= 0.345\ngini = 0.124\nsamples = 1572\nvalue = [54, 2323, 0, 109]\nclass = b'),
Text(2678.4, 906.0, 'CO <= 0.165\ngini = 0.47\nsamples = 260\nvalue = [41, 285, 0, 87]\nclass = b'),
Text(2529.600000000004, 543.599999999999, 'OXY <= 1.67\ngini = 0.083\nsamples = 43\nvalue = [3, 66, 0, 0]
\nclass = b'),
Text(2455.200000000003, 181.19999999999982, 'gini = 0.0\nsamples = 28\nvalue = [0, 48, 0, 0]\nclass = b'),
Text(2604.0, 181.19999999999982, 'gini = 0.245\nsamples = 15\nvalue = [3, 18, 0, 0]\nclass = b'),
Text(2827.200000000003, 543.599999999999, 'PXY <= 0.73\ngini = 0.519\nsamples = 217\nvalue = [38, 219, 0,
87]\nclass = b'),
Text(2752.8, 181.19999999999982, 'gini = 0.14\nsamples = 74\nvalue = [5, 112, 0, 4]\nclass = b'),
Text(2901.600000000004, 181.19999999999982, 'gini = 0.609\nsamples = 143\nvalue = [33, 107, 0, 83]\nclass
= b'),
Text(3050.4, 906.0, 'SO_2 <= 5.005\ngini = 0.033\nsamples = 1312\nvalue = [13, 2038, 0, 22]\nclass = b'),
Text(2976.0, 543.599999999999, 'gini = 0.0\nsamples = 1081\nvalue = [0, 1694, 0, 0]\nclass = b'),
Text(3124.8, 543.599999999999, 'NMHC <= 0.095\ngini = 0.172\nsamples = 231\nvalue = [13, 344, 0, 22]\nclass
= b'),
Text(3050.4, 181.19999999999982, 'gini = 0.632\nsamples = 22\nvalue = [13, 6, 0, 14]\nclass = d'),
Text(3199.200000000003, 181.19999999999982, 'gini = 0.045\nsamples = 209\nvalue = [0, 338, 0, 8]\nclass
= b'),
Text(3868.8, 1268.4, 'OXY <= 4.185\ngini = 0.694\nsamples = 10356\nvalue = [4726, 793, 5295, 5590]\nclass
= d'),
Text(3571.200000000003, 906.0, 'SO_2 <= 7.145\ngini = 0.661\nsamples = 6896\nvalue = [1678, 669, 3814, 472
3]\nclass = d'),
Text(3422.4, 543.599999999999, 'NMHC <= 0.035\ngini = 0.657\nsamples = 848\nvalue = [215, 321, 138, 677]\nclass
= d'),
Text(3348.000000000005, 181.19999999999982, 'gini = 0.158\nsamples = 142\nvalue = [214, 0, 18, 2]\nclass
= a'),
Text(3496.8, 181.19999999999982, 'gini = 0.541\nsamples = 706\nvalue = [1, 321, 120, 675]\nclass = d'),
Text(3720.000000000005, 543.599999999999, 'O_3 <= 7.035\ngini = 0.646\nsamples = 6048\nvalue = [1463, 34
8, 3676, 4046]\nclass = d'),
Text(3645.600000000004, 181.19999999999982, 'gini = 0.598\nsamples = 872\nvalue = [99, 165, 792, 333]\nclass
= c'),
```

```

Text(3794.4, 181.1999999999982, 'gini = 0.638\nsamples = 5176\nvalue = [1364, 183, 2884, 3713]\nklass = d'),
Text(4166.400000000001, 906.0, 'OXY <= 5.565\ngini = 0.598\nsamples = 3460\nvalue = [3048, 124, 1481, 867]\nklass = a'),
Text(4017.600000000004, 543.5999999999999, 'O_3 <= 11.56\ngini = 0.675\nsamples = 1292\nvalue = [845, 89, 631, 459]\nklass = a'),
Text(3943.200000000003, 181.1999999999982, 'gini = 0.703\nsamples = 565\nvalue = [197, 75, 325, 279]\nklass = c'),
Text(4092.000000000005, 181.1999999999982, 'gini = 0.586\nsamples = 727\nvalue = [648, 14, 306, 180]\nklass = a'),
Text(4315.200000000001, 543.5999999999999, 'O_3 <= 8.505\ngini = 0.53\nsamples = 2168\nvalue = [2203, 35, 850, 408]\nklass = a'),
Text(4240.8, 181.1999999999982, 'gini = 0.653\nsamples = 870\nvalue = [555, 32, 553, 267]\nklass = a'),
Text(4389.6, 181.1999999999982, 'gini = 0.353\nsamples = 1298\nvalue = [1648, 3, 297, 141]\nklass = a')

```



## Conclusion

## Accuracy

*Linear Regression:* 0.20226473515319232

*Ridge Regression:* 0.20204408096897053

*Lasso Regression:* 0.05896723554942673

*ElasticNet Regression:* 0.09906557806958294

*Logistic Regression:* 0.8480899292795158

*Random Forest:* 0.7758757610517957

**Logistic Regression is suitable for this dataset**