

Importing Libraries

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```

1 df=pd.read_csv("madrid_2003")
2 df

```

Out[2]:

| | | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PXY | SO_2 | TCH | TOL | stat |
|--------|-----|---------------------|------|------|------|------|------|-----------|------------|------|-----------|-----------|------|-----------|------|------|--------|
| 0 | | 2003-03-01 01:00:00 | NaN | 1.72 | NaN | NaN | NaN | 73.900002 | 316.299988 | NaN | 10.550000 | 55.209999 | NaN | 24.299999 | NaN | NaN | 280790 |
| 1 | | 2003-03-01 01:00:00 | NaN | 1.45 | NaN | NaN | 0.26 | 72.110001 | 250.000000 | 0.73 | 6.720000 | 52.389999 | NaN | 14.230000 | 1.55 | NaN | 280790 |
| 2 | | 2003-03-01 01:00:00 | NaN | 1.57 | NaN | NaN | NaN | 80.559998 | 224.199997 | NaN | 21.049999 | 63.240002 | NaN | 17.879999 | NaN | NaN | 280790 |
| 3 | | 2003-03-01 01:00:00 | NaN | 2.45 | NaN | NaN | NaN | 78.370003 | 450.399994 | NaN | 4.220000 | 67.839996 | NaN | 24.900000 | NaN | NaN | 280790 |
| 4 | | 2003-03-01 01:00:00 | NaN | 3.26 | NaN | NaN | NaN | 96.250000 | 479.100006 | NaN | 8.460000 | 95.779999 | NaN | 18.750000 | NaN | NaN | 280790 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 243979 | | 2003-10-01 00:00:00 | 0.20 | 0.16 | 2.01 | 3.17 | 0.02 | 31.799999 | 32.299999 | 1.68 | 34.049999 | 7.380000 | 1.20 | 4.870000 | 1.27 | 1.00 | 280790 |
| 243980 | | 2003-10-01 00:00:00 | 0.32 | 0.08 | 0.36 | 0.72 | NaN | 10.450000 | 14.760000 | 1.00 | 34.610001 | 7.400000 | 0.50 | 8.360000 | NaN | 0.88 | 280790 |
| 243981 | | 2003-10-01 00:00:00 | NaN | NaN | NaN | NaN | 0.07 | 34.639999 | 50.810001 | NaN | 32.160000 | 16.830000 | NaN | 5.330000 | 1.55 | NaN | 280790 |
| 243982 | | 2003-10-01 00:00:00 | NaN | NaN | NaN | NaN | 0.07 | 32.580002 | 41.020000 | NaN | NaN | 13.570000 | NaN | 6.830000 | 1.27 | NaN | 280790 |
| 243983 | | 2003-10-01 00:00:00 | 1.00 | 0.29 | 2.15 | 6.41 | 0.07 | 37.150002 | 56.849998 | 2.28 | 21.480000 | 12.350000 | 2.43 | 6.060000 | 1.32 | 5.67 | 280790 |

243984 rows × 16 columns



Data Cleaning and Data Preprocessing

```
In [3]: 1 df=df.dropna()
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
               dtype='object')
```

```
In [5]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 33010 entries, 5 to 243983
Data columns (total 16 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      33010 non-null   object 
 1   BEN       33010 non-null   float64
 2   CO        33010 non-null   float64
 3   EBE       33010 non-null   float64
 4   MXY       33010 non-null   float64
 5   NMHC      33010 non-null   float64
 6   NO_2      33010 non-null   float64
 7   NOx       33010 non-null   float64
 8   OXY       33010 non-null   float64
 9   O_3        33010 non-null   float64
 10  PM10      33010 non-null   float64
 11  PXY       33010 non-null   float64
 12  SO_2      33010 non-null   float64
 13  TCH       33010 non-null   float64
 14  TOL       33010 non-null   float64
 15  station    33010 non-null   int64  
dtypes: float64(14), int64(1), object(1)
memory usage: 4.3+ MB
```

```
In [6]: 1 data=df[['CO' , 'station']]  
2 data
```

Out[6]:

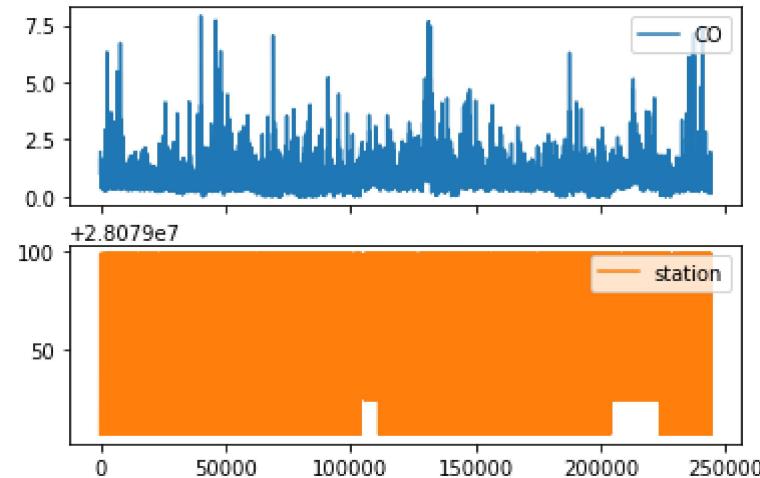
| | CO | station |
|--------|------|----------|
| 5 | 1.94 | 28079006 |
| 23 | 1.27 | 28079024 |
| 27 | 1.79 | 28079099 |
| 33 | 1.47 | 28079006 |
| 51 | 1.29 | 28079024 |
| ... | ... | ... |
| 243955 | 0.41 | 28079099 |
| 243957 | 0.60 | 28079035 |
| 243961 | 0.82 | 28079006 |
| 243979 | 0.16 | 28079024 |
| 243983 | 0.29 | 28079099 |

33010 rows × 2 columns

Line chart

```
In [7]: 1 data.plot.line(subplots=True)
```

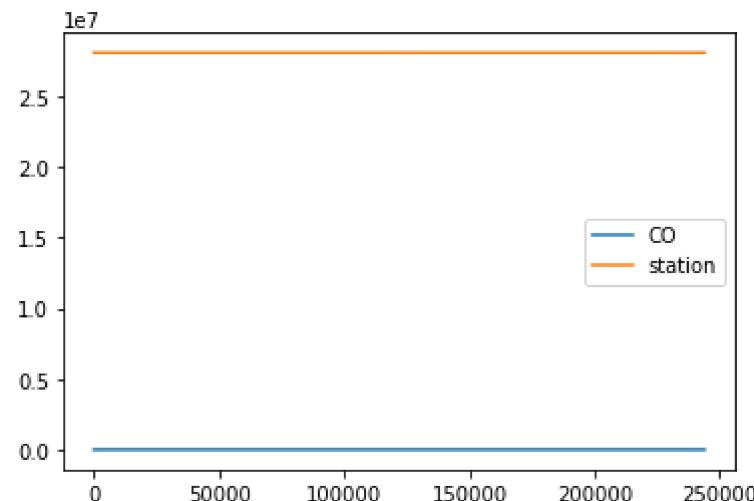
```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



Line chart

```
In [8]: 1 data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

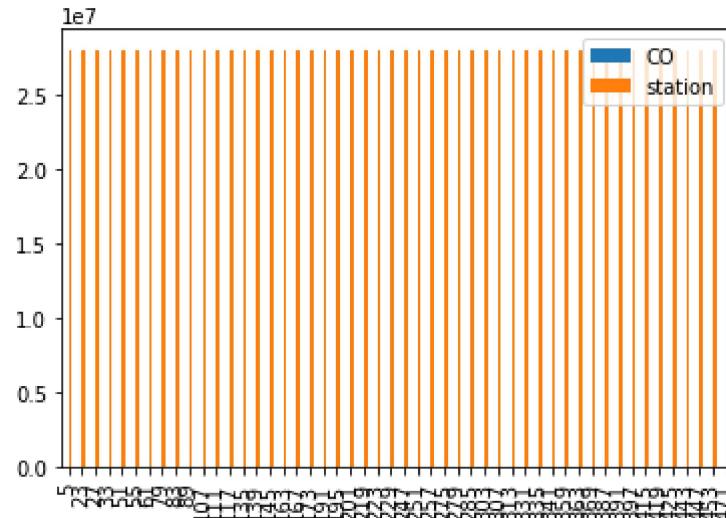


Bar chart

```
In [9]: 1 b=data[0:50]
```

```
In [10]: 1 b.plot.bar()
```

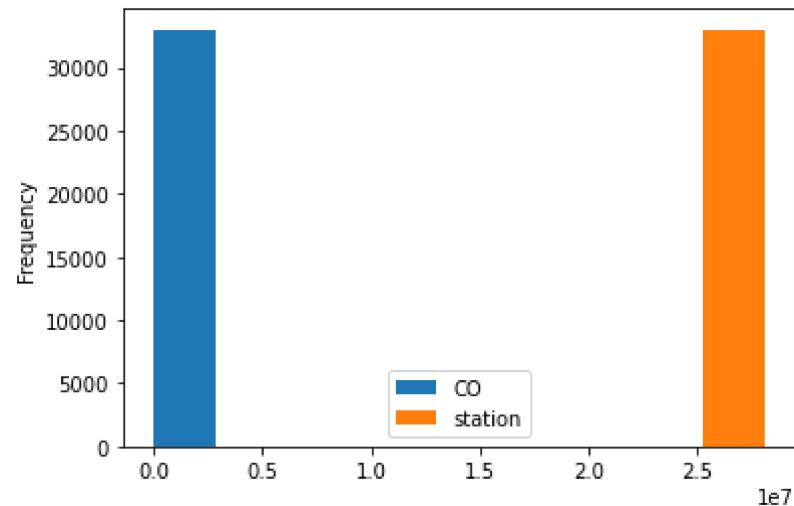
```
Out[10]: <AxesSubplot:>
```



Histogram

```
In [11]: 1 data.plot.hist()
```

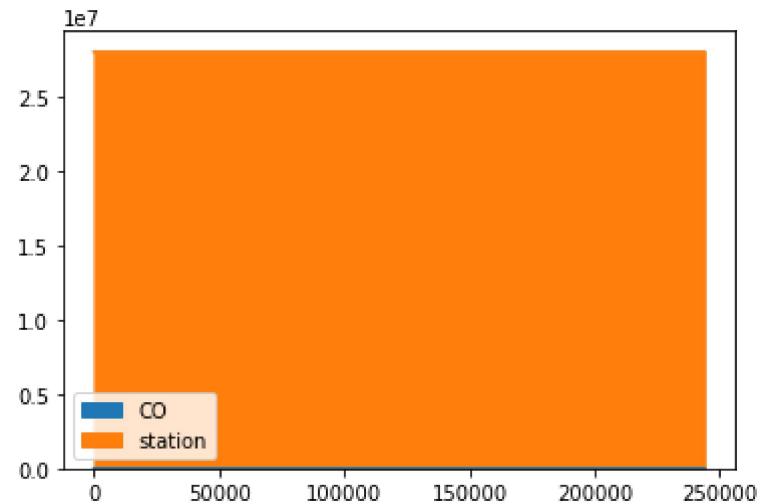
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

In [12]: 1 data.plot.area()

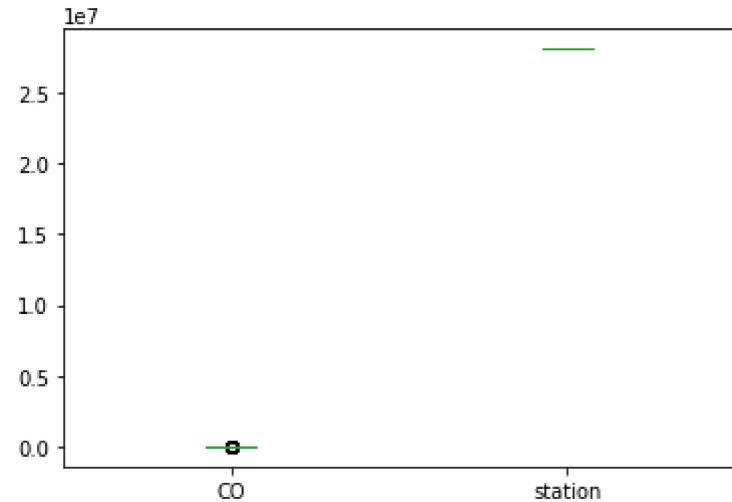
Out[12]: <AxesSubplot:>



Box chart

In [13]: 1 data.plot.box()

Out[13]: <AxesSubplot:>



Pie chart

```
In [14]: 1 b.plot.pie(y='station' )
```

```
Out[14]: <AxesSubplot:ylabel='station'>
```

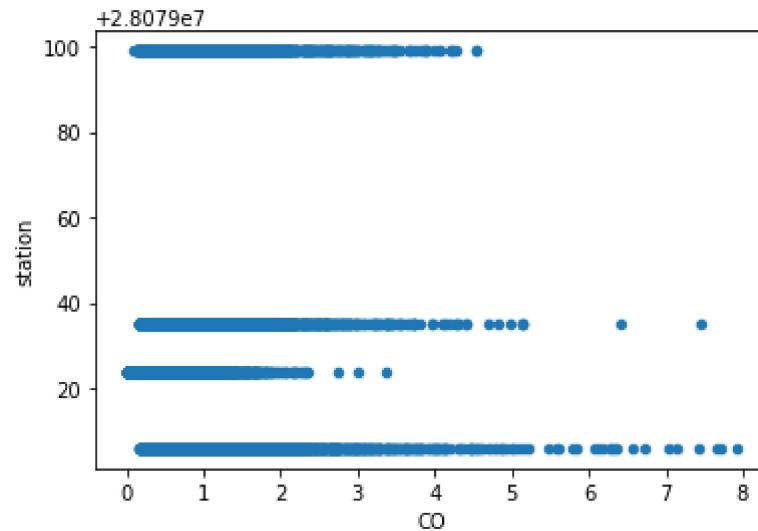





Scatter chart

```
In [15]: 1 data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 33010 entries, 5 to 243983
Data columns (total 16 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   date      33010 non-null   object  
 1   BEN        33010 non-null   float64 
 2   CO         33010 non-null   float64 
 3   EBE        33010 non-null   float64 
 4   MXY        33010 non-null   float64 
 5   NMHC       33010 non-null   float64 
 6   NO_2       33010 non-null   float64 
 7   NOx        33010 non-null   float64 
 8   OXY        33010 non-null   float64 
 9   O_3         33010 non-null   float64 
 10  PM10       33010 non-null   float64 
 11  PXY        33010 non-null   float64 
 12  SO_2       33010 non-null   float64 
 13  TCH        33010 non-null   float64 
 14  TOL        33010 non-null   float64
```

```
In [17]: 1 df.describe()
```

Out[17]:

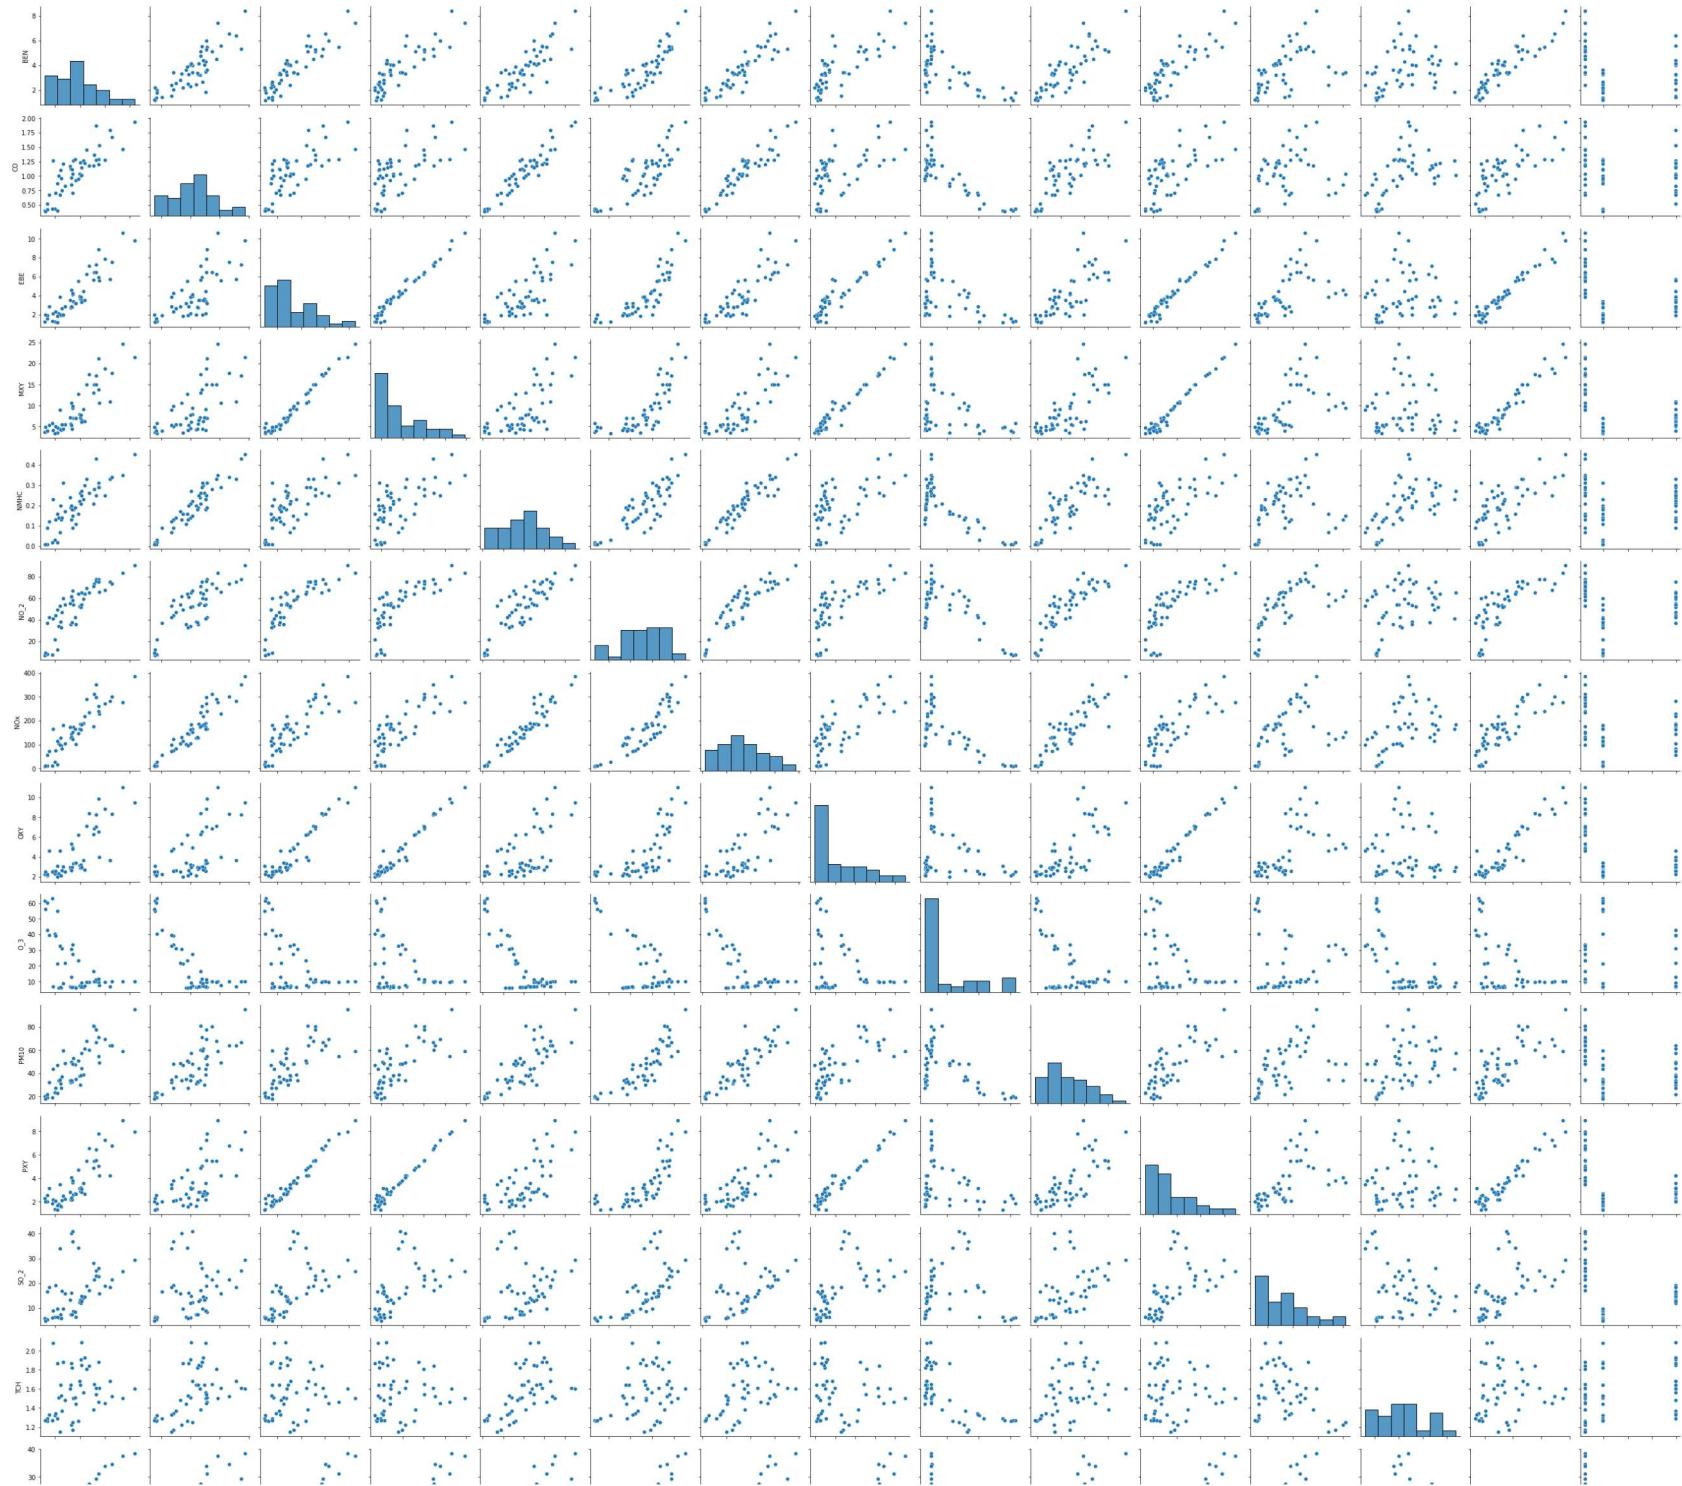
| | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|--------------|
| count | 33010.000000 | 33010.000000 | 33010.000000 | 33010.000000 | 33010.000000 | 33010.000000 | 33010.000000 | 33010.000000 | 33010.000000 |
| mean | 2.192633 | 0.759868 | 2.639726 | 5.838414 | 0.137177 | 57.328049 | 120.153676 | 2.684084 | 36.914485 |
| std | 2.064160 | 0.545999 | 2.825194 | 6.267296 | 0.127863 | 31.811082 | 104.521700 | 2.717832 | 28.988487 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.130000 |
| 25% | 0.900000 | 0.430000 | 1.010000 | 1.880000 | 0.060000 | 34.529999 | 49.070000 | 1.000000 | 12.230000 |
| 50% | 1.610000 | 0.620000 | 1.890000 | 4.070000 | 0.110000 | 55.105000 | 92.779999 | 1.790000 | 30.400000 |
| 75% | 2.810000 | 0.930000 | 3.300000 | 7.530000 | 0.170000 | 76.160004 | 160.100006 | 3.340000 | 54.349998 |
| max | 66.389999 | 7.920000 | 92.589996 | 177.600006 | 2.180000 | 342.700012 | 1246.000000 | 88.180000 | 178.699997 |

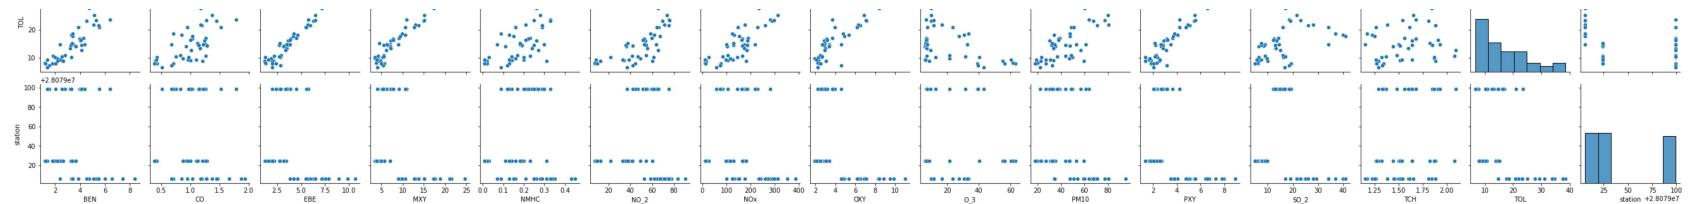
```
In [18]: 1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

```
In [19]: 1 sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x29a6651d670>
```

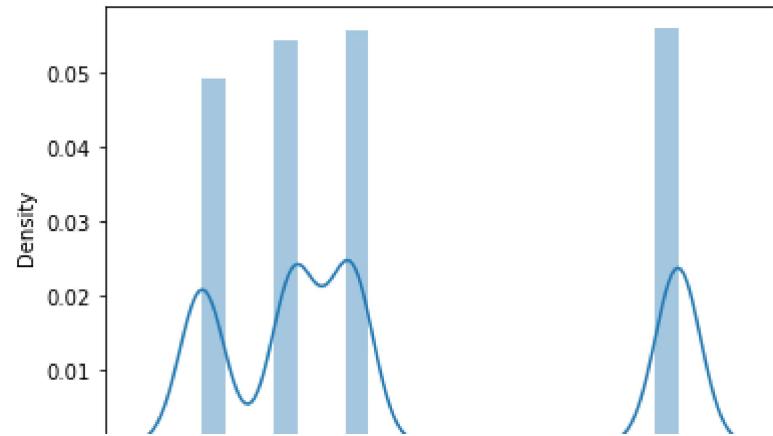





```
In [20]: 1 sns.distplot(df1['station'])
```

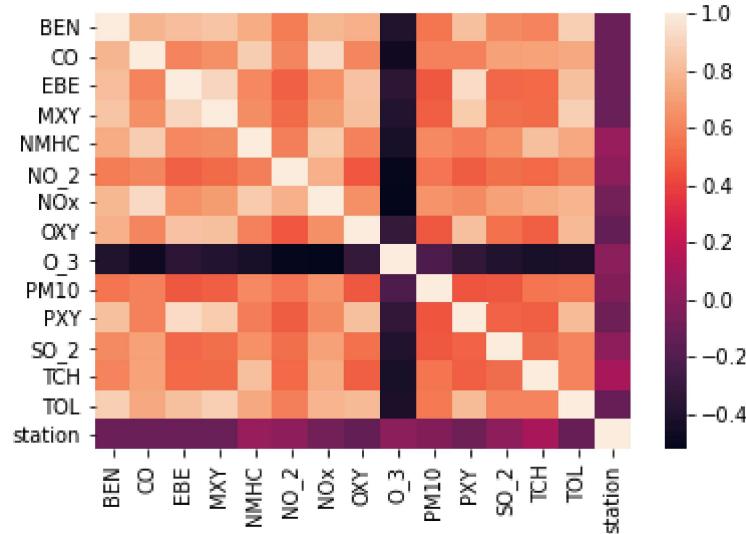
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: 1 sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
```

```
In [23]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: 1 from sklearn.linear_model import LinearRegression  
2 lr=LinearRegression()  
3 lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: 1 lr.intercept_
```

```
Out[25]: 28079002.289702866
```

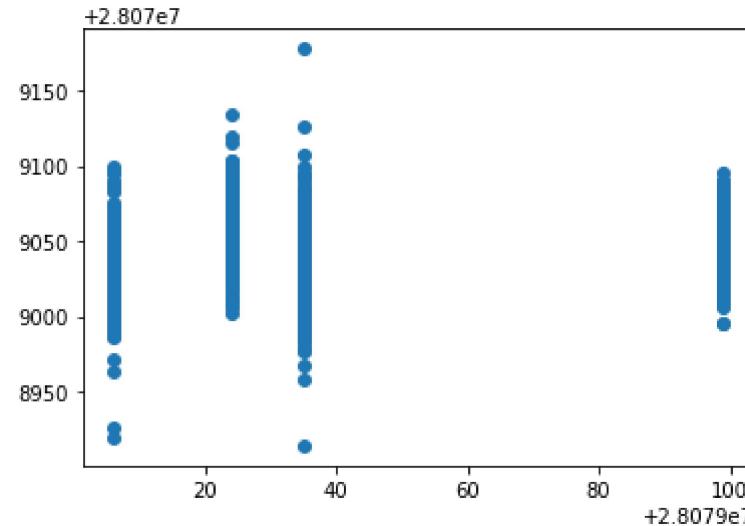
```
In [26]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
2 coeff
```

```
Out[26]:
```

| | Co-efficient |
|------|--------------|
| BEN | 2.207835 |
| CO | -37.683984 |
| EBE | -1.881216 |
| MXY | 0.139755 |
| NMHC | 150.686970 |
| NO_2 | 0.175650 |
| NOx | -0.075895 |
| OXY | -1.435456 |
| O_3 | -0.009738 |
| PM10 | -0.069049 |
| PXY | 1.811652 |
| SO_2 | 0.809381 |
| TCH | 34.508579 |
| TOL | -0.875362 |

```
In [27]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x29a7620a670>
```



ACCURACY

```
In [28]: 1 lr.score(x_test,y_test)
```

```
Out[28]: 0.1809341834858923
```

```
In [29]: 1 lr.score(x_train,y_train)
```

```
Out[29]: 0.17328056958079485
```

Ridge and Lasso

```
In [30]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: 1 rr.score(x_test,y_test)
```

```
Out[32]: 0.17879287952357292
```

```
In [33]: 1 rr.score(x_train,y_train)
```

```
Out[33]: 0.17227501231797626
```

```
In [34]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

```
In [35]: 1 la.score(x_train,y_train)
```

```
Out[35]: 0.03578590238383095
```

Accuracy(Lasso)

```
In [36]: 1 la.score(x_test,y_test)
```

```
Out[36]: 0.033703301186020385
```

```
In [37]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: 1 en.coef_
```

```
Out[38]: array([ 0.          , -0.19683471,  0.          , -0.04925999,  0.14178781,
   0.16321412, -0.06889313, -1.27697124, -0.03629348,  0.05801668,
   0.27493181,  0.73121501,  1.57332176, -0.40068216])
```

```
In [39]: 1 en.intercept_
```

```
Out[39]: 28079037.33122153
```

```
In [40]: 1 prediction=en.predict(x_test)
```

```
In [41]: 1 en.score(x_test,y_test)
```

```
Out[41]: 0.04813959207478369
```

Evaluation Metrics

```
In [42]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
29.002865032872688
```

```
1173.984806997495
```

```
34.263461690224695
```

Logistic Regression

```
In [43]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [44]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df['station']
```

```
In [45]: 1 feature_matrix.shape
```

```
Out[45]: (33010, 14)
```

```
In [46]: 1 target_vector.shape
```

```
Out[46]: (33010,)
```

```
In [47]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [48]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: 1 logr=LogisticRegression(max_iter=10000)  
2 logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: 1 prediction=logr.predict(observation)  
2 print(prediction)
```

```
[28079035]
```

```
In [52]: 1 logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079035, 28079099], dtype=int64)
```

```
In [53]: 1 logr.score(fs,target_vector)
```

```
Out[53]: 0.7584974250227204
```

```
In [54]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[54]: 2.3306153265290618e-23
```

```
In [55]: 1 logr.predict_proba(observation)
```

```
Out[55]: array([[2.33061533e-23, 1.44436075e-55, 1.00000000e+00, 6.68457491e-16]])
```

Random Forest

```
In [56]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: 1 rfc=RandomForestClassifier()
2 rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: 1 parameters={'max_depth':[1,2,3,4,5],
2                 'min_samples_leaf':[5,10,15,20,25],
3                 'n_estimators':[10,20,30,40,50]
4 }
```

```
In [59]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')
```

```
In [60]: 1 grid_search.best_score_
```

```
Out[60]: 0.7276148693348015
```

```
In [61]: 1 rfc_best=grid_search.best_estimator_
```

In [62]:

```
1 from sklearn.tree import plot_tree
2
3 plt.figure(figsize=(80,40))
4 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

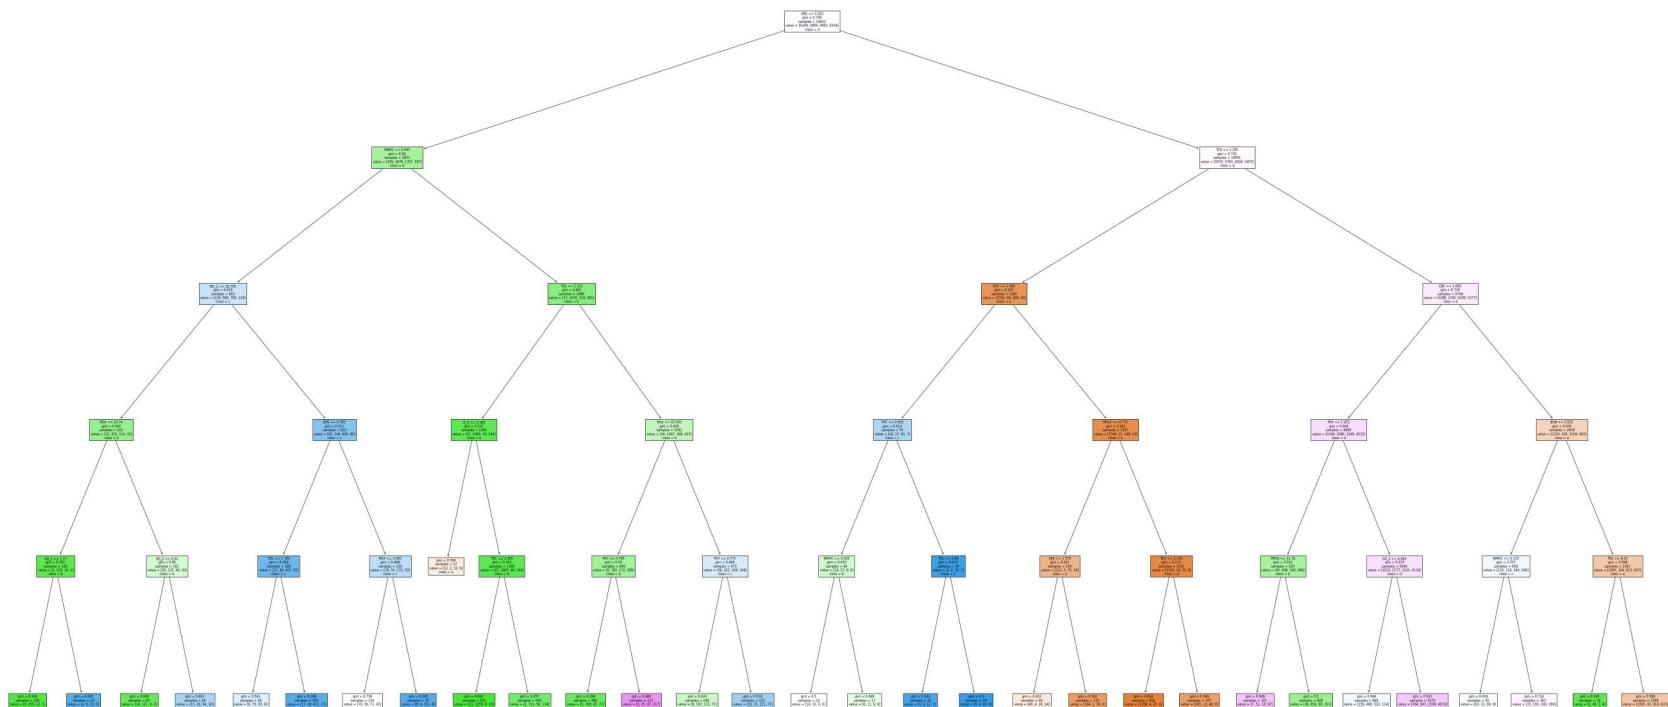
Out[62]: [Text(2166.9, 1993.2, 'EBE <= 1.015\ngini = 0.749\nsamples = 14642\nvalue = [5269, 5869, 5865, 6104]\nnclass = d'),
Text(1060.2, 1630.8000000000002, 'NMHC <= 0.045\ngini = 0.56\nsamples = 3651\nvalue = [199, 3476, 1357, 697]\nnclass = b'),
Text(595.2, 1268.4, 'NO_2 <= 18.795\ngini = 0.616\nsamples = 963\nvalue = [126, 500, 783, 116]\nnclass = c'),
Text(297.6, 906.0, 'NOx <= 14.74\ngini = 0.502\nsamples = 341\nvalue = [31, 352, 114, 31]\nnclass = b'),
Text(148.8, 543.5999999999999, 'SO_2 <= 7.77\ngini = 0.201\nsamples = 160\nvalue = [2, 215, 24, 1]\nnclass = b'),
Text(74.4, 181.1999999999982, 'gini = 0.108\nsamples = 150\nvalue = [0, 215, 12, 1]\nnclass = b'),
Text(223.2000000000002, 181.1999999999982, 'gini = 0.245\nsamples = 10\nvalue = [2, 0, 12, 0]\nnclass = c'),
Text(446.4000000000003, 543.5999999999999, 'SO_2 <= 5.52\ngini = 0.65\nsamples = 181\nvalue = [29, 137, 90, 30]\nnclass = b'),
Text(372.0, 181.1999999999982, 'gini = 0.269\nsamples = 87\nvalue = [14, 111, 6, 0]\nnclass = b'),
Text(520.800000000001, 181.1999999999982, 'gini = 0.631\nsamples = 94\nvalue = [15, 26, 84, 30]\nnclass = c'),
Text(892.800000000001, 906.0, 'EBE <= 0.785\ngini = 0.511\nsamples = 622\nvalue = [95, 148, 669, 85]\nnclass = c'),
Text(744.0, 543.5999999999999, 'TOL <= 1.745\ngini = 0.383\nsamples = 398\nvalue = [17, 98, 497, 35]\nnclass = c'),
Text(669.6, 181.1999999999982, 'gini = 0.541\nsamples = 94\nvalue = [0, 70, 85, 8]\nnclass = c'),
Text(818.400000000001, 181.1999999999982, 'gini = 0.268\nsamples = 304\nvalue = [17, 28, 412, 27]\nnclass = c'),
Text(1041.600000000001, 543.5999999999999, 'MXY <= 2.045\ngini = 0.668\nsamples = 224\nvalue = [78, 50, 172, 50]\nnclass = c'),
Text(967.2, 181.1999999999982, 'gini = 0.738\nsamples = 154\nvalue = [70, 50, 71, 42]\nnclass = c'),
Text(1116.0, 181.1999999999982, 'gini = 0.245\nsamples = 70\nvalue = [8, 0, 101, 8]\nnclass = c'),
Text(1525.2, 1268.4, 'TOL <= 2.315\ngini = 0.461\nsamples = 2688\nvalue = [73, 2976, 574, 581]\nnclass = b'),
Text(1264.800000000002, 906.0, 'O_3 <= 5.185\ngini = 0.212\nsamples = 1406\nvalue = [37, 1969, 76, 144]\nnclass = b'),
Text(1190.4, 543.5999999999999, 'gini = 0.569\nsamples = 17\nvalue = [12, 2, 10, 0]\nnclass = a'),
Text(1339.2, 543.5999999999999, 'TOL <= 1.455\ngini = 0.197\nsamples = 1389\nvalue = [25, 1967, 66, 144]\nnclass = b'),
Text(1264.800000000002, 181.1999999999982, 'gini = 0.062\nsamples = 823\nvalue = [23, 1256, 8, 10]\nnclass = b'),
Text(1413.600000000001, 181.1999999999982, 'gini = 0.357\nsamples = 566\nvalue = [2, 711, 58, 134]\nnclass = b'),
Text(1785.600000000001, 906.0, 'NOx <= 63.995\ngini = 0.628\nsamples = 1282\nvalue = [36, 1007, 498, 437]\nnclass = b'),
Text(1636.800000000002, 543.5999999999999, 'PXY <= 0.795\ngini = 0.53\nsamples = 809\nvalue = [0, 785, 172, 289]\nnclass = b')]

```
Text(1562.4, 181.1999999999982, 'gini = 0.298\nsamples = 586\nvalue = [0, 760, 85, 72]\nclass = b'),  
Text(1711.2, 181.1999999999982, 'gini = 0.489\nsamples = 223\nvalue = [0, 25, 87, 217]\nclass = d'),  
Text(1934.4, 543.5999999999999, 'PXY <= 0.775\nngini = 0.666\nsamples = 473\nvalue = [36, 222, 326, 148]\nclass = c'),  
Text(1860.000000000002, 181.1999999999982, 'gini = 0.633\nsamples = 248\nvalue = [8, 187, 115, 71]\nclass = b'),  
Text(2008.800000000002, 181.1999999999982, 'gini = 0.574\nsamples = 225\nvalue = [28, 35, 211, 77]\nclass = c'),  
Text(3273.600000000004, 1630.800000000002, 'TCH <= 1.255\nngini = 0.732\nsamples = 10991\nvalue = [5070, 2393, 4508, 5407]\nclass = d'),  
Text(2678.4, 1268.4, 'OXY <= 1.005\nngini = 0.247\nsamples = 1283\nvalue = [1782, 48, 209, 30]\nclass = a'),  
Text(2380.8, 906.0, 'PXY <= 0.905\nngini = 0.614\nsamples = 76\nvalue = [16, 27, 61, 7]\nclass = c'),  
Text(2232.0, 543.5999999999999, 'NMHC <= 0.035\nngini = 0.672\nsamples = 40\nvalue = [16, 27, 9, 6]\nclass = b'),  
Text(2157.600000000004, 181.1999999999982, 'gini = 0.5\nsamples = 23\nvalue = [16, 16, 0, 0]\nclass = a'),  
Text(2306.4, 181.1999999999982, 'gini = 0.648\nsamples = 17\nvalue = [0, 11, 9, 6]\nclass = b'),  
Text(2529.600000000004, 543.5999999999999, 'TOL <= 4.96\nngini = 0.037\nsamples = 36\nvalue = [0, 0, 52, 1]\nclass = c'),  
Text(2455.200000000003, 181.1999999999982, 'gini = 0.142\nsamples = 10\nvalue = [0, 0, 12, 1]\nclass = c'),  
Text(2604.0, 181.1999999999982, 'gini = 0.0\nsamples = 26\nvalue = [0, 0, 40, 0]\nclass = c'),  
Text(2976.0, 906.0, 'PM10 <= 9.775\nngini = 0.181\nsamples = 1207\nvalue = [1766, 21, 148, 23]\nclass = a'),  
Text(2827.200000000003, 543.5999999999999, 'EBE <= 1.575\nngini = 0.457\nsamples = 194\nvalue = [213, 5, 75, 14]\nclass = a'),  
Text(2752.8, 181.1999999999982, 'gini = 0.632\nsamples = 69\nvalue = [49, 4, 39, 14]\nclass = a'),  
Text(2901.600000000004, 181.1999999999982, 'gini = 0.302\nsamples = 125\nvalue = [164, 1, 36, 0]\nclass = a'),  
Text(3124.8, 543.5999999999999, 'TCH <= 1.235\nngini = 0.113\nsamples = 1013\nvalue = [1553, 16, 73, 9]\nclass = a'),  
Text(3050.4, 181.1999999999982, 'gini = 0.043\nsamples = 816\nvalue = [1288, 4, 25, 0]\nclass = a'),  
Text(3199.200000000003, 181.1999999999982, 'gini = 0.348\nsamples = 197\nvalue = [265, 12, 48, 9]\nclass = a'),  
Text(3868.8, 1268.4, 'EBE <= 3.685\nngini = 0.728\nsamples = 9708\nvalue = [3288, 2345, 4299, 5377]\nclass = d'),  
Text(3571.200000000003, 906.0, 'PXY <= 1.055\nngini = 0.694\nsamples = 6869\nvalue = [1068, 2085, 3240, 4532]\nclass = d'),  
Text(3422.4, 543.5999999999999, 'PM10 <= 11.14\nngini = 0.532\nsamples = 929\nvalue = [45, 908, 109, 398]\nclass = b'),  
Text(3348.000000000005, 181.1999999999982, 'gini = 0.608\nsamples = 101\nvalue = [7, 52, 19, 87]\nclass = d'),  
Text(3496.8, 181.1999999999982, 'gini = 0.5\nsamples = 828\nvalue = [38, 856, 90, 311]\nclass = b'),  
Text(3720.000000000005, 543.5999999999999, 'SO_2 <= 6.905\nngini = 0.673\nsamples = 5940\nvalue = [1023, 11
```

```

77, 3131, 4134]\nclass = d'),
Text(3645.600000000004, 181.1999999999982, 'gini = 0.694\nsamples = 864\nvalue = [259, 480, 532, 124]\nklass = c'),
Text(3794.4, 181.1999999999982, 'gini = 0.633\nsamples = 5076\nvalue = [764, 697, 2599, 4010]\nklass = d'),
Text(4166.400000000001, 906.0, 'BEN <= 3.025\ngini = 0.645\nsamples = 2839\nvalue = [2220, 260, 1059, 845]\nklass = a'),
Text(4017.600000000004, 543.5999999999999, 'NMHC <= 0.115\ngini = 0.727\nsamples = 458\nvalue = [135, 116, 246, 208]\nklass = c'),
Text(3943.200000000003, 181.1999999999982, 'gini = 0.624\nsamples = 95\nvalue = [63, 13, 66, 9]\nklass = c'),
Text(4092.000000000005, 181.1999999999982, 'gini = 0.714\nsamples = 363\nvalue = [72, 103, 180, 199]\nklass = d'),
Text(4315.200000000001, 543.5999999999999, 'TOL <= 8.32\ngini = 0.598\nsamples = 2381\nvalue = [2085, 144, 813, 637]\nklass = a'),
Text(4240.8, 181.1999999999982, 'gini = 0.226\nsamples = 32\nvalue = [0, 49, 3, 4]\nklass = b'),
Text(4389.6, 181.1999999999982, 'gini = 0.588\nsamples = 2349\nvalue = [2085, 95, 810, 633]\nklass = a')])

```



Conclusion

Accuracy

Linear Regression:0.17328056958079485

Ridge Regression:0.17227501231797626

Lasso Regression:0.03578590238383095

ElasticNet Regression:0.04813959207478369

Logistic Regression:0.7584974250227204

Random Forest:0.7276148693348015

Logistic Regression is suitable for this dataset