

## Importing Libraries

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

## Importing Datasets

In [2]:

```

1 df=pd.read_csv("madrid_2004")
2 df

```

Out[2]:

		date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2	TCH	TO
0		2004-08-01 01:00:00	NaN	0.66	NaN	NaN	NaN	89.550003	118.900002	NaN	40.020000	39.990002	25.860001	NaN	12.20	NaN	Na
1		2004-08-01 01:00:00	2.66	0.54	2.99	6.08	0.18	51.799999	53.860001	3.28	51.689999	22.950001	NaN	3.38	6.12	1.57	11.3
2		2004-08-01 01:00:00	NaN	1.02	NaN	NaN	NaN	93.389999	138.600006	NaN	20.860001	49.480000	NaN	NaN	8.99	NaN	Na
3		2004-08-01 01:00:00	NaN	0.53	NaN	NaN	NaN	87.290001	105.000000	NaN	36.730000	31.070000	NaN	NaN	8.82	NaN	Na
4		2004-08-01 01:00:00	NaN	0.17	NaN	NaN	NaN	34.910000	35.349998	NaN	86.269997	54.080002	NaN	NaN	8.71	NaN	Na
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
245491		2004-06-01 00:00:00	0.75	0.21	0.85	1.55	0.07	59.580002	64.389999	0.66	33.029999	30.900000	14.860000	0.52	6.62	1.28	2.5
245492		2004-06-01 00:00:00	2.49	0.75	2.44	4.57	NaN	97.139999	146.899994	2.34	7.740000	37.689999	NaN	2.35	6.92	NaN	11.9
245493		2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.13	102.699997	132.600006	NaN	17.809999	22.840000	12.040000	NaN	7.82	1.52	Na
245494		2004-06-01 00:00:00	NaN	NaN	NaN	NaN	0.09	82.599998	102.599998	NaN	NaN	45.630001	NaN	NaN	5.53	1.33	Na
245495		2004-06-01 00:00:00	3.01	0.67	2.78	5.12	0.20	92.550003	141.000000	2.60	11.460000	24.389999	17.959999	2.29	8.68	1.47	12.9

245496 rows × 17 columns



# Data Cleaning and Data Preprocessing

```
In [3]: 1 df=df.dropna()
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
               dtype='object')
```

```
In [5]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19397 entries, 5 to 245495
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      19397 non-null   object 
 1   BEN       19397 non-null   float64
 2   CO        19397 non-null   float64
 3   EBE       19397 non-null   float64
 4   MXY       19397 non-null   float64
 5   NMHC      19397 non-null   float64
 6   NO_2      19397 non-null   float64
 7   NOx       19397 non-null   float64
 8   OXY       19397 non-null   float64
 9   O_3        19397 non-null   float64
 10  PM10      19397 non-null   float64
 11  PM25      19397 non-null   float64
 12  PXY       19397 non-null   float64
 13  SO_2      19397 non-null   float64
 14  TCH       19397 non-null   float64
 15  TOL       19397 non-null   float64
 16  station    19397 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 2.7+ MB
```

```
In [6]: 1 data=df[['CO' , 'station']]  
2 data
```

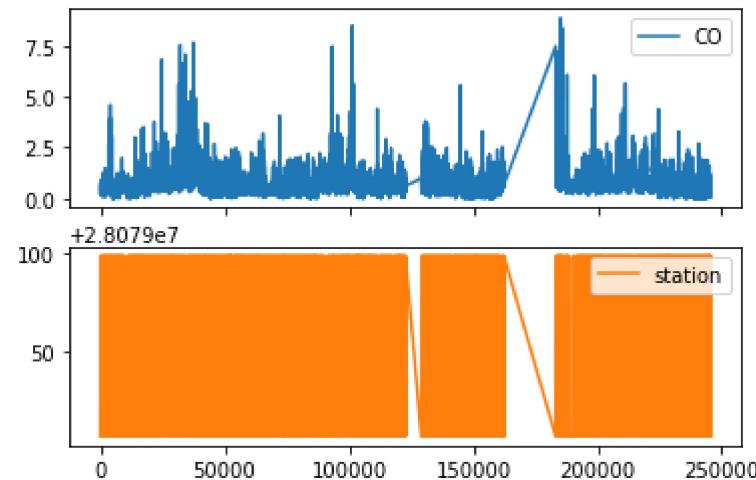
```
Out[6]:      CO    station  
_____  
 5  0.63  28079006  
22  0.36  28079024  
26  0.46  28079099  
32  0.67  28079006  
49  0.30  28079024  
...   ...     ...  
245463  0.08  28079024  
245467  0.67  28079099  
245473  1.12  28079006  
245491  0.21  28079024  
245495  0.67  28079099
```

19397 rows × 2 columns

## Line chart

```
In [7]: 1 data.plot.line(subplots=True)
```

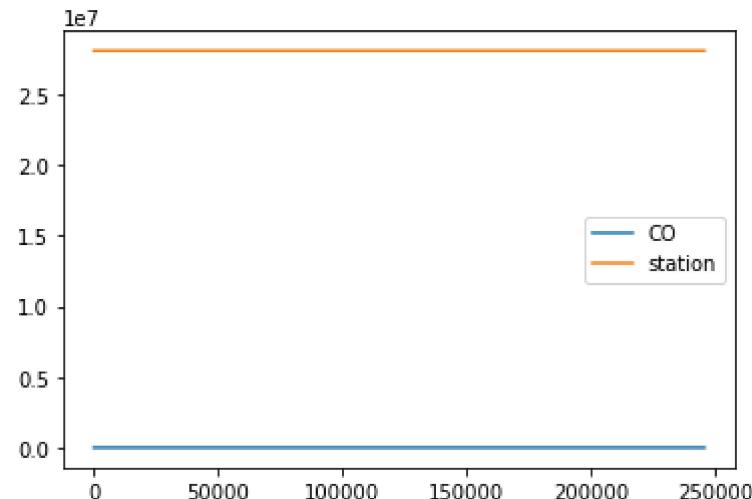
```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



## Line chart

```
In [8]: 1 data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```



## Bar chart

```
In [9]: 1 b=data[0:50]
```

```
In [10]: 1 b.plot.bar()
```

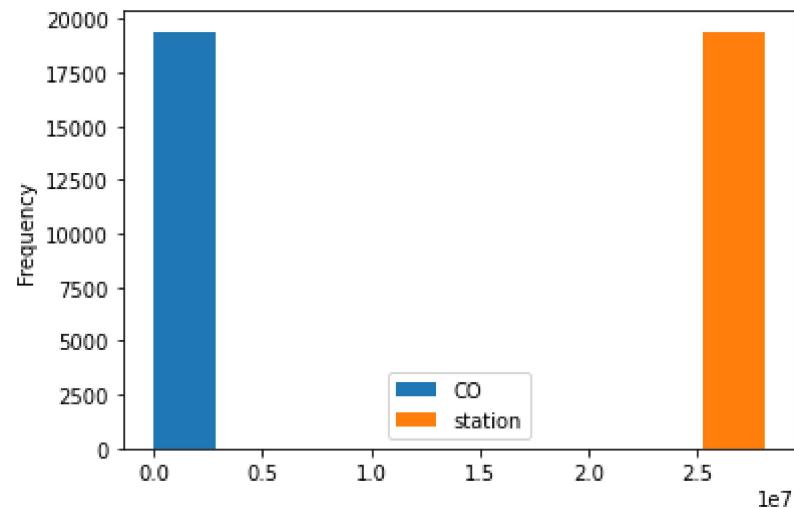
```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: 1 data.plot.hist()
```

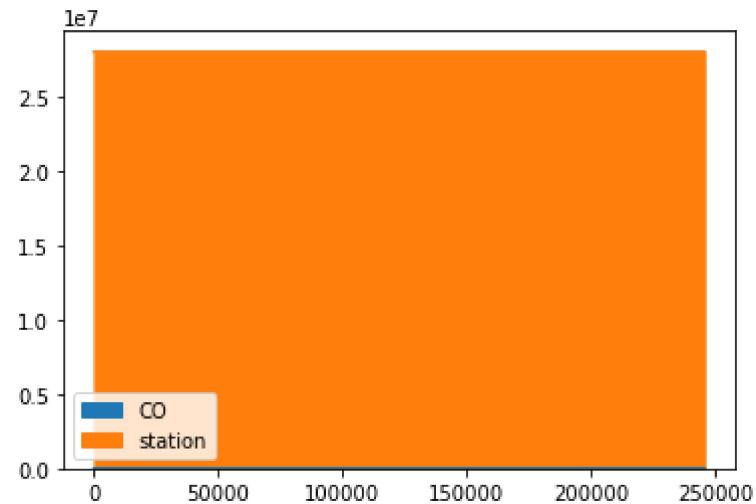
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

In [12]: 1 data.plot.area()

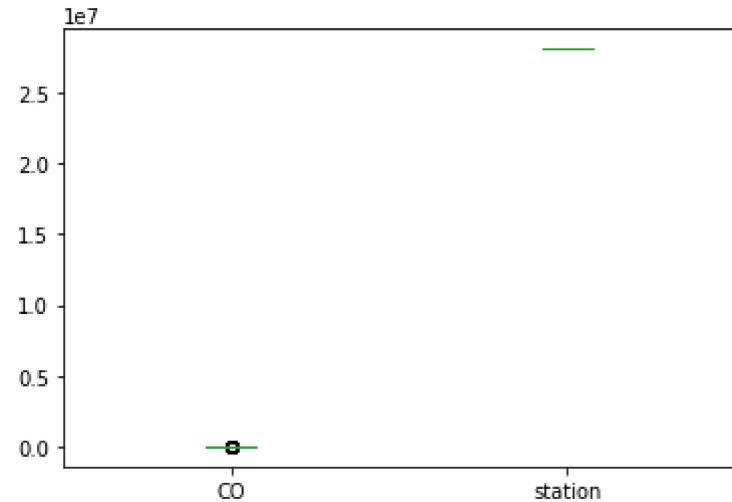
Out[12]: <AxesSubplot:>



## Box chart

In [13]: 1 data.plot.box()

Out[13]: <AxesSubplot:>



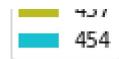
## Pie chart

```
In [14]: 1 b.plot.pie(y='station' )
```

```
Out[14]: <AxesSubplot:ylabel='station'>
```



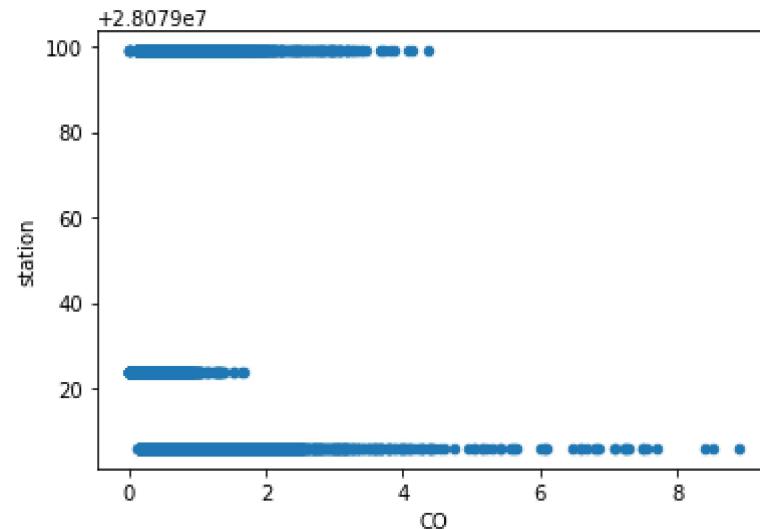




## Scatter chart

```
In [15]: 1 data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19397 entries, 5 to 245495
Data columns (total 17 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   date      19397 non-null    object  
 1   BEN        19397 non-null    float64 
 2   CO         19397 non-null    float64 
 3   EBE        19397 non-null    float64 
 4   MXY        19397 non-null    float64 
 5   NMHC       19397 non-null    float64 
 6   NO_2       19397 non-null    float64 
 7   NOx        19397 non-null    float64 
 8   OXY        19397 non-null    float64 
 9   O_3         19397 non-null    float64 
 10  PM10       19397 non-null    float64 
 11  PM25       19397 non-null    float64 
 12  PXY        19397 non-null    float64 
 13  SO_2       19397 non-null    float64 
 14  TSP        19397 non-null    float64
```

```
In [17]: 1 df.describe()
```

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
count	19397.000000	19397.000000	19397.000000	19397.000000	19397.000000	19397.000000	19397.000000	19397.000000	19397.000000
mean	2.250781	0.675347	2.775913	5.424809	0.151024	62.887023	128.554023	2.71421	38.292501
std	2.184724	0.591026	2.729622	5.554358	0.158603	37.952255	127.912411	2.54485	29.808431
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.090000	2.410000	0.00000	0.770000
25%	0.870000	0.320000	1.020000	1.780000	0.060000	35.150002	45.209999	1.00000	12.710000
50%	1.620000	0.520000	1.970000	3.800000	0.110000	58.310001	93.220001	1.93000	31.830000
75%	2.910000	0.860000	3.580000	7.260000	0.200000	85.730003	174.300003	3.55000	56.139999
max	34.180000	8.900000	41.880001	91.599998	4.810000	355.100006	1700.000000	45.34000	192.500000

In [18]:

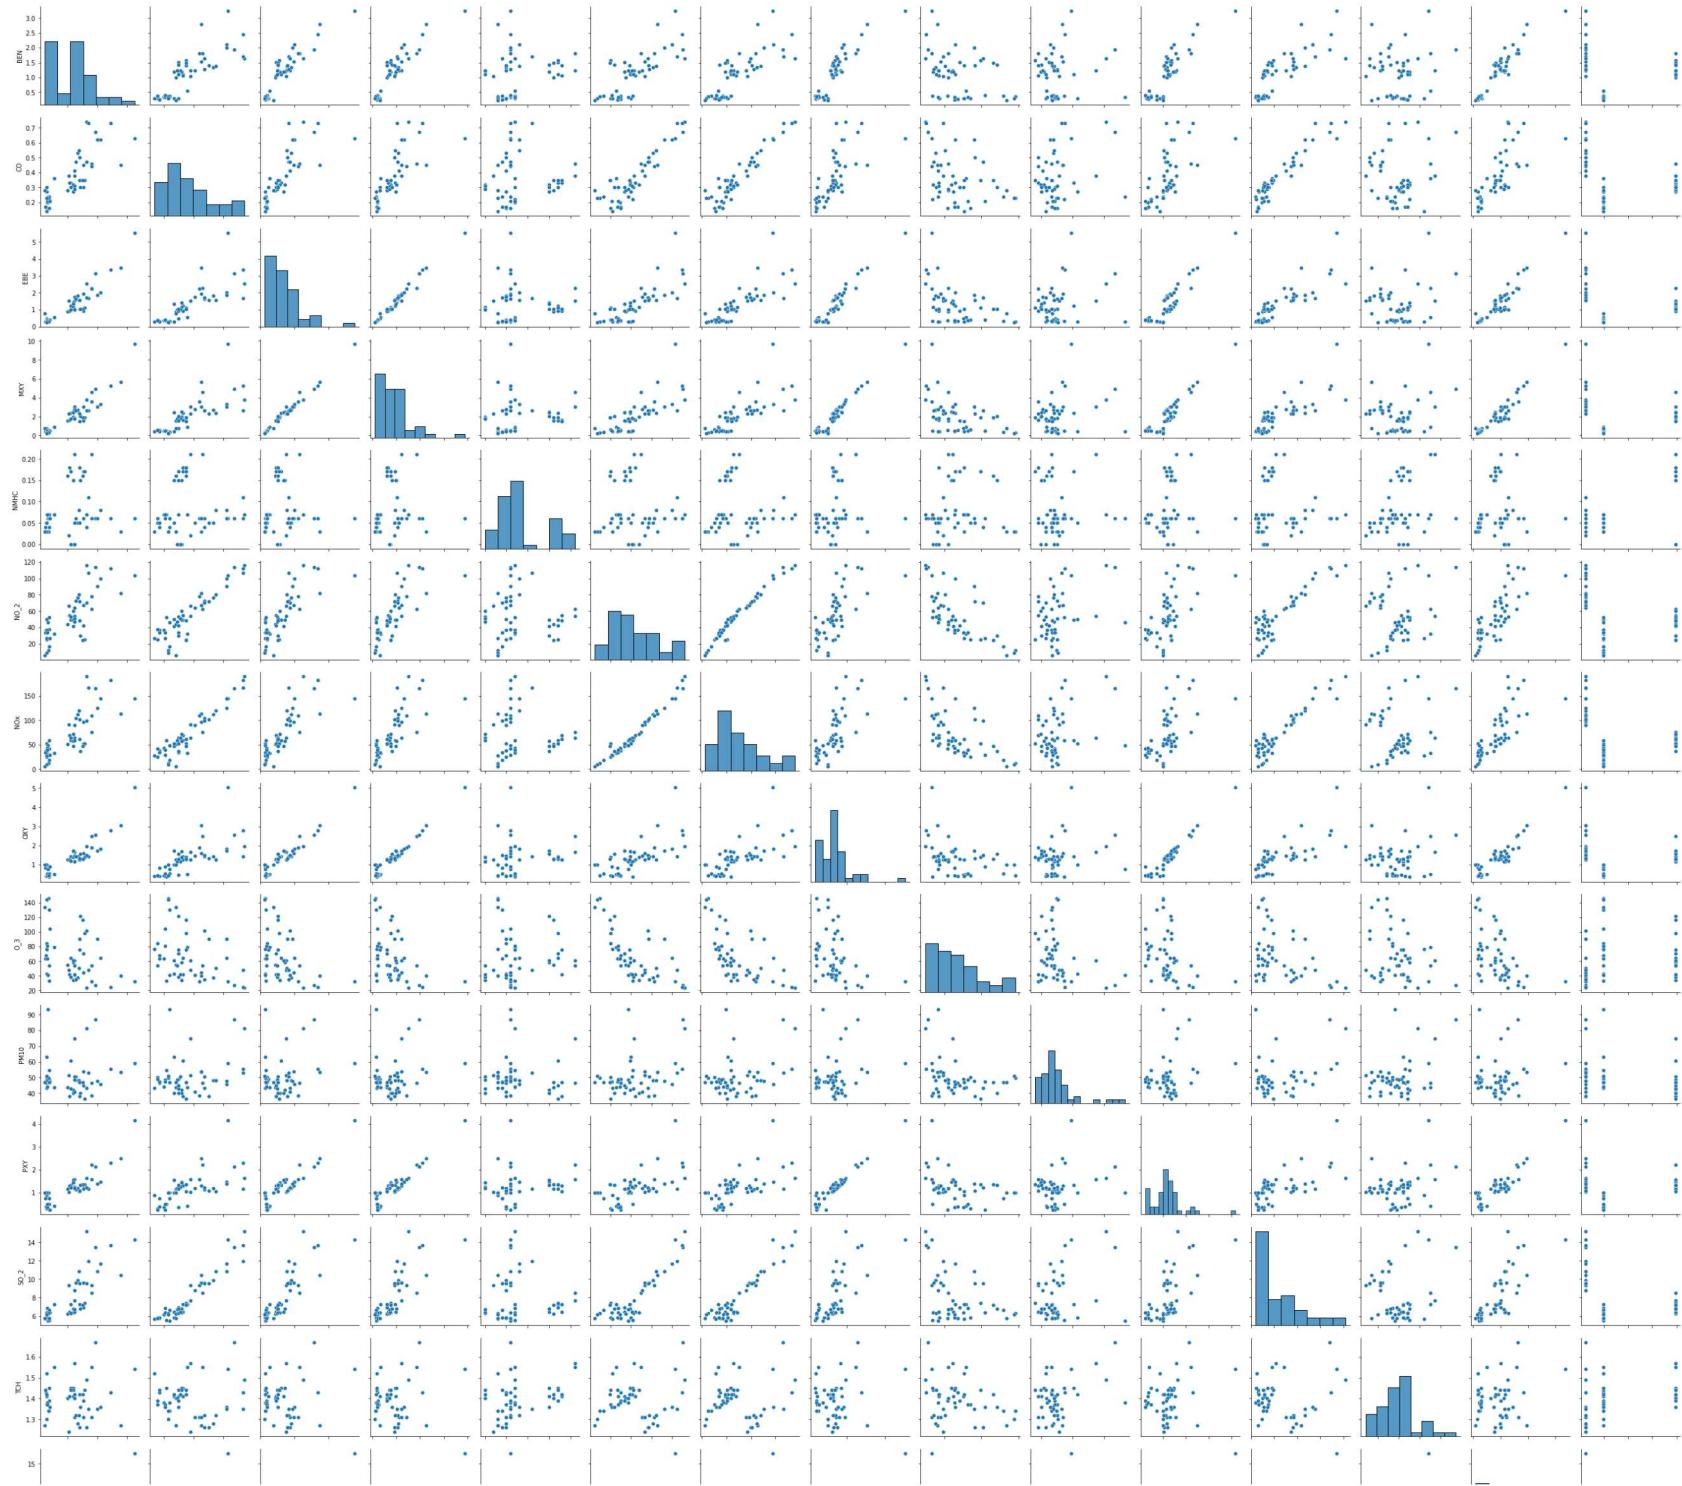
```
1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2         'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

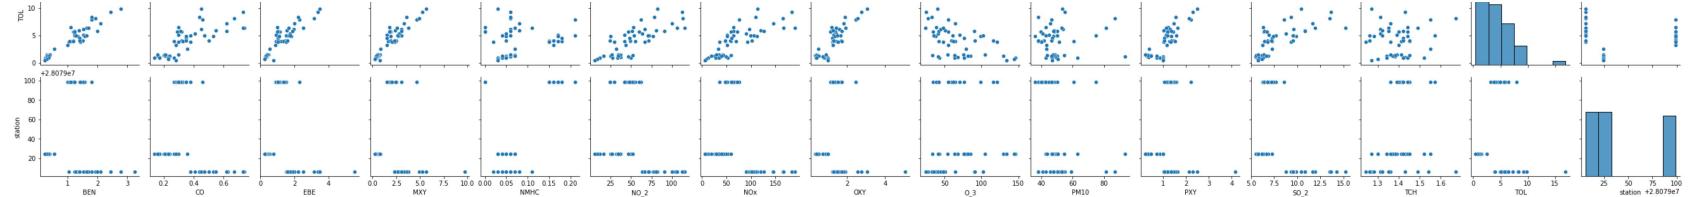
## EDA AND VISUALIZATION

```
In [19]: 1 sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x1e04b4fd8b0>
```



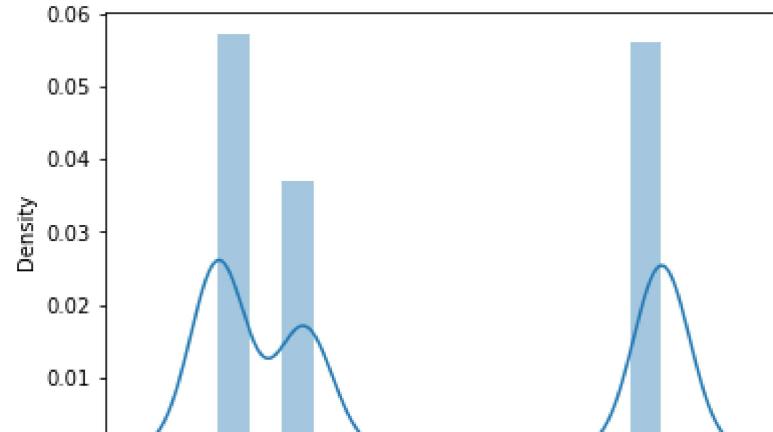




```
In [20]: 1 sns.distplot(df1['station'])
```

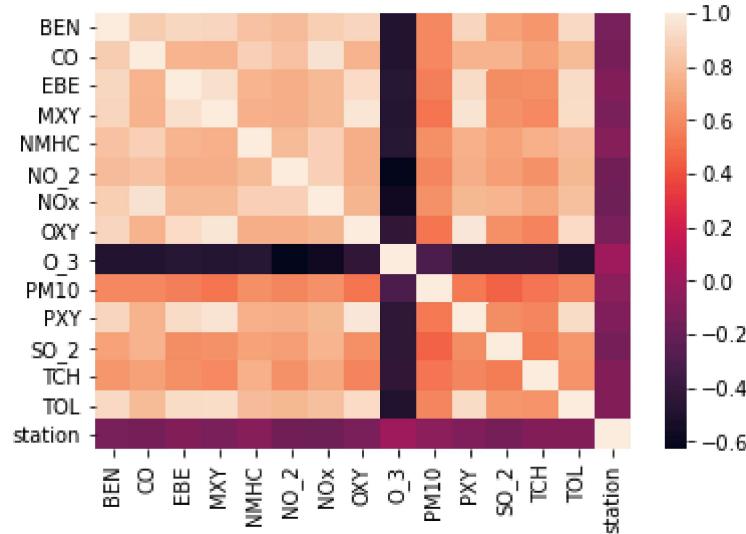
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: 1 sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
```

```
In [23]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

```
In [24]: 1 from sklearn.linear_model import LinearRegression  
2 lr=LinearRegression()  
3 lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: 1 lr.intercept_
```

```
Out[25]: 28079076.59600632
```

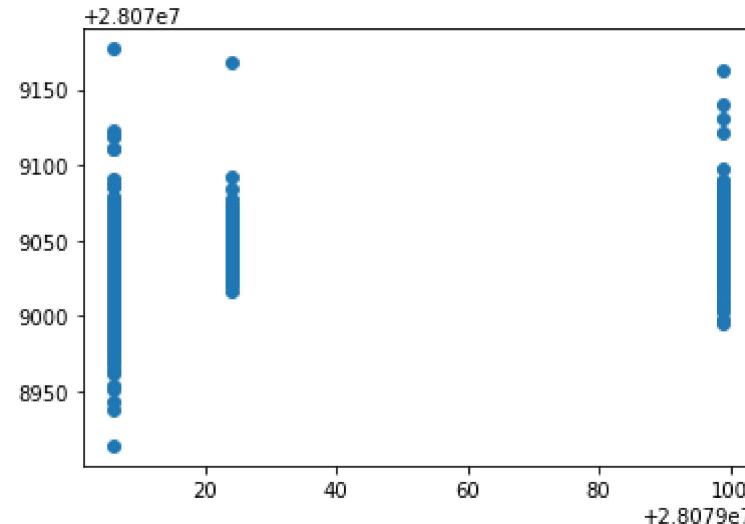
```
In [26]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
2 coeff
```

```
Out[26]:
```

Co-efficient	
BEN	-4.773847
CO	27.162786
EBE	4.108199
MXY	-3.078540
NMHC	78.272374
NO_2	-0.168998
NOx	-0.246871
OXY	-2.892686
O_3	-0.289682
PM10	0.068920
PXY	5.847097
SO_2	-0.132592
TCH	-7.286821
TOL	1.197092

```
In [27]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x1e059fc1190>
```



## ACCURACY

```
In [28]: 1 lr.score(x_test,y_test)
```

```
Out[28]: 0.1112746228512006
```

```
In [29]: 1 lr.score(x_train,y_train)
```

```
Out[29]: 0.10390101399679486
```

## Ridge and Lasso

```
In [30]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: 1 rr.score(x_test,y_test)
```

```
Out[32]: 0.10899529791144424
```

```
In [33]: 1 rr.score(x_train,y_train)
```

```
Out[33]: 0.10358897784064869
```

```
In [34]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

## Accuracy(Lasso)

```
In [35]: 1 la.score(x_train,y_train)
```

```
Out[35]: 0.054967136004089534
```

```
In [36]: 1 la.score(x_test,y_test)
```

```
Out[36]: 0.05002837680426531
```

```
In [37]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: 1 en.coef_
```

```
Out[38]: array([-0.28915406,  0.36158915,  1.39965368, -1.68984222,  0. ,
 -0.18932432, -0.08633701, -0.06288545, -0.22891864,  0.09380176,
 0.3580717 , -0.05759837,  0.         ,  1.1580202 ])
```

```
In [39]: 1 en.intercept_
```

```
Out[39]: 28079067.816384424
```

```
In [40]: 1 prediction=en.predict(x_test)
```

```
In [41]: 1 en.score(x_test,y_test)
```

```
Out[41]: 0.06584479039848501
```

## Evaluation Metrics

```
In [42]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
38.56286219795878
1660.9196913918136
40.75438248080584
```

## Logistic Regression

```
In [43]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [44]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df['station']
```

```
In [45]: 1 feature_matrix.shape
```

```
Out[45]: (19397, 14)
```

```
In [46]: 1 target_vector.shape
```

```
Out[46]: (19397,)
```

```
In [47]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [48]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: 1 logr=LogisticRegression(max_iter=10000)  
2 logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: 1 prediction=logr.predict(observation)  
2 print(prediction)
```

```
[28079006]
```

```
In [52]: 1 logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [53]: 1 logr.score(fs,target_vector)
```

```
Out[53]: 0.7360416559261741
```

```
In [54]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[54]: 0.9999978255573396
```

```
In [55]: 1 logr.predict_proba(observation)
```

```
Out[55]: array([[9.99997826e-01, 7.75018107e-20, 2.17444266e-06]])
```

# Random Forest

```
In [56]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: 1 rfc=RandomForestClassifier()
2 rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: 1 parameters={'max_depth':[1,2,3,4,5],
2                 'min_samples_leaf':[5,10,15,20,25],
3                 'n_estimators':[10,20,30,40,50]
4 }
```

```
In [59]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')
```

```
In [60]: 1 grid_search.best_score_
```

```
Out[60]: 0.7751348675493556
```

```
In [61]: 1 rfc_best=grid_search.best_estimator_
```

In [62]:

```
1 from sklearn.tree import plot_tree
2
3 plt.figure(figsize=(80,40))
4 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

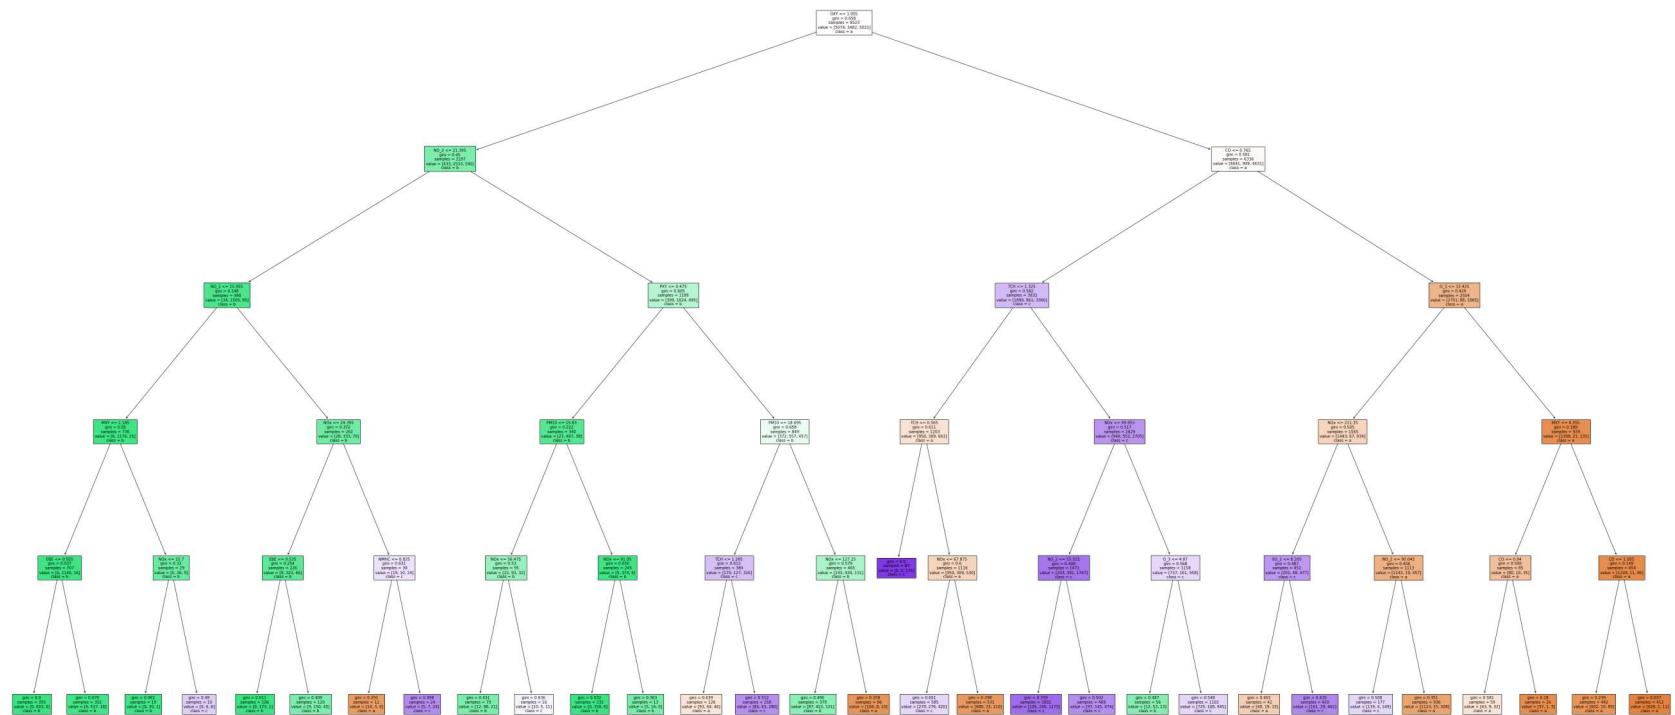
```
Out[62]: [Text(2241.3, 1993.2, 'OXY <= 1.005\ngini = 0.658\nsamples = 8523\nvalue = [5074, 3482, 5021]\nnclass = a'),  
Text(1190.4, 1630.800000000002, 'NO_2 <= 21.395\ngini = 0.45\nsamples = 2187\nvalue = [433, 2533, 590]\nnclass = b'),  
Text(595.2, 1268.4, 'NO_2 <= 15.955\ngini = 0.148\nsamples = 998\nvalue = [34, 1509, 95]\nnclass = b'),  
Text(297.6, 906.0, 'MXY <= 1.185\ngini = 0.05\nsamples = 736\nvalue = [6, 1176, 25]\nnclass = b'),  
Text(148.8, 543.599999999999, 'EBE <= 0.505\ngini = 0.037\nsamples = 707\nvalue = [6, 1140, 16]\nnclass = b'),  
Text(74.4, 181.1999999999982, 'gini = 0.0\nsamples = 355\nvalue = [0, 603, 0]\nnclass = b'),  
Text(223.200000000002, 181.1999999999982, 'gini = 0.076\nsamples = 352\nvalue = [6, 537, 16]\nnclass = b'),  
Text(446.400000000003, 543.599999999999, 'NOx <= 15.7\ngini = 0.32\nsamples = 29\nvalue = [0, 36, 9]\nnclass = b'),  
Text(372.0, 181.1999999999982, 'gini = 0.062\nsamples = 19\nvalue = [0, 30, 1]\nnclass = b'),  
Text(520.800000000001, 181.1999999999982, 'gini = 0.49\nsamples = 10\nvalue = [0, 6, 8]\nnclass = c'),  
Text(892.800000000001, 906.0, 'NOx <= 24.395\ngini = 0.372\nsamples = 262\nvalue = [28, 333, 70]\nnclass = b'),  
Text(744.0, 543.599999999999, 'EBE <= 0.535\ngini = 0.254\nsamples = 226\nvalue = [9, 323, 46]\nnclass = b'),  
Text(669.6, 181.1999999999982, 'gini = 0.011\nsamples = 106\nvalue = [0, 173, 1]\nnclass = b'),  
Text(818.400000000001, 181.1999999999982, 'gini = 0.409\nsamples = 120\nvalue = [9, 150, 45]\nnclass = b'),  
Text(1041.600000000001, 543.599999999999, 'NMHC <= 0.025\ngini = 0.631\nsamples = 36\nvalue = [19, 10, 24]\nnclass = c'),  
Text(967.2, 181.1999999999982, 'gini = 0.291\nsamples = 12\nvalue = [14, 3, 0]\nnclass = a'),  
Text(1116.0, 181.1999999999982, 'gini = 0.498\nsamples = 24\nvalue = [5, 7, 24]\nnclass = c'),  
Text(1785.600000000001, 1268.4, 'PXY <= 0.475\ngini = 0.605\nsamples = 1189\nvalue = [399, 1024, 495]\nnclass = b'),  
Text(1488.0, 906.0, 'PM10 <= 15.83\ngini = 0.222\nsamples = 340\nvalue = [27, 467, 38]\nnclass = b'),  
Text(1339.2, 543.599999999999, 'NOx <= 56.475\ngini = 0.53\nsamples = 95\nvalue = [22, 93, 32]\nnclass = b'),  
Text(1264.800000000002, 181.1999999999982, 'gini = 0.431\nsamples = 79\nvalue = [12, 88, 21]\nnclass = b'),  
Text(1413.600000000001, 181.1999999999982, 'gini = 0.636\nsamples = 16\nvalue = [10, 5, 11]\nnclass = c'),  
Text(1636.800000000002, 543.599999999999, 'NOx <= 91.05\ngini = 0.056\nsamples = 245\nvalue = [5, 374, 6]\nnclass = b'),  
Text(1562.4, 181.1999999999982, 'gini = 0.032\nsamples = 232\nvalue = [0, 358, 6]\nnclass = b'),  
Text(1711.2, 181.1999999999982, 'gini = 0.363\nsamples = 13\nvalue = [5, 16, 0]\nnclass = b'),  
Text(2083.200000000003, 906.0, 'PM10 <= 18.695\ngini = 0.658\nsamples = 849\nvalue = [372, 557, 457]\nnclass = b'),  
Text(1934.4, 543.599999999999, 'TCH <= 1.285\ngini = 0.613\nsamples = 384\nvalue = [179, 127, 326]\nnclass = c'),  
Text(1860.000000000002, 181.1999999999982, 'gini = 0.639\nsamples = 126\nvalue = [93, 64, 46]\nnclass = a'),
```

```
Text(2008.800000000002, 181.19999999999982, 'gini = 0.512\nsamples = 258\nvalue = [86, 63, 280]\nclass = c'),
Text(2232.0, 543.599999999999, 'NOx <= 127.25\ngini = 0.579\nsamples = 465\nvalue = [193, 430, 131]\nclass = b'),
Text(2157.600000000004, 181.19999999999982, 'gini = 0.495\nsamples = 379\nvalue = [87, 422, 121]\nclass = b'),
Text(2306.4, 181.19999999999982, 'gini = 0.259\nsamples = 86\nvalue = [106, 8, 10]\nclass = a'),
Text(3292.200000000003, 1630.800000000002, 'CO <= 0.765\ngini = 0.581\nsamples = 6336\nvalue = [4641, 949, 4431]\nclass = a'),
Text(2715.600000000004, 1268.4, 'TCH <= 1.325\ngini = 0.582\nsamples = 3832\nvalue = [1890, 861, 3366]\nclass = c'),
Text(2455.200000000003, 906.0, 'TCH <= 0.565\ngini = 0.611\nsamples = 1203\nvalue = [950, 309, 661]\nclass = a'),
Text(2380.8, 543.599999999999, 'gini = 0.0\nsamples = 87\nvalue = [0, 0, 131]\nclass = c'),
Text(2529.600000000004, 543.599999999999, 'NOx <= 67.875\ngini = 0.6\nsamples = 1116\nvalue = [950, 309, 530]\nclass = a'),
Text(2455.200000000003, 181.19999999999982, 'gini = 0.651\nsamples = 585\nvalue = [270, 276, 420]\nclass = c'),
Text(2604.0, 181.19999999999982, 'gini = 0.298\nsamples = 531\nvalue = [680, 33, 110]\nclass = a'),
Text(2976.0, 906.0, 'NOx <= 99.955\ngini = 0.517\nsamples = 2629\nvalue = [940, 552, 2705]\nclass = c'),
Text(2827.200000000003, 543.599999999999, 'NO_2 <= 55.025\ngini = 0.408\nsamples = 1471\nvalue = [203, 391, 1747]\nclass = c'),
Text(2752.8, 181.19999999999982, 'gini = 0.359\nsamples = 1002\nvalue = [106, 246, 1273]\nclass = c'),
Text(2901.600000000004, 181.19999999999982, 'gini = 0.502\nsamples = 469\nvalue = [97, 145, 474]\nclass = c'),
Text(3124.8, 543.599999999999, 'O_3 <= 4.87\ngini = 0.568\nsamples = 1158\nvalue = [737, 161, 958]\nclass = c'),
Text(3050.4, 181.19999999999982, 'gini = 0.487\nsamples = 56\nvalue = [12, 53, 13]\nclass = b'),
Text(3199.200000000003, 181.19999999999982, 'gini = 0.548\nsamples = 1102\nvalue = [725, 108, 945]\nclass = c'),
Text(3868.8, 1268.4, 'O_3 <= 13.425\ngini = 0.429\nsamples = 2504\nvalue = [2751, 88, 1065]\nclass = a'),
Text(3571.200000000003, 906.0, 'NOx <= 211.35\ngini = 0.505\nsamples = 1565\nvalue = [1443, 67, 934]\nclass = a'),
Text(3422.4, 543.599999999999, 'SO_2 <= 8.205\ngini = 0.487\nsamples = 452\nvalue = [201, 48, 477]\nclass = c'),
Text(3348.000000000005, 181.19999999999982, 'gini = 0.601\nsamples = 42\nvalue = [40, 19, 15]\nclass = a'),
Text(3496.8, 181.19999999999982, 'gini = 0.435\nsamples = 410\nvalue = [161, 29, 462]\nclass = c'),
Text(3720.000000000005, 543.599999999999, 'NO_2 <= 90.045\ngini = 0.406\nsamples = 1113\nvalue = [1242, 19, 457]\nclass = a'),
Text(3645.600000000004, 181.19999999999982, 'gini = 0.508\nsamples = 177\nvalue = [119, 4, 149]\nclass = c'),
Text(3794.4, 181.19999999999982, 'gini = 0.351\nsamples = 936\nvalue = [1123, 15, 308]\nclass = a'),
```

```

Text(4166.400000000001, 906.0, 'MXY <= 4.355\ngini = 0.189\nsamples = 939\nvalue = [1308, 21, 131]\nclass = a'),
Text(4017.600000000004, 543.5999999999999, 'CO <= 0.94\ngini = 0.506\nsamples = 85\nvalue = [80, 10, 35]\nclass = a'),
Text(3943.200000000003, 181.1999999999982, 'gini = 0.581\nsamples = 59\nvalue = [43, 9, 32]\nclass = a'),
Text(4092.000000000005, 181.1999999999982, 'gini = 0.18\nsamples = 26\nvalue = [37, 1, 3]\nclass = a'),
Text(4315.200000000001, 543.5999999999999, 'CO <= 1.005\ngini = 0.149\nsamples = 854\nvalue = [1228, 11, 96]\nclass = a'),
Text(4240.8, 181.1999999999982, 'gini = 0.239\nsamples = 442\nvalue = [602, 10, 85]\nclass = a'),
Text(4389.6, 181.1999999999982, 'gini = 0.037\nsamples = 412\nvalue = [626, 1, 11]\nclass = a')]

```



## Conclusion

## Accuracy

*Linear Regression:* 0.10390101399679486

*Ridge Regression:* 0.10358897784064869

*Lasso Regression:* 0.054967136004089534

*ElasticNet Regression:* 0.06584479039848501

*Logistic Regression:* 0.7360416559261741

*Random Forest:* 0.7751348675493556

**Random Forest is suitable for this dataset**

In [ ]: 1