

## Importing Libraries

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

## Importing Datasets

In [2]:

```

1 df=pd.read_csv("madrid_2005")
2 df

```

Out[2]:

		date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2	TCH	TOL	static
0		2005-11-01 01:00:00	NaN	0.77	NaN	NaN	NaN	57.130001	128.699997	NaN	14.720000	14.91	10.65	NaN	4.62	NaN	NaN	2807900
1		2005-11-01 01:00:00	1.52	0.65	1.49	4.57	0.25	86.559998	181.699997	1.27	11.680000	30.93	NaN	1.59	7.80	1.35	7.98	2807900
2		2005-11-01 01:00:00	NaN	0.40	NaN	NaN	NaN	46.119999	53.000000	NaN	30.469999	14.60	NaN	NaN	5.76	NaN	NaN	2807900
3		2005-11-01 01:00:00	NaN	0.42	NaN	NaN	NaN	37.220001	52.009998	NaN	21.379999	15.16	NaN	NaN	6.60	NaN	NaN	2807900
4		2005-11-01 01:00:00	NaN	0.57	NaN	NaN	NaN	32.160000	36.680000	NaN	33.410000	5.00	NaN	NaN	3.00	NaN	NaN	2807900
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	
236995		2006-01-01 00:00:00	1.08	0.36	1.01	NaN	0.11	21.990000	23.610001	NaN	43.349998	5.00	NaN	NaN	6.68	1.37	2.75	2807902
236996		2006-01-01 00:00:00	0.39	0.54	1.00	1.00	0.11	2.200000	4.220000	1.00	69.639999	4.95	1.49	1.00	7.06	1.28	0.38	2807902
236997		2006-01-01 00:00:00	0.19	NaN	0.26	NaN	0.08	26.730000	30.809999	NaN	43.840000	4.31	2.93	NaN	13.20	1.28	0.56	2807902
236998		2006-01-01 00:00:00	0.14	NaN	1.00	NaN	0.06	13.770000	17.770000	NaN	NaN	5.00	NaN	NaN	5.81	1.25	0.22	2807902
236999		2006-01-01 00:00:00	0.50	0.40	0.73	1.84	0.13	20.940001	26.950001	1.49	48.259998	5.67	2.11	1.09	11.07	1.30	1.29	2807902

237000 rows × 17 columns



# Data Cleaning and Data Preprocessing

```
In [3]: 1 df=df.dropna()
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
      dtype='object')
```

```
In [5]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20070 entries, 5 to 236999
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      20070 non-null   object 
 1   BEN        20070 non-null   float64
 2   CO         20070 non-null   float64
 3   EBE        20070 non-null   float64
 4   MXY        20070 non-null   float64
 5   NMHC       20070 non-null   float64
 6   NO_2       20070 non-null   float64
 7   NOx        20070 non-null   float64
 8   OXY        20070 non-null   float64
 9   O_3         20070 non-null   float64
 10  PM10       20070 non-null   float64
 11  PM25       20070 non-null   float64
 12  PXY        20070 non-null   float64
 13  SO_2       20070 non-null   float64
 14  TCH        20070 non-null   float64
 15  TOL        20070 non-null   float64
 16  station    20070 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 2.8+ MB
```

```
In [6]: 1 data=df[['CO' , 'station']]  
2 data
```

Out[6]:

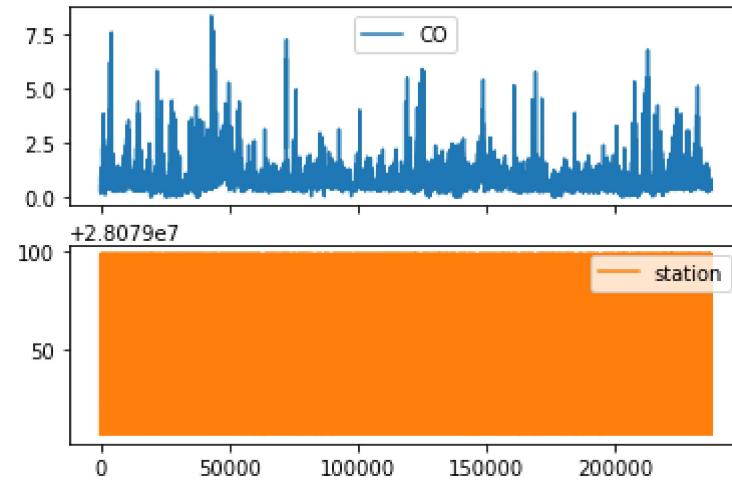
	CO	station
5	0.88	28079006
22	0.22	28079024
25	0.49	28079099
31	0.84	28079006
48	0.20	28079024
...	...	...
236970	0.39	28079024
236973	0.45	28079099
236979	0.38	28079006
236996	0.54	28079024
236999	0.40	28079099

20070 rows × 2 columns

## Line chart

```
In [7]: 1 data.plot.line(subplots=True)
```

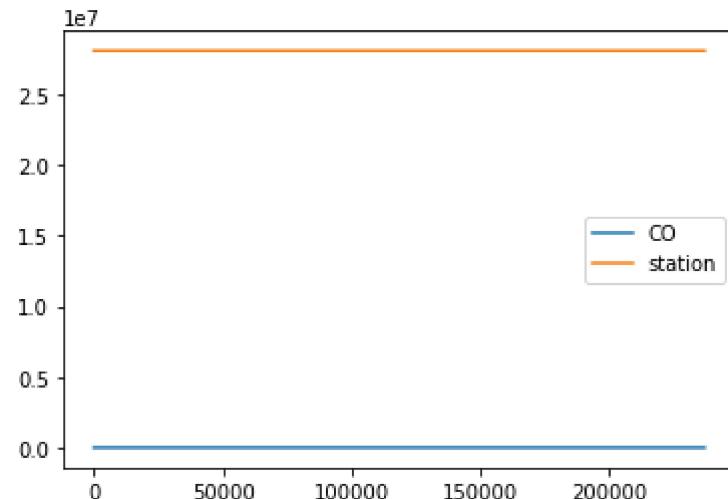
```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



## Line chart

```
In [8]: 1 data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

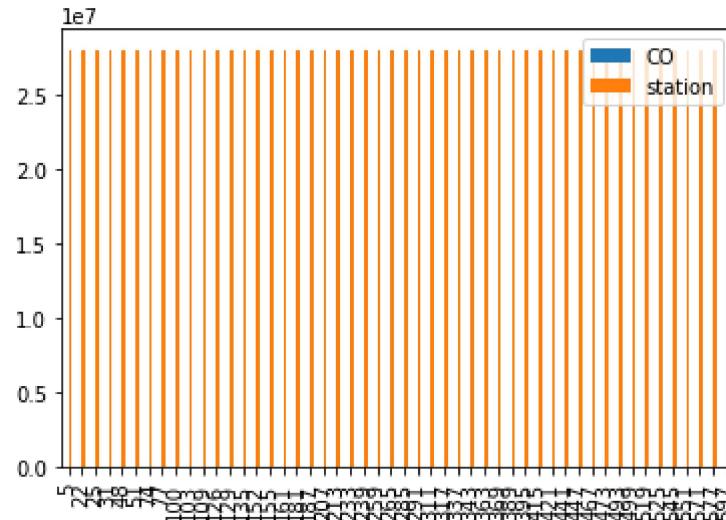


## Bar chart

```
In [9]: 1 b=data[0:50]
```

```
In [10]: 1 b.plot.bar()
```

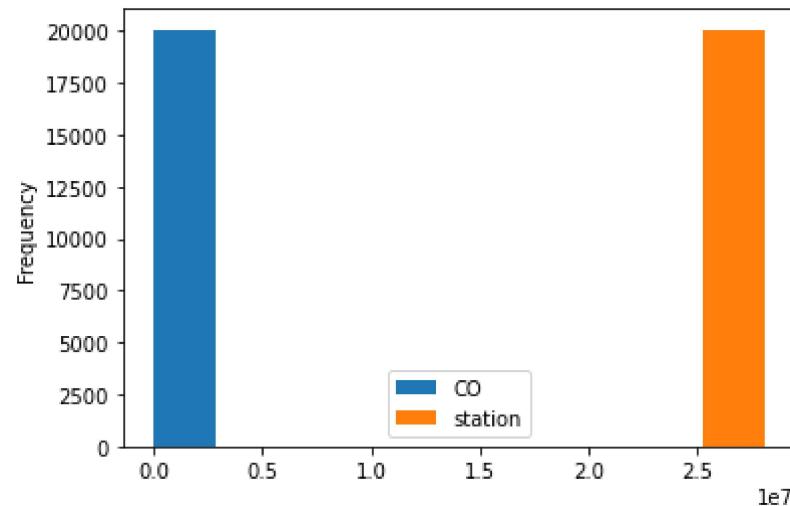
```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: 1 data.plot.hist()
```

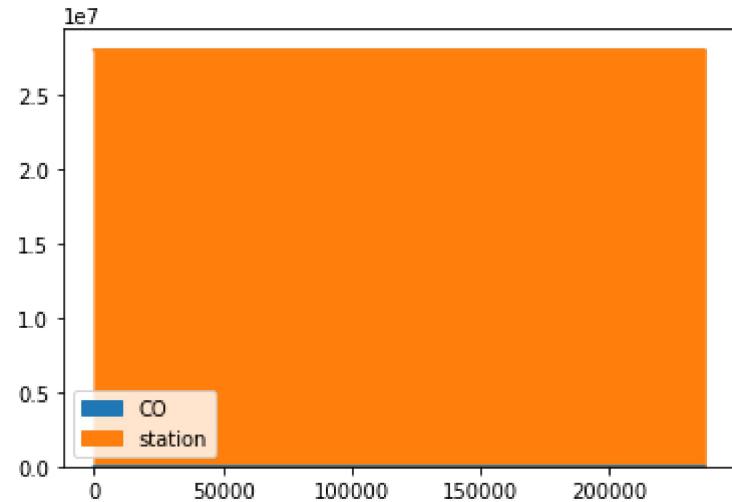
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

In [12]: 1 data.plot.area()

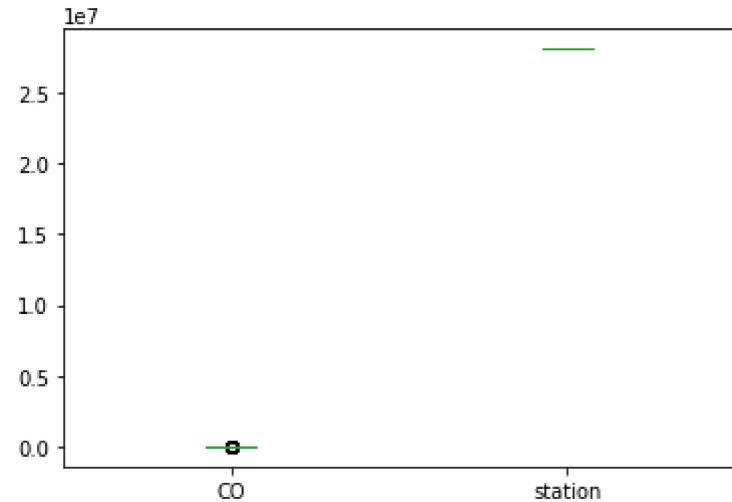
Out[12]: <AxesSubplot:>



## Box chart

In [13]: 1 data.plot.box()

Out[13]: <AxesSubplot:>

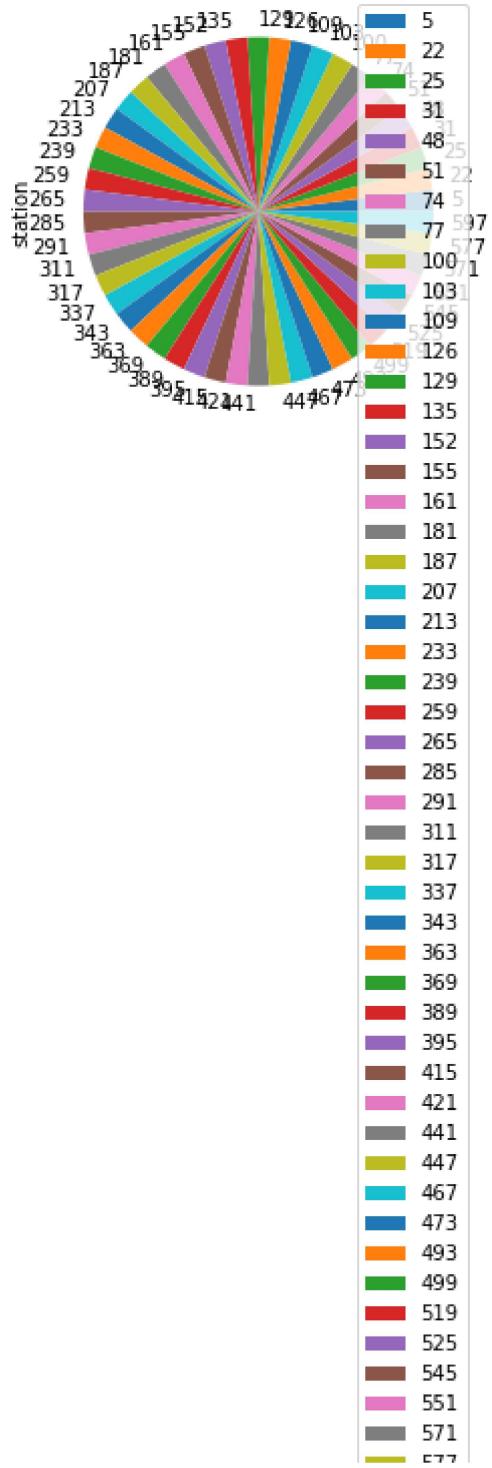


## Pie chart

```
In [14]: 1 b.plot.pie(y='station' )
```

```
Out[14]: <AxesSubplot:ylabel='station'>
```



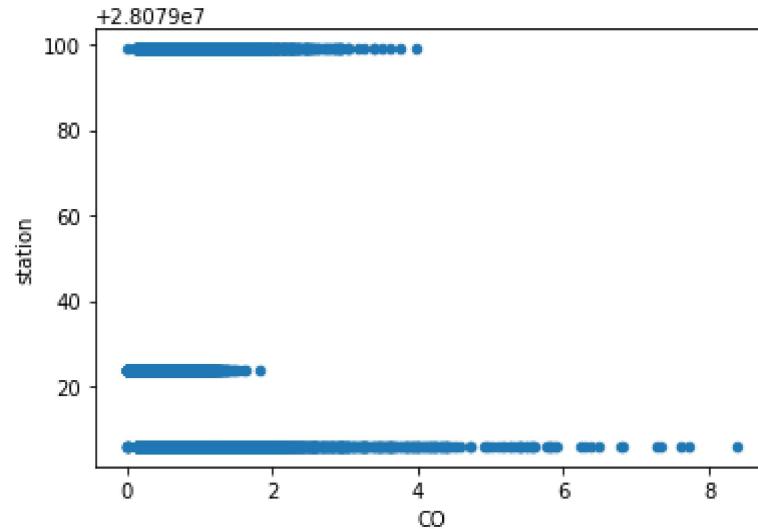




## Scatter chart

In [15]: 1 data.plot.scatter(x='CO' ,y='station')

Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>



```
In [16]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 20070 entries, 5 to 236999
Data columns (total 17 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   date      20070 non-null   object  
 1   BEN        20070 non-null   float64 
 2   CO         20070 non-null   float64 
 3   EBE        20070 non-null   float64 
 4   MXY        20070 non-null   float64 
 5   NMHC       20070 non-null   float64 
 6   NO_2       20070 non-null   float64 
 7   NOx        20070 non-null   float64 
 8   OXY        20070 non-null   float64 
 9   O_3         20070 non-null   float64 
 10  PM10       20070 non-null   float64 
 11  PM25       20070 non-null   float64 
 12  PXY        20070 non-null   float64 
 13  SO_2       20070 non-null   float64 
 14  TCO        20070 non-null   float64
```

```
In [17]: 1 df.describe()
```

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
count	20070.000000	20070.000000	20070.000000	20070.000000	20070.000000	20070.000000	20070.000000	20070.000000	20070.000000
mean	1.923656	0.720657	2.345423	5.457855	0.179282	66.226924	143.046536	2.774935	32.121323
std	2.019061	0.549723	2.379219	5.495147	0.152783	40.568197	136.582521	2.705508	23.954714
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.690000	0.400000	0.950000	1.930000	0.090000	36.602499	56.102499	1.000000	11.540000
50%	1.260000	0.580000	1.480000	3.800000	0.150000	60.525000	105.699997	1.890000	26.255000
75%	2.510000	0.880000	2.950000	7.210000	0.220000	89.317499	190.100006	3.620000	47.150002
max	26.570000	8.380000	29.870001	71.050003	1.880000	419.500000	1774.000000	38.680000	149.800003

In [18]:

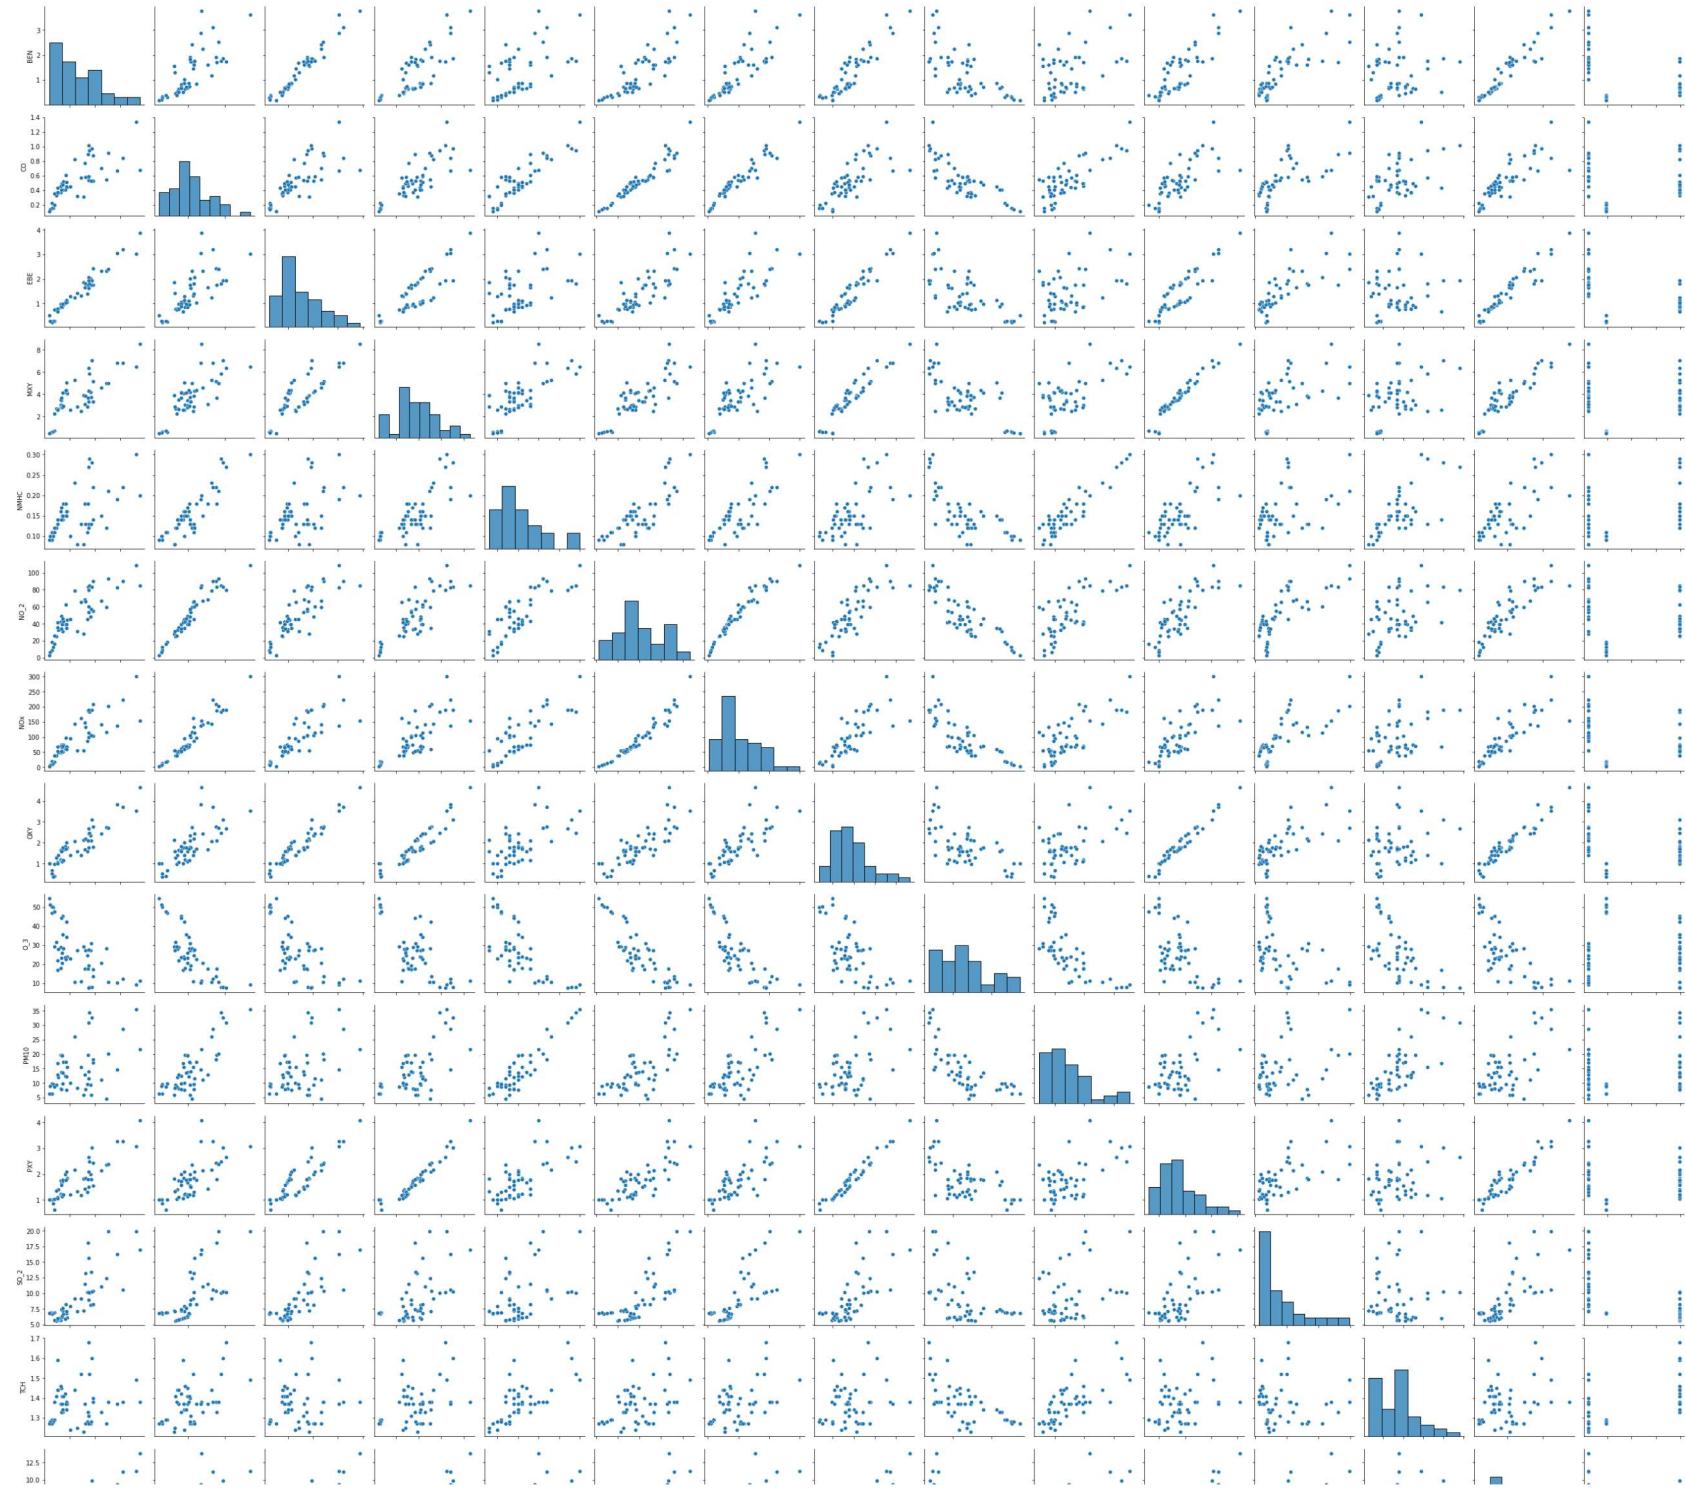
```
1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2         'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

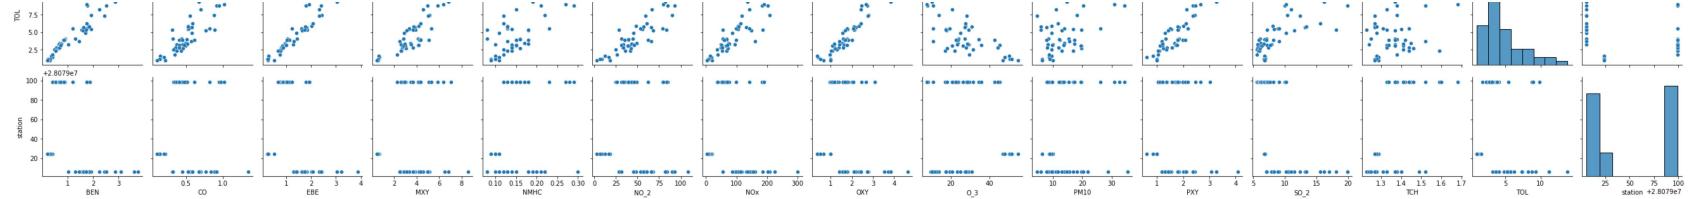
## EDA AND VISUALIZATION

```
In [19]: 1 sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x2338ba80af0>
```



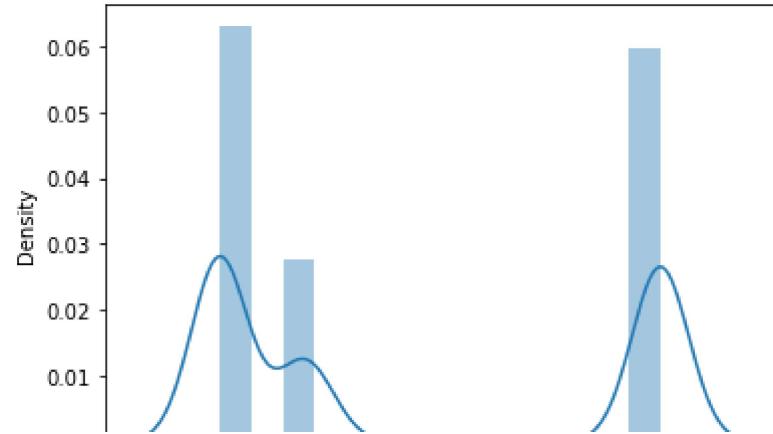




```
In [20]: 1 sns.distplot(df1['station'])
```

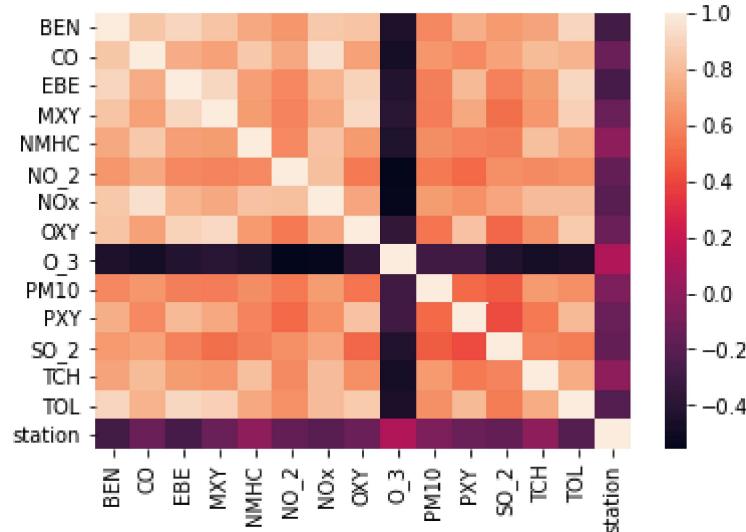
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: 1 sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
```

```
In [23]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

```
In [24]: 1 from sklearn.linear_model import LinearRegression  
2 lr=LinearRegression()  
3 lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: 1 lr.intercept_
```

```
Out[25]: 28078959.43807614
```

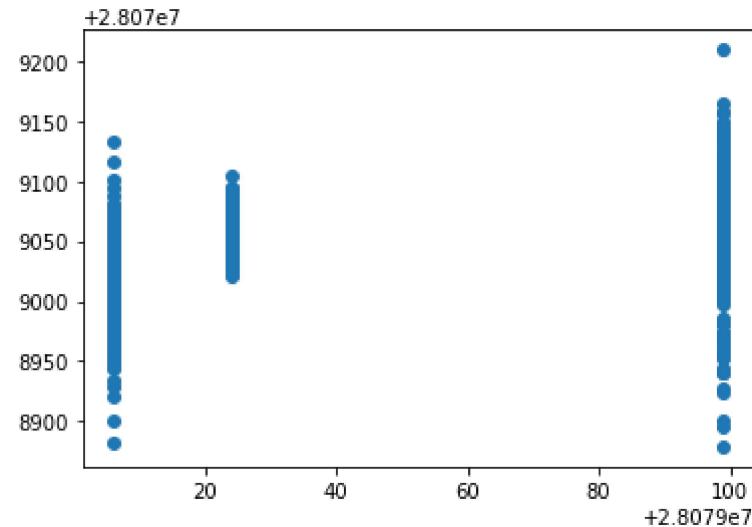
```
In [26]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
2 coeff
```

```
Out[26]:
```

	Co-efficient
BEN	-8.957694
CO	38.918415
EBE	-13.198253
MXY	3.642700
NMHC	77.780575
NO_2	0.125812
NOx	-0.274325
OXY	3.167298
O_3	0.008482
PM10	0.048127
PXY	2.543436
SO_2	0.180207
TCH	62.510949
TOL	-0.544119

```
In [27]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x2339a38f580>
```



## ACCURACY

```
In [28]: 1 lr.score(x_test,y_test)
```

```
Out[28]: 0.31417673670713897
```

```
In [29]: 1 lr.score(x_train,y_train)
```

```
Out[29]: 0.2998431212306839
```

## Ridge and Lasso

```
In [30]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: 1 rr.score(x_test,y_test)
```

```
Out[32]: 0.31355614379783414
```

```
In [33]: 1 rr.score(x_train,y_train)
```

```
Out[33]: 0.2996231647869222
```

```
In [34]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

## Accuracy(Lasso)

```
In [35]: 1 la.score(x_train,y_train)
```

```
Out[35]: 0.062201095160138054
```

```
In [36]: 1 la.score(x_test,y_test)
```

```
Out[36]: 0.0691094813594253
```

```
In [37]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: 1 en.coef_
```

```
Out[38]: array([-5.44544857,  1.50574277, -7.32302084,  2.61297051,  0.92964019,
 -0.0573147 , -0.00983528,  1.89670049, -0.02333062,  0.22792507,
 1.40928848,  0.13714068,  1.59571901, -0.77222586])
```

```
In [39]: 1 en.intercept_
```

```
Out[39]: 28079050.10880432
```

```
In [40]: 1 prediction=en.predict(x_test)
```

```
In [41]: 1 en.score(x_test,y_test)
```

```
Out[41]: 0.18421578404776207
```

## Evaluation Metrics

```
In [42]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
36.82616234889504
1527.3078091416137
39.080785677128006
```

## Logistic Regression

```
In [43]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [44]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df['station']
```

```
In [45]: 1 feature_matrix.shape
```

```
Out[45]: (20070, 14)
```

```
In [46]: 1 target_vector.shape
```

```
Out[46]: (20070,)
```

```
In [47]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [48]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: 1 logr=LogisticRegression(max_iter=10000)  
2 logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: 1 prediction=logr.predict(observation)  
2 print(prediction)
```

```
[28079006]
```

```
In [52]: 1 logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [53]: 1 logr.score(fs,target_vector)
```

```
Out[53]: 0.879023418036871
```

```
In [54]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[54]: 0.9998967601812779
```

```
In [55]: 1 logr.predict_proba(observation)
```

```
Out[55]: array([[9.99896760e-01, 3.21124597e-30, 1.03239819e-04]])
```

# Random Forest

```
In [56]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: 1 rfc=RandomForestClassifier()
2 rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: 1 parameters={'max_depth':[1,2,3,4,5],
2                 'min_samples_leaf':[5,10,15,20,25],
3                 'n_estimators':[10,20,30,40,50]
4 }
```

```
In [59]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')
```

```
In [60]: 1 grid_search.best_score_
```

```
Out[60]: 0.8659692442383612
```

```
In [61]: 1 rfc_best=grid_search.best_estimator_
```

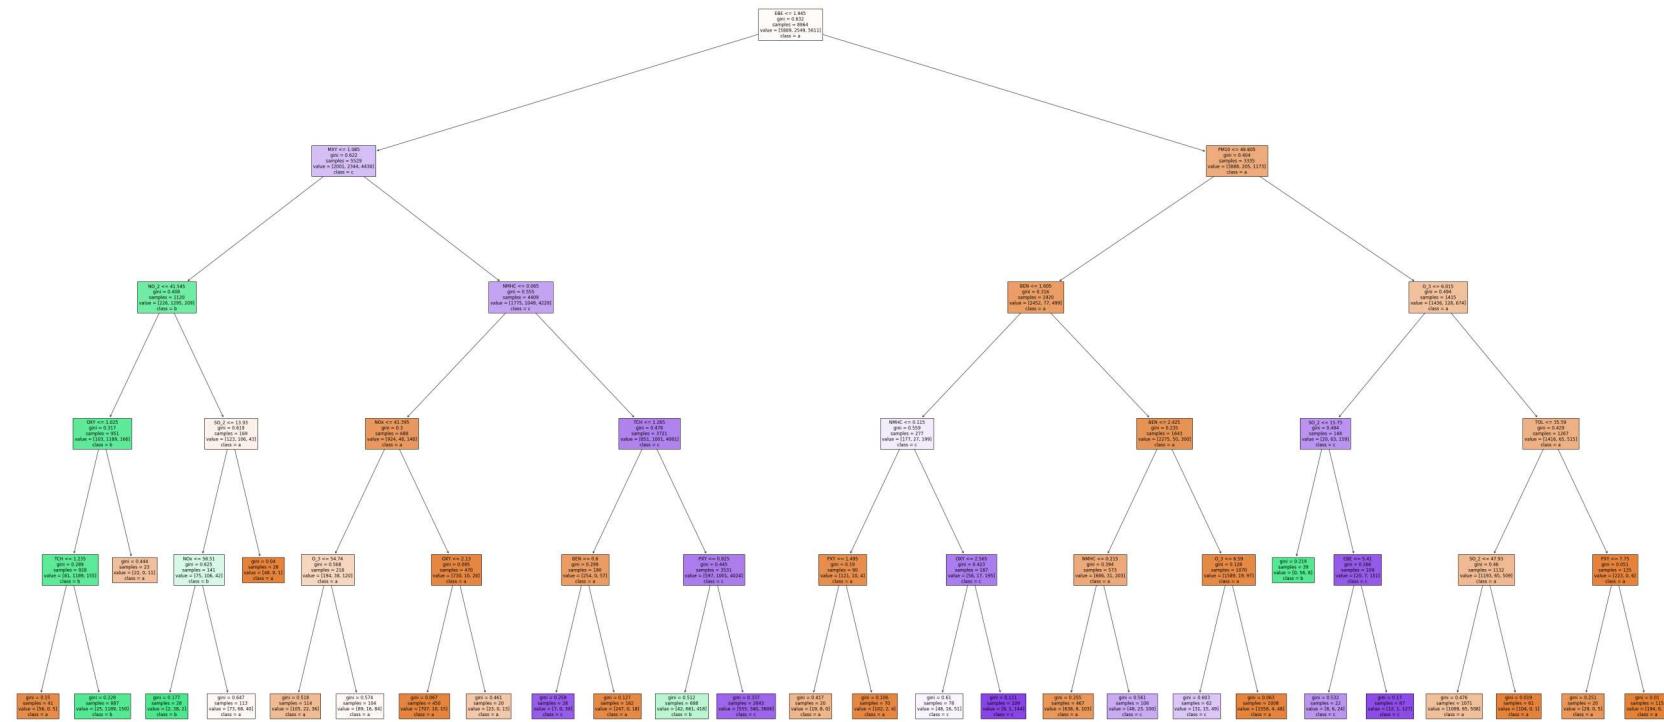
In [62]:

```
1 from sklearn.tree import plot_tree
2
3 plt.figure(figsize=(80,40))
4 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

```
Out[62]: [Text(2092.5, 1993.2, 'EBE <= 1.945\ngini = 0.632\nsamples = 8864\nvalue = [5889, 2549, 5611]\nclass = a'),  
Text(901.3846153846154, 1630.800000000002, 'MXY <= 1.085\ngini = 0.622\nsamples = 5529\nvalue = [2001, 234  
4, 4438]\nclass = c'),  
Text(429.23076923076917, 1268.4, 'NO_2 <= 41.545\ngini = 0.408\nsamples = 1120\nvalue = [226, 1295, 209]\nnc  
lass = b'),  
Text(257.53846153846155, 906.0, 'OXY <= 1.025\ngini = 0.317\nsamples = 951\nvalue = [103, 1189, 166]\nnclass  
= b'),  
Text(171.69230769230768, 543.5999999999999, 'TCH <= 1.235\ngini = 0.289\nsamples = 928\nvalue = [81, 1189,  
155]\nclass = b'),  
Text(85.84615384615384, 181.1999999999982, 'gini = 0.15\ngsamples = 41\nvalue = [56, 0, 5]\nclass = a'),  
Text(257.53846153846155, 181.1999999999982, 'gini = 0.228\nsamples = 887\nvalue = [25, 1189, 150]\nclass =  
b'),  
Text(343.38461538461536, 543.5999999999999, 'gini = 0.444\nsamples = 23\nvalue = [22, 0, 11]\nclass = a'),  
Text(600.9230769230769, 906.0, 'SO_2 <= 13.93\ngini = 0.619\nsamples = 169\nvalue = [123, 106, 43]\nclass =  
a'),  
Text(515.0769230769231, 543.5999999999999, 'NOx <= 58.51\ngini = 0.625\nsamples = 141\nvalue = [75, 106, 4  
2]\nclass = b'),  
Text(429.23076923076917, 181.1999999999982, 'gini = 0.177\nsamples = 28\nvalue = [2, 38, 2]\nclass = b'),  
Text(600.9230769230769, 181.1999999999982, 'gini = 0.647\nsamples = 113\nvalue = [73, 68, 40]\nclass =  
a'),  
Text(686.7692307692307, 543.5999999999999, 'gini = 0.04\nsamples = 28\nvalue = [48, 0, 1]\nclass = a'),  
Text(1373.5384615384614, 1268.4, 'NMHC <= 0.065\ngini = 0.555\nsamples = 4409\nvalue = [1775, 1049, 4229]\nnc  
lass = c'),  
Text(1030.1538461538462, 906.0, 'NOx <= 41.395\ngini = 0.3\nsamples = 688\nvalue = [924, 48, 148]\nclass =  
a'),  
Text(858.4615384615383, 543.5999999999999, 'O_3 <= 54.74\ngini = 0.568\nsamples = 218\nvalue = [194, 38, 12  
0]\nclass = a'),  
Text(772.6153846153845, 181.1999999999982, 'gini = 0.518\nsamples = 114\nvalue = [105, 22, 36]\nclass =  
a'),  
Text(944.3076923076923, 181.1999999999982, 'gini = 0.574\nsamples = 104\nvalue = [89, 16, 84]\nclass =  
a'),  
Text(1201.8461538461538, 543.5999999999999, 'OXY <= 2.13\ngini = 0.095\nsamples = 470\nvalue = [730, 10, 2  
8]\nclass = a'),  
Text(1116.0, 181.1999999999982, 'gini = 0.067\nsamples = 450\nvalue = [707, 10, 15]\nclass = a'),  
Text(1287.6923076923076, 181.1999999999982, 'gini = 0.461\nsamples = 20\nvalue = [23, 0, 13]\nclass = a'),  
Text(1716.9230769230767, 906.0, 'TCH <= 1.265\ngini = 0.478\nsamples = 3721\nvalue = [851, 1001, 4081]\nclass =  
c'),  
Text(1545.230769230769, 543.5999999999999, 'BEN <= 0.6\ngini = 0.299\nsamples = 190\nvalue = [254, 0, 57]\nclass = a'),  
Text(1459.3846153846152, 181.1999999999982, 'gini = 0.258\nsamples = 28\nvalue = [7, 0, 39]\nclass = c'),  
Text(1631.0769230769229, 181.1999999999982, 'gini = 0.127\nsamples = 162\nvalue = [247, 0, 18]\nclass =  
a'),  
Text(1888.6153846153845, 543.5999999999999, 'PXY <= 0.825\ngini = 0.445\nsamples = 3531\nvalue = [597, 100
```

```
1, 4024]\nclass = c'),
Text(1802.7692307692307, 181.19999999999982, 'gini = 0.512\nsamples = 688\nvalue = [42, 661, 418]\nclass =
b'),
Text(1974.4615384615383, 181.19999999999982, 'gini = 0.337\nsamples = 2843\nvalue = [555, 340, 3606]\nclass
= c'),
Text(3283.6153846153843, 1630.8000000000002, 'PM10 <= 48.805\ngini = 0.404\nsamples = 3335\nvalue = [3888,
205, 1173]\nclass = a'),
Text(2747.076923076923, 1268.4, 'BEN <= 1.605\ngini = 0.316\nsamples = 1920\nvalue = [2452, 77, 499]\nclass
= a'),
Text(2403.6923076923076, 906.0, 'NMHC <= 0.115\ngini = 0.559\nsamples = 277\nvalue = [177, 27, 199]\nclass
= c'),
Text(2232.0, 543.5999999999999, 'PXY <= 1.495\ngini = 0.19\nsamples = 90\nvalue = [121, 10, 4]\nclass =
a'),
Text(2146.153846153846, 181.19999999999982, 'gini = 0.417\nsamples = 20\nvalue = [19, 8, 0]\nclass = a'),
Text(2317.846153846154, 181.19999999999982, 'gini = 0.106\nsamples = 70\nvalue = [102, 2, 4]\nclass = a'),
Text(2575.3846153846152, 543.5999999999999, 'OXY <= 2.565\ngini = 0.423\nsamples = 187\nvalue = [56, 17, 19
5]\nclass = c'),
Text(2489.5384615384614, 181.19999999999982, 'gini = 0.61\nsamples = 78\nvalue = [48, 16, 51]\nclass = c'),
Text(2661.230769230769, 181.19999999999982, 'gini = 0.111\nsamples = 109\nvalue = [8, 1, 144]\nclass = c'),
Text(3090.461538461538, 906.0, 'BEN <= 2.425\ngini = 0.235\nsamples = 1643\nvalue = [2275, 50, 300]\nclass
= a'),
Text(2918.7692307692305, 543.5999999999999, 'NMHC <= 0.215\ngini = 0.394\nsamples = 573\nvalue = [686, 31,
203]\nclass = a'),
Text(2832.9230769230767, 181.19999999999982, 'gini = 0.255\nsamples = 467\nvalue = [638, 8, 103]\nclass =
a'),
Text(3004.6153846153843, 181.19999999999982, 'gini = 0.561\nsamples = 106\nvalue = [48, 23, 100]\nclass
= c'),
Text(3262.1538461538457, 543.5999999999999, 'O_3 <= 6.59\ngini = 0.128\nsamples = 1070\nvalue = [1589, 19,
97]\nclass = a'),
Text(3176.307692307692, 181.19999999999982, 'gini = 0.603\nsamples = 62\nvalue = [31, 15, 49]\nclass = c'),
Text(3347.9999999999995, 181.19999999999982, 'gini = 0.063\nsamples = 1008\nvalue = [1558, 4, 48]\nclass
= a'),
Text(3820.1538461538457, 1268.4, 'O_3 <= 6.015\ngini = 0.494\nsamples = 1415\nvalue = [1436, 128, 674]\ncla
ss = a'),
Text(3519.6923076923076, 906.0, 'SO_2 <= 15.75\ngini = 0.494\nsamples = 148\nvalue = [20, 63, 159]\nclass =
c'),
Text(3433.8461538461534, 543.5999999999999, 'gini = 0.219\nsamples = 39\nvalue = [0, 56, 8]\nclass = b'),
Text(3605.5384615384614, 543.5999999999999, 'EBE <= 5.41\ngini = 0.266\nsamples = 109\nvalue = [20, 7, 151]
\nclass = c'),
Text(3519.6923076923076, 181.19999999999982, 'gini = 0.532\nsamples = 22\nvalue = [8, 6, 24]\nclass = c'),
Text(3691.3846153846152, 181.19999999999982, 'gini = 0.17\nsamples = 87\nvalue = [12, 1, 127]\nclass = c'),
Text(4120.615384615385, 906.0, 'TOL <= 35.59\ngini = 0.429\nsamples = 1267\nvalue = [1416, 65, 515]\nclass
= a'),
```

```
Text(3948.9230769230767, 543.5999999999999, 'SO_2 <= 47.93\n gini = 0.46\n samples = 1132\n value = [1193, 65, 509]\n class = a'),  
Text(3863.076923076923, 181.1999999999982, 'gini = 0.476\n samples = 1071\n value = [1089, 65, 508]\n class = a'),  
Text(4034.7692307692305, 181.1999999999982, 'gini = 0.019\n samples = 61\n value = [104, 0, 1]\n class = a'),  
Text(4292.307692307692, 543.5999999999999, 'PXY <= 7.75\n gini = 0.051\n samples = 135\n value = [223, 0, 6]\n class = a'),  
Text(4206.461538461538, 181.1999999999982, 'gini = 0.251\n samples = 20\n value = [29, 0, 5]\n class = a'),  
Text(4378.153846153846, 181.1999999999982, 'gini = 0.01\n samples = 115\n value = [194, 0, 1]\n class = a')]
```



# Conclusion

## Accuracy

*Linear Regression:* 0.2998431212306839

*Ridge Regression:* 0.2996231647869222

*Lasso Regression:* 0.062201095160138054

*ElasticNet Regression:* 0.18421578404776207

*Logistic Regression:* 0.879023418036871

*Random Forest:* 0.8659692442383612

**Logistic Regression is suitable for this dataset**