

Importing Libraries

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

Importing Datasets

```
In [2]: 1 df=pd.read_csv("madrid_2006")
2 df
```

Out[2]:

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2	TCH
0	2006-02-01 01:00:00	NaN	1.84	NaN	NaN	NaN	155.100006	490.100006	NaN	4.880000	97.570000	40.259998	NaN	33.779999	NaN
1	2006-02-01 01:00:00	1.68	1.01	2.38	6.36	0.32	94.339996	229.699997	3.04	7.100000	25.820000	NaN	2.48	11.890000	1.59
2	2006-02-01 01:00:00	NaN	1.25	NaN	NaN	NaN	66.800003	192.000000	NaN	4.430000	34.419998	NaN	NaN	19.719999	NaN
3	2006-02-01 01:00:00	NaN	1.68	NaN	NaN	NaN	103.000000	407.799988	NaN	4.830000	28.260000	NaN	NaN	21.129999	NaN
4	2006-02-01 01:00:00	NaN	1.31	NaN	NaN	NaN	105.400002	269.200012	NaN	6.990000	54.180000	NaN	NaN	11.050000	NaN
...
230563	2006-05-01 00:00:00	5.88	0.83	6.23	NaN	0.20	112.500000	218.000000	NaN	24.389999	93.120003	NaN	NaN	7.400000	1.50
230564	2006-05-01 00:00:00	0.76	0.32	0.48	1.09	0.08	51.900002	54.820000	0.61	48.410000	29.469999	15.640000	0.50	8.840000	1.32
230565	2006-05-01 00:00:00	0.96	NaN	0.69	NaN	0.19	135.100006	179.199997	NaN	11.460000	64.680000	35.000000	NaN	12.110000	1.51
230566	2006-05-01 00:00:00	0.50	NaN	0.67	NaN	0.10	82.599998	105.599998	NaN	NaN	94.360001	NaN	NaN	4.890000	1.47
230567	2006-05-01 00:00:00	1.95	0.74	1.99	4.00	0.24	107.300003	160.199997	2.01	17.730000	52.490002	27.920000	1.70	9.170000	1.50

230568 rows × 17 columns

Data Cleaning and Data Preprocessing

```
In [3]: 1 df=df.dropna()
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
       dtype='object')
```

```
In [5]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24758 entries, 5 to 230567
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      24758 non-null   object 
 1   BEN        24758 non-null   float64
 2   CO         24758 non-null   float64
 3   EBE        24758 non-null   float64
 4   MXY        24758 non-null   float64
 5   NMHC       24758 non-null   float64
 6   NO_2       24758 non-null   float64
 7   NOx        24758 non-null   float64
 8   OXY        24758 non-null   float64
 9   O_3         24758 non-null   float64
 10  PM10       24758 non-null   float64
 11  PM25       24758 non-null   float64
 12  PXY        24758 non-null   float64
 13  SO_2       24758 non-null   float64
 14  TCH        24758 non-null   float64
 15  TOL        24758 non-null   float64
 16  station    24758 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.4+ MB
```

```
In [6]: 1 data=df[['CO' , 'station']]  
2 data
```

Out[6]:

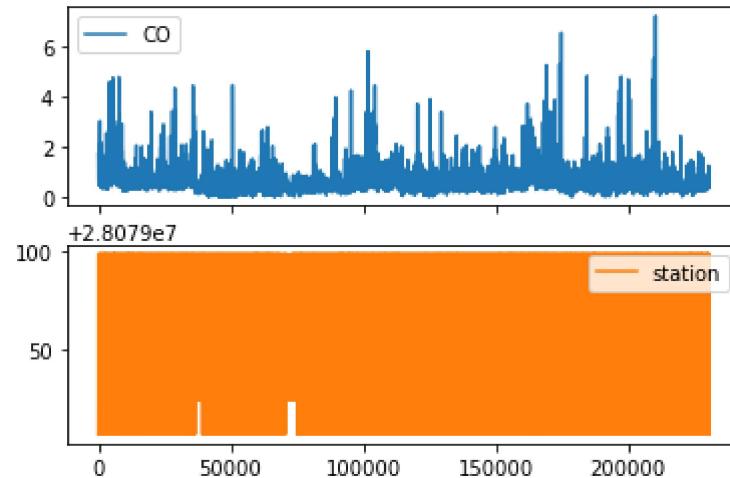
	CO	station
5	1.69	28079006
22	0.79	28079024
25	1.47	28079099
31	0.85	28079006
48	0.79	28079024
...
230538	0.40	28079024
230541	0.94	28079099
230547	1.06	28079006
230564	0.32	28079024
230567	0.74	28079099

24758 rows × 2 columns

Line chart

```
In [7]: 1 data.plot.line(subplots=True)
```

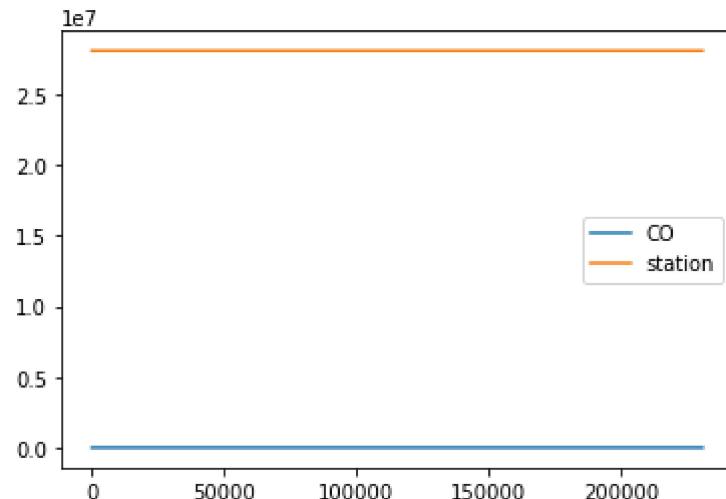
```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



Line chart

```
In [8]: 1 data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

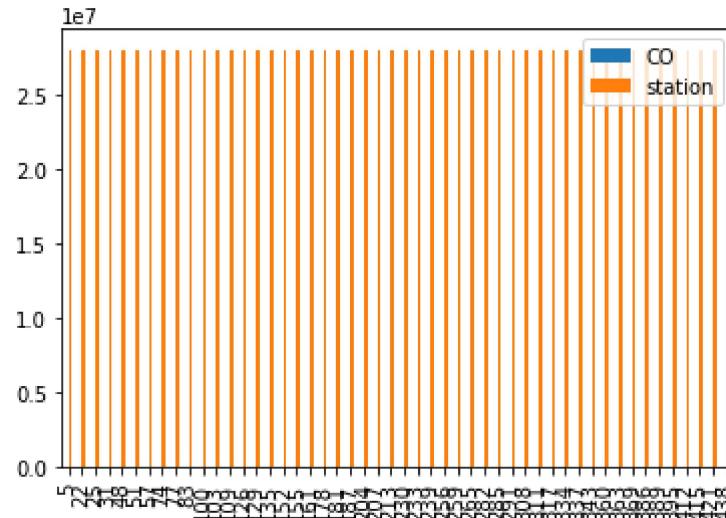


Bar chart

```
In [9]: 1 b=data[0:50]
```

```
In [10]: 1 b.plot.bar()
```

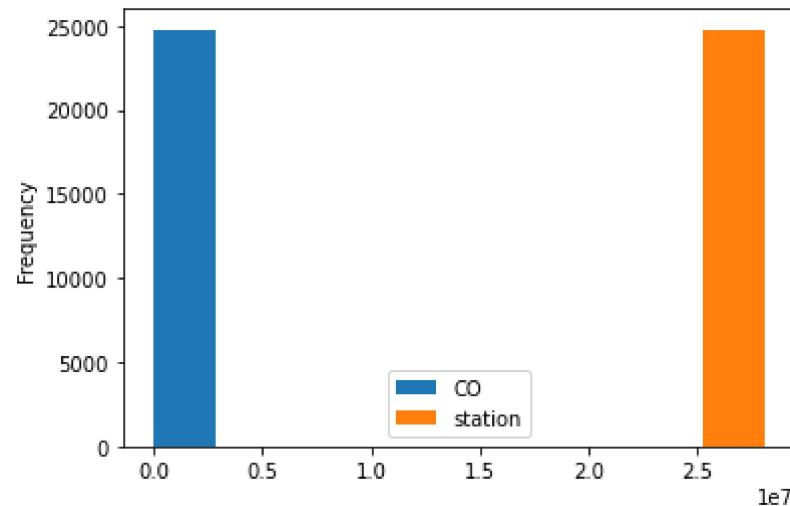
```
Out[10]: <AxesSubplot:>
```



Histogram

```
In [11]: 1 data.plot.hist()
```

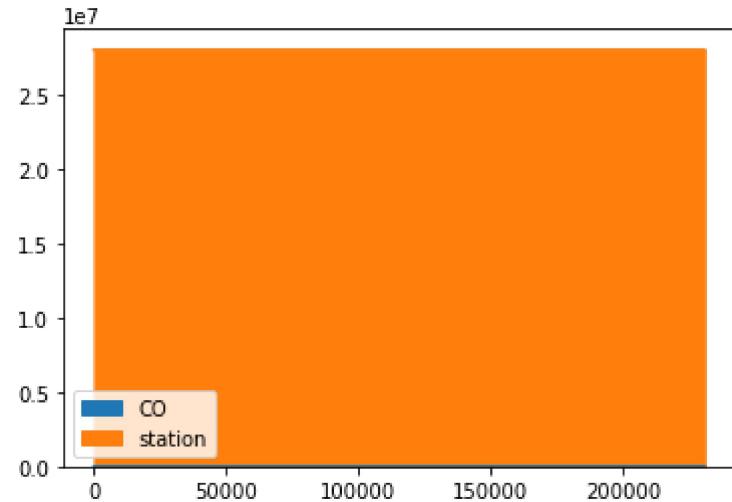
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

In [12]: 1 data.plot.area()

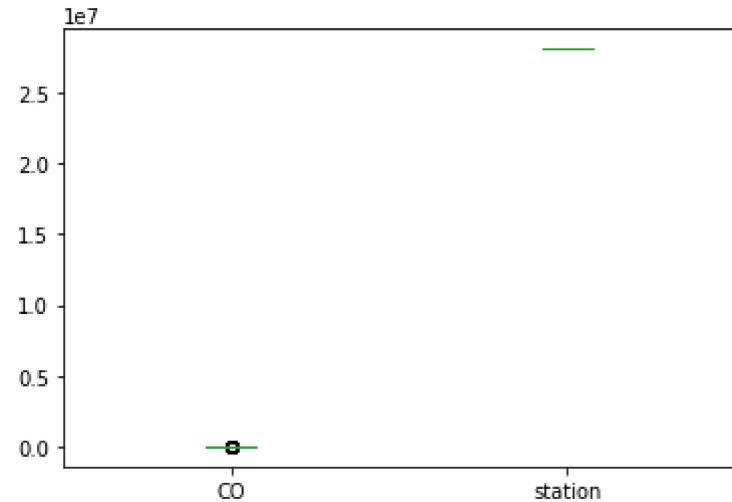
Out[12]: <AxesSubplot:>



Box chart

In [13]: 1 data.plot.box()

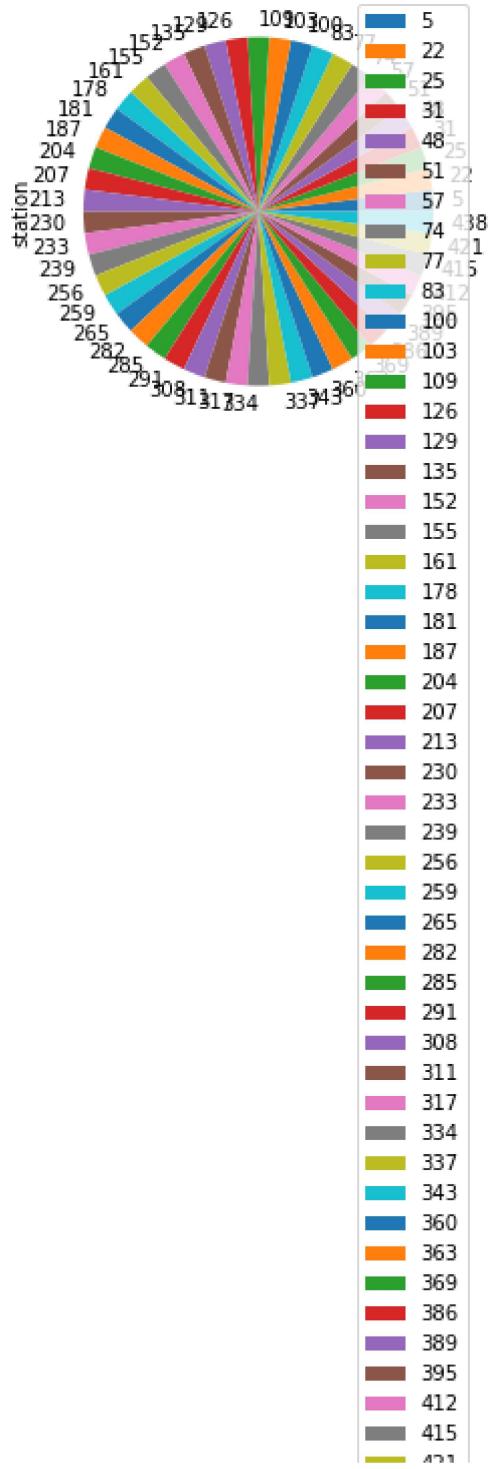
Out[13]: <AxesSubplot:>



Pie chart

```
In [14]: 1 b.plot.pie(y='station' )
```

```
Out[14]: <AxesSubplot:ylabel='station'>
```

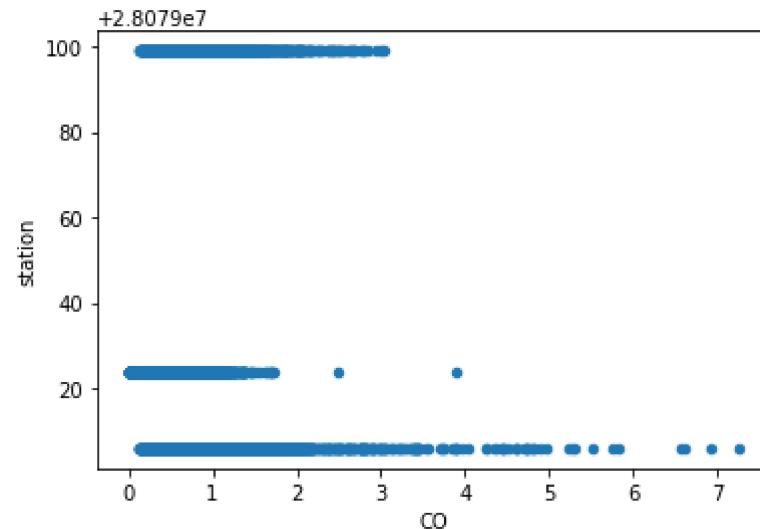





Scatter chart

```
In [15]: 1 data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24758 entries, 5 to 230567
Data columns (total 17 columns):
 #   Column   Non-Null Count  Dtype  
--- 
 0   date      24758 non-null   object  
 1   BEN        24758 non-null   float64 
 2   CO         24758 non-null   float64 
 3   EBE        24758 non-null   float64 
 4   MXY        24758 non-null   float64 
 5   NMHC       24758 non-null   float64 
 6   NO_2       24758 non-null   float64 
 7   NOx        24758 non-null   float64 
 8   OXY        24758 non-null   float64 
 9   O_3         24758 non-null   float64 
 10  PM10       24758 non-null   float64 
 11  PM25       24758 non-null   float64 
 12  PXY        24758 non-null   float64 
 13  SO_2       24758 non-null   float64 
 14  TOL        24758 non-null   float64
```

```
In [17]: 1 df.describe()
```

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
count	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000	24758.000000
mean	1.350624	0.600713	1.824534	3.835034	0.176546	58.333481	116.419090	1.990347	37.155685
std	1.541636	0.419048	1.868939	4.069036	0.126683	40.529382	117.557064	1.931620	31.127681
min	0.110000	0.000000	0.170000	0.150000	0.000000	1.680000	2.020000	0.190000	0.310000
25%	0.450000	0.360000	0.810000	1.060000	0.100000	28.450001	36.882501	0.960000	12.110000
50%	0.850000	0.500000	1.130000	2.500000	0.150000	52.959999	85.180000	1.260000	28.675000
75%	1.680000	0.720000	2.160000	5.090000	0.220000	79.347498	158.300003	2.470000	53.029999
max	45.430000	7.250000	57.799999	66.900002	2.020000	461.299988	1680.000000	63.000000	178.899994

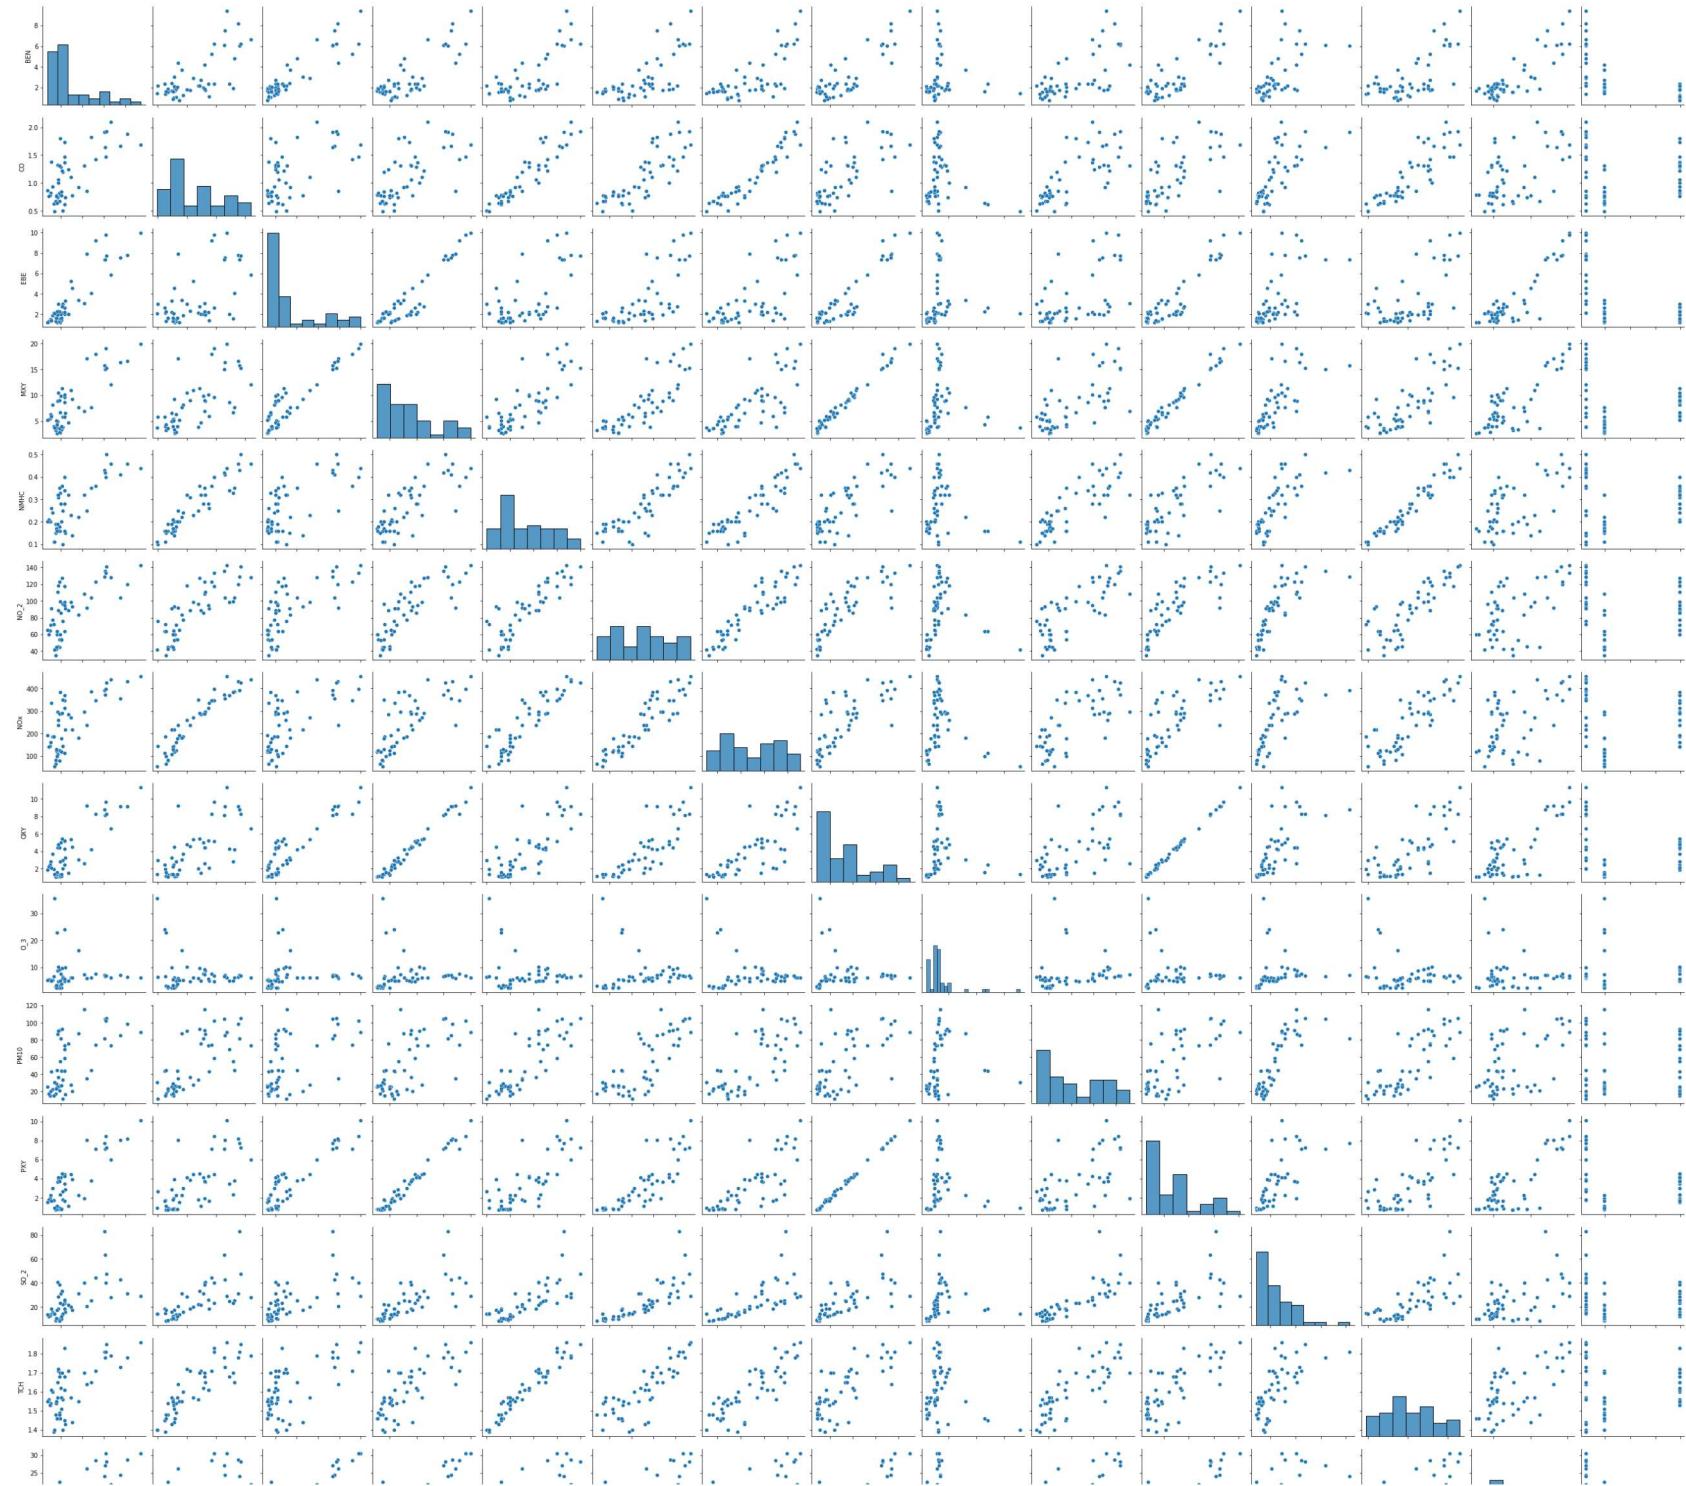
In [18]:

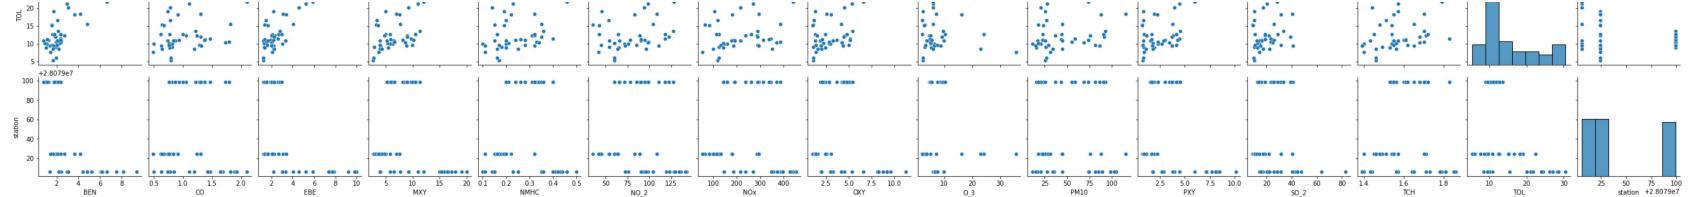
```
1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2         'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

```
In [19]: 1 sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x11e9d425880>
```

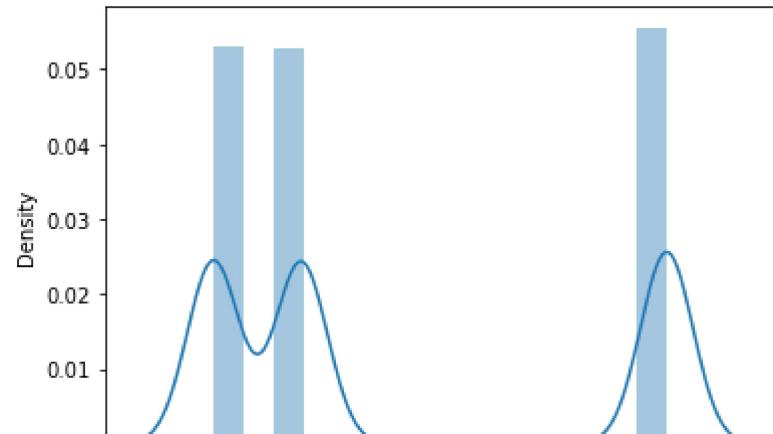





```
In [20]: 1 sns.distplot(df1['station'])
```

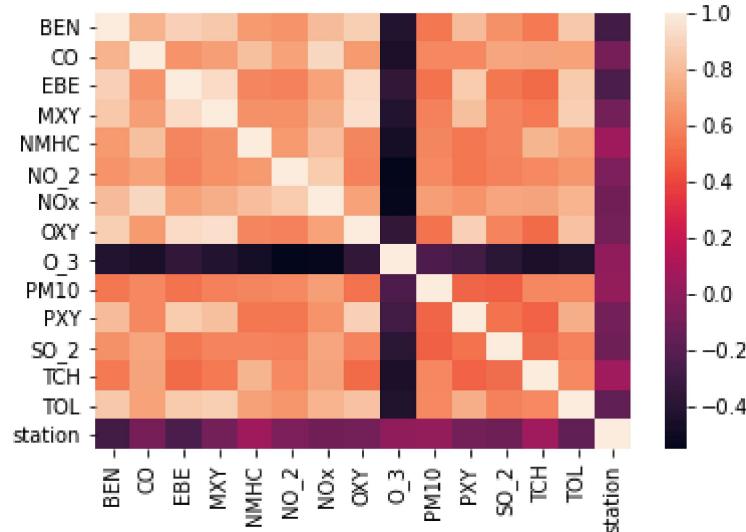
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: 1 sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
```

```
In [23]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: 1 from sklearn.linear_model import LinearRegression  
2 lr=LinearRegression()  
3 lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: 1 lr.intercept_
```

```
Out[25]: 28079017.64075772
```

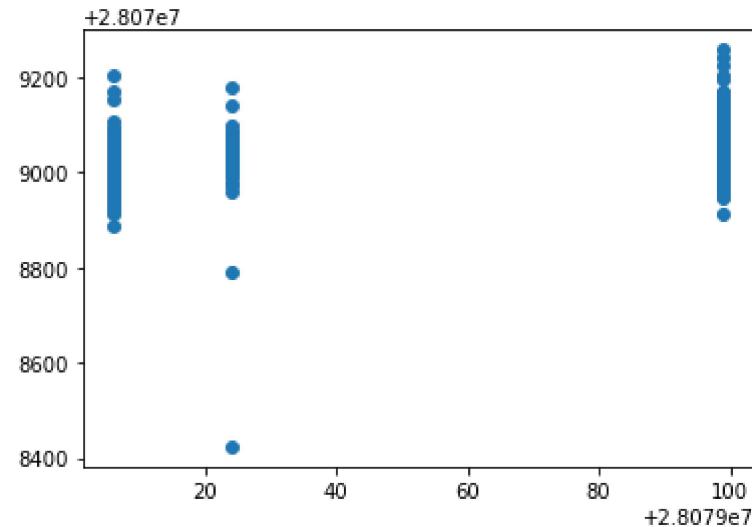
```
In [26]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
2 coeff
```

```
Out[26]:
```

	Co-efficient
BEN	-18.174056
CO	-12.057900
EBE	-23.022451
MXY	4.652349
NMHC	131.095210
NO_2	-0.034385
NOx	0.000609
OXY	14.122589
O_3	-0.055439
PM10	0.144916
PXY	6.404906
SO_2	-0.630762
TCH	20.843452
TOL	-0.566375

```
In [27]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x11eabde0b80>
```



ACCURACY

```
In [28]: 1 lr.score(x_test,y_test)
```

```
Out[28]: 0.3788685188488047
```

```
In [29]: 1 lr.score(x_train,y_train)
```

```
Out[29]: 0.3989029834268095
```

Ridge and Lasso

```
In [30]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: 1 rr.score(x_test,y_test)
```

```
Out[32]: 0.3786409396350162
```

```
In [33]: 1 rr.score(x_train,y_train)
```

```
Out[33]: 0.3982129859207212
```

```
In [34]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

Accuracy(Lasso)

```
In [35]: 1 la.score(x_train,y_train)
```

```
Out[35]: 0.06309829772793729
```

```
In [36]: 1 la.score(x_test,y_test)
```

```
Out[36]: 0.05729762636752733
```

```
In [37]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: 1 en.coef_
```

```
Out[38]: array([-8.54250558,  0.          , -9.02976539,  3.25611115,  0.42688624,
 -0.01771263,  0.01021728,  3.36745815, -0.12640861,  0.31114115,
 2.74860128, -0.4380658 ,  0.56883655, -1.01138549])
```

```
In [39]: 1 en.intercept_
```

```
Out[39]: 28079052.05014169
```

```
In [40]: 1 prediction=en.predict(x_test)
```

```
In [65]: 1 en.score(x_test,y_test)
```

```
Out[65]: 0.2319550412034177
```

Evaluation Metrics

```
In [42]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
32.39661552047374
```

```
1274.3280704097153
```

```
35.697732006525506
```

Logistic Regression

```
In [43]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [44]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df['station']
```

```
In [45]: 1 feature_matrix.shape
```

```
Out[45]: (24758, 14)
```

```
In [46]: 1 target_vector.shape
```

```
Out[46]: (24758,)
```

```
In [47]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [48]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: 1 logr=LogisticRegression(max_iter=10000)  
2 logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: 1 prediction=logr.predict(observation)  
2 print(prediction)
```

```
[28079099]
```

```
In [52]: 1 logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [53]: 1 logr.score(fs,target_vector)
```

```
Out[53]: 0.8741416915744405
```

```
In [54]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[54]: 3.5557727473608076e-15
```

```
In [55]: 1 logr.predict_proba(observation)
```

```
Out[55]: array([[3.55577275e-15, 7.80743173e-29, 1.00000000e+00]])
```

Random Forest

```
In [56]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: 1 rfc=RandomForestClassifier()
2 rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: 1 parameters={'max_depth':[1,2,3,4,5],
2                 'min_samples_leaf':[5,10,15,20,25],
3                 'n_estimators':[10,20,30,40,50]
4 }
```

```
In [59]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')
```

```
In [60]: 1 grid_search.best_score_
```

```
Out[60]: 0.8706866705135603
```

```
In [61]: 1 rfc_best=grid_search.best_estimator_
```

In [62]:

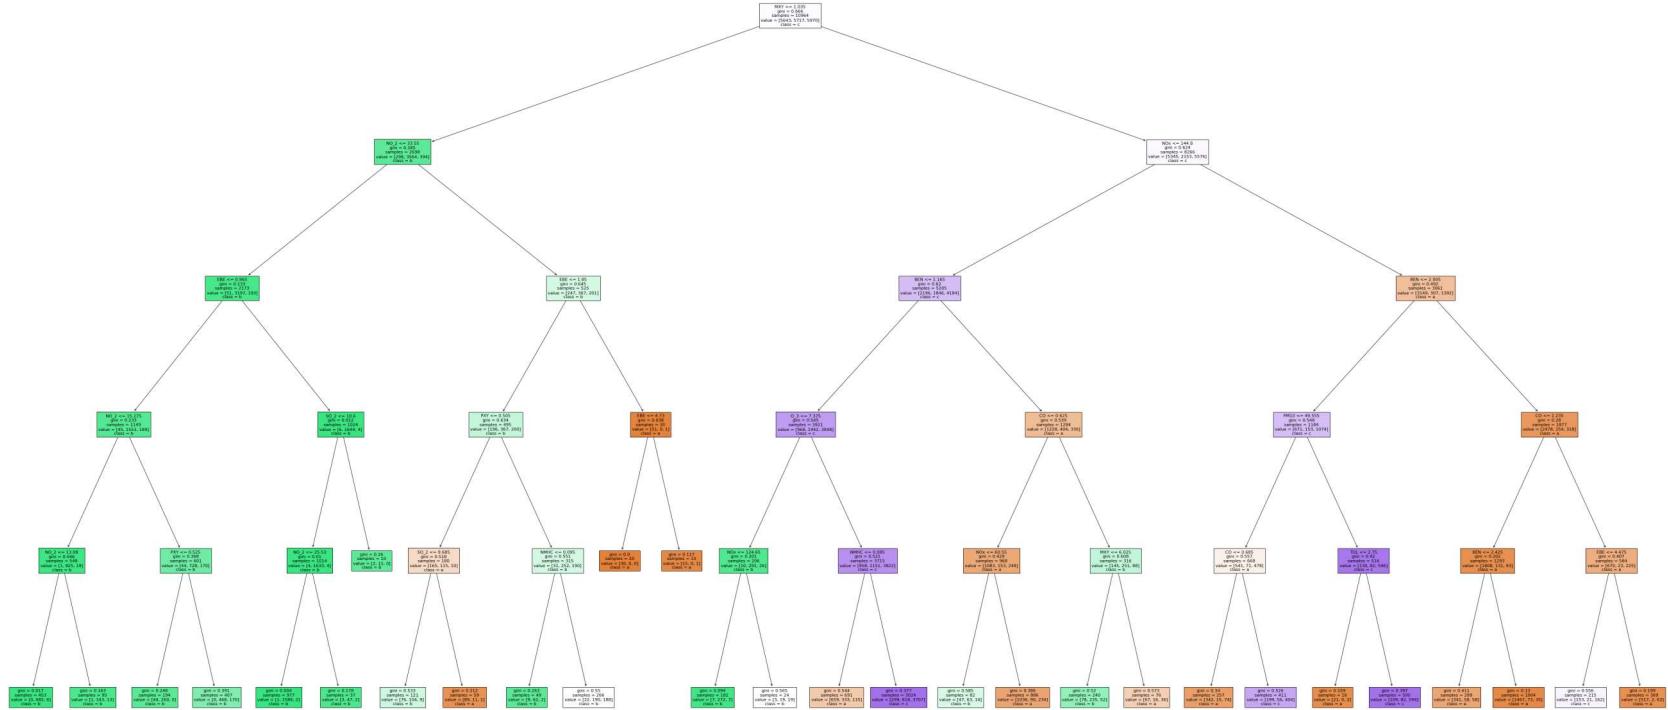
```
1 from sklearn.tree import plot_tree
2
3 plt.figure(figsize=(80,40))
4 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

Out[62]: [Text(2108.0, 1993.2, 'MXY <= 1.035\ngini = 0.666\nsamples = 10964\nvalue = [5643, 5717, 5970]\nnclass = c'),
Text(1074.6666666666667, 1630.800000000002, 'NO_2 <= 33.55\ngini = 0.285\nsamples = 2698\nvalue = [298, 35
64, 394]\nnclass = b'),
Text(620.0, 1268.4, 'EBE <= 0.965\ngini = 0.133\nsamples = 2173\nvalue = [51, 3197, 193]\nnclass = b'),
Text(330.6666666666667, 906.0, 'NO_2 <= 15.275\ngini = 0.233\nsamples = 1149\nvalue = [45, 1553, 189]\nclas
s = b'),
Text(165.3333333333334, 543.5999999999999, 'NO_2 <= 13.08\ngini = 0.046\nsamples = 548\nvalue = [1, 825, 1
9]\nnclass = b'),
Text(82.66666666666667, 181.1999999999982, 'gini = 0.017\nsamples = 453\nvalue = [0, 682, 6]\nnclass = b'),
Text(248.0, 181.1999999999982, 'gini = 0.163\nsamples = 95\nvalue = [1, 143, 13]\nnclass = b'),
Text(496.0, 543.5999999999999, 'PXY <= 0.525\ngini = 0.368\nsamples = 601\nvalue = [44, 728, 170]\nnclass =
b'),
Text(413.3333333333337, 181.1999999999982, 'gini = 0.248\nsamples = 194\nvalue = [44, 260, 0]\nnclass =
b'),
Text(578.6666666666667, 181.1999999999982, 'gini = 0.391\nsamples = 407\nvalue = [0, 468, 170]\nnclass =
b'),
Text(909.333333333334, 906.0, 'SO_2 <= 10.6\ngini = 0.012\nsamples = 1024\nvalue = [6, 1644, 4]\nnclass =
b'),
Text(826.6666666666667, 543.5999999999999, 'NO_2 <= 25.53\ngini = 0.01\nsamples = 1014\nvalue = [4, 1633,
4]\nnclass = b'),
Text(744.0, 181.1999999999982, 'gini = 0.004\nsamples = 977\nvalue = [1, 1586, 2]\nnclass = b'),
Text(909.333333333334, 181.1999999999982, 'gini = 0.178\nsamples = 37\nvalue = [3, 47, 2]\nnclass = b'),
Text(992.0, 543.5999999999999, 'gini = 0.26\nsamples = 10\nvalue = [2, 11, 0]\nnclass = b'),
Text(1529.333333333335, 1268.4, 'EBE <= 1.85\ngini = 0.645\nsamples = 525\nvalue = [247, 367, 201]\nnclass
= b'),
Text(1322.6666666666667, 906.0, 'PXY <= 0.505\ngini = 0.634\nsamples = 495\nvalue = [196, 367, 200]\nnclass
= b'),
Text(1157.333333333335, 543.5999999999999, 'SO_2 <= 9.685\ngini = 0.518\nsamples = 180\nvalue = [165, 115,
10]\nnclass = a'),
Text(1074.6666666666667, 181.1999999999982, 'gini = 0.533\nsamples = 121\nvalue = [76, 104, 9]\nnclass =
b'),
Text(1240.0, 181.1999999999982, 'gini = 0.212\nsamples = 59\nvalue = [89, 11, 1]\nnclass = a'),
Text(1488.0, 543.5999999999999, 'NMHC <= 0.095\ngini = 0.551\nsamples = 315\nvalue = [31, 252, 190]\nnclass
= b'),
Text(1405.333333333335, 181.1999999999982, 'gini = 0.263\nsamples = 49\nvalue = [9, 62, 2]\nnclass = b'),
Text(1570.6666666666667, 181.1999999999982, 'gini = 0.55\nsamples = 266\nvalue = [22, 190, 188]\nnclass =
b'),
Text(1736.0, 906.0, 'EBE <= 4.73\ngini = 0.038\nsamples = 30\nvalue = [51, 0, 1]\nnclass = a'),
Text(1653.333333333335, 543.5999999999999, 'gini = 0.0\nsamples = 20\nvalue = [36, 0, 0]\nnclass = a'),
Text(1818.6666666666667, 543.5999999999999, 'gini = 0.117\nsamples = 10\nvalue = [15, 0, 1]\nnclass = a'),
Text(3141.333333333335, 1630.800000000002, 'NOx <= 144.8\ngini = 0.624\nsamples = 8266\nvalue = [5345, 21
53, 5576]\nnclass = c'),
Text(2480.0, 1268.4, 'BEN <= 1.165\ngini = 0.62\nsamples = 5205\nvalue = [2196, 1846, 4184]\nnclass = c')]

```
Text(2149.333333333335, 906.0, 'O_3 <= 7.375\n gini = 0.545\n samples = 3921\n value = [968, 1442, 3848]\n class = c'),  
Text(1984.0, 543.599999999999, 'NOx <= 124.65\n gini = 0.201\n samples = 206\n value = [10, 291, 26]\n class = b'),  
Text(1901.333333333335, 181.1999999999982, 'gini = 0.094\n samples = 182\n value = [7, 272, 7]\n class = b'),  
Text(2066.666666666667, 181.1999999999982, 'gini = 0.565\n samples = 24\n value = [3, 19, 19]\n class = b'),  
Text(2314.666666666667, 543.599999999999, 'NMHC <= 0.085\n gini = 0.521\n samples = 3715\n value = [958, 1151, 3822]\n class = c'),  
Text(2232.0, 181.1999999999982, 'gini = 0.544\n samples = 691\n value = [659, 333, 115]\n class = a'),  
Text(2397.333333333335, 181.1999999999982, 'gini = 0.377\n samples = 3024\n value = [299, 818, 3707]\n class = c'),  
Text(2810.666666666667, 906.0, 'CO <= 0.625\n gini = 0.539\n samples = 1284\n value = [1228, 404, 336]\n class = a'),  
Text(2645.333333333335, 543.599999999999, 'NOx <= 60.55\n gini = 0.429\n samples = 968\n value = [1083, 153, 248]\n class = a'),  
Text(2562.666666666667, 181.1999999999982, 'gini = 0.585\n samples = 82\n value = [47, 63, 14]\n class = b'),  
Text(2728.0, 181.1999999999982, 'gini = 0.386\n samples = 886\n value = [1036, 90, 234]\n class = a'),  
Text(2976.0, 543.599999999999, 'MXY <= 6.025\n gini = 0.608\n samples = 316\n value = [145, 251, 88]\n class = b'),  
Text(2893.333333333335, 181.1999999999982, 'gini = 0.52\n samples = 240\n value = [78, 235, 52]\n class = b'),  
Text(3058.666666666667, 181.1999999999982, 'gini = 0.573\n samples = 76\n value = [67, 16, 36]\n class = a'),  
Text(3802.666666666667, 1268.4, 'BEN <= 2.005\n gini = 0.492\n samples = 3061\n value = [3149, 307, 1392]\n class = a'),  
Text(3472.0, 906.0, 'PM10 <= 49.555\n gini = 0.548\n samples = 1184\n value = [671, 153, 1074]\n class = c'),  
Text(3306.666666666667, 543.599999999999, 'CO <= 0.685\n gini = 0.557\n samples = 668\n value = [541, 71, 478]\n class = a'),  
Text(3224.0, 181.1999999999982, 'gini = 0.34\n samples = 257\n value = [342, 15, 74]\n class = a'),  
Text(3389.333333333335, 181.1999999999982, 'gini = 0.526\n samples = 411\n value = [199, 56, 404]\n class = c'),  
Text(3637.333333333335, 543.599999999999, 'TOL <= 2.75\n gini = 0.42\n samples = 516\n value = [130, 82, 596]\n class = c'),  
Text(3554.666666666667, 181.1999999999982, 'gini = 0.159\n samples = 16\n value = [21, 0, 2]\n class = a'),  
Text(3720.0, 181.1999999999982, 'gini = 0.397\n samples = 500\n value = [109, 82, 594]\n class = c'),  
Text(4133.33333333334, 906.0, 'CO <= 1.235\n gini = 0.28\n samples = 1877\n value = [2478, 154, 318]\n class = a'),  
Text(3968.0, 543.599999999999, 'BEN <= 2.425\n gini = 0.202\n samples = 1293\n value = [1808, 131, 93]\n class = a'),  
Text(3885.333333333335, 181.1999999999982, 'gini = 0.411\n samples = 289\n value = [341, 58, 58]\n class = a'),  
Text(4050.666666666667, 181.1999999999982, 'gini = 0.13\n samples = 1004\n value = [1467, 73, 35]\n class = a'),
```

```

Text(4298.666666666667, 543.5999999999999, 'EBE <= 4.475\nngini = 0.407\nnsamples = 584\nvalue = [670, 23, 22
5]\nnclass = a'),
Text(4216.0, 181.1999999999982, 'gini = 0.556\nnsamples = 215\nvalue = [153, 21, 162]\nnclass = c'),
Text(4381.333333333334, 181.1999999999982, 'gini = 0.199\nnsamples = 369\nvalue = [517, 2, 63]\nnclass =
a')]
```



Conclusion

Accuracy

Linear Regression: 0.3989029834268095

Ridge Regression: 0.3982129859207212

Lasso Regression:0.06309829772793729

ElasticNet Regression:0.2319550412034177

Logistic Regression:0.8741416915744405

Random Forest:0.8706866705135603

Logistic Regression is suitable for this dataset