

Importing Libraries

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

Importing Datasets

```
In [2]: 1 df=pd.read_csv("madrid_2007")
2 df
```

Out[2]:

		date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PM25	PXY	SO_2	TCH
0		2007-12-01 01:00:00	NaN	2.86	NaN	NaN	NaN	282.200012	1054.000000	NaN	4.030000	156.199997	97.43	NaN	64.519997	NaN
1		2007-12-01 01:00:00	NaN	1.82	NaN	NaN	NaN	86.419998	354.600006	NaN	3.260000	80.809998	NaN	NaN	35.419998	NaN
2		2007-12-01 01:00:00	NaN	1.47	NaN	NaN	NaN	94.639999	319.000000	NaN	5.310000	53.099998	NaN	NaN	19.080000	NaN
3		2007-12-01 01:00:00	NaN	1.64	NaN	NaN	NaN	127.900002	476.700012	NaN	4.500000	105.300003	NaN	NaN	17.670000	NaN
4		2007-12-01 01:00:00	4.64	1.86	4.26	7.98	0.57	145.100006	573.900024	3.49	52.689999	106.500000	15.90	3.56	40.230000	1.94
...
225115		2007-03-01 00:00:00	0.30	0.45	1.00	0.30	0.26	8.690000	11.690000	1.00	42.209999	6.760000	5.14	1.00	7.420000	1.44
225116		2007-03-01 00:00:00	NaN	0.16	NaN	NaN	NaN	46.820000	51.480000	NaN	22.150000	5.700000	NaN	NaN	7.130000	NaN
225117		2007-03-01 00:00:00	0.24	NaN	0.20	NaN	0.09	51.259998	66.809998	NaN	18.540001	13.010000	6.95	NaN	8.740000	1.30
225118		2007-03-01 00:00:00	0.11	NaN	1.00	NaN	0.05	24.240000	36.930000	NaN	NaN	6.610000	NaN	NaN	9.890000	1.29
225119		2007-03-01 00:00:00	0.53	0.40	1.00	1.70	0.12	32.360001	47.860001	1.37	24.150000	10.260000	7.08	1.23	9.890000	1.32

225120 rows × 17 columns

Data Cleaning and Data Preprocessing

```
In [3]: 1 df=df.dropna()
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
               'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
               dtype='object')
```

```
In [5]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25443 entries, 4 to 225119
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      25443 non-null   object 
 1   BEN       25443 non-null   float64
 2   CO        25443 non-null   float64
 3   EBE       25443 non-null   float64
 4   MXY       25443 non-null   float64
 5   NMHC      25443 non-null   float64
 6   NO_2      25443 non-null   float64
 7   NOx       25443 non-null   float64
 8   OXY       25443 non-null   float64
 9   O_3        25443 non-null   float64
 10  PM10      25443 non-null   float64
 11  PM25      25443 non-null   float64
 12  PXY       25443 non-null   float64
 13  SO_2      25443 non-null   float64
 14  TCH       25443 non-null   float64
 15  TOL       25443 non-null   float64
 16  station    25443 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.5+ MB
```

```
In [6]: 1 data=df[['CO' , 'station']]  
2 data
```

Out[6]:

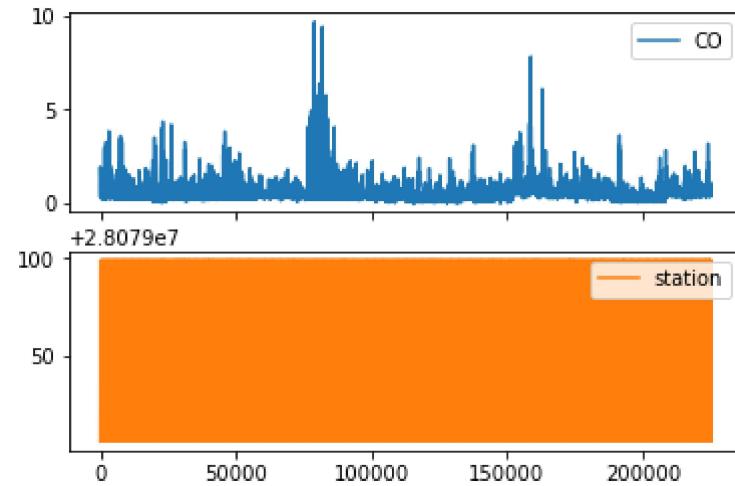
	CO	station
4	1.86	28079006
21	0.31	28079024
25	1.42	28079099
30	1.89	28079006
47	0.30	28079024
...
225073	0.47	28079006
225094	0.45	28079099
225098	0.41	28079006
225115	0.45	28079024
225119	0.40	28079099

25443 rows × 2 columns

Line chart

```
In [7]: 1 data.plot.line(subplots=True)
```

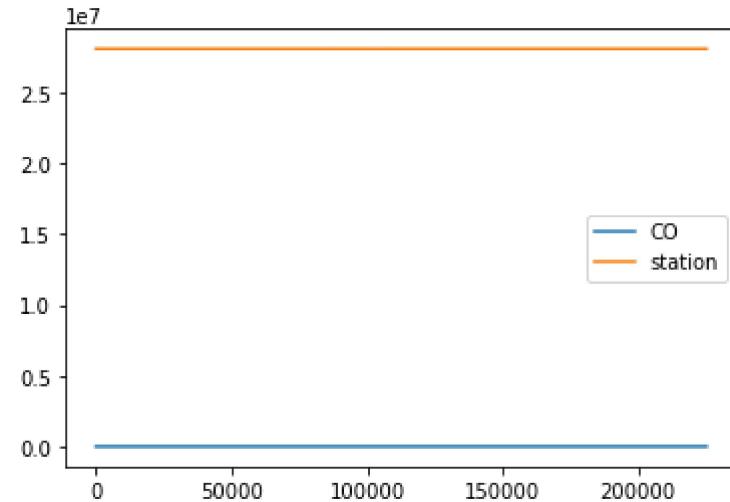
```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



Line chart

```
In [8]: 1 data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

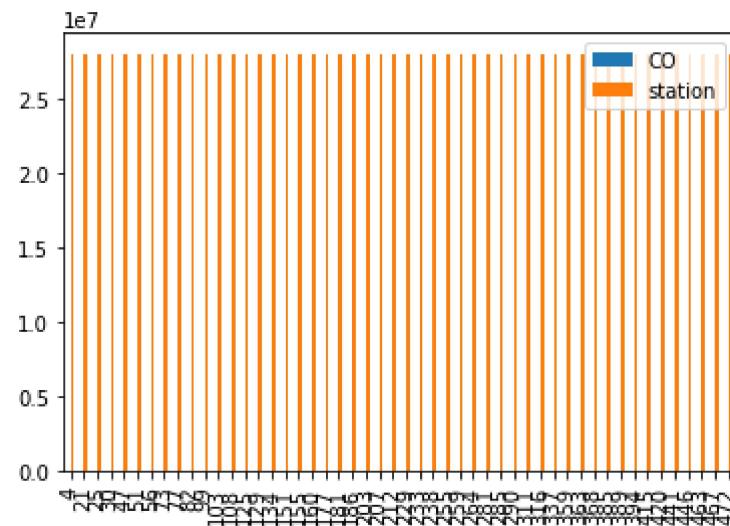


Bar chart

```
In [9]: 1 b=data[0:50]
```

In [10]: 1 b.plot.bar()

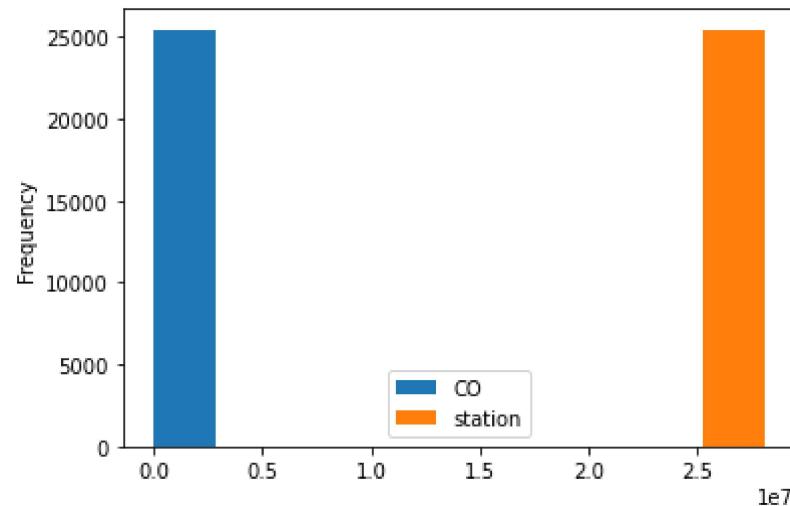
Out[10]: <AxesSubplot:>



Histogram

```
In [11]: 1 data.plot.hist()
```

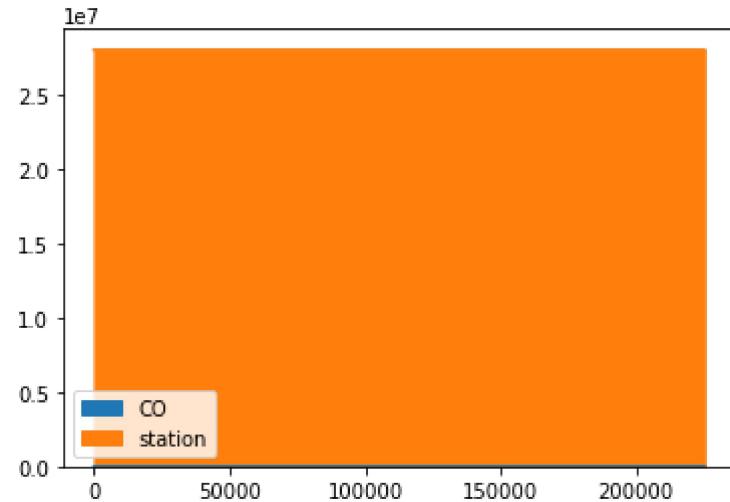
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

In [12]: 1 data.plot.area()

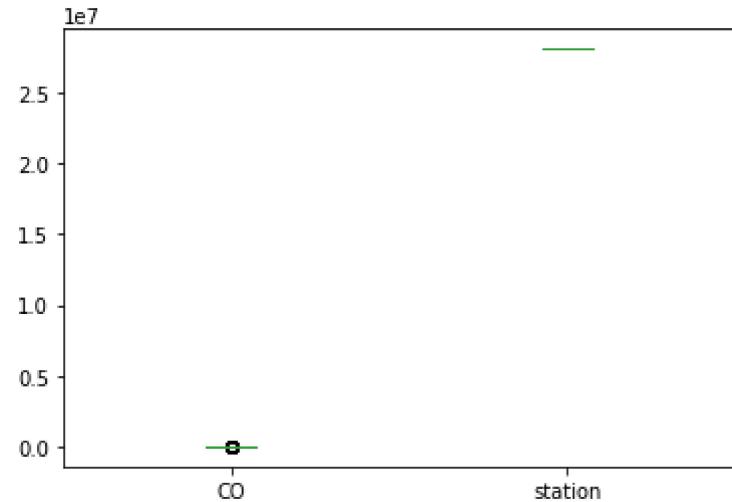
Out[12]: <AxesSubplot:>



Box chart

In [13]: 1 data.plot.box()

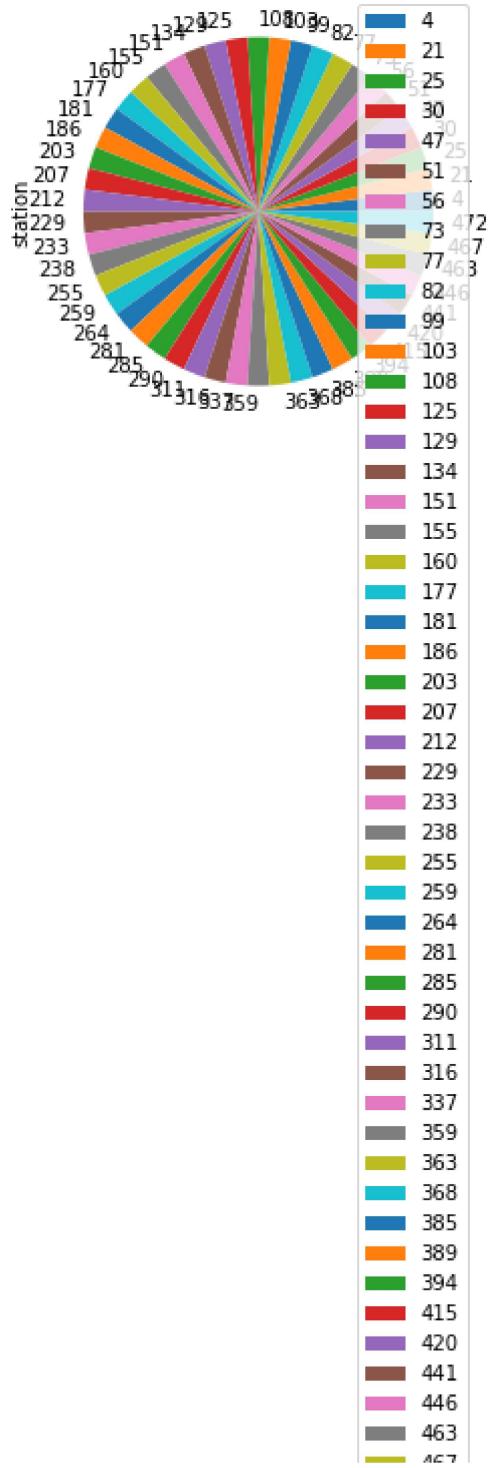
Out[13]: <AxesSubplot:>



Pie chart

```
In [14]: 1 b.plot.pie(y='station' )
```

```
Out[14]: <AxesSubplot:ylabel='station'>
```

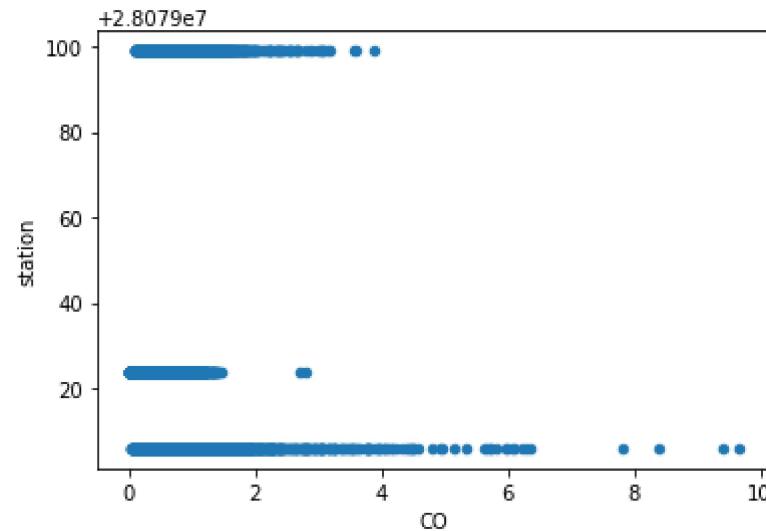





Scatter chart

```
In [15]: 1 data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25443 entries, 4 to 225119
Data columns (total 17 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   date      25443 non-null    object  
 1   BEN        25443 non-null    float64 
 2   CO         25443 non-null    float64 
 3   EBE        25443 non-null    float64 
 4   MXY        25443 non-null    float64 
 5   NMHC       25443 non-null    float64 
 6   NO_2       25443 non-null    float64 
 7   NOx        25443 non-null    float64 
 8   OXY        25443 non-null    float64 
 9   O_3         25443 non-null    float64 
 10  PM10       25443 non-null    float64 
 11  PM25       25443 non-null    float64 
 12  PXY        25443 non-null    float64 
 13  SO_2       25443 non-null    float64 
 14  TOL        25443 non-null    float64
```

```
In [17]: 1 df.describe()
```

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3
count	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000	25443.000000
mean	1.146744	0.505120	1.394071	2.392008	0.249967	58.532683	112.741861	1.270278	35.653534
std	1.278733	0.423231	1.268265	2.784302	0.142627	37.755029	115.527006	1.143188	26.911492
min	0.130000	0.000000	0.120000	0.150000	0.000000	1.690000	1.780000	0.110000	0.280000
25%	0.450000	0.260000	0.780000	0.960000	0.160000	31.285001	39.910000	0.740000	11.680000
50%	0.770000	0.400000	1.000000	1.500000	0.220000	54.080002	82.809998	1.000000	30.980000
75%	1.390000	0.640000	1.580000	2.855000	0.300000	79.230003	149.300003	1.360000	53.310001
max	30.139999	9.660000	31.680000	65.480003	2.570000	430.299988	1893.000000	32.540001	149.100006

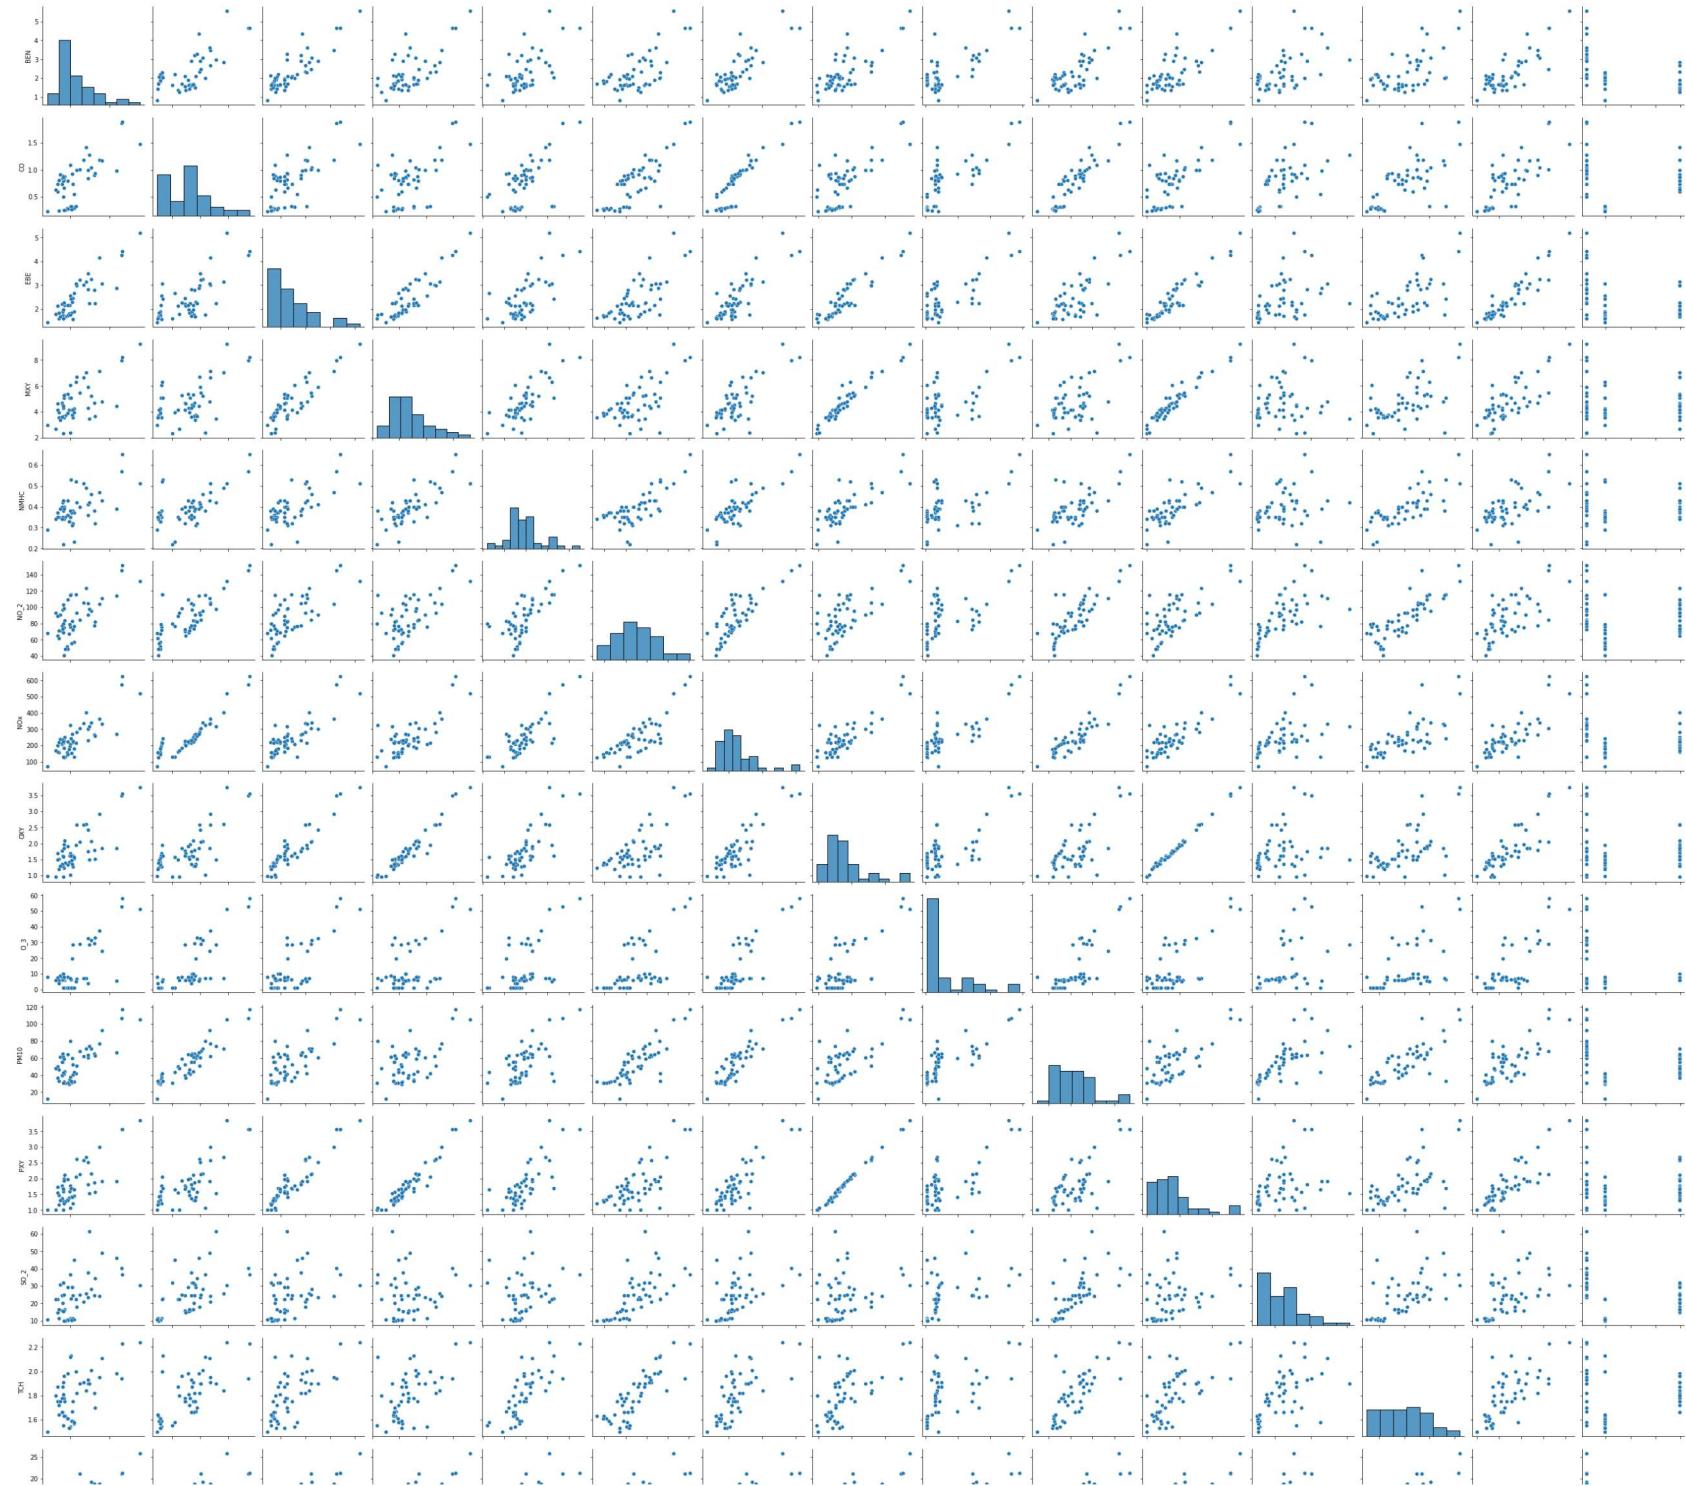
In [18]:

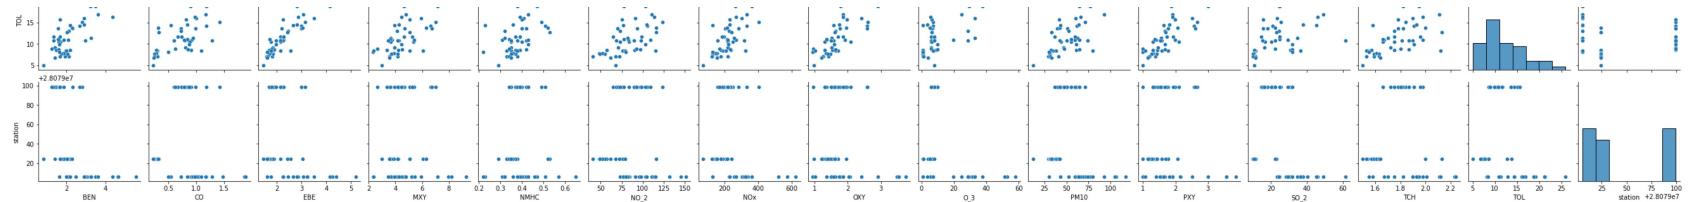
```
1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2         'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

```
In [19]: 1 sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x245ef4504f0>
```

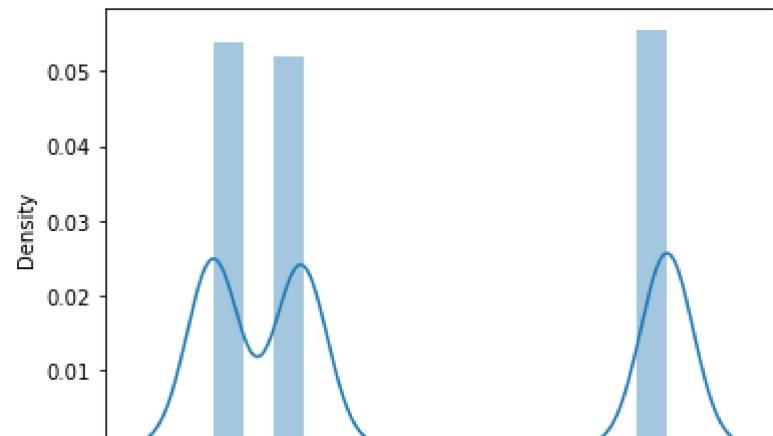





```
In [20]: 1 sns.distplot(df1['station'])
```

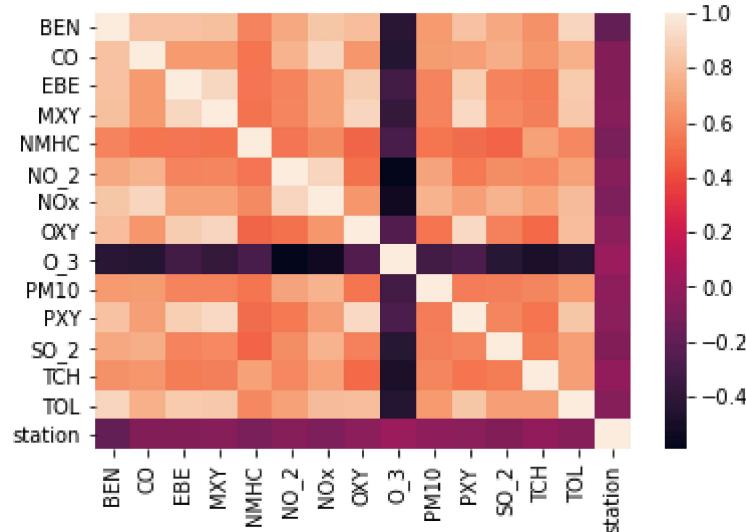
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: 1 sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
```

```
In [23]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: 1 from sklearn.linear_model import LinearRegression  
2 lr=LinearRegression()  
3 lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: 1 lr.intercept_
```

```
Out[25]: 28079009.799331877
```

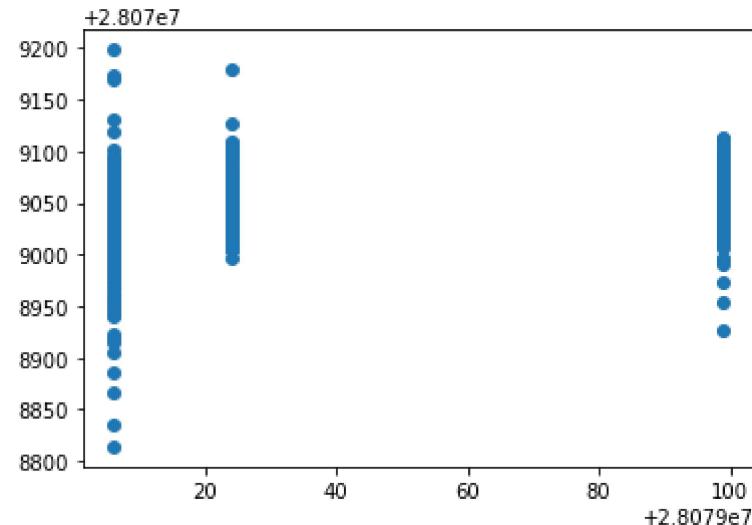
```
In [26]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
2 coeff
```

```
Out[26]:
```

	Co-efficient
BEN	-32.676195
CO	17.012253
EBE	0.846044
MXY	-1.163146
NMHC	-41.008556
NO_2	0.100799
NOx	-0.035712
OXY	5.301298
O_3	-0.037435
PM10	0.137878
PXY	6.509756
SO_2	0.233654
TCH	26.243788
TOL	3.097605

```
In [27]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x245fddb6f70>
```



ACCURACY

```
In [28]: 1 lr.score(x_test,y_test)
```

```
Out[28]: 0.15888618593861437
```

```
In [29]: 1 lr.score(x_train,y_train)
```

```
Out[29]: 0.15930303157272896
```

Ridge and Lasso

```
In [30]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: 1 rr.score(x_test,y_test)
```

```
Out[32]: 0.15883525144058097
```

```
In [33]: 1 rr.score(x_train,y_train)
```

```
Out[33]: 0.15925082658615208
```

```
In [34]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

Accuracy(Lasso)

```
In [35]: 1 la.score(x_train,y_train)
```

```
Out[35]: 0.012581248316343108
```

```
In [36]: 1 la.score(x_test,y_test)
```

```
Out[36]: 0.014373582045566824
```

```
In [37]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: 1 en.coef_
```

```
Out[38]: array([-8.15496477e+00,  0.00000000e+00,  0.00000000e+00,  1.52019723e-01,
 -0.00000000e+00,  3.80128154e-02, -4.42837414e-02,  5.71943162e-01,
 -6.15289655e-02,  1.57697948e-01,  5.63096780e-01,  5.54937274e-04,
 0.00000000e+00,  9.18614145e-01])
```

```
In [39]: 1 en.intercept_
```

```
Out[39]: 28079046.59911058
```

```
In [40]: 1 prediction=en.predict(x_test)
```

```
In [41]: 1 en.score(x_test,y_test)
```

```
Out[41]: 0.06773347039258326
```

Evaluation Metrics

```
In [42]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
36.616311701215956
```

```
1537.9390731126289
```

```
39.21656630956653
```

Logistic Regression

```
In [43]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [44]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df[ 'station']
```

```
In [45]: 1 feature_matrix.shape
```

```
Out[45]: (25443, 14)
```

```
In [46]: 1 target_vector.shape
```

```
Out[46]: (25443,)
```

```
In [47]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [48]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: 1 logr=LogisticRegression(max_iter=10000)
2 logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: 1 prediction=logr.predict(observation)
2 print(prediction)
```

```
[28079099]
```

```
In [52]: 1 logr.classes_
```

```
Out[52]: array([28079006, 28079024, 28079099], dtype=int64)
```

```
In [53]: 1 logr.score(fs,target_vector)
```

```
Out[53]: 0.8146838030106512
```

```
In [54]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[54]: 1.082753977181323e-19
```

```
In [55]: 1 logr.predict_proba(observation)
```

```
Out[55]: array([[1.08275398e-19, 1.80383815e-19, 1.00000000e+00]])
```

Random Forest

```
In [56]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: 1 rfc=RandomForestClassifier()
2 rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: 1 parameters={'max_depth':[1,2,3,4,5],
2                 'min_samples_leaf':[5,10,15,20,25],
3                 'n_estimators':[10,20,30,40,50]
4 }
```

```
In [59]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')
```

```
In [60]: 1 grid_search.best_score_
```

```
Out[60]: 0.8211117349803481
```

```
In [61]: 1 rfc_best=grid_search.best_estimator_
```

In [62]:

```
1 from sklearn.tree import plot_tree
2
3 plt.figure(figsize=(80,40))
4 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

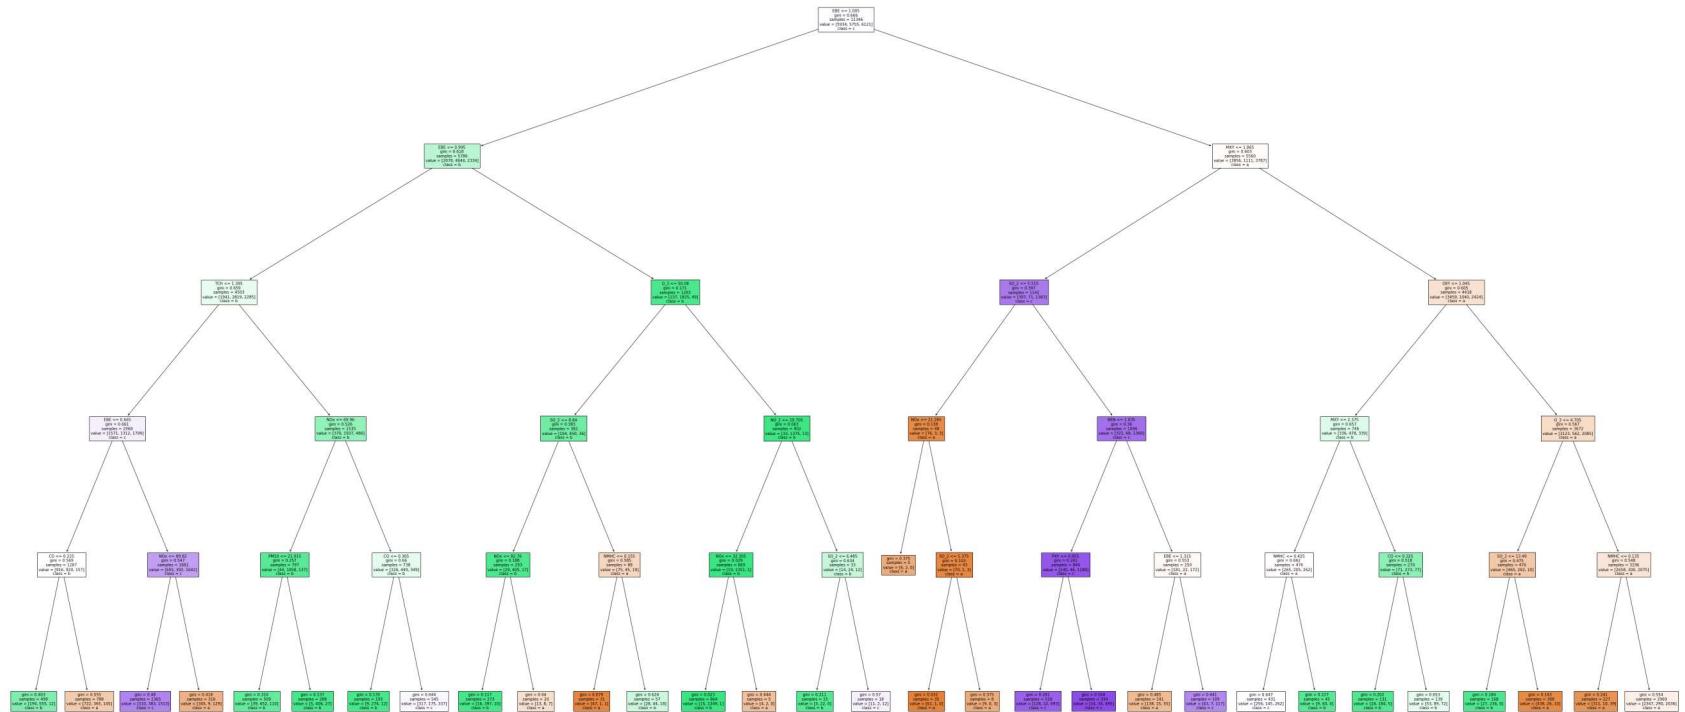
```
Out[62]: [Text(2241.3, 1993.2, 'EBE <= 1.005\ngini = 0.666\nsamples = 11346\nvalue = [5934, 5755, 6121]\nnclass = c'),  
Text(1190.4, 1630.800000000002, 'EBE <= 0.995\ngini = 0.618\nsamples = 5786\nvalue = [2078, 4644, 2334]\nnclass = b'),  
Text(595.2, 1268.4, 'TCH <= 1.395\ngini = 0.659\nsamples = 4503\nvalue = [1941, 2819, 2285]\nnclass = b'),  
Text(297.6, 906.0, 'EBE <= 0.665\ngini = 0.661\nsamples = 2968\nvalue = [1571, 1312, 1799]\nnclass = c'),  
Text(148.8, 543.5999999999999, 'CO <= 0.225\ngini = 0.569\nsamples = 1287\nvalue = [916, 920, 157]\nnclass = b'),  
Text(74.4, 181.1999999999982, 'gini = 0.403\nsamples = 499\nvalue = [194, 555, 12]\nnclass = b'),  
Text(223.200000000002, 181.1999999999982, 'gini = 0.555\nsamples = 788\nvalue = [722, 365, 145]\nnclass = a'),  
Text(446.400000000003, 543.5999999999999, 'NOx <= 89.82\ngini = 0.547\nsamples = 1681\nvalue = [655, 392, 1642]\nnclass = c'),  
Text(372.0, 181.1999999999982, 'gini = 0.48\nsamples = 1365\nvalue = [310, 383, 1513]\nnclass = c'),  
Text(520.800000000001, 181.1999999999982, 'gini = 0.418\nsamples = 316\nvalue = [345, 9, 129]\nnclass = a'),  
Text(892.800000000001, 906.0, 'NOx <= 69.96\ngini = 0.526\nsamples = 1535\nvalue = [370, 1507, 486]\nnclass = b'),  
Text(744.0, 543.5999999999999, 'PM10 <= 21.015\ngini = 0.257\nsamples = 797\nvalue = [44, 1058, 137]\nnclass = b'),  
Text(669.6, 181.1999999999982, 'gini = 0.316\nsamples = 509\nvalue = [39, 652, 110]\nnclass = b'),  
Text(818.400000000001, 181.1999999999982, 'gini = 0.137\nsamples = 288\nvalue = [5, 406, 27]\nnclass = b'),  
Text(1041.600000000001, 543.5999999999999, 'CO <= 0.305\ngini = 0.66\nsamples = 738\nvalue = [326, 449, 349]\nnclass = b'),  
Text(967.2, 181.1999999999982, 'gini = 0.135\nsamples = 193\nvalue = [9, 274, 12]\nnclass = b'),  
Text(1116.0, 181.1999999999982, 'gini = 0.644\nsamples = 545\nvalue = [317, 175, 337]\nnclass = c'),  
Text(1785.600000000001, 1268.4, 'O_3 <= 50.08\ngini = 0.171\nsamples = 1283\nvalue = [137, 1825, 49]\nnclass = b'),  
Text(1488.0, 906.0, 'SO_2 <= 8.84\ngini = 0.383\nsamples = 381\nvalue = [104, 450, 36]\nnclass = b'),  
Text(1339.2, 543.5999999999999, 'NOx <= 82.76\ngini = 0.188\nsamples = 293\nvalue = [29, 405, 17]\nnclass = b'),  
Text(1264.800000000002, 181.1999999999982, 'gini = 0.117\nsamples = 273\nvalue = [16, 397, 10]\nnclass = b'),  
Text(1413.600000000001, 181.1999999999982, 'gini = 0.64\nsamples = 20\nvalue = [13, 8, 7]\nnclass = a'),  
Text(1636.800000000002, 543.5999999999999, 'NMHC <= 0.155\ngini = 0.585\nsamples = 88\nvalue = [75, 45, 19]\nnclass = a'),  
Text(1562.4, 181.1999999999982, 'gini = 0.079\nsamples = 31\nvalue = [47, 1, 1]\nnclass = a'),  
Text(1711.2, 181.1999999999982, 'gini = 0.624\nsamples = 57\nvalue = [28, 44, 18]\nnclass = b'),  
Text(2083.200000000003, 906.0, 'NO_2 <= 29.705\ngini = 0.063\nsamples = 902\nvalue = [33, 1375, 13]\nnclass = b'),  
Text(1934.4, 543.5999999999999, 'NOx <= 32.305\ngini = 0.029\nsamples = 869\nvalue = [19, 1351, 1]\nnclass = b'),  
Text(1860.000000000002, 181.1999999999982, 'gini = 0.023\nsamples = 864\nvalue = [15, 1349, 1]\nnclass =
```

b'),
Text(2008.800000000002, 181.19999999999982, 'gini = 0.444\nsamples = 5\nvalue = [4, 2, 0]\nclass = a'),
Text(2232.0, 543.599999999999, 'SO_2 <= 6.485\ngini = 0.634\nsamples = 33\nvalue = [14, 24, 12]\nclass = b'),
Text(2157.600000000004, 181.19999999999982, 'gini = 0.211\nsamples = 15\nvalue = [3, 22, 0]\nclass = b'),
Text(2306.4, 181.19999999999982, 'gini = 0.57\ngsamples = 18\nvalue = [11, 2, 12]\nclass = c'),
Text(3292.200000000003, 1630.800000000002, 'MXY <= 1.865\ngini = 0.603\nsamples = 5560\nvalue = [3856, 11
11, 3787]\nclass = a'),
Text(2715.600000000004, 1268.4, 'SO_2 <= 5.515\ngini = 0.397\nsamples = 1142\nvalue = [397, 71, 1363]\ncla
ss = c'),
Text(2455.200000000003, 906.0, 'NOx <= 21.285\ngini = 0.138\nsamples = 48\nvalue = [76, 3, 3]\nclass =
a'),
Text(2380.8, 543.599999999999, 'gini = 0.375\nsamples = 5\nvalue = [6, 2, 0]\nclass = a'),
Text(2529.600000000004, 543.599999999999, 'SO_2 <= 5.375\ngini = 0.103\nsamples = 43\nvalue = [70, 1, 3]
\nclass = a'),
Text(2455.200000000003, 181.19999999999982, 'gini = 0.032\nsamples = 35\nvalue = [61, 1, 0]\nclass = a'),
Text(2604.0, 181.19999999999982, 'gini = 0.375\nsamples = 8\nvalue = [9, 0, 3]\nclass = a'),
Text(2976.0, 906.0, 'BEN <= 1.035\ngini = 0.36\nsamples = 1094\nvalue = [321, 68, 1360]\nclass = c'),
Text(2827.200000000003, 543.599999999999, 'PXY <= 0.855\ngini = 0.241\nsamples = 844\nvalue = [140, 46, 1
188]\nclass = c'),
Text(2752.8, 181.19999999999982, 'gini = 0.281\nsamples = 510\nvalue = [126, 12, 693]\nclass = c'),
Text(2901.600000000004, 181.19999999999982, 'gini = 0.164\nsamples = 334\nvalue = [14, 34, 495]\nclass =
c'),
Text(3124.8, 543.599999999999, 'EBE <= 1.315\ngini = 0.553\nsamples = 250\nvalue = [181, 22, 172]\nclass =
a'),
Text(3050.4, 181.19999999999982, 'gini = 0.485\nsamples = 141\nvalue = [138, 15, 55]\nclass = a'),
Text(3199.200000000003, 181.19999999999982, 'gini = 0.441\nsamples = 109\nvalue = [43, 7, 117]\nclass =
c'),
Text(3868.8, 1268.4, 'OXY <= 1.045\ngini = 0.605\nsamples = 4418\nvalue = [3459, 1040, 2424]\nclass = a'),
Text(3571.200000000003, 906.0, 'MXY <= 2.375\ngini = 0.657\nsamples = 746\nvalue = [336, 478, 339]\nclass =
b'),
Text(3422.4, 543.599999999999, 'NMHC <= 0.425\ngini = 0.662\nsamples = 476\nvalue = [265, 205, 262]\nclass =
a'),
Text(3348.000000000005, 181.19999999999982, 'gini = 0.647\nsamples = 431\nvalue = [256, 145, 262]\nclass =
c'),
Text(3496.8, 181.19999999999982, 'gini = 0.227\nsamples = 45\nvalue = [9, 60, 0]\nclass = b'),
Text(3720.000000000005, 543.599999999999, 'CO <= 0.325\ngini = 0.518\nsamples = 270\nvalue = [71, 273, 7
7]\nclass = b'),
Text(3645.600000000004, 181.19999999999982, 'gini = 0.202\nsamples = 131\nvalue = [18, 184, 5]\nclass =
b'),
Text(3794.4, 181.19999999999982, 'gini = 0.653\nsamples = 139\nvalue = [53, 89, 72]\nclass = b'),
Text(4166.400000000001, 906.0, 'O_3 <= 4.705\ngini = 0.567\nsamples = 3672\nvalue = [3123, 562, 2085]\ncl
ass = a'),

```

Text(4017.600000000004, 543.5999999999999, 'SO_2 <= 13.49\ngini = 0.475\nsamples = 476\nvalue = [465, 262, 10]\nclass = a'),
Text(3943.200000000003, 181.1999999999982, 'gini = 0.184\nsamples = 168\nvalue = [27, 236, 0]\nclass = b'),
Text(4092.000000000005, 181.1999999999982, 'gini = 0.143\nsamples = 308\nvalue = [438, 26, 10]\nclass = a'),
Text(4315.200000000001, 543.5999999999999, 'NMHC <= 0.135\ngini = 0.548\nsamples = 3196\nvalue = [2658, 300, 2075]\nclass = a'),
Text(4240.8, 181.1999999999982, 'gini = 0.241\nsamples = 227\nvalue = [311, 10, 39]\nclass = a'),
Text(4389.6, 181.1999999999982, 'gini = 0.554\nsamples = 2969\nvalue = [2347, 290, 2036]\nclass = a')]

```



Conclusion

Accuracy

Linear Regression: 0.15930303157272896

Ridge Regression: 0.15925082658615208

Lasso Regression: 0.012581248316343108

ElasticNet Regression: 0.06773347039258326

Logistic Regression: 0.8146838030106512

Random Forest: 0.8211117349803481

Random Forest is suitable for this dataset