

Importing Libraries

In [1]:

```
1 import numpy as np
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
```

Importing Datasets

In [2]:

```

1 df=pd.read_csv("madrid_2010")
2 df

```

Out[2]:

| | | date | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM10 | PM25 | PXY | SO_2 | TCH | TOL |
|--------|-----|---------------------|------|------|------|-----|------|------------|------------|-----|-----------|-----------|-----------|-----|-------|------|------|
| 0 | | 2010-03-01 01:00:00 | NaN | 0.29 | NaN | NaN | NaN | 25.090000 | 29.219999 | NaN | 68.930000 | NaN | NaN | NaN | 10.15 | NaN | NaN |
| 1 | | 2010-03-01 01:00:00 | NaN | 0.27 | NaN | NaN | NaN | 24.879999 | 30.040001 | NaN | NaN | NaN | NaN | NaN | 12.24 | NaN | NaN |
| 2 | | 2010-03-01 01:00:00 | NaN | 0.28 | NaN | NaN | NaN | 17.410000 | 20.540001 | NaN | 72.120003 | NaN | NaN | NaN | NaN | NaN | NaN |
| 3 | | 2010-03-01 01:00:00 | 0.38 | 0.24 | 1.74 | NaN | 0.05 | 15.610000 | 21.080000 | NaN | 72.970001 | 19.410000 | 7.870000 | NaN | 10.06 | 1.52 | 1.49 |
| 4 | | 2010-03-01 01:00:00 | 0.79 | NaN | 1.32 | NaN | NaN | 21.430000 | 26.070000 | NaN | NaN | 24.670000 | 22.030001 | NaN | 10.68 | NaN | 2.88 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 209443 | | 2010-08-01 00:00:00 | NaN | 0.55 | NaN | NaN | NaN | 125.000000 | 219.899994 | NaN | 25.379999 | NaN | NaN | NaN | NaN | NaN | NaN |
| 209444 | | 2010-08-01 00:00:00 | NaN | 0.27 | NaN | NaN | NaN | 45.709999 | 47.410000 | NaN | NaN | 51.259998 | NaN | NaN | 7.26 | NaN | NaN |
| 209445 | | 2010-08-01 00:00:00 | NaN | NaN | NaN | NaN | 0.24 | 46.560001 | 49.040001 | NaN | 46.250000 | NaN | NaN | NaN | NaN | 1.47 | NaN |
| 209446 | | 2010-08-01 00:00:00 | NaN | NaN | NaN | NaN | NaN | 46.770000 | 50.119999 | NaN | 77.709999 | NaN | NaN | NaN | NaN | NaN | NaN |
| 209447 | | 2010-08-01 00:00:00 | 0.92 | 0.43 | 0.71 | NaN | 0.25 | 76.330002 | 88.190002 | NaN | 52.259998 | 47.150002 | 26.860001 | NaN | 7.03 | 1.44 | 2.87 |

209448 rows × 17 columns



Data Cleaning and Data Preprocessing

```
In [3]: 1 df=df.dropna()
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'],
       dtype='object')
```

```
In [5]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6666 entries, 11 to 191927
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      6666 non-null   object 
 1   BEN        6666 non-null   float64
 2   CO         6666 non-null   float64
 3   EBE        6666 non-null   float64
 4   MXY        6666 non-null   float64
 5   NMHC       6666 non-null   float64
 6   NO_2       6666 non-null   float64
 7   NOx        6666 non-null   float64
 8   OXY        6666 non-null   float64
 9   O_3         6666 non-null   float64
 10  PM10       6666 non-null   float64
 11  PM25       6666 non-null   float64
 12  PXY        6666 non-null   float64
 13  SO_2       6666 non-null   float64
 14  TCH        6666 non-null   float64
 15  TOL        6666 non-null   float64
 16  station    6666 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 937.4+ KB
```

```
In [6]: 1 data=df[['CO' , 'station']]  
2 data
```

Out[6]:

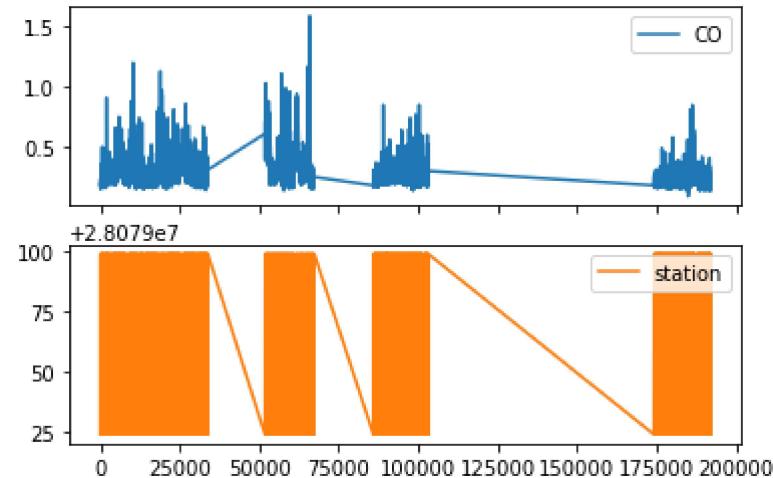
| | CO | station |
|--------|------|----------|
| 11 | 0.18 | 28079024 |
| 23 | 0.23 | 28079099 |
| 35 | 0.17 | 28079024 |
| 47 | 0.21 | 28079099 |
| 59 | 0.16 | 28079024 |
| ... | ... | ... |
| 191879 | 0.26 | 28079099 |
| 191891 | 0.16 | 28079024 |
| 191903 | 0.28 | 28079099 |
| 191915 | 0.16 | 28079024 |
| 191927 | 0.25 | 28079099 |

6666 rows × 2 columns

Line chart

```
In [7]: 1 data.plot.line(subplots=True)
```

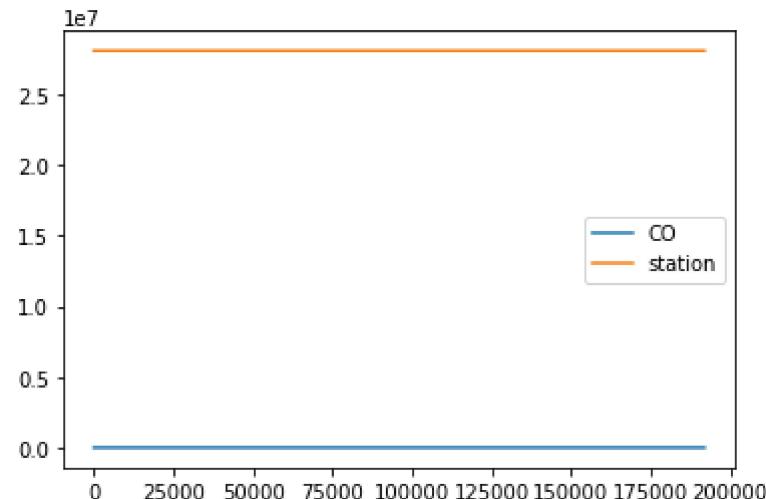
```
Out[7]: array([<AxesSubplot:>, <AxesSubplot:>], dtype=object)
```



Line chart

```
In [8]: 1 data.plot.line()
```

```
Out[8]: <AxesSubplot:>
```

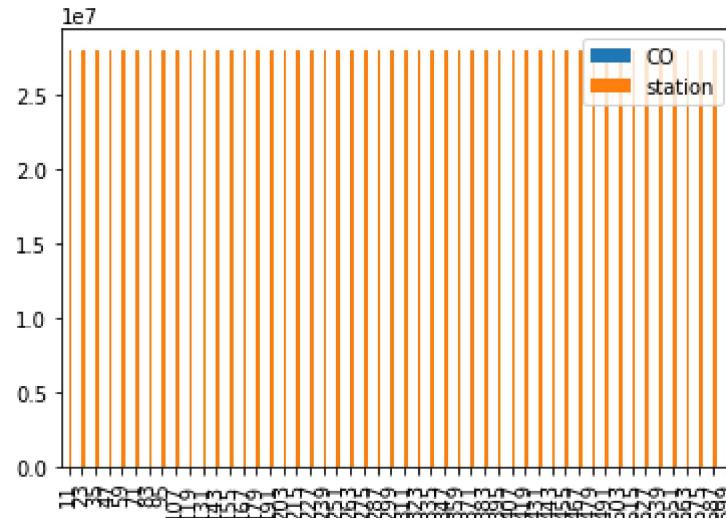


Bar chart

```
In [9]: 1 b=data[0:50]
```

```
In [10]: 1 b.plot.bar()
```

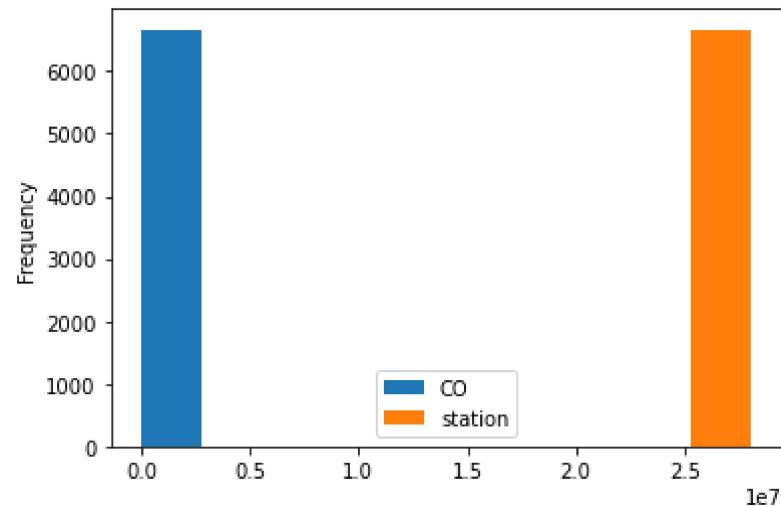
```
Out[10]: <AxesSubplot:>
```



Histogram

```
In [11]: 1 data.plot.hist()
```

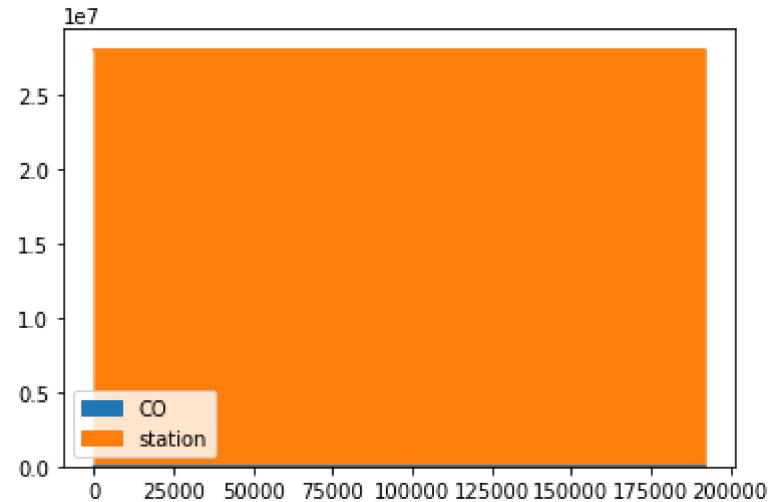
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



Area chart

In [12]: 1 data.plot.area()

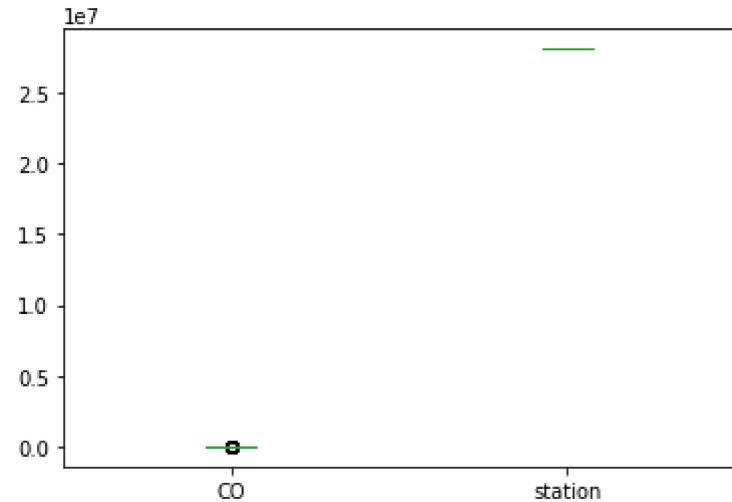
Out[12]: <AxesSubplot:>



Box chart

In [13]: 1 data.plot.box()

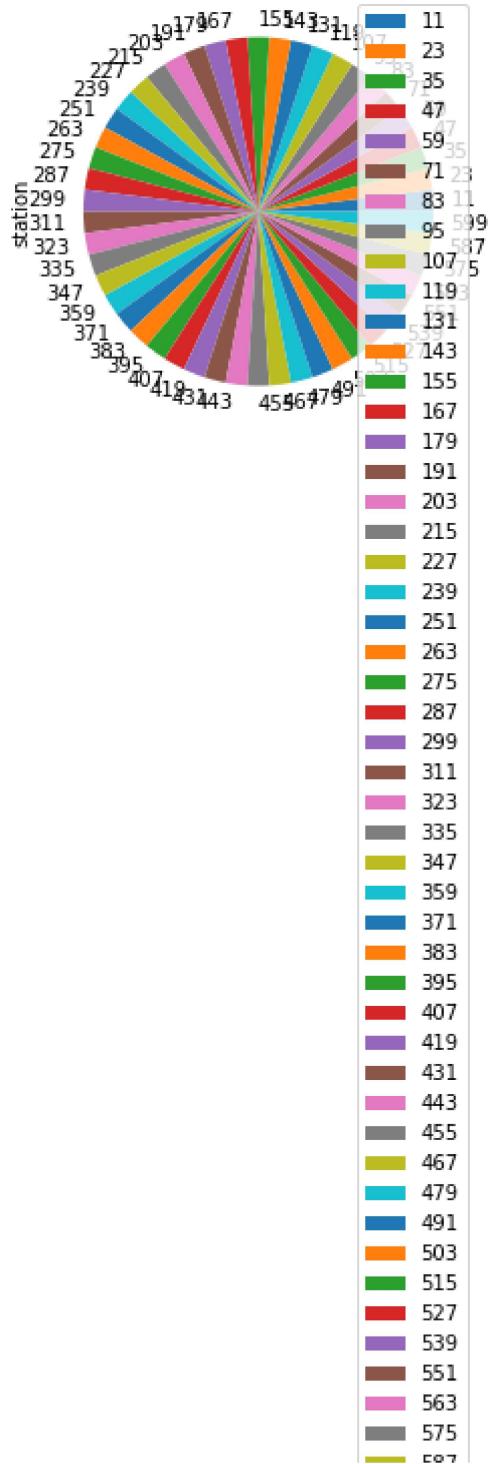
Out[13]: <AxesSubplot:>

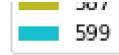


Pie chart

```
In [14]: 1 b.plot.pie(y='station' )
```

```
Out[14]: <AxesSubplot:ylabel='station'>
```

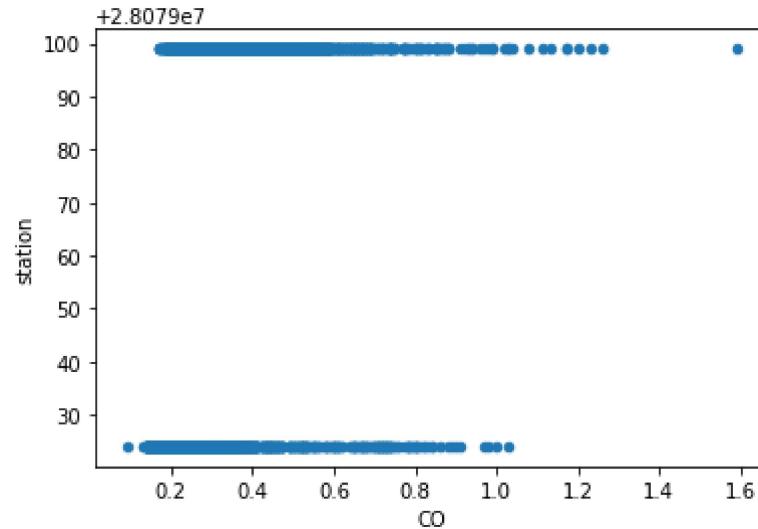





Scatter chart

```
In [15]: 1 data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



```
In [16]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6666 entries, 11 to 191927
Data columns (total 17 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   date      6666 non-null    object  
 1   BEN        6666 non-null    float64 
 2   CO         6666 non-null    float64 
 3   EBE        6666 non-null    float64 
 4   MXY        6666 non-null    float64 
 5   NMHC       6666 non-null    float64 
 6   NO_2       6666 non-null    float64 
 7   NOx        6666 non-null    float64 
 8   OXY        6666 non-null    float64 
 9   O_3         6666 non-null    float64 
 10  PM10       6666 non-null    float64 
 11  PM25       6666 non-null    float64 
 12  PXY        6666 non-null    float64 
 13  SO_2       6666 non-null    float64 
 14  TCO        6666 non-null    float64
```

```
In [17]: 1 df.describe()
```

Out[17]:

| | BEN | CO | EBE | MXY | NMHC | NO_2 | NOx | OXY | O_3 | PM |
|-------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|
| count | 6666.000000 | 6666.000000 | 6666.000000 | 6666.000000 | 6666.000000 | 6666.000000 | 6666.000000 | 6666.000000 | 6666.000000 | 6666.000000 |
| mean | 0.648425 | 0.296280 | 0.840585 | 0.839959 | 0.243378 | 33.888744 | 47.540617 | 0.916668 | 56.246101 | 17.1473 |
| std | 0.395346 | 0.133296 | 0.508031 | 0.382263 | 0.115730 | 23.465169 | 41.230578 | 0.192521 | 30.380535 | 13.6790 |
| min | 0.170000 | 0.090000 | 0.140000 | 0.110000 | 0.000000 | 1.290000 | 2.760000 | 0.200000 | 0.600000 | 1.3200 |
| 25% | 0.380000 | 0.200000 | 0.470000 | 0.590000 | 0.180000 | 15.752500 | 19.442501 | 1.000000 | 31.740000 | 8.9700 |
| 50% | 0.540000 | 0.260000 | 0.755000 | 1.000000 | 0.220000 | 29.320000 | 36.770000 | 1.000000 | 57.750000 | 14.0000 |
| 75% | 0.810000 | 0.340000 | 1.000000 | 1.000000 | 0.280000 | 47.657500 | 62.102501 | 1.000000 | 79.489998 | 21.1500 |
| max | 5.110000 | 1.590000 | 5.190000 | 6.810000 | 0.930000 | 133.399994 | 409.299988 | 2.760000 | 161.100006 | 174.6999 |

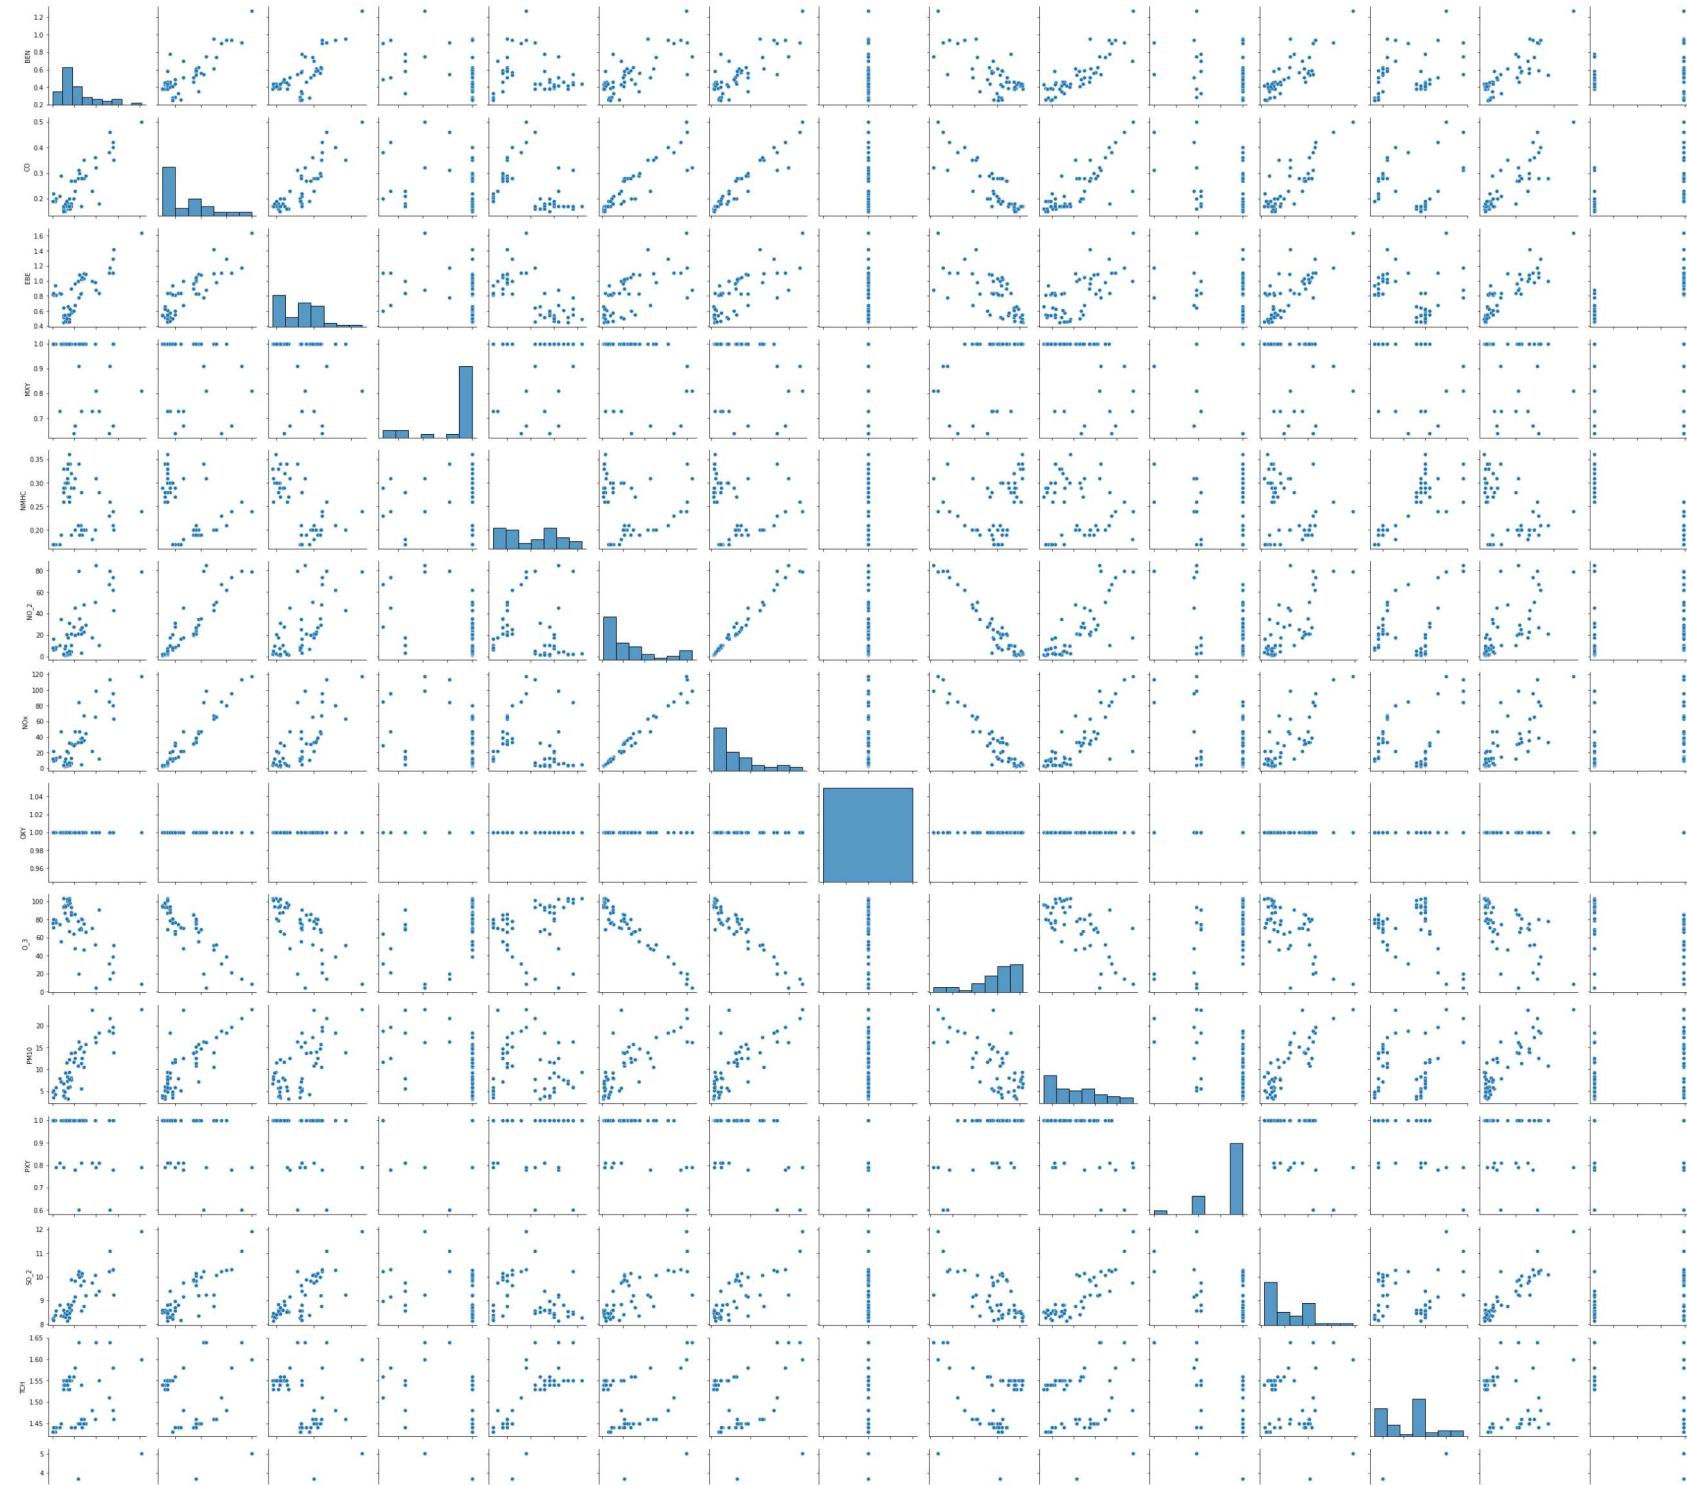
In [18]:

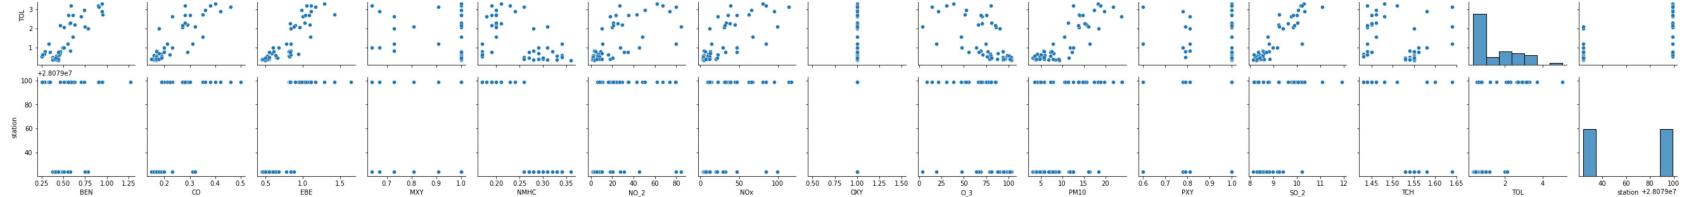
```
1 df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2         'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

EDA AND VISUALIZATION

```
In [19]: 1 sns.pairplot(df1[0:50])
```

```
Out[19]: <seaborn.axisgrid.PairGrid at 0x2868f6c7220>
```

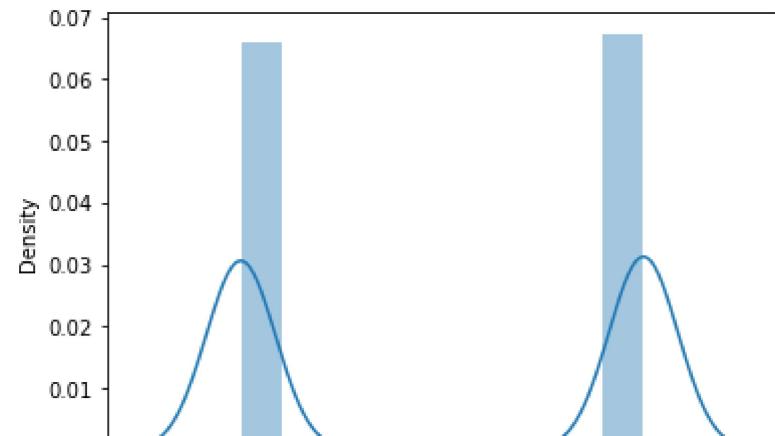





```
In [20]: 1 sns.distplot(df1['station'])
```

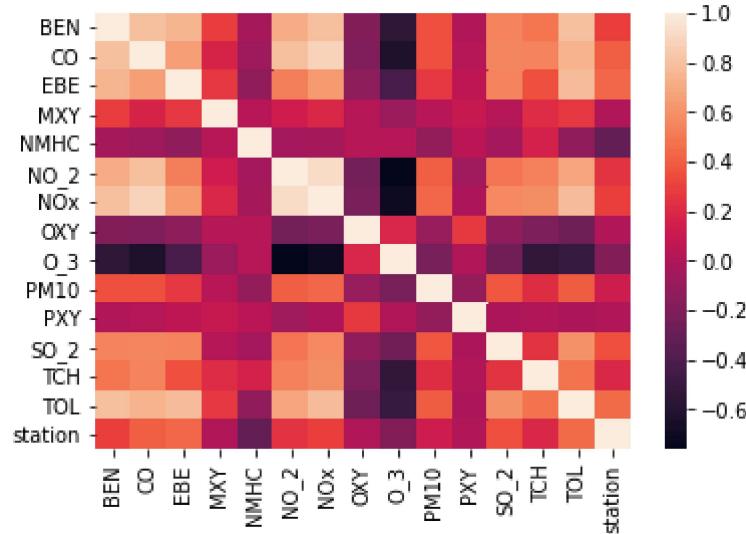
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>
```



```
In [21]: 1 sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: 1 x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2           'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 y=df['station']
```

```
In [23]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

Linear Regression

```
In [24]: 1 from sklearn.linear_model import LinearRegression  
2 lr=LinearRegression()  
3 lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: 1 lr.intercept_
```

```
Out[25]: 28078931.02099493
```

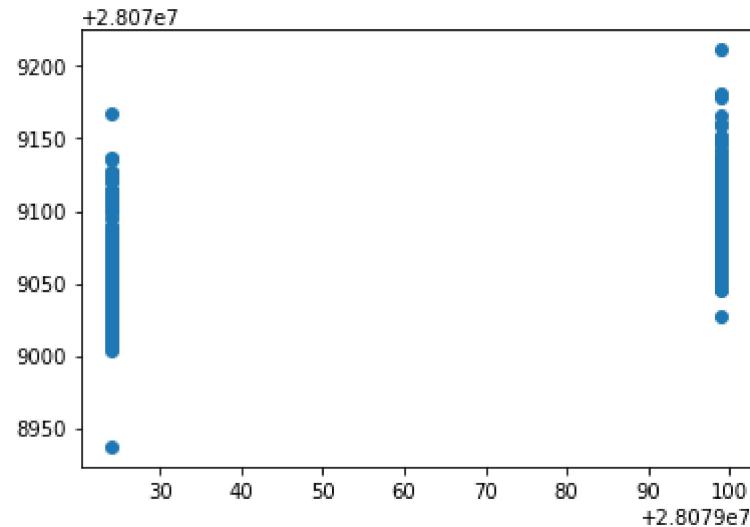
```
In [26]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
2 coeff
```

```
Out[26]:
```

| | Co-efficient |
|------|--------------|
| BEN | -35.980070 |
| CO | 174.138112 |
| EBE | 16.081525 |
| MXY | -9.700544 |
| NMHC | -70.949113 |
| NO_2 | 0.278852 |
| NOx | -0.637720 |
| OXY | 33.041587 |
| O_3 | 0.067798 |
| PM10 | -0.165901 |
| PXY | -5.456597 |
| SO_2 | 1.700364 |
| TCH | 48.133074 |
| TOL | 10.323340 |

```
In [27]: 1 prediction =lr.predict(x_test)
2 plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x2869d6ee880>
```



ACCURACY

```
In [28]: 1 lr.score(x_test,y_test)
```

```
Out[28]: 0.41834875044261377
```

```
In [29]: 1 lr.score(x_train,y_train)
```

```
Out[29]: 0.4313187068979043
```

Ridge and Lasso

```
In [30]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: 1 rr=Ridge(alpha=10)
          2 rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

Accuracy(Ridge)

```
In [32]: 1 rr.score(x_test,y_test)
```

```
Out[32]: 0.41071757471844017
```

```
In [33]: 1 rr.score(x_train,y_train)
```

```
Out[33]: 0.4183715947196103
```

```
In [34]: 1 la=Lasso(alpha=10)
          2 la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

Accuracy(Lasso)

```
In [35]: 1 la.score(x_train,y_train)
```

```
Out[35]: 0.1793088808576867
```

```
In [36]: 1 la.score(x_test,y_test)
```

```
Out[36]: 0.18729059757800426
```

```
In [37]: 1 from sklearn.linear_model import ElasticNet
          2 en=ElasticNet()
          3 en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: 1 en.coef_
```

```
Out[38]: array([-0.          ,  0.20397676,  2.73065235, -1.09460886, -1.23331305,
                 0.          , -0.097484  ,  0.65885519, -0.03918052, -0.11680219,
                 -0.         ,  2.44426388,  0.          ,  7.22599835])
```

```
In [39]: 1 en.intercept_
```

```
Out[39]: 28079028.297450814
```

```
In [40]: 1 prediction=en.predict(x_test)
```

```
In [41]: 1 en.score(x_test,y_test)
```

```
Out[41]: 0.2426219931577691
```

Evaluation Metrics

```
In [42]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
30.542128097679466
1064.8540698087513
32.63210182946773
```

Logistic Regression

```
In [43]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [44]: 1 feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
2                  'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
3 target_vector=df['station']
```

```
In [45]: 1 feature_matrix.shape
```

```
Out[45]: (6666, 14)
```

```
In [46]: 1 target_vector.shape
```

```
Out[46]: (6666,)
```

```
In [47]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [48]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: 1 logr=LogisticRegression(max_iter=10000)  
2 logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: 1 observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]
```

```
In [51]: 1 prediction=logr.predict(observation)  
2 print(prediction)
```

```
[28079099]
```

```
In [52]: 1 logr.classes_
```

```
Out[52]: array([28079024, 28079099], dtype=int64)
```

```
In [53]: 1 logr.score(fs,target_vector)
```

```
Out[53]: 0.8660366036603661
```

```
In [54]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[54]: 0.0
```

```
In [55]: 1 logr.predict_proba(observation)
```

```
Out[55]: array([[0., 1.]])
```

Random Forest

```
In [56]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: 1 rfc=RandomForestClassifier()
2 rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: 1 parameters={'max_depth':[1,2,3,4,5],
2                 'min_samples_leaf':[5,10,15,20,25],
3                 'n_estimators':[10,20,30,40,50]
4 }
```

```
In [59]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
param_grid={'max_depth': [1, 2, 3, 4, 5],
'min_samples_leaf': [5, 10, 15, 20, 25],
'n_estimators': [10, 20, 30, 40, 50]},
scoring='accuracy')
```

```
In [60]: 1 grid_search.best_score_
```

```
Out[60]: 0.9258465495070725
```

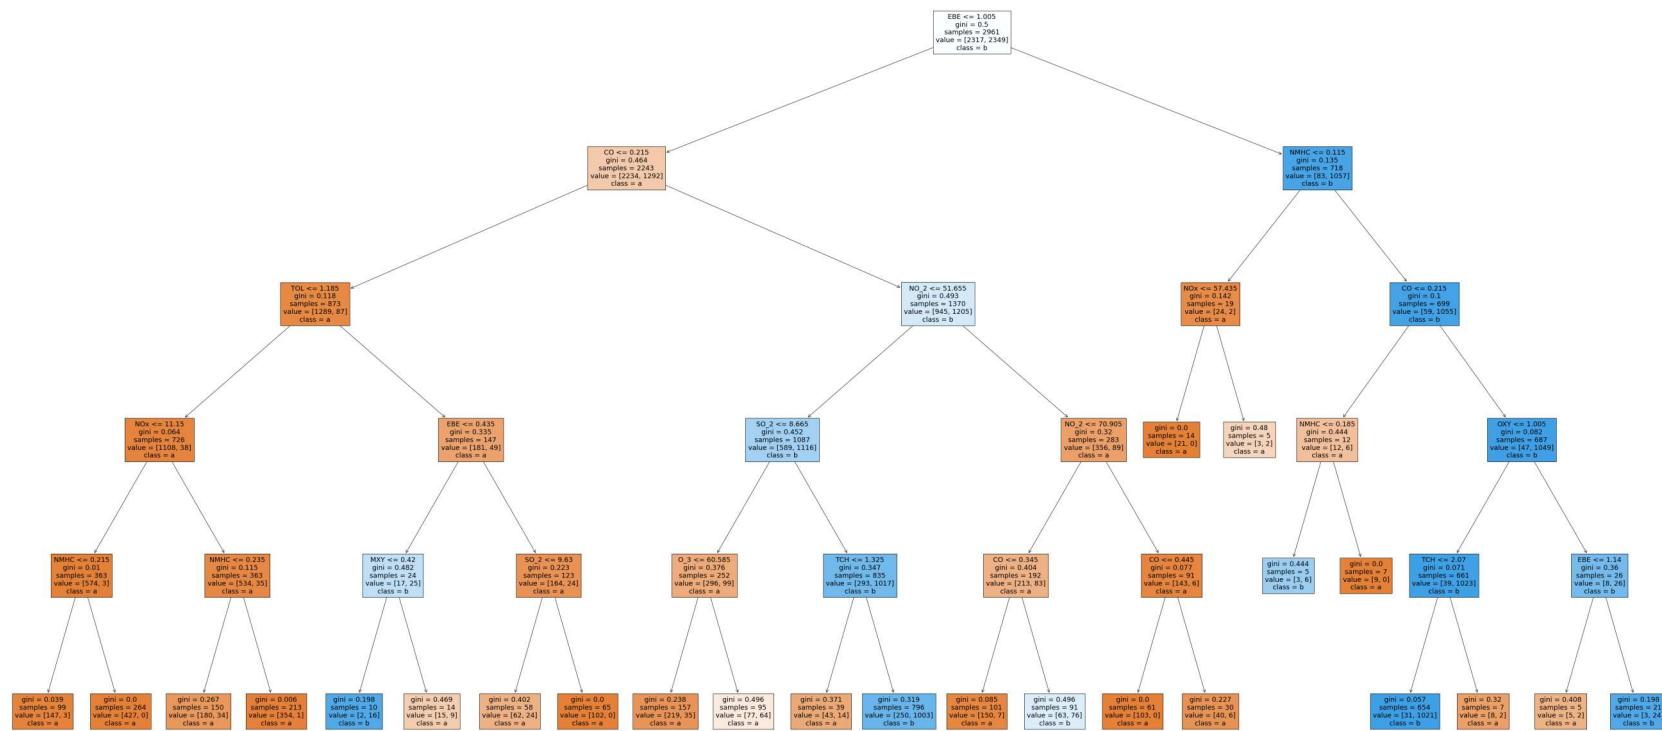
```
In [61]: 1 rfc_best=grid_search.best_estimator_
```

In [62]:

```
1 from sklearn.tree import plot_tree
2
3 plt.figure(figsize=(80,40))
4 plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'],filled=True)
```

Out[62]: [Text(2582.3720930232557, 1993.2, 'EBE <= 1.005\ngini = 0.5\nsamples = 2961\nvalue = [2317, 2349]\nclass = b'),
Text(1661.0232558139535, 1630.8000000000002, 'CO <= 0.215\ngini = 0.464\nsamples = 2243\nvalue = [2234, 1292]\nclass = a'),
Text(830.5116279069767, 1268.4, 'TOL <= 1.185\ngini = 0.118\nsamples = 873\nvalue = [1289, 87]\nclass = a'),
Text(415.25581395348837, 906.0, 'NOx <= 11.15\ngini = 0.064\nsamples = 726\nvalue = [1108, 38]\nclass = a'),
Text(207.62790697674419, 543.5999999999999, 'NMHC <= 0.215\ngini = 0.01\nsamples = 363\nvalue = [574, 3]\nclass = a'),
Text(103.81395348837209, 181.19999999999982, 'gini = 0.039\nsamples = 99\nvalue = [147, 3]\nclass = a'),
Text(311.4418604651163, 181.19999999999982, 'gini = 0.0\nsamples = 264\nvalue = [427, 0]\nclass = a'),
Text(622.8837209302326, 543.5999999999999, 'NMHC <= 0.235\ngini = 0.115\nsamples = 363\nvalue = [534, 35]\nclass = a'),
Text(519.0697674418604, 181.19999999999982, 'gini = 0.267\nsamples = 150\nvalue = [180, 34]\nclass = a'),
Text(726.6976744186046, 181.19999999999982, 'gini = 0.006\nsamples = 213\nvalue = [354, 1]\nclass = a'),
Text(1245.7674418604652, 906.0, 'EBE <= 0.435\ngini = 0.335\nsamples = 147\nvalue = [181, 49]\nclass = a'),
Text(1038.139534883721, 543.5999999999999, 'MXY <= 0.42\ngini = 0.482\nsamples = 24\nvalue = [17, 25]\nclass = b'),
Text(934.3255813953489, 181.19999999999982, 'gini = 0.198\nsamples = 10\nvalue = [2, 16]\nclass = b'),
Text(1141.953488372093, 181.19999999999982, 'gini = 0.469\nsamples = 14\nvalue = [15, 9]\nclass = a'),
Text(1453.3953488372092, 543.5999999999999, 'SO_2 <= 9.63\ngini = 0.223\nsamples = 123\nvalue = [164, 24]\nclass = a'),
Text(1349.5813953488373, 181.19999999999982, 'gini = 0.402\nsamples = 58\nvalue = [62, 24]\nclass = a'),
Text(1557.2093023255813, 181.19999999999982, 'gini = 0.0\nsamples = 65\nvalue = [102, 0]\nclass = a'),
Text(2491.5348837209303, 1268.4, 'NO_2 <= 51.655\ngini = 0.493\nsamples = 1370\nvalue = [945, 1205]\nclass = b'),
Text(2076.279069767442, 906.0, 'SO_2 <= 8.665\ngini = 0.452\nsamples = 1087\nvalue = [589, 1116]\nclass = b'),
Text(1868.6511627906978, 543.5999999999999, 'O_3 <= 60.585\ngini = 0.376\nsamples = 252\nvalue = [296, 99]\nclass = a'),
Text(1764.8372093023256, 181.19999999999982, 'gini = 0.238\nsamples = 157\nvalue = [219, 35]\nclass = a'),
Text(1972.4651162790697, 181.19999999999982, 'gini = 0.496\nsamples = 95\nvalue = [77, 64]\nclass = a'),
Text(2283.906976744186, 543.5999999999999, 'TCH <= 1.325\ngini = 0.347\nsamples = 835\nvalue = [293, 1017]\nclass = b'),
Text(2180.093023255814, 181.19999999999982, 'gini = 0.371\nsamples = 39\nvalue = [43, 14]\nclass = a'),
Text(2387.720930232558, 181.19999999999982, 'gini = 0.319\nsamples = 796\nvalue = [250, 1003]\nclass = b'),
Text(2906.7906976744184, 906.0, 'NO_2 <= 70.905\ngini = 0.32\nsamples = 283\nvalue = [356, 89]\nclass = a'),
Text(2699.1627906976746, 543.5999999999999, 'CO <= 0.345\ngini = 0.404\nsamples = 192\nvalue = [213, 83]\nclass = a'),
Text(2595.3488372093025, 181.19999999999982, 'gini = 0.085\nsamples = 101\nvalue = [150, 7]\nclass = a'),
Text(2802.9767441860463, 181.19999999999982, 'gini = 0.496\nsamples = 91\nvalue = [63, 76]\nclass = b')]

```
Text(3114.4186046511627, 543.5999999999999, 'CO <= 0.445\ngini = 0.077\nsamples = 91\nvalue = [143, 6]\nclass = a'),  
Text(3010.6046511627906, 181.19999999999982, 'gini = 0.0\nsamples = 61\nvalue = [103, 0]\nclass = a'),  
Text(3218.232558139535, 181.19999999999982, 'gini = 0.227\nsamples = 30\nvalue = [40, 6]\nclass = a'),  
Text(3503.720930232558, 1630.8000000000002, 'NMHC <= 0.115\ngini = 0.135\nsamples = 718\nvalue = [83, 1057]\nclass = b'),  
Text(3218.232558139535, 1268.4, 'NOx <= 57.435\ngini = 0.142\nsamples = 19\nvalue = [24, 2]\nclass = a'),  
Text(3114.4186046511627, 906.0, 'gini = 0.0\nsamples = 14\nvalue = [21, 0]\nclass = a'),  
Text(3322.046511627907, 906.0, 'gini = 0.48\nsamples = 5\nvalue = [3, 2]\nclass = a'),  
Text(3789.2093023255816, 1268.4, 'CO <= 0.215\ngini = 0.1\nsamples = 699\nvalue = [59, 1055]\nclass = b'),  
Text(3529.6744186046512, 906.0, 'NMHC <= 0.185\ngini = 0.444\nsamples = 12\nvalue = [12, 6]\nclass = a'),  
Text(3425.860465116279, 543.5999999999999, 'gini = 0.444\nsamples = 5\nvalue = [3, 6]\nclass = b'),  
Text(3633.4883720930234, 543.5999999999999, 'gini = 0.0\nsamples = 7\nvalue = [9, 0]\nclass = a'),  
Text(4048.7441860465115, 906.0, 'OXY <= 1.005\ngini = 0.082\nsamples = 687\nvalue = [47, 1049]\nclass = b'),  
Text(3841.1162790697676, 543.5999999999999, 'TCH <= 2.07\ngini = 0.071\nsamples = 661\nvalue = [39, 1023]\nclass = b'),  
Text(3737.3023255813955, 181.19999999999982, 'gini = 0.057\nsamples = 654\nvalue = [31, 1021]\nclass = b'),  
Text(3944.9302325581393, 181.19999999999982, 'gini = 0.32\nsamples = 7\nvalue = [8, 2]\nclass = a'),  
Text(4256.372093023256, 543.5999999999999, 'EBE <= 1.14\ngini = 0.36\nsamples = 26\nvalue = [8, 26]\nclass = b'),  
Text(4152.558139534884, 181.19999999999982, 'gini = 0.408\nsamples = 5\nvalue = [5, 2]\nclass = a'),  
Text(4360.186046511628, 181.19999999999982, 'gini = 0.198\nsamples = 21\nvalue = [3, 24]\nclass = b')]
```



Conclusion

Accuracy

Linear Regression

Ridge Regression: 0.4183715947196103

Lasso Regression: 0.1793088808576867

ElasticNet Regression: 0.2426219931577691

Logistic Regression:0.8660366036603661

Random Forest:0.9258465495070725

Random Forest is suitable for this dataset