

## Importing Libraries

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
```

## Importing Datasets

```
In [2]: 1 df=pd.read_csv("stations1")
        2 df
```

Out[2]:

	id	name	address	lon	lat	elevation
0	28079004	Pza. de España	Plaza de España	-3.712247	40.423853	635
1	28079008	Escuelas Aguirre	Entre C/ Alcalá y C/ O' Donell	-3.682319	40.421564	670
2	28079011	Avda. Ramón y Cajal	Avda. Ramón y Cajal esq. C/ Príncipe de Vergara	-3.677356	40.451475	708
3	28079016	Arturo Soria	C/ Arturo Soria esq. C/ Vizconde de los Asilos	-3.639233	40.440047	693
4	28079017	Villaverde	C/. Juan Peñalver	-3.713322	40.347139	604
5	28079018	Farolillo	Calle Farolillo - C/Ervigio	-3.731853	40.394781	630
6	28079024	Casa de Campo	Casa de Campo (Terminal del Teleférico)	-3.747347	40.419356	642
7	28079027	Barajas Pueblo	C/. Júpiter, 21 (Barajas)	-3.580031	40.476928	621
8	28079035	Pza. del Carmen	Plaza del Carmen esq. Tres Cruces.	-3.703172	40.419208	659
9	28079036	Moratalaz	Avd. Moratalaz esq. Camino de los Vinateros	-3.645306	40.407947	685
10	28079038	Cuatro Caminos	Avda. Pablo Iglesias esq. C/ Marqués de Lema	-3.707128	40.445544	698
11	28079039	Barrio del Pilar	Avd. Betanzos esq. C/ Monforte de Lemos	-3.711542	40.478228	674
12	28079040	Vallecas	C/ Arroyo del Olivar esq. C/ Río Grande.	-3.651522	40.388153	677
13	28079047	Mendez Alvaro	C/ Juan de Mariana / Pza. Amanecer Mendez Alvaro	-3.686825	40.398114	599
14	28079048	Castellana	C/ Jose Gutierrez Abascal	-3.690367	40.439897	676
15	28079049	Parque del Retiro	Paseo Venezuela- Casa de Vacas	-3.682583	40.414444	662
16	28079050	Plaza Castilla	Plaza Castilla (Canal)	-3.688769	40.465572	728
17	28079054	Ensanche de Vallecas	Avda La Gavia / Avda. Las Suertes	-3.612117	40.372933	627
18	28079055	Urb. Embajada	C/ Riaño (Barajas)	-3.580747	40.462531	618
19	28079056	Pza. Fernández Ladreda	Pza. Fernández Ladreda - Avda. Oporto	-3.718728	40.384964	604
20	28079057	Sanchinarro	C/ Princesa de Eboji esq C/ Maria Tudor	-3.660503	40.494208	700
21	28079058	El Pardo	Avda. La Guardia	-3.774611	40.518058	615
22	28079059	Juan Carlos I	Parque Juan Carlos I (frente oficinas mantenim...	-3.609072	40.465250	660
23	28079060	Tres Olivos	Plaza Tres Olivos	-3.689761	40.500589	715

## Data Cleaning and Data Preprocessing

```
In [3]: 1 df=df.dropna()
```

```
In [4]: 1 df.columns
```

```
Out[4]: Index(['id', 'name', 'address', 'lon', 'lat', 'elevation'], dtype='object')
```

```
In [5]: 1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24 entries, 0 to 23
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   id          24 non-null    int64
1   name        24 non-null    object
2   address     24 non-null    object
3   lon         24 non-null    float64
4   lat         24 non-null    float64
5   elevation   24 non-null    int64
dtypes: float64(2), int64(2), object(2)
memory usage: 1.3+ KB
```

```
In [6]: 1 data=df[['lon', 'lat', 'elevation']]
        2 data
```

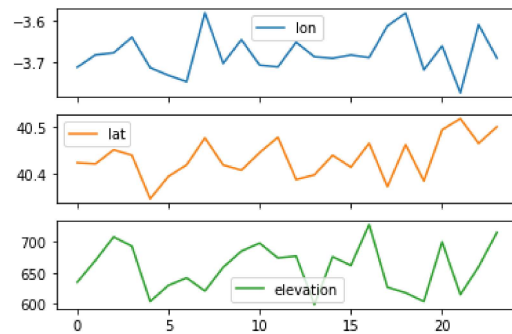
Out[6]:

	lon	lat	elevation
0	-3.712247	40.423853	635
1	-3.682319	40.421564	670
2	-3.677356	40.451475	708
3	-3.639233	40.440047	693
4	-3.713322	40.347139	604
5	-3.731853	40.394781	630
6	-3.747347	40.419356	642
7	-3.580031	40.476928	621
8	-3.703172	40.419208	659
9	-3.645306	40.407947	685
10	-3.707128	40.445544	698
11	-3.711542	40.478228	674
12	-3.651522	40.388153	677
13	-3.686825	40.398114	599
14	-3.690367	40.439897	676
15	-3.682583	40.414444	662
16	-3.688769	40.465572	728
17	-3.612117	40.372933	627
18	-3.580747	40.462531	618
19	-3.718728	40.384964	604
20	-3.660503	40.494208	700
21	-3.774611	40.518058	615
22	-3.609072	40.465250	660
23	-3.689761	40.500589	715

Line chart

```
In [7]: 1 data.plot.line(subplots=True)
```


```
Out[7]: array([<AxesSubplot:~>, <AxesSubplot:~>, <AxesSubplot:~>], dtype=object)
```



## Line chart

```
In [8]: 1 data.plot.line()
```

```
Out[8]: <AxesSubplot:~>
```

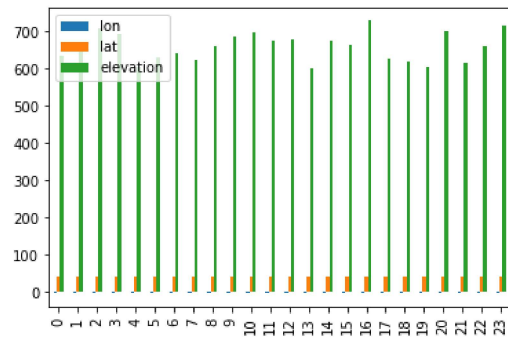


## Bar chart

```
In [9]: 1 b=data[0:50]
```

```
In [10]: 1 b.plot.bar()
```


```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: 1 data.plot.hist()
```

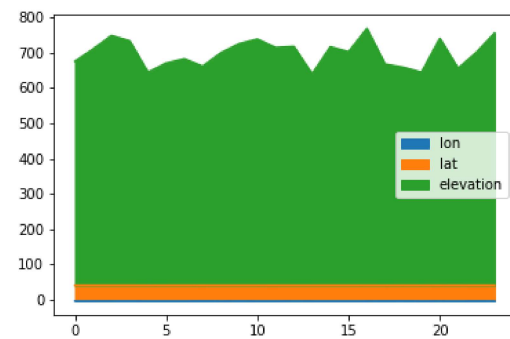
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

```
In [12]: 1 data.plot.area()
```

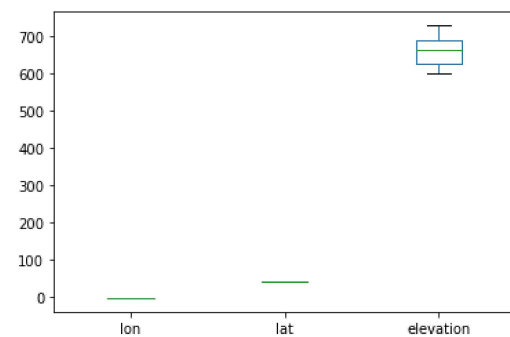
```
Out[12]: <AxesSubplot:>
```



## Box chart

```
In [13]: 1 data.plot.box()
```


```
Out[13]: <AxesSubplot:>
```



## Pie chart

```
In [14]: 1 b.plot.pie(y= 'elevation' )
```


```
Out[14]: <AxesSubplot:ylabel='elevation'>
```



## Scatter chart

```
In [15]: 1 data.plot.scatter(x='lat',y='lon')
```

```
Out[15]: <AxesSubplot:xlabel='lat', ylabel='lon'>
```



In [16]:

```
1 df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 24 entries, 0 to 23
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype  
---  ------  -
0    id          24 non-null    int64   
1    name        24 non-null    object  
2    address     24 non-null    object  
3    lon         24 non-null    float64  
4    lat         24 non-null    float64  
5    elevation   24 non-null    int64   
dtypes: float64(2), int64(2), object(2)
memory usage: 1.3+ KB
```

In [17]:

```
1 df.describe()
```

Out[17]:

	id	lon	lat	elevation
count	2.400000e+01	24.000000	24.000000	24.000000
mean	2.807904e+07	-3.679019	40.434616	658.333333
std	1.799094e+01	0.049324	0.043022	38.295949
min	2.807900e+07	-3.774611	40.347139	599.000000
25%	2.807902e+07	-3.711718	40.405489	625.500000
50%	2.807904e+07	-3.687797	40.431875	661.000000
75%	2.807905e+07	-3.649968	40.465331	687.000000
max	2.807906e+07	-3.580031	40.518058	728.000000

In [18]:

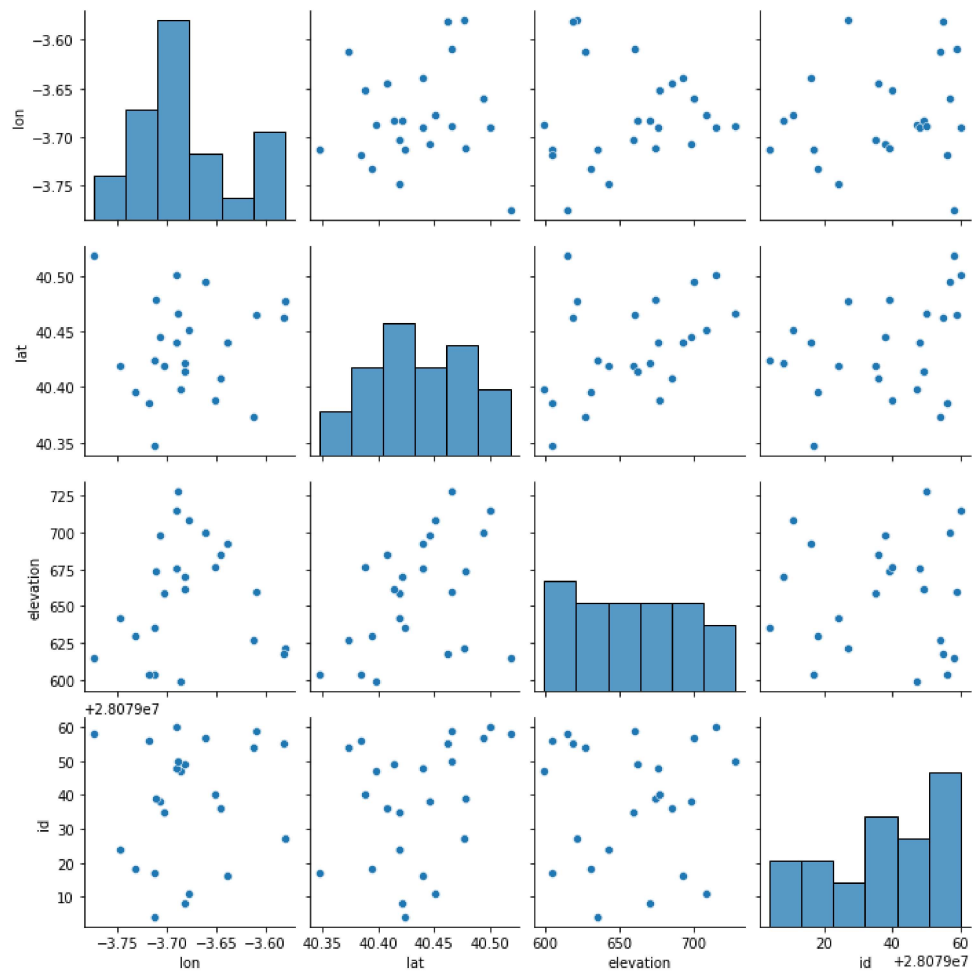
```
1 df1=df[['lon', 'lat', 'elevation', 'id']]
```

## EDA AND VISUALIZATION



```
In [19]: 1 sns.pairplot(df1[0:50])
```


```
Out[19]: <seaborn.axisgrid.PairGrid at 0x18019d7d1f0>
```



```
In [20]: 1 sns.distplot(df1['lon'])
```


C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).  
warnings.warn(msg, FutureWarning)

```
Out[20]: <AxesSubplot:xlabel='lon', ylabel='Density'>
```



```
In [21]: 1 sns.heatmap(df1.corr())
```

```
Out[21]: <AxesSubplot:>
```



## TO TRAIN THE MODEL AND MODEL BUILDING

```
In [22]: 1 x=df[['id']]
2 y=df['elevation']
```

```
In [23]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

```
In [24]: 1 from sklearn.linear_model import LinearRegression
2 lr=LinearRegression()
3 lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: 1 lr.intercept_
```

```
Out[25]: -2129342.4097207235
```


```
In [26]: 1 coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
2 coeff
```

```
Out[26]:
```

Co-efficient	
id	0.075857

```
In [27]: 1 prediction =lr.predict(x_test)  
2 plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x1801b37ca00>
```



## ACCURACY

```
In [28]: 1 lr.score(x_test,y_test)
```

```
Out[28]: -0.09486822413919249
```

```
In [29]: 1 lr.score(x_train,y_train)
```

```
Out[29]: 0.0009583837933925254
```

## Ridge and Lasso

```
In [30]: 1 from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: 1 rr=Ridge(alpha=10)  
2 rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: 1 rr.score(x_test,y_test)
```

```
Out[32]: -0.09485306914926017
```

```
In [33]: 1 rr.score(x_train,y_train)
```

```
Out[33]: 0.0009583793573364474
```

```
In [34]: 1 la=Lasso(alpha=10)
2 la.fit(x_train,y_train)
```

```
Out[34]: Lasso(alpha=10)
```

```
In [35]: 1 la.score(x_train,y_train)
```

```
Out[35]: 0.0007600890445066399
```

## Accuracy(Lasso)

```
In [36]: 1 la.score(x_test,y_test)
```

```
Out[36]: -0.09226892618842597
```

```
In [37]: 1 from sklearn.linear_model import ElasticNet
2 en=ElasticNet()
3 en.fit(x_train,y_train)
```

```
Out[37]: ElasticNet()
```

```
In [38]: 1 en.coef_
```

```
Out[38]: array([0.07400431])
```

```
In [39]: 1 en.intercept_
```

```
Out[39]: -2077313.9320505918
```

```
In [40]: 1 prediction=en.predict(x_test)
```

```
In [41]: 1 en.score(x_test,y_test)
```

```
Out[41]: -0.09469779583781057
```

## Evaluation Metrics

```
In [42]: 1 from sklearn import metrics
2 print(metrics.mean_absolute_error(y_test,prediction))
3 print(metrics.mean_squared_error(y_test,prediction))
4 print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

```
23.697995156078832
```

```
759.97684010734
```

```
27.5676774521783
```

# Logistic Regression

```
In [43]: 1 from sklearn.linear_model import LogisticRegression
```

```
In [44]: 1 feature_matrix=df[['lon', 'lat', 'elevation']]
2 target_vector=df[ 'id']
```

```
In [45]: 1 feature_matrix.shape
```

```
Out[45]: (24, 3)
```

```
In [46]: 1 target_vector.shape
```

```
Out[46]: (24,)
```

```
In [47]: 1 from sklearn.preprocessing import StandardScaler
```

```
In [48]: 1 fs=StandardScaler().fit_transform(feature_matrix)
```

```
In [49]: 1 logr=LogisticRegression(max_iter=10000)
2 logr.fit(fs,target_vector)
```

```
Out[49]: LogisticRegression(max_iter=10000)
```

```
In [50]: 1 observation=[[1,2,3]]
```

```
In [51]: 1 prediction=logr.predict(observation)
2 print(prediction)
```

```
[28079060]
```

```
In [52]: 1 logr.classes_
```

```
Out[52]: array([28079004, 28079008, 28079011, 28079016, 28079017, 28079018,
28079024, 28079027, 28079035, 28079036, 28079038, 28079039,
28079040, 28079047, 28079048, 28079049, 28079050, 28079054,
28079055, 28079056, 28079057, 28079058, 28079059, 28079060],
dtype=int64)
```

```
In [53]: 1 logr.score(fs,target_vector)
```

```
Out[53]: 0.625
```

```
In [54]: 1 logr.predict_proba(observation)[0][0]
```

```
Out[54]: 0.00267009423713907
```

```
In [55]: 1 logr.predict_proba(observation)
```

```
Out[55]: array([[2.67009424e-03, 1.31718443e-02, 9.28886937e-02, 7.05390759e-02,
8.34704249e-05, 8.61239187e-04, 1.92424615e-03, 9.42530382e-03,
6.74552285e-03, 2.41930772e-02, 3.99013815e-02, 3.17626697e-02,
1.03584997e-02, 3.64133186e-04, 2.15799332e-02, 8.44889971e-03,
1.94987201e-01, 1.02771134e-03, 5.74880974e-03, 2.41257564e-04,
1.88161280e-01, 2.7774487e-03, 4.15162302e-02, 2.30621680e-01]])
```

# Random Forest

```
In [56]: 1 from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: 1 rfc=RandomForestClassifier()
        2 rfc.fit(x_train,y_train)
```

Out[57]: RandomForestClassifier()

```
In [58]: 1 parameters={'max_depth':[1,2,3,4,5],
        2               'min_samples_leaf':[5,10,15,20,25],
        3               'n_estimators':[10,20,30,40,50]}
        4 }
```

```
In [59]: 1 from sklearn.model_selection import GridSearchCV
        2 grid_search =GridSearchCV(estimator=rfc,param_grid=parameters,cv=2,scoring="accuracy")
        3 grid_search.fit(x_train,y_train)
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\model\_selection\\_split.py:666: UserWarning: The least populated class in y has only 1 members, which is less than n\_splits=2.  
warnings.warn(("The least populated class in y has only %d"

Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),  
param\_grid={'max\_depth': [1, 2, 3, 4, 5],  
'min\_samples\_leaf': [5, 10, 15, 20, 25],  
'n\_estimators': [10, 20, 30, 40, 50]},  
scoring='accuracy')

```
In [60]: 1 grid_search.best_score_
```

Out[60]: 0.125

```
In [61]: 1 rfc_best=grid_search.best_estimator_
```

```
In [66]: rn.tree import plot_tree      1
          (figsize=(80,40))            2
          rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d','e','f','g','h','i','j','k','l','m','n','o','p','q','r','s','t','u','v','w','x'],filled=True)  3

Out[66]: [Text(2232.0, 1087.2, 'gini = 0.852\nsamples = 10\nvalue = [0, 2, 0, 3, 1, 0, 4, 2, 1, 0, 0, 1, 0, 1\n1]\nnclass = g')]
```

gini = 0.852  
samples = 10  
value = [0, 2, 0, 3, 1, 0, 4, 2, 1, 0, 0, 1, 0, 1  
1]  
class = g

## Conclusion

### Accuracy

*Linear Regression:*0.0009583837933925254

*Ridge Regression:*0.0009583793573364474 ¶

*Lasso Regression:*0.09226892618842597

*ElasticNet Regression:*0.09469779583781057

*Logistic Regression:0.625*

*Random Forest:0.125*

**Logistic Regression is suitable for this dataset**